



See the light - agile, industrial strength, rapid web application development made easy

The Grails Framework - Reference Documentation

Authors: Graeme Rocher, Peter Ledbrook, Marc Palmer, Jeff Brown, Luke Daley, Burt Beckwith

Version: null

Table of Contents

1 Introduction

1.1 What's new in Grails 2.0?

1.1.1 Development Environment Features

1.1.2 Core Features

1.1.3 Web Features

1.1.4 Persistence Features

1.1.5 Testing Features

2 Getting Started

2.1 Downloading and Installing

2.2 Upgrading from previous versions of Grails

2.3 Creating an Application

2.4 A Hello World Example

2.5 Using Interactive Mode

2.6 Getting Set Up in an IDE

2.7 Convention over Configuration

2.8 Running an Application

2.9 Testing an Application

2.10 Deploying an Application

2.11 Supported Java EE Containers

2.12 Generating an Application

2.13 Creating Artefacts

3 Configuration

3.1 Basic Configuration

3.1.1 Built in options

3.1.2 Logging

3.1.3 GORM

3.2 Environments

3.3 The DataSource

3.3.1 DataSources and Environments

3.3.2 JNDI DataSources

3.3.3 Automatic Database Migration

3.3.4 Transaction-aware DataSource Proxy

- 3.3.5** Database Console
 - 3.3.6** Multiple Datasources
 - 3.4** Externalized Configuration
 - 3.5** Versioning
 - 3.6** Project Documentation
 - 3.7** Dependency Resolution
 - 3.7.1** Configurations and Dependencies
 - 3.7.2** Dependency Repositories
 - 3.7.3** Debugging Resolution
 - 3.7.4** Inherited Dependencies
 - 3.7.5** Providing Default Dependencies
 - 3.7.6** Dependency Reports
 - 3.7.7** Plugin JAR Dependencies
 - 3.7.8** Maven Integration
 - 3.7.9** Deploying to a Maven Repository
 - 3.7.10** Plugin Dependencies
- 4** The Command Line
 - 4.1** Interactive Mode
 - 4.2** Creating Gant Scripts
 - 4.3** Re-using Grails scripts
 - 4.4** Hooking into Events
 - 4.5** Customising the build
 - 4.6** Ant and Maven
- 5** Object Relational Mapping (GORM)
 - 5.1** Quick Start Guide
 - 5.1.1** Basic CRUD
 - 5.2** Domain Modelling in GORM
 - 5.2.1** Association in GORM
 - 5.2.1.1** Many-to-one and one-to-one
 - 5.2.1.2** One-to-many
 - 5.2.1.3** Many-to-many
 - 5.2.1.4** Basic Collection Types
 - 5.2.2** Composition in GORM
 - 5.2.3** Inheritance in GORM
 - 5.2.4** Sets, Lists and Maps
 - 5.3** Persistence Basics
 - 5.3.1** Saving and Updating

5.3.2 Deleting Objects

5.3.3 Understanding Cascading Updates and Deletes

5.3.4 Eager and Lazy Fetching

5.3.5 Pessimistic and Optimistic Locking

5.3.6 Modification Checking

5.4 Querying with GORM

5.4.1 Dynamic Finders

5.4.2 Criteria

5.4.3 Hibernate Query Language (HQL)

5.5 Advanced GORM Features

5.5.1 Events and Auto Timestamping

5.5.2 Custom ORM Mapping

5.5.2.1 Table and Column Names

5.5.2.2 Caching Strategy

5.5.2.3 Inheritance Strategies

5.5.2.4 Custom Database Identity

5.5.2.5 Composite Primary Keys

5.5.2.6 Database Indices

5.5.2.7 Optimistic Locking and Versioning

5.5.2.8 Eager and Lazy Fetching

5.5.2.9 Custom Cascade Behaviour

5.5.2.10 Custom Hibernate Types

5.5.2.11 Derived Properties

5.5.2.12 Custom Naming Strategy

5.5.3 Default Sort Order

5.6 Programmatic Transactions

5.7 GORM and Constraints

6 The Web Layer

6.1 Controllers

6.1.1 Understanding Controllers and Actions

6.1.2 Controllers and Scopes

6.1.3 Models and Views

6.1.4 Redirects and Chaining

6.1.5 Controller Interceptors

6.1.6 Data Binding

6.1.7 XML and JSON Responses

6.1.8 More on JSONBuilder

- 6.1.9** Uploading Files
- 6.1.10** Command Objects
- 6.1.11** Handling Duplicate Form Submissions
- 6.1.12** Simple Type Converters
- 6.1.13** Asynchronous Request Processing
- 6.2** Groovy Server Pages
 - 6.2.1** GSP Basics
 - 6.2.1.1** Variables and Scopes
 - 6.2.1.2** Logic and Iteration
 - 6.2.1.3** Page Directives
 - 6.2.1.4** Expressions
 - 6.2.2** GSP Tags
 - 6.2.2.1** Variables and Scopes
 - 6.2.2.2** Logic and Iteration
 - 6.2.2.3** Search and Filtering
 - 6.2.2.4** Links and Resources
 - 6.2.2.5** Forms and Fields
 - 6.2.2.6** Tags as Method Calls
 - 6.2.3** Views and Templates
 - 6.2.4** Layouts with Sitemesh
 - 6.2.5** Static Resources
 - 6.2.5.1** Including resources using the resource tags
 - 6.2.5.2** Other resource tags
 - 6.2.5.3** Declaring resources
 - 6.2.5.4** Overriding plugin resources
 - 6.2.5.5** Optimizing your resources
 - 6.2.5.6** Debugging
 - 6.2.5.7** Preventing processing of resources
 - 6.2.5.8** Other Resources-aware plugins
 - 6.2.6** Sitemesh Content Blocks
 - 6.2.7** Making Changes to a Deployed Application
 - 6.2.8** GSP Debugging
- 6.3** Tag Libraries
 - 6.3.1** Variables and Scopes
 - 6.3.2** Simple Tags
 - 6.3.3** Logical Tags
 - 6.3.4** Iterative Tags

- 6.3.5** Tag Namespaces
- 6.3.6** Using JSP Tag Libraries
- 6.3.7** Tag return value
- 6.4** URL Mappings
 - 6.4.1** Mapping to Controllers and Actions
 - 6.4.2** Embedded Variables
 - 6.4.3** Mapping to Views
 - 6.4.4** Mapping to Response Codes
 - 6.4.5** Mapping to HTTP methods
 - 6.4.6** Mapping Wildcards
 - 6.4.7** Automatic Link Re-Writing
 - 6.4.8** Applying Constraints
 - 6.4.9** Named URL Mappings
 - 6.4.10** Customizing URL Formats
- 6.5** Web Flow
 - 6.5.1** Start and End States
 - 6.5.2** Action States and View States
 - 6.5.3** Flow Execution Events
 - 6.5.4** Flow Scopes
 - 6.5.5** Data Binding and Validation
 - 6.5.6** Subflows and Conversations
- 6.6** Filters
 - 6.6.1** Applying Filters
 - 6.6.2** Filter Types
 - 6.6.3** Variables and Scopes
 - 6.6.4** Filter Dependencies
- 6.7** Ajax
 - 6.7.1** Ajax Support
 - 6.7.1.1** Remoting Linking
 - 6.7.1.2** Updating Content
 - 6.7.1.3** Remote Form Submission
 - 6.7.1.4** Ajax Events
 - 6.7.2** Ajax with Prototype
 - 6.7.3** Ajax with Dojo
 - 6.7.4** Ajax with GWT
 - 6.7.5** Ajax on the Server
- 6.8** Content Negotiation

- 7 Validation**
 - 7.1 Declaring Constraints**
 - 7.2 Validating Constraints**
 - 7.3 Validation on the Client**
 - 7.4 Validation and Internationalization**
 - 7.5 Validation Non Domain and Command Object Classes**
- 8 The Service Layer**
 - 8.1 Declarative Transactions**
 - 8.1.1 Transactions Rollback and the Session**
 - 8.2 Scoped Services**
 - 8.3 Dependency Injection and Services**
 - 8.4 Using Services from Java**
- 9 Testing**
 - 9.1 Unit Testing**
 - 9.1.1 Unit Testing Controllers**
 - 9.1.2 Unit Testing Tag Libraries**
 - 9.1.3 Unit Testing Domains**
 - 9.1.4 Unit Testing Filters**
 - 9.1.5 Unit Testing URL Mappings**
 - 9.1.6 Mocking Collaborators**
 - 9.2 Integration Testing**
 - 9.3 Functional Testing**
- 10 Internationalization**
 - 10.1 Understanding Message Bundles**
 - 10.2 Changing Locales**
 - 10.3 Reading Messages**
 - 10.4 Scaffolding and i18n**
- 11 Security**
 - 11.1 Securing Against Attacks**
 - 11.2 Encoding and Decoding Objects**
 - 11.3 Authentication**
 - 11.4 Security Plugins**
 - 11.4.1 Spring Security**
 - 11.4.2 Shiro**
- 12 Plugins**
 - 12.1 Creating and Installing Plugins**
 - 12.2 Plugin Repositories**

- 12.3** Understanding a Plugin's Structure
- 12.4** Providing Basic Artefacts
- 12.5** Evaluating Conventions
- 12.6** Hooking into Build Events
- 12.7** Hooking into Runtime Configuration
- 12.8** Adding Dynamic Methods at Runtime
- 12.9** Participating in Auto Reload Events
- 12.10** Understanding Plugin Load Order
- 12.11** The Artefact API
 - 12.11.1** Asking About Available Artefacts
 - 12.11.2** Adding Your Own Artefact Types
- 12.12** Binary Plugins
- 13** Web Services
 - 13.1** REST
 - 13.2** SOAP
 - 13.3** RSS and Atom
- 14** Grails and Spring
 - 14.1** The Underpinnings of Grails
 - 14.2** Configuring Additional Beans
 - 14.3** Runtime Spring with the Beans DSL
 - 14.4** The BeanBuilder DSL Explained
 - 14.5** Property Placeholder Configuration
 - 14.6** Property Override Configuration
- 15** Grails and Hibernate
 - 15.1** Using Hibernate XML Mapping Files
 - 15.2** Mapping with Hibernate Annotations
 - 15.3** Adding Constraints
- 16** Scaffolding
- 17** Deployment

1 Introduction

Java WebJavaWebDon't Repeat YourselfDRY

RailsDjangoTurboGearsWebGrailsJavaWebGrailsSpringHibernateJava

Grailsfull-stackplug-inWeb

*[Hibernate](#)-(ORM) *Groovy Server Pages (GSP) *[Spring](#) MVC *[Gant](#) *[Tomcat](#) *Spring *SpringMessag

GroovyDomain Specific LanguageDSL

GrailsWeb


1.1 What's new in Grails 2.0?

This section covers the new features that are present in 2.0 and is broken down into sections covering the persistence enhancements and improvements in testing. Note there are many more small enhancements cover some of the highlights.

1.1.1 Development Environment Features


Interactive Mode and Console Enhancements

Grails 2.0 features brand new console output that is more concise and user friendly to consume. An example tests can be seen below:

A screenshot of a terminal window titled 'bookstore — java — 74x22'. The terminal shows the output of running unit tests. It starts with 'Running 2 unit tests... 1 of 2'. Then, a failure is reported: 'Failure: testFindBook(bookstore.BookControllerTests)'. The error message is 'Assertion failed: assert model.book.title == "The Shining"', followed by a comparison table showing 'false' instead of 'The Shining'. The stack trace points to 'bookstore.BookControllerTests.testFindBook'. The final output is 'Completed 2 unit tests, 1 failed in 103ms' and 'Tests FAILED - view reports in target/test-reports'. The prompt 'grails>' is visible at the bottom.

```
bash java bash bash
| Running 2 unit tests... 1 of 2
| Failure: testFindBook(bookstore.BookControllerTests)
| Assertion failed:
|
| assert model.book.title == "The Shining"
|      |      |      |
|      |      |      false
|      |      The Stand
|      bookstore.Book : 1
| [book:bookstore.Book : 1]
|
| at bookstore.BookControllerTests.testFindBook(BookControllerTests.
| groovy:23)
| Completed 2 unit tests, 1 failed in 103ms
| Tests FAILED - view reports in target/test-reports
grails>
```

In general Grails makes its best effort to display update information on a single line and only present the that while in previous versions of Grails the [war](#) command produced many lines of output, in Grails 2.0 on

A screenshot of a terminal window titled 'bookstore — java — 74x22'. The window has four tabs labeled 'bash', 'java', 'bash', and 'bash'. The terminal output shows a single line: 'Done creating WAR target/bookstore-0.1.war'. Below this line, the prompt 'grails>' is visible with a cursor. The terminal background is dark, and the text is light-colored.

In addition simply typing 'grails' at the command line activates the new interactive mode which features `⌘` keeps the JVM running to ensure commands execute much quicker than otherwise

bookstore

1 Introduction - Reference Documentation

Authors:**Version:** 0.1

1 Introduction

This is an example documentation template. The syntax format is similar to [Textile](#)

You can apply formatting such as **bold**, *italic* and `code`. Bullets are possible too:

- Bullet 1
- Bullet 2

As well as numbered lists:

1. Number 1
2. Number 2

The documentation also handles links to [guide items](#) as well as [reference](#)

Quick Reference (hide)

Items

See the section on the [documentation engine](#) for more usage info.

Use a TOC for Project Docs

The old documentation engine relied on you putting section numbers into the gdoc filenames. Although difficult to restructure your user guide by inserting new chapters and sections. In addition, any such result resulted in breaking changes to the URLs.

You can now use logical names for your gdoc files and define the structure and section titles in a YAMl the section on the [documentation engine](#). The logical names appear in the URLs, so as long as you don't remain the same no matter how much restructuring or changing of titles you do.

Grails 2.0 even provides a [migrate-docs](#) command to aid you in migrating existing gdoc user guides.

Enhanced Error Reporting and Diagnosis

Error reporting and problem diagnosis has been greatly improved with a new errors view that analyses problem areas in your code:

Error 500: Internal Server Error

URI: /bookstore/book/find

Class: groovy.lang.MissingPropertyException

Message: No such property: titl for class: bookstore.BookService

Around line 6 of *grails-app/services/bookstore/BookService.groovy*

```
3: class BookService {
4:
5:     Book findByTitle(String title) {
6:         Book.findByTitle(titl)
7:     }
8: }
```

Around line 10 of *grails-app/controllers/bookstore/BookController.groovy*

```
7:     def bookService
8:     def find() {
9:
10:         def b = bookService.findByTitle(params.title)
11:
12:         [book:b]
13:     }
```

Trace

Line	Method
->> 6	findByTitle in BookService.groovy

10	find in BookController.groovy
886	runTask . . in java.util.concurrent.ThreadPoolExecutor\$Worker
908	run in ''
680	run in java.lang.Thread


In addition stack trace filtering has been further enhanced to display only relevant trace information:

Line	Method
->> 9	getValue in Book.groovy

7	getBookValue in BookService.groovy
886	runTask . . in ThreadPoolExecutor.java
908	run in ''
662	run in Thread.java

H2 Database and Console

Grails 2.0 now uses the H2 database instead of HSQLDB, and enables the H2 database console in development that the in-memory database can be easily queried from the browser:



Plugin Usage Tracking

To enhance community awareness of the most popular plugins an opt-in plugin usage tracking system will be introduced. This will help drive the roadmap and increase support of key plugins while reducing the need to support less popular ones. We encourage all plugin developers to participate in providing feedback to the plugin community on which plugins are most popular.

This will help drive the roadmap and increase support of key plugins while reducing the need to support less popular ones. We encourage all plugin developers to participate in providing feedback to the plugin community on which plugins are most popular.

Dependency Resolution Improvements

There are numerous improvements to dependency resolution handling via Ivy including:

- Grails now makes a best effort to cache the previous resolve and avoid resolving again unless you change the dependencies
- Plugins dependencies now appear in the dependency report generated by `grails dependency-report`
- Plugins published with the release plugin now publish their transitive plugin dependencies in the generated POM
- It is now possible to customize the ivy cache directory via `BuildConfig.groovy`

```
grails.project.dependency.resolution = {
    cacheDir "target/ivy-cache"
}
```

- It is now possible to completely disable resolution from inherited repositories (repositories defined by

```
grails.project.dependency.resolution = {
  repositories {
    inherits false // Whether to inherit repository definitions from plugins
    ...
  }
  ...
}
```

- It is now possible to easily disable checksum validation errors:

```
grails.project.dependency.resolution = {
  checksums false // whether to verify checksums or not
}
```

1.1.2 Core Features

Binary Plugins

Grails plugins can now be packaged as JAR files and published to standard maven repositories. This evolution (with resources plugin 1.0.1). See the section on [Binary plugins](#) for more information.

Groovy 1.8

Grails 2.0 comes with Groovy 1.8 which includes many new [features and enhancements](#)

Spring 3.1 Profile Support

Grails' existing environment support has been bridged into the Spring 3.1 profile support. For example, an environment called "production", a Spring profile of "production" is activated so that you can use Spring beans for a specific profile.

1.1.3 Web Features

Controller Actions as Methods

It is now possible to define controller actions as methods instead of using closures as in previous versions. The preferred way of expressing an action. For example:

```
// action as a method
def index() {
}
// action as a closure
def index = {
}
```

Binding Primitive Method Action Arguments

It is now possible to bind form parameters to action arguments where the name of the form element matches the following form:

```
<g:form name="myForm" action="save">
  <input name="name" />
  <input name="age" />
</g:form>
```

You can define an action that declares arguments for each input and automatically converts the parameters

```
def save(String name, int age) {
  // remaining
}
```

Static Resource Abstraction

A new [static resource abstraction](#) is included that allows declarative handling of JavaScript, CSS and ordering, compression, caching and gzip handling.

Servlet 3.0 Async Features

Grails now supports Servlet 3.0 including the Asynchronous programming model defined by the specification

```
def index() {
  def ctx = startAsync()
  ctx.start {
    new Book(title:"The Stand").save()
    render template:"books", model:[books:Book.list()]
    ctx.complete()
  }
}
```

Link Generation API

A general purpose LinkGenerator class is now available that is usable anywhere within a Grails application or controller. For example if you need to generate links in a service or an asynchronous background job out

```
LinkGenerator grailsLinkGenerator

def generateLink() {
  grailsLinkGenerator.link(controller:"book", action:"list")
}
```

Page Rendering API

Like the LinkGenerator the new PageRenderer can be used to render GSP pages outside the scheduled job or web service. The PageRenderer class features a very similar API to the render method


```
grails.gsp.PageRenderer groovyPageRenderer

void welcomeUser(User user) {
    def contents = groovyPageRenderer.render(view: "/emails/welcomeLetter", model:
    sendEmail {
        to user.email
        body contents
    }
}
```

The PageRenderer service also allows you to pre-process GSPs into HTML templates:

```
new File("/path/to/welcome.html").withWriter { w ->
    groovyPageRenderer.renderTo(view: "/page/content", w)
}
```

Filter Exclusions

Filters may now express controller, action and uri exclusions to offer more options for expressing to which filters are applied.

```
filter1(actionExclude: 'log*') {
    before = {
        // ...
    }
}
filter2(controllerExclude: 'auth') {
    before = {
        // ...
    }
}
filter3(uriExclude: '/secure*') {
    before = {
        // ...
    }
}
```

Performance Improvements

Performance of GSP page rendering has once again been improved by optimizing the GSP compiler to inli

HTML5 Scaffolding

There is a new HTML5-based scaffolding UI:

[Home](#)[Book List](#)

Create Book

Title *

[Create](#)

jQuery by Default

The jQuery plugin is now the default JavaScript library installed into a Grails application. For backward compatibility, the Prototype plugin is also available. Refer to the [documentation](#) on the Prototype plugin for installation instructions.

1.1.4 Persistence Features

The GORM API

The GORM API has been formalized into a set of classes (`GormStaticApi`, `GormInstanceApi`) statically wired into every domain class at the byte code level. The result is better code completion for IDEs and the potential for more GORM implementations for other types of data stores.

New `findOrCreate` and `findOrCreateSave` Methods

Domain classes have support for the `findOrCreateWhere`, `findOrCreateSaveWhere`, `findOrCreateBy` and `findOrCreateSaveBy` just like `findWhere` and `findBy` methods except that they should never return null. If a matching instance is found, the instance is returned. If no matching instance is found, a new instance is created, populated with values represented in the query parameters and returned. For `findOrCreateSaveBy`, the instance is saved before being returned.

```
def book = Book.findOrCreateWhere(author: 'Douglas Adams', title: "The Hitchhiker's Guide to the Galaxy")
def book = Book.findOrCreateSaveWhere(author: 'Daniel Suarez', title: 'Daemon')
def book = Book.findOrCreateByAuthorAndTitle('Daniel Suarez', 'Daemon')
def book = Book.findOrCreateSaveByAuthorAndTitle('Daniel Suarez', 'Daemon')
```

Abstract Inheritance

GORM now supports abstract inheritance trees which means you can define queries and associations linking

```

abstract class Media {
    String title
    ...
}
class Book extends Media {
}
class Album extends Media {
}
class Account {
    static hasMany = [purchasedMedia:Media]
}

..

def allMedia = Media.list()

```

Multiple Data Sources Support

It is now possible to define multiple datasources in `DataSource.groovy` and declare one or more default:

```

class ZipCode {
    String code
    static mapping = {
        datasource 'ZIP_CODES'
    }
}

```

If multiple datasources are specified for a domain then you can use the name of a particular datasource GORM method:

```

def zipCode = ZipCode.auditing.get(42)

```

For more information see the section on [Multiple Data Sources](#) in the user guide.

Database Migrations

A new [database migration plugin](#) has been designed and built for Grails 2.0 allowing you to apply migrations and diff your domain model with the current state of the database.

Database Reverse Engineering

A new [database reverse engineering](#) plugin has been designed and built for Grails 2.0 that allows you to reverse engineer an existing database schema.

Hibernate 3.6

Grails 2.0 is now built on Hibernate 3.6

Bag Collections

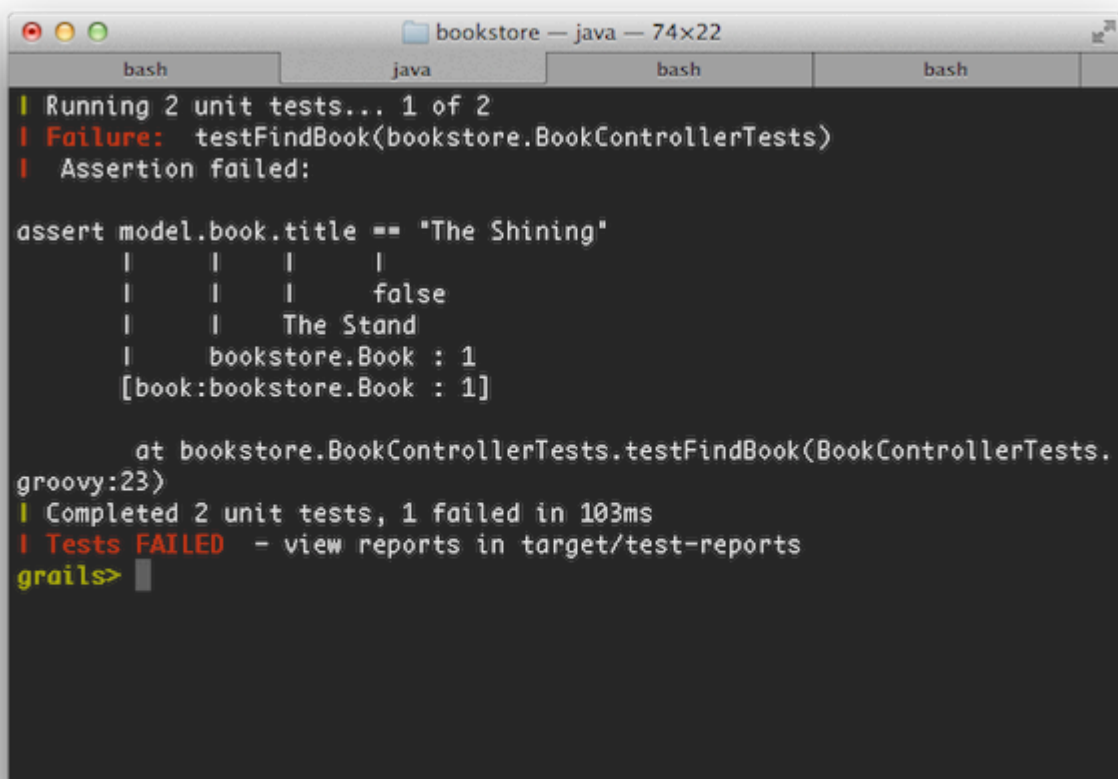
You can now use Hibernate [Bags](#) for mapped collections to avoid the memory and performance issues of `Set` uniqueness or `List` order.

For more information see the section on [Sets, Lists and Maps](#) in the user guide.

1.1.5 Testing Features

New Unit Testing Console Output

Test output from the `test-app` command has been improved:



```
bookstore — java — 74x22
bash java bash bash
| Running 2 unit tests... 1 of 2
| Failure: testFindBook(bookstore.BookControllerTests)
| Assertion failed:
|
| assert model.book.title == "The Shining"
|      |      |      |
|      |      |      false
|      |      The Stand
|      bookstore.Book : 1
|      [book:bookstore.Book : 1]
|
|      at bookstore.BookControllerTests.testFindBook(BookControllerTests.
groovy:23)
| Completed 2 unit tests, 1 failed in 103ms
| Tests FAILED - view reports in target/test-reports
grails>
```

New Unit Testing API

There is a new unit testing API based on mixins that supports JUnit 3, 4 and Spock style tests (with Spock

```
import grails.test.mixin.TestFor
@TestFor(SimpleController)
class SimpleControllerTests {
    void testIndex() {
        controller.home()

        assert view == "/simple/homePage"
        assert model.title == "Hello World"
    }
}
```

The [documentation on testing](#) has also been re-written around this new framework.

Unit Testing GORM

A new in-memory GORM implementation is present that supports many more features of the GORM API named queries and other previously unsupported methods possible.

Faster Unit Testing with Interactive Mode

The new interactive mode (activated by typing 'grails') greatly improves the execution time of running unit

Unit Test Scaffolding

A unit test is now generated for scaffolded controllers.

2 Getting Started

2.1 Downloading and Installing

The first step to getting up and running with Grails is to install the distribution. To do so follow these steps

- [Download](#) a binary distribution of Grails and extract the resulting zip file to a location of your choice
- Set the `GRAILS_HOME` environment variable to the location where you extracted the zip
 - On Unix/Linux based systems this is typically a matter of adding something like `GRAILS_HOME=/path/to/grails` to your profile
 - On Windows this is typically a matter of setting an environment variable under My Computer > System > Environment Variables
- Then add the bin directory to your `PATH` variable:
 - On Unix/Linux based systems this can be done by adding `export PATH="$PATH:$GRAILS_HOME/bin"`
 - On Windows this is done by modifying the Path environment variable under My Computer > System > Environment Variables

If Grails is working correctly you should now be able to type `grails -version` in the terminal window

```
Grails version: 2.0.0
```

2.2 Upgrading from previous versions of Grails

Although the Grails development team have tried to keep breakages to a minimum there are a number of things you need to be aware of when upgrading from Grails 1.0.x, 1.1.x, 1.2.x, or 1.3.x applications to Grails 2.0. The major changes are described in detail below.

Upgrading from Grails 1.3.x

HSQldb Has Been Replaced With H2

HSQldb is still bundled with Grails but is not configured as a default runtime dependency. Upgrade your application by replacing references in `DataSource.groovy` with H2 references or adding HSQldb as a runtime dependency for the application.

If you want to run an application with different versions of Grails, it's simplest to add HSQldb as a runtime dependency in `BuildConfig.groovy`:

```

grails.project.dependency.resolution = {
    inherits("global") {
    }
    repositories {
        grailsPlugins()
        grailsHome()
        grailsCentral()
    }
    dependencies {
        // Add HSQLDB as a runtime dependency
        runtime 'hsqldb:hsqldb:1.8.0.10'
    }
}

```

A default DataSource.groovy which is compatible with H2 looks like this:

```

dataSource {
    driverClassName = "org.h2.Driver"
    username = "sa"
    password = ""
}
// environment specific settings
environments {
    development {
        dataSource {
            dbCreate = "create-drop" // one of 'create', 'create-drop','update'
            url = "jdbc:h2:mem:devDb"
        }
    }
    test {
        dataSource {
            dbCreate = "update"
            url = "jdbc:h2:mem:testDb"
        }
    }
    production {
        dataSource {
            dbCreate = "update"
            url = "jdbc:h2:prodDb"
        }
    }
}

```

Another significant difference between H2 and HSQLDB is in the handling of `byte[]` domain class properties. In HSQLDB, the `byte[]` domain class property is large and so you typically don't need to specify a maximum size. But H2 defaults to a maximum size of 255 bytes for a `byte[]` column in the database, the saves are likely to fail because of this. The easy fix is to add a `maxSize` constraint to the `byte[]` property.

```

class MyDomain {
    byte[] data

    static constraints = {
        data maxSize: 1024 * 1024 * 2 // 2MB
    }
}

```

This constraint influences schema generation, so in the above example H2 will have the `data` column set to `VARBINARY(2048)`.

Abstract Inheritance Changes

In previous versions of Grails abstract classes in `grails-app/domain` were not treated as persistent and this had a significant impact on upgrading your application. For example consider the following domain model in a C

```
abstract class Sellable {  
    }  
class Book extends Sellable {  
    }  
}
```

In Grails 1.3.x you would get a `BOOK` table and the properties from the `Sellable` class would be stored in the `BOOK` table. In Grails 2.0.x you will get `SELLABLE` table and the default table-per-hierarchy inheritance rules apply with `SELLABLE` table.

You have two options when upgrading in this scenario:

1. Move the abstract `Sellable` class into the `src/groovy` package. If the `Sellable` class is in the `src/groovy` package it will be regarded as non-persistent.
2. Use the [database migration](#) plugin to apply the appropriate changes to the database (typically renaming the `BOOK` table to `SELLABLE` and updating the inheritance tree)

Criteria Queries Default to INNER JOIN

The previous default of `LEFT JOIN` for criteria queries across associations is now `INNER JOIN`.

Logging By Convention Changes

The packages that you should use for Grails artifacts have mostly changed. In particular:

- `service` -> `services`
- `controller` -> `controllers`
- `tagLib` -> `taglib` (case change)
- `bootstrap` -> `conf`
- `dataSource` -> `conf`

You can find out more about logging by convention in the [main part](#) of the user guide, under "Configuring Logging".

jQuery Replaces Prototype

The Prototype Javascript library has been removed from Grails core and now new Grails applications default to jQuery. This will only impact you if you are using Prototype with the adaptive AJAX tags in your application. Those tags will break as soon as you upgrade.

To resolve this issue, simply install the [jQuery plugin](#) in your application. You can also rename the `web-app/js/prototype` directory if you want.

Access Control and Resources

The Resources plugin is a great new feature of Grails, but you do need to be aware that it adds an extra control in your application, this may mean that the static resources require an authenticated user to load t account of the `/static` URL.

Controller Public Methods

As of Grails 2.0, public methods of controllers are now treated as actions in addition to actions define relying on the use of methods for privacy controls or as helper methods then this could result in unexpect should mark all methods of your application that are not to be exposed as actions as `private` methods.

The redirect Method

The [redirect](#) method no longer commits the response. The result of this is code that relies of this behavior v

```
redirect action: "next"
if (response.committed) {
    // do something
}
```

In this case in Grails 1.3.x and below the `response.committed` property would return true and the i is no longer the case and you should instead use the new `isRedirected()` method of the request ob

```
redirect action: "next"
if (request.redirected) {
    // do something
}
```

Another side-effect of the changes to the redirect method is that it now always uses the `grails.serve` Previous versions of Grails included default values for all the environments, but when upgrading to Grai break redirection. So, we recommend you remove the development and test settings for `grails.serve` appropriate for your application.

Content Negotiation

As of Grails 2.0 the [withFormat](#) method of controllers no longer takes into account the request content t header), but instead deals exclusively with the response content type (dictated by the ACCEPT header or application has code that relies on reading XML from the request using `withFormat` this will no longer

```
def processBook() {
    withFormat {
        xml {
            // read request XML
        }
        html {
            // read request parameters
        }
    }
}
```

Instead you use the `withFormat` method provided on the request object:

```
def processBook() {
  request.withFormat {
    xml {
      // read request XML
    }
    html {
      // read request parameters
    }
  }
}
```

Command Line Output

Ant output is now hidden by default to keep the noise in the terminal to a minimum. That means if you want to communicate messages to the user, we recommend switching to an alternative mechanism.

For status related messages, you can use the event system:

```
event "StatusUpdate", [ "Some message" ]
event "StatusFinal", [ "Some message" ]
event "StatusError", [ "Some message" ]
```

For more control you can use the `grailsConsole` script variable, which gives you access to an instance of `GrailsConsole`. You can log information messages with `log()` or `info()`, errors and warnings with `error()` and `warn()`, and user input with `userInput()`.

Updated Underlying APIs

Grails 2.0 contains updated dependencies including Servlet 3.0, Tomcat 7, Spring 3.1, Hibernate 3.6 and plugins and applications that depend on earlier versions of these APIs may no longer work. The `HttpServletRequest` interface includes new methods, so if a plugin implements this interface for Servlet 2.4, said plugin will break. The same can be said of any Spring interface.

Removal of release-plugin

The built in `release-plugin` command for releases plugins to the central Grails plugin repository has should be used instead which provides an equivalent `publish-plugin` command.

Removal of Deprecated Classes

The following deprecated classes have been removed: `grails.web.JsonBuilder`, `grails.web.C`

Upgrading from Grails 1.2.x

Plugin Repositories

As of Grails 1.3, Grails no longer natively supports resolving plugins against secured SVN repositories. Grails 1.2 and below has been replaced by one built on [Ivy](#), the upside of which is that you can now resolve repositories as well as regular Grails repositories.

Ivy supports a much richer set of repository resolvers for resolving plugins, including support for WebAssembly in the section on [resolvers](#) in the Ivy docs for all the available options and the section of [plugin repositories](#) in the Ivy docs to configure additional resolvers.

If you still need support for resolving plugins against secured SVN repositories then the [IvySvn](#) project provides support for secured SVN repositories.

Upgrading from Grails 1.1.x

Plugin paths

In Grails 1.1.x typically a `pluginContextPath` variable was used to establish paths to plugin resources.

```
<g:resource dir="${pluginContextPath}/images" file="foo.jpg" />
```

In Grails 1.2 views have been made plugin aware and this is no longer necessary:

```
<g:resource dir="images" file="foo.jpg" />
```

Additionally the above example will no longer link to an application image from a plugin view. To do so correctly you need to specify the context path:

```
<g:resource contextPath="" dir="images" file="foo.jpg" />
```

The same rules apply to the [javascript](#) and [render](#) tags.

Tag and Body return values

Tags no longer return `java.lang.String` instances but instead return a Grails `StreamCharBuffer` class which implements all the same methods as `String` but doesn't extend `String`, so code like this will break:

```
def foo = body()
if (foo instanceof String) {
    // do something
}
```

In these cases you should check for the `java.lang.CharSequence` interface, which both `String` and `StreamCharBuffer` implement:

```
def foo = body()
if (foo instanceof CharSequence) {
    // do something
}
```

New JSONBuilder

There is a new version of JSONBuilder which is semantically different from the one used in earlier application depends on the older semantics you can still use the deprecated implementation by setting Config.groovy:

```
grails.json.legacy.builder=true
```

Validation on Flush

Grails now executes validation routines when the underlying Hibernate session is flushed to ensure that no your constraints (such as a custom validator) executes a query then this can cause an additional flush, resulting in a StackOverflowError. For example:

```
static constraints = {
    author validator: { a ->
        assert a != Book.findByTitle("My Book").author
    }
}
```

The above code can lead to a StackOverflowError in Grails 1.2. The solution is to run the query in a new session (this is recommended in general as doing Hibernate work during flushing can cause other issues):

```
static constraints = {
    author validator: { a ->
        Book.withNewSession {
            assert a != Book.findByTitle("My Book").author
        }
    }
}
```

Upgrading from Grails 1.0.x

Groovy 1.6

Grails 1.1 and above ship with Groovy 1.6 and no longer supports code compiled against Groovy 1.5. If you are using Groovy 1.5 you must recompile it against Groovy 1.6 or higher before using it with Grails 1.1.

Java 5.0

Grails 1.1 now no longer supports JDK 1.4, if you wish to continue using Grails then it is recommended that you are able to upgrade your JDK.

Configuration Changes

- 1) The setting `grails.testing.reports.destDir` has been renamed to `grails.project.test.reports.destDir`.
- 2) The following settings have been moved from `grails-app/conf/BuildConfig.groovy` to `grails-app/conf/BuildConfig.groovy`:

- `grails.config.base.webXml`
- `grails.project.war.file` (renamed from `grails.war.destFile`)
- `grails.war.dependencies`
- `grails.war.copyToWebApp`
- `grails.war.resources`

3) The `grails.war.java5.dependencies` option is no longer supported, since Java 5.0 is now the

4) The use of `jsessionid` (now considered harmful) is disabled by default. If your application requires `jsessionid`, adding the following to `grails-app/conf/Config.groovy`:

```
grails.views.enable.jsessionid=true
```

5) The syntax used to configure Log4j has changed. See the user guide section on [Logging](#) for more information.

Plugin Changes

As of version 1.1, Grails no longer stores plugins inside your `PROJECT_HOME/plugins` directory by default. If you have errors in your application unless you either re-install all your plugins or set `grails-app/conf/BuildConfig.groovy`:

```
grails.project.plugins.dir="./plugins"
```

Script Changes

1) If you were previously using Grails 1.0.3 or below the following syntax is no longer supported for importing scripts:

```
Ant.property(environment:"env")
grailsHome = Ant.antProject.properties."env.GRAILS_HOME"
includeTargets << new File("${grailsHome}/scripts/Bootstrap.groovy")
```

Instead you should use the new `grailsScript` method to import a named script:

```
includeTargets << grailsScript("_GrailsBootstrap")
```

2) Due to an upgrade of Gant all references to the variable `Ant` should be changed to `ant`.

3) The root directory of the project is no longer on the classpath, so loading a resource like this will no longer work:

```
def stream = getClass().classLoader.getResourceAsStream(
    "grails-app/conf/my-config.xml")
```

Instead you should use the Java File APIs with the `basedir` property:

```
new File("${basedir}/grails-app/conf/my-config.xml").withInputStream { stream ->
    // read the file
}
```

Command Line Changes

The `run-app-https` and `run-war-https` commands no longer exist and have been replaced by an `https` flag:

```
grails run-app -https
```

Data Mapping Changes

1) Enum types are now mapped using their String value rather than the ordinal value. You can revert mapping as follows:

```
static mapping = {
    someEnum enumType: "ordinal"
}
```

2) Bidirectional one-to-one associations are now mapped with a single column on the owning side and need to change anything; however you should drop column on the inverse side as it contains duplicate data

REST Support

Incoming XML requests are now no longer automatically parsed. To enable parsing of REST requests you need to add the `parseRequest` argument inside a URL mapping:

```
"/book"(controller: "book", parseRequest: true)
```

Alternatively, you can use the new `resource` argument, which enables parsing by default:

```
"/book"(resource: "book")
```

2.3 Creating an Application

To create a Grails application you first need to familiarize yourself with the usage of the `grails` command in the following manner:

```
grails [command name]
```

Run [create-app](#) to create an application:

```
grails create-app helloworld
```

This will create a new directory inside the current one that contains the project. Navigate to this directory in your terminal by running `cd helloworld`.

2.4 A Hello World Example

To implement the typical "hello world!" example run the [create-controller](#) command:

```
grails create-controller hello
```

This will create a new controller (Refer to the section on [Controllers](#) for more information) in the `grails` directory called `helloworld/HelloController.groovy`.



If no package is specified with `create-controller` script, Grails automatically uses the application package name. This default is configurable with the `grails.project.groupId` attribute.

Controllers are capable of dealing with web requests and to fulfil the "hello world!" use case our implementation will look like this:

```
package helloworld

class HelloController {
    def world() {
        render "Hello World!"
    }
}
```

Job done. Now start-up the container with another new command called [run-app](#):

```
grails run-app
```

This will start-up a server on port 8080 and you should now be able to access your application at `http://localhost:8080/helloworld`.

The result will look something like the following screenshot:

APPLICATION STATUS

App version: 0.1
Grails version: 2.0.0.BUILD-SNAPSHOT
Groovy version: 1.8.3-SNAPSHOT
JVM version: 1.6.0_26
Controllers: 1
Domains: 0
Services: 2
Tag Libraries: 12

INSTALLED PLUGINS

logging - 2.0.0.BUILD-SNAPSHOT

Welcome to Grails

Congratulations, you have successfully started your first Grails application! At the moment this is the default page, feel free to modify it to either redirect to a controller or display whatever content you may choose. Below is a list of controllers that are currently deployed in this application, click on each to execute its default action:

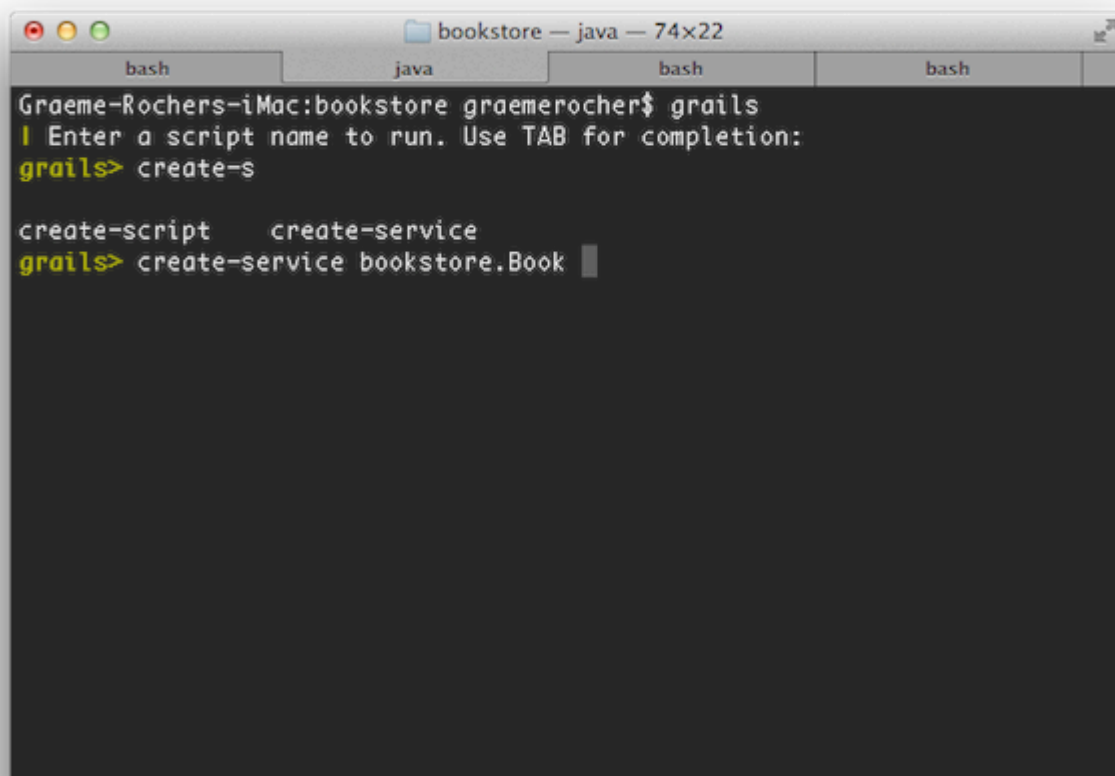
Available Controllers:

- [helloworld.HelloController](#)

This is the Grails intro page which is rendered by the `web-app/index.gsp` file. You will note a link to a controller and clicking on the link to our controller we can see the text "Hello World!" printed to the browser.

2.5 Using Interactive Mode

Grails 2.0 features an interactive mode which makes command execution faster since the JVM doesn't have to start for every command. To use interactive mode simply type 'grails' from the root of any project and use TAB completion to get suggestions. Screenshot below for an example:



```
bookstore — java — 74x22
bash      java      bash      bash
Graeme-Rochers-iMac:bookstore graemerocher$ grails
| Enter a script name to run. Use TAB for completion:
grails> create-s

create-script    create-service
grails> create-service bookstore.Book
```

For more information on the capabilities of interactive mode refer to the section on [Interactive Mode](#) in the Grails User Guide.

2.6 Getting Set Up in an IDE

IntelliJ IDEA

[IntelliJ IDEA](#) and the [JetGroovy](#) plugin offer good support for Groovy and Grails developers. Refer to the [JetBrains website](#) for a feature overview.

To integrate Grails with IntelliJ run the following command to generate appropriate project files:

```
grails integrate-with --intellij
```

Eclipse

We recommend that users of [Eclipse](#) looking to develop Grails application take a look at [SpringSource Tool Suite](#) for Grails including automatic classpath management, a GSP editor and quick access to Grails command overview.

NetBeans

NetBeans provides a Groovy/Grails plugin that automatically recognizes Grails projects and provides the IDE, code completion and integration with the Glassfish server. For an overview of features see the [NetBeans website](#) which was written by the NetBeans team.

TextMate

Since Grails' focus is on simplicity it is often possible to utilize more simple editors and [TextMate](#) on the [Mac OS X](#) bundle available from the [Textmate bundles SVN](#).

To integrate Grails with TextMate run the following command to generate appropriate project files:

```
grails integrate-with --textmate
```

Alternatively TextMate can easily open any project with its command line integration by issuing the following command:

```
mate .
```

2.7 Convention over Configuration

Grails uses "convention over configuration" to configure itself. This typically means that the name and location of configuration files is determined by convention, hence you need to familiarize yourself with the directory structure provided by Grails.

Here is a breakdown and links to the relevant sections:

- `grails-app` - top level directory for Groovy sources
 - `conf` - [Configuration sources](#).
 - `controllers` - [Web controllers](#) - The C in MVC.
 - `domain` - The [application domain](#).
 - `i18n` - Support for [internationalization \(i18n\)](#).
 - `services` - The [service layer](#).
 - `taglib` - [Tag libraries](#).
 - `utils` - Grails specific utilities.
 - `views` - [Groovy Server Pages](#) - The V in MVC.
- `scripts` - [Gant scripts](#).
- `src` - Supporting sources
 - `groovy` - Other Groovy sources
 - `java` - Other Java sources
- `test` - [Unit and integration tests](#).

2.8 Running an Application

Grails applications can be run with the built in Tomcat server using the [run-app](#) command which will load

```
grails run-app
```

You can specify a different port by using the `server.port` argument:

```
grails -Dserver.port=8090 run-app
```

More information on the [run-app](#) command can be found in the reference guide.

2.9 Testing an Application

The `create-*` commands in Grails automatically create unit or integration tests for you within the `test` directory. It is of course up to you to populate these tests with valid test logic, information on which can be

To execute tests you run the [test-app](#) command as follows:

```
grails test-app
```

2.10 Deploying an Application

Grails applications are deployed as Web Application Archives (WAR files), and Grails includes the [war](#) co

```
grails war
```

This will produce a WAR file under the `target` directory which can then be deployed as per your contain

Unlike most scripts which default to the `development` environment unless overridden, the `war` environment by default. You can override this like any script by specifying the environment name, for exa

```
grails dev war
```



NEVER deploy Grails using the [run-app](#) command as this command sets Grails up for aut which has a severe performance and scalability implications

When deploying Grails you should always run your containers JVM with the `-server` option and with set of VM flags would be:

```
-server -Xmx512M -XX:MaxPermSize=256m
```

2.11 Supported Java EE Containers

Grails runs on any container that supports Servlet 2.5 and above and is known to work on the following sp

- Tomcat 7
- Tomcat 6
- SpringSource tc Server
- Eclipse Virgo
- GlassFish 3
- GlassFish 2
- Resin 4
- Resin 3
- JBoss 6
- JBoss 5
- Jetty 7
- Jetty 6
- IBM Websphere 7.0
- IBM Websphere 6.1
- Oracle Weblogic 10.3
- Oracle Weblogic 10
- Oracle Weblogic 9

Some containers have bugs however, which in most cases can be worked around. A [list of known deploy wiki](#).

2.12 Generating an Application

To get started quickly with Grails it is often useful to use a feature called [Scaffolding](#) to generate the skeleton of one of the `generate-*` commands such as [generate-all](#), which will generate a [controller](#) (and its unit tests).

```
grails generate-all Book
```

2.13 Creating Artefacts

Grails ships with a few convenience targets such as [create-controller](#), [create-domain-class](#) and so on for different artefact types for you.



These are just for your convenience and you can just as easily use an IDE or your favourite text editor.

For example to create the basis of an application you typically need a [domain model](#):

```
grails create-domain-class book
```

This will result in the creation of a domain class at `grails-app/domain/Book.groovy` such as:

```
class Book {  
}
```

There are many such `create-*` commands that can be explored in the command line reference guide.



To decrease the amount of time it takes to run Grails scripts, use the [interactive](#) mode.

3 Configuration

It may seem odd that in a framework that embraces "convention-over-configuration" that we tackle this. There is typically a one-off, it is best to get it out the way.

With Grails' default settings you can actually develop an application without doing any configuration whatsoever. Grails provides a servlet container and in-memory H2 database, so there isn't even a database to set up.

However, typically you should configure a more robust database at some point and that is described in the

3.1 Basic Configuration

For general configuration Grails provides a file called `grails-app/conf/Config.groovy`. This file is very similar to Java properties files except it is pure Groovy hence you can reuse variables and use properties.

You can add your own configuration in here, for example:

```
foo.bar.hello = "world"
```

Then later in your application you can access these settings in one of two ways. The most common is from within a controller or tag library:

```
assert "world" == grailsApplication.config.foo.bar.hello
```

The other way involves getting a reference to the [ConfigurationHolder](#) class that holds a reference to the configuration.

```
import org.codehaus.groovy.grails.commons.*
...
def config = ConfigurationHolder.config
assert "world" == config.foo.bar.hello
```



`ConfigurationHolder` and `ApplicationHolder` are deprecated and will be removed in a future version. It is highly preferable to access the `GrailsApplication` and `config` from the `grailsApp`.

3.1.1 Built in options

Grails also provides the following configuration options:

- `grails.config.locations` - The location of properties files or additional Grails configuration
- `grails.enable.native2ascii` - Set this to false if you do not require native2ascii conversion
- `grails.views.default.codec` - Sets the default encoding regime for GSPs - can be one of 'n' or 'u'. To reduce risk of XSS attacks, set this to 'html'.
- `grails.views.gsp.encoding` - The file encoding used for GSP source files (default is 'utf-8')
- `grails.mime.file.extensions` - Whether to use the file extension to dictate the mime type
- `grails.mime.types` - A map of supported mime types used for [Content Negotiation](#)
- `grails.serverURL` - A string specifying the server URL portion of absolute links. `grails.serverURL="http://my.yourportal.com"`. See [createLink](#).

War generation

- `grails.project.war.file` - Sets the name and location of the WAR file generated by the `war` task
- `grails.war.dependencies` - A closure containing Ant builder syntax or a list of JAR filenames included in the WAR file.
- `grails.war.copyToWebApp` - A closure containing Ant builder syntax that is legal inside an Ant script. You control what gets included in the WAR file from the "web-app" directory.
- `grails.war.resources` - A closure containing Ant builder syntax. Allows the application to define resources for the final WAR file

For more information on using these options, see the section on [deployment](#).

3.1.2 Logging

The Basics

Grails uses its common configuration mechanism to provide the settings for the underlying [Log4j](#) logging. The `log4j` setting is placed in the file `grails-app/conf/Config.groovy`.

So what does this `log4j` setting look like? Here's a basic example:

```
log4j = {
    error 'org.codehaus.groovy.grails.web.servlet', // controllers
         'org.codehaus.groovy.grails.web.pages' // GSP
    warn  'org.apache.catalina'
}
```

This says that for loggers whose name starts with 'org.codehaus.groovy.grails.web.servlet' or 'org.codehaus.groovy.grails.web.pages', messages logged at 'error' level and above will be shown. Loggers with names starting with 'org.apache.catalina' will have messages logged at the 'warn' level and above. What does that mean? First of all, you have to understand how levels work.

Logging levels

There are several standard logging levels, which are listed here in order of descending priority:

1. off
2. fatal
3. error
4. warn
5. info
6. debug
7. trace
8. all

When you log a message, you implicitly give that message a level. For example, the method `log.error(msg)` will log it at 'error' level. Likewise, `log.debug(msg)` will log it at 'debug'. Each of the above levels apart from 'off' and 'all' has the same name.

The logging system uses that *message* level combined with the configuration for the logger (see next section) to decide what gets written out. For example, if you have an 'org.example.domain' logger configured like so:

```
warn 'org.example.domain'
```

then messages with a level of 'warn', 'error', or 'fatal' will be written out. Messages at other levels will be ignored.

Before we go on to loggers, a quick note about those 'off' and 'all' levels. These are special in that they can't log messages at these levels. So if you configure a logger with a level of 'off', then no messages will be written out. If you configure a logger with a level of 'all', then you will see all messages. Simple.

Loggers

Loggers are fundamental to the logging system, but they are a source of some confusion. For a start, what are they and how do you configure them?

A logger is the object you log messages to, so in the call `log.debug(msg)`, `log` is a logger instance (created by the logging system) and uniquely identified by name, so if two separate classes use loggers with the same name, those loggers will be the same.

There are two main ways to get hold of a logger:

1. use the `log` instance injected into artifacts such as domain classes, controllers and services;
2. use the Commons Logging API directly.

If you use the dynamic `log` property, then the name of the logger is 'grails.app.<type>.<className>', where <type> is either 'controller' or 'service', and <className> is the fully qualified name of the artifact. For example, if you have a service class named `org.example.domain.MyService`, then the logger name would be 'grails.app.service.org.example.domain.MyService'.

```
package org.example

class MyService {
    ...
}
```


then the name of the logger will be 'grails.app.service.org.example.MyService'.

For other classes, the typical approach is to store a logger based on the class name in a constant static field

```
package org.other

import org.apache.commons.logging.LogFactory

class MyClass {
    private static final log = LogFactory.getLog(this)
    ...
}
```

This will create a logger with the name 'org.other.MyClass' - note the lack of a 'grails.app.' prefix since pass a name to the `getLog()` method, such as "myLogger", but this is less common because the logging special way.

Configuring loggers

You have already seen how to configure loggers in Grails:

```
log4j = {
    error 'org.codehaus.groovy.grails.web.servlet'
}
```

This example configures loggers with names starting with 'org.codehaus.groovy.grails.web.servlet' to ignore of 'warn' or lower. But is there a logger with this name in the application? No. So why have a configuration to any logger whose name *begins with* 'org.codehaus.groovy.grails.web.servlet.' as well. For example `org.codehaus.groovy.grails.web.servlet.GrailsDispatcherServlet` or `org.codehaus.groovy.grails.web.servlet.mvc.GrailsWebRequest` one.

In other words, loggers are hierarchical. This makes configuring them by package much simpler than it would be otherwise.

The most common things that you will want to capture log output from are your controllers, services, and domain classes. As mentioned earlier to do that: `grails.app.<artifactType>.<className>`. In particular the class name must exist if there is one:

```
log4j = {
    // Set level for all application artifacts
    info "grails.app"

    // Set for a specific controller in the default package
    debug "grails.app.controllers.YourController"

    // Set for a specific domain class
    debug "grails.app.domain.org.example.Book"

    // Set for all taglibs
    info "grails.app.taglib"
}
```

The standard artifact names used in the logging configuration are:

- `conf` - For anything under `grails-app/conf` such as `BootStrap.groovy` (but excluding `filters`)
- `filters` - For filters
- `taglib` - For tag libraries
- `services` - For service classes
- `controllers` - For controllers
- `domain` - For domain entities

Grails itself generates plenty of logging information and it can sometimes be helpful to see that. Here are some loggers that you can use, especially when tracking down problems with your application:

- `org.codehaus.groovy.grails.commons` - Core artifact information such as class loading
- `org.codehaus.groovy.grails.web` - Grails web request processing
- `org.codehaus.groovy.grails.web.mapping` - URL mapping debugging
- `org.codehaus.groovy.grails.plugins` - Log plugin activity
- `grails.spring` - See what Spring beans Grails and plugins are defining
- `org.springframework` - See what Spring is doing
- `org.hibernate` - See what Hibernate is doing

So far, we've only looked at explicit configuration of loggers. But what about all those loggers that *don't* have explicit configuration? The answer lies with the root logger.

The Root Logger

All logger objects inherit their configuration from the root logger, so if no explicit configuration is present, all log messages that go to that logger are subject to the rules defined for the root logger. In other words, the root logger defines the default configuration for the logging system.

Grails automatically configures the root logger to only handle messages at 'error' level and above, and all messages are sent to `stdout` for those with a C background). You can customise this behaviour by specifying a 'root' section in `log4j.properties`:

```
log4j = {
    root {
        info()
    }
    ...
}
```

The above example configures the root logger to log messages at 'info' level and above to the default console appender. You can also configure the root logger to log to one or more named appenders (which we'll talk more about shortly):

```
log4j = {
  appenders {
    file name:'file', file: '/var/logs/mylog.log'
  }
  root {
    debug 'stdout', 'file'
  }
}
```

In the above example, the root logger will log to two appenders - the default 'stdout' (console) appender and the 'file' appender. For power users there is an alternative syntax for configuring the root logger: the root org.apache.log4j.Level.DEBUG argument to the log4j closure. This lets you work with the logger directly:

```
log4j = { root ->
  root.level = org.apache.log4j.Level.DEBUG
  ...
}
```


For more information on what you can do with this Logger instance, refer to the Log4j API documentation.

Those are the basics of logging pretty well covered and they are sufficient if you're happy to only send log messages to the console. You want to send them to a file? How do you make sure that messages from a particular logger go to a file and more will be answered as we look into appenders.

Appenders

Loggers are a useful mechanism for filtering messages, but they don't physically write the messages anywhere. There are various types. For example, there is the default one that writes messages to the console, several others. You can even create your own appender implementations!

This diagram shows how they fit into the logging pipeline:



As you can see, a single logger may have several appenders attached to it. In a standard Grails configuration, the root logger is attached to all loggers through the default root logger configuration. But that's the only one. Adding more appenders is done in the 'appenders' block:

```
log4j = {
    appenders {
        rollingFile name: "myAppender",
                    maxFileSize: 1024,
                    file: "/tmp/logs/myApp.log"
    }
}
```

The following appenders are available by default:

Name	Class	Description
jdbc	JDBCAppender	Logs to a JDBC connection.
console	ConsoleAppender	Logs to the console.
file	FileAppender	Logs to a single file.
rollingFile	RollingFileAppender	Logs to rolling files, for example a new file each day.

Each named argument passed to an appender maps to a property of the underlying [Appender](#) implementation, name, maxFileSize and file properties of the RollingFileAppender instance.

You can have as many appenders as you like - just make sure that they all have unique names. You can even use different appender types, for example several file appenders that log to different files.

If you prefer to create the appender programmatically or if you want to use an appender implementation that is not available by default, you can simply declare an appender entry with an instance of the appender you want:

```
import org.apache.log4j.*

log4j = {
    appenders {
        appender new RollingFileAppender(
            name: "myAppender",
            maxFileSize: 1024,
            file: "/tmp/logs/myApp.log")
    }
}
```

This approach can be used to configure JMSAppender, SocketAppender, SMTPAppender, and more.

Once you have declared your extra appenders, you can attach them to specific loggers by passing the logger name to the methods from the previous section:

```
error myAppender: "grails.app.controller.BookController"
```

This will ensure that the 'grails.app.controller.BookController' logger sends log messages to 'myAppender' for the root logger. To add more than one appender to the logger, then add them to the same level declaration.

```

error myAppender:      "grails.app.controller.BookController",
myFileAppender:      ["grails.app.controller.BookController",
                      "grails.app.service.BookService"],
rollingFile:          "grails.app.controller.BookController"

```

The above example also shows how you can configure more than one logger at a time for a given appender.

Be aware that you can only configure a single level for a logger, so if you tried this code:

```

error myAppender:      "grails.app.controller.BookController"
debug myFileAppender:  "grails.app.controller.BookController"
fatal rollingFile:     "grails.app.controller.BookController"

```

you'd find that only 'fatal' level messages get logged for 'grails.app.controller.BookController'. That's because the highest level logger wins. What you probably want to do is limit what level of messages an appender writes.

An appender that is attached to a logger configured with the 'all' level will generate a lot of logging information, making working at the console difficult. So we configure the console appender to only write out messages at a certain level.

```

log4j = {
  appenders {
    console name: "stdout", threshold: org.apache.log4j.Level.INFO
  }
}

```

The key here is the `threshold` argument which determines the cut-off for log messages. This argument must be a `Level` instance - a string such as "info" will not work.

Custom Layouts

By default the Log4j DSL assumes that you want to use a [PatternLayout](#). However, there are other layouts

- `xml` - Create an XML log file
- `html` - Creates an HTML log file
- `simple` - A simple textual log
- `pattern` - A Pattern layout

You can specify custom patterns to an appender using the `layout` setting:

```

log4j = {
  appenders {
    console name: "customAppender",
            layout: pattern(conversionPattern: "%c{2} %m%n")
  }
}

```

This also works for the built-in appender "stdout", which logs to the console:

```
log4j = {
  appenders {
    console name: "stdout",
            layout: pattern(conversionPattern: "%c{2} %m%n")
  }
}
```

Environment-specific configuration

Since the logging configuration is inside `Config.groovy`, you can put it inside an environment-specific block with this approach: you have to provide the full logging configuration each time you define the `log4j` block, selectively override parts of the configuration - it's all or nothing.

To get around this, the logging DSL provides its own environment blocks that you can put anywhere in the

```
log4j = {
  appenders {
    console name: "stdout",
            layout: pattern(conversionPattern: "%c{2} %m%n")
  }

  environments {
    production {
      rollingFile name: "myAppender", maxFileSize: 1024,
                  file: "/tmp/logs/myApp.log"
    }
  }

  root {
    //...
  }

  // other shared config
  info "grails.app.controller"

  environments {
    production {
      // Override previous setting for 'grails.app.controller'
      error "grails.app.controller"
    }
  }
}
```

The one place you can't put an environment block is *inside* the `root` definition, but you can put the `root` block.

Full stacktraces

When exceptions occur, there can be an awful lot of noise in the stacktrace from Java and Groovy classes, so you can restrict traces to non-core Grails/Groovy class packages.

When this happens, the full trace is always logged to the `StackTrace` logger, which by default is `stacktrace.log`. As with other loggers though, you can change its behaviour in the configuration. For example, to go to the console, add this entry:

```
error stdout: "StackTrace"
```

This won't stop Grails from attempting to create the `stacktrace.log` file - it just redirects where stack traces is to change the location of the 'stacktrace' appender's file:

```
log4j = {
  appenders {
    rollingFile name: "stacktrace", maxFileSize: 1024,
               file:  "/var/tmp/logs/myApp-stacktrace.log"
  }
}
```

or, if you don't want to the 'stacktrace' appender at all, configure it as a 'null' appender:

```
log4j = {
  appenders {
    'null' name: "stacktrace"
  }
}
```

You can of course combine this with attaching the 'stdout' appender to the 'StackTrace' logger if you want. Finally, you can completely disable stacktrace filtering by setting the `grails.full.stacktrace` VM

```
grails -Dgrails.full.stacktrace=true run-app
```

Masking Request Parameters From Stacktrace Logs

When Grails logs a stacktrace, the log message may include the names and values of all of the request mask out the values of secure request parameters, specify the parameter names in the `grails.exception` config property:

```
grails.exceptionresolver.params.exclude = ['password', 'creditCard']
```

Request parameter logging may be turned off altogether by setting the `grails.exceptionresolver` property to `false`. The default value is `true` when the application is running in `DEVELOPMENT` mode

```
grails.exceptionresolver.logRequestParameters=false
```

Logger inheritance

Earlier, we mentioned that all loggers inherit from the root logger and that loggers are hierarchical based is that unless you override a parent setting, a logger retains the level and the appenders configured for that

```
log4j = {
    appenders {
        file name:'file', file: '/var/logs/mylog.log'
    }
    root {
        debug 'stdout', 'file'
    }
}
```

all loggers in the application will have a level of 'debug' and will log to both the 'stdout' and 'file' appender. How can you change the level of 'debug' for a particular logger? Change the 'additivity' for a logger in that case.

Additivity simply determines whether a logger inherits the configuration from its parent. If additivity is false for all loggers is true, i.e. they inherit the configuration. So how do you change this setting? Here's an example:

```
log4j = {
    appenders {
        ...
    }
    root {
        ...
    }
}

info additivity: false
    stdout: [ "grails.app.controller.BookController",
              "grails.app.service.BookService" ]
}
```

So when you specify a log level, add an 'additivity' named argument. Note that you when you specify loggers for a named appender. The following syntax will *not* work:

```
info additivity: false, [ "grails.app.controller.BookController",
                          "grails.app.service.BookService" ]
```

Customizing stack trace printing and filtering

Stacktraces in general and those generated when using Groovy in particular are quite verbose and can be interesting when diagnosing problems. So Grails uses a `org.codehaus.groovy.grails.exceptions.StackTraceFilterer` interface to filter out the approach used for filtering, implement that interface in a class in `src/groovy` or `src/java` and register it in `Config.groovy`:

```
grails.logging.stackTraceFiltererClass =
    'com.yourcompany.yourapp.MyStackTraceFilterer'
```

In addition, Grails customizes the display of the filtered stacktrace to make the information more readable. To customize this, implement the `org.codehaus.groovy.grails.exceptions.StackTracePrinter` interface in a class in `src/groovy` or `src/java` and register it in `Config.groovy`:


```
grails.logging.stackTracePrinterClass =  
    'com.yourcompany.yourapp.MyStackTracePrinter'
```

Finally, to render error information in the error GSP, an HTML-generating printer implementation is `org.codehaus.groovy.grails.web.errors.ErrorsViewStackTracePrinter` and its own implementation, either implement the `org.codehaus.groovy.grails.exceptions.ErrorsViewStackTracePrinter` or subclass `ErrorsViewStackTracePrinter` and register it in `grails-app/conf/spring/resources.groovy`.

```
import com.yourcompany.yourapp.MyErrorsViewStackTracePrinter  
  
beans = {  
    errorsViewStackTracePrinter(MyErrorsViewStackTracePrinter,  
                                ref('grailsResourceLocator'))  
}
```

3.1.3 GORM

Grails provides the following GORM configuration options:

- `grails.gorm.failOnError` - If set to `true`, causes the `save()` method of `grails.validation.ValidationException` if [validation](#) fails during a save. This option represents package names. If the value is a list of Strings then the `failOnError` behavior will only apply to the specified packages (including sub-packages). See the [save](#) method docs for more information.

For example, to enable `failOnError` for all domain classes:

```
grails.gorm.failOnError=true
```

and to enable `failOnError` for domain classes by package:

```
grails.gorm.failOnError = ['com.companyname.somepackage',  
                           'com.companyname.someotherpackage']
```

- `grails.gorm.autoFlush` - If set to `true`, causes the [merge](#), [save](#) and [delete](#) methods to flush explicitly using `save(flush: true)`.

3.2 Environments

Per Environment Configuration

Grails supports the concept of per environment configuration. The `Config.groovy`, `DataSource.groovy`, and `application.yml` files in the `grails-app/conf` directory can use per-environment configuration using the syntax provided by Grails. Consider the following default `DataSource` definition provided by Grails:

```

dataSource {
    pooled = false
    driverClassName = "org.h2.Driver"
    username = "sa"
    password = ""
}
environments {
    development {
        dataSource {
            dbCreate = "create-drop"
            url = "jdbc:h2:mem:devDb"
        }
    }
    test {
        dataSource {
            dbCreate = "update"
            url = "jdbc:h2:mem:testDb"
        }
    }
    production {
        dataSource {
            dbCreate = "update"
            url = "jdbc:h2:prodDb"
        }
    }
}

```

Notice how the common configuration is provided at the top level and then an `environments` block specifies the `dbCreate` and `url` properties of the `DataSource`.

Packaging and Running for Different Environments

Grails' [command line](#) has built in capabilities to execute any command within the context of a specific environment.

```
grails [environment] [command name]
```

In addition, there are 3 preset environments known to Grails: `dev`, `prod`, and `test` for development, production, and testing respectively. For example to create a WAR for the test environment you would run:

```
grails test war
```

To target other environments you can pass a `grails.env` variable to any command:

```
grails -Dgrails.env=UAT run-app
```

Programmatic Environment Detection

Within your code, such as in a Gant script or a bootstrap class you can detect the environment using the [Grails](#) `environment` property.

```
import grails.util.Environment

...

switch (Environment.current) {
    case Environment.DEVELOPMENT:
        configureForDevelopment()
        break
    case Environment.PRODUCTION:
        configureForProduction()
        break
}
```

Per Environment Bootstrapping

It's often desirable to run code when your application starts up on a per-environment basis. The `grails-app/conf/BootStrap.groovy` file's support for per-environment execution:

```
def init = { ServletContext ctx ->
    environments {
        production {
            ctx.setAttribute("env", "prod")
        }
        development {
            ctx.setAttribute("env", "dev")
        }
    }
    ctx.setAttribute("foo", "bar")
}
```

Generic Per Environment Execution

The previous `BootStrap` example uses the `grails.util.Environment` class internally to execute your own environment specific logic:

```
Environment.executeForCurrentEnvironment {
    production {
        // do something in production
    }
    development {
        // do something only in development
    }
}
```

3.3 The DataSource

Since Grails is built on Java technology setting up a data source requires some knowledge of JDBC (the Database Connectivity).

If you use a database other than H2 you need a JDBC driver. For example for MySQL you would need [Co](#)

Drivers typically come in the form of a JAR archive. It's best to use Ivy to resolve the jar if it's available. You could add a dependency for the MySQL driver like this:

```

grails.project.dependency.resolution = {
    inherits("global")
    log "warn"
    repositories {
        grailsPlugins()
        grailsHome()
        grailsCentral()
        mavenCentral()
    }
    dependencies {
        runtime 'mysql:mysql-connector-java:5.1.16'
    }
}

```

Note that the built-in `mavenCentral()` repository is included here since that's a reliable location for this

If you can't use Ivy then just put the JAR in your project's `lib` directory.

Once you have the JAR resolved you need to get familiar Grails' `DataSource` `grails-app/conf/DataSource.groovy`. This file contains the `dataSource` definition which includes

- `driverClassName` - The class name of the JDBC driver
- `username` - The username used to establish a JDBC connection
- `password` - The password used to establish a JDBC connection
- `url` - The JDBC URL of the database
- `dbCreate` - Whether to auto-generate the database from the domain model - one of 'create-drop', 'create', 'drop'
- `pooled` - Whether to use a pool of connections (defaults to true)
- `logSql` - Enable SQL logging to stdout
- `formatSql` - Format logged SQL
- `dialect` - A String or Class that represents the Hibernate dialect used to communicate with the package for available dialects.
- `readOnly` - If true makes the `DataSource` read-only, which results in the connection pool call `Connection`
- `properties` - Extra properties to set on the `DataSource` bean. See the [Commons DBCP BasicDataSource](#)

A typical configuration for MySQL may be something like:

```

dataSource {
    pooled = true
    dbCreate = "update"
    url = "jdbc:mysql://localhost/yourDB"
    driverClassName = "com.mysql.jdbc.Driver"
    dialect = org.hibernate.dialect.MySQL5InnoDBDialect
    username = "yourUser"
    password = "yourPassword"
}

```



When configuring the DataSource do not include the type or the def keyword before any settings as Groovy will treat these as local variable definitions and they will not be processed. The following is invalid:

```
dataSource {
    boolean pooled = true // type declaration results in ignored local variable
    ...
}
```

Example of advanced configuration using extra properties:

```
dataSource {
    pooled = true
    dbCreate = "update"
    url = "jdbc:mysql://localhost/yourDB"
    driverClassName = "com.mysql.jdbc.Driver"
    dialect = org.hibernate.dialect.MySQL5InnoDBDialect
    username = "yourUser"
    password = "yourPassword"
    properties {
        maxActive = 50
        maxIdle = 25
        minIdle = 5
        initialSize = 5
        minEvictableIdleTimeMillis = 60000
        timeBetweenEvictionRunsMillis = 60000
        maxWait = 10000
        validationQuery = "/* ping */"
    }
}
```

More on dbCreate

Hibernate can automatically create the database tables required for your domain model. You have some control over this through the dbCreate property, which can take these values:

- **create** - Drops the existing schema and creates the schema on startup, dropping existing tables, indexes, etc.
- **create-drop** - Same as **create**, but also drops the tables when the application shuts down cleanly.
- **update** - Creates missing tables and indexes, and updates the current schema without dropping existing tables. It properly handles many schema changes like column renames (you're left with the old column containing nulls).
- **validate** - Makes no changes to your database. Compares the configuration with the existing database.
- any other value - does nothing

You can also remove the dbCreate setting completely, which is recommended once your schema is released into production. Database changes are then managed through proper database migration tools like [Liquibase](#) (the [Database Migration](#) plugin uses Liquibase and is tightly integrated with it).

3.3.1 DataSources and Environments

The previous example configuration assumes you want the same config for all environments: production, test, etc.

Grails' DataSource definition is "environment aware", however, so you can do:

```
dataSource {
    pooled = true
    driverClassName = "com.mysql.jdbc.Driver"
    dialect = org.hibernate.dialect.MySQL5InnoDBDialect
    // other common settings here
}

environments {
    production {
        dataSource {
            url = "jdbc:mysql://liveip.com/liveDb"
            // other environment-specific settings here
        }
    }
}
```

3.3.2 JNDI DataSources

Referring to a JNDI DataSource

Most Java EE containers supply DataSource instances via [Java Naming and Directory Interface](#) (JNDI) data sources as follows:

```
dataSource {
    jndiName = "java:comp/env/myDataSource"
}
```

The format on the JNDI name may vary from container to container, but the way you define the DataSource

Configuring a Development time JNDI resource

The way in which you configure JNDI data sources at development time is plugin dependent. Using the resources using the `grails.naming.entries` setting in `grails-app/conf/Config.groovy`:

```

grails.naming.entries = [
  "bean/MyBeanFactory": [
    auth: "Container",
    type: "com.mycompany.MyBean",
    factory: "org.apache.naming.factory.BeanFactory",
    bar: "23"
  ],
  "jdbc/EmployeeDB": [
    type: "javax.sql.DataSource", //required
    auth: "Container", // optional
    description: "Data source for Foo", //optional
    driverClassName: "org.h2.Driver",
    url: "jdbc:h2:mem:database",
    username: "dbusername",
    password: "dbpassword",
    maxActive: "8",
    maxIdle: "4"
  ],
  "mail/session": [
    type: "javax.mail.Session",
    auth: "Container",
    "mail.smtp.host": "localhost"
  ]
]

```

3.3.3 Automatic Database Migration

The `dbCreate` property of the `DataSource` definition is important as it dictates what Grails does when automatically generating the database tables from [GORM](#) classes. The options are described in the [DataSource](#) section.

- create
- create-drop
- update
- validate
- no value

In [development](#) mode `dbCreate` is by default set to "create-drop", but at some point in development (and in production) you'll need to stop dropping and re-creating the database every time you start up your server.

It's tempting to switch to `update` so you retain existing data and only update the schema when your support is very conservative. It won't make any changes that could result in data loss, and doesn't detect if the old schema is left with the old one and will also have the new one.

Grails supports Rails-style migrations via the [Database Migration](#) plugin which can be installed by running

```

grails install-plugin database-migration

```

The plugin uses [Liquibase](#) and provides access to all of its functionality, and also has support for GC (by comparing your domain classes to a database).

3.3.4 Transaction-aware DataSource Proxy

The actual `DataSource` bean is wrapped in a transaction-aware proxy so you will be given the connection or Hibernate `Session` if one is active.

If this were not the case, then retrieving a connection from the `dataSource` would be a new connection changes that haven't been committed yet (assuming you have a sensible transaction isolation setting, e.g. `RepeatableRead`). The "real" unproxied `dataSource` is still available to you if you need access to it; its bean name is `dataSourceUnproxied`. You can access this bean like any other Spring bean, i.e. using dependency injection:

```
class MyService {
  def dataSourceUnproxied
  ...
}
```

or by pulling it from the `ApplicationContext`:

```
def dataSourceUnproxied = ctx.dataSourceUnproxied
```

3.3.5 Database Console

The [H2 database console](#) is a convenient feature of H2 that provides a web-based interface to any database. It's very useful to view the database you're developing against. It's especially useful when running against a remote database.

You can access the console by navigating to **`http://localhost:8080/appname/dbconsole`** in a browser. The `grails.dbconsole.urlRoot` attribute in `Config.groovy` and defaults to `'/dbconsole'`.

The console is enabled by default in development mode and can be disabled or enabled in the `grails.dbconsole.enabled` attribute in `Config.groovy`. For example you could enable the console in production:

```
environments {
  production {
    grails.serverURL = "http://www.changeme.com"
    grails.dbconsole.enabled = true
    grails.dbconsole.urlRoot = '/admin/dbconsole'
  }
  development {
    grails.serverURL = "http://localhost:8080/${appName}"
  }
  test {
    grails.serverURL = "http://localhost:8080/${appName}"
  }
}
```



If you enable the console in production be sure to guard access to it using a trusted security framework.

Configuration

By default the console is configured for an H2 database which will work with the default settings if you have an H2 database. If you just need to change the JDBC URL to `jdbc:h2:mem:devDB`. If you've configured an external database you can use the Saved Settings dropdown to choose a settings template and fill in the url and username. See `DataSource.groovy`.

3.3.6 Multiple Datasources

By default all domain classes share a single `DataSource` and a single database, but you have the option of two or more `DataSource`s.

Configuring Additional DataSources

The default `DataSource` configuration in `grails-app/conf/DataSource.groovy` looks something like this:

```
dataSource {
    pooled = true
    driverClassName = "org.h2.Driver"
    username = "sa"
    password = ""
}
hibernate {
    cache.use_second_level_cache = true
    cache.use_query_cache = true
    cache.provider_class = 'net.sf.ehcache.hibernate.EhCacheProvider'
}
environments {
    development {
        dataSource {
            dbCreate = "create-drop"
            url = "jdbc:h2:mem:devDb"
        }
    }
    test {
        dataSource {
            dbCreate = "update"
            url = "jdbc:h2:mem:testDb"
        }
    }
    production {
        dataSource {
            dbCreate = "update"
            url = "jdbc:h2:prodDb"
        }
    }
}
```

This configures a single `DataSource` with the Spring bean named `dataSource`. To configure multiple `DataSource` blocks (at the top level, in an environment block, or both, just like the standard `DataSource`), you separate them by an underscore. For example, this configuration adds a second `DataSource`, using MySQL in production:

```

environments {
    development {
        dataSource {
            dbCreate = "create-drop"
            url = "jdbc:h2:mem:devDb"
        }
        dataSource_lookup {
            dialect = org.hibernate.dialect.MySQLInnoDBDialect
            driverClassName = 'com.mysql.jdbc.Driver'
            username = 'lookup'
            password = 'secret'
            url = 'jdbc:mysql://localhost/lookup'
            dbCreate = 'update'
        }
    }
    test {
        dataSource {
            dbCreate = "update"
            url = "jdbc:h2:mem:testDb"
        }
    }
    production {
        dataSource {
            dbCreate = "update"
            url = "jdbc:h2:prodDb"
        }
        dataSource_lookup {
            dialect = org.hibernate.dialect.Oracle10gDialect
            driverClassName = 'oracle.jdbc.driver.OracleDriver'
            username = 'lookup'
            password = 'secret'
            url = 'jdbc:oracle:thin:@localhost:1521:lookup'
            dbCreate = 'update'
        }
    }
}

```

You can use the same or different databases as long as they're supported by Hibernate.

Configuring Domain Classes

If a domain class has no `DataSource` configuration, it defaults to the standard `'dataSource'`. S mapping block to configure a non-default `DataSource`. For example, if you want to use the `ZipDataSource`, configure it like this;

```

class ZipCode {
    String code
    static mapping = {
        datasource 'lookup'
    }
}

```

A domain class can also use two or more `DataSources`. Use the `datasources` property with a list of example:

```
class ZipCode {
  String code
  static mapping = {
    datasources(['lookup', 'auditing'])
  }
}
```

If a domain class uses the default DataSource and one or more others, use the special name DataSource:

```
class ZipCode {
  String code
  static mapping = {
    datasources(['lookup', 'DEFAULT'])
  }
}
```

If a domain class uses all configured DataSources use the special value 'ALL':

```
class ZipCode {
  String code
  static mapping = {
    datasource 'ALL'
  }
}
```

Namespaces and GORM Methods

If a domain class uses more than one DataSource then you can use the namespace implied by each DataSource for a particular DataSource. For example, consider this class which uses two DataSources:

```
class ZipCode {
  String code
  static mapping = {
    datasources(['lookup', 'auditing'])
  }
}
```

The first DataSource specified is the default when not using an explicit namespace, so in this case we can use GORM methods on the 'auditing' DataSource with the DataSource name, for example:

```
def zipCode = ZipCode.auditing.get(42)
...
zipCode.auditing.save()
```

As you can see, you add the `DataSource` to the method call in both the static case and the instance case.

Services

Like Domain classes, by default Services use the default `DataSource` and `PlatformTransactionManager`. If you want to use a different `DataSource`, use the static `datasource` property, for example:

```
class DataService {
    static datasource = 'lookup'
    void someMethod(...) {
        ...
    }
}
```

A transactional service can only use a single `DataSource`, so be sure to only make changes for domains that are the same as the Service.

Note that the `datasource` specified in a service has no bearing on which `datasources` are used for domains. It's used to declare which transaction manager to use.

What you'll see is that if you have a `Foo` domain class in `dataSource1` and a `Bar` domain class in `dataSource2`, a service method that saves a new `Foo` and a new `Bar` will only be transactional for `Foo`. A transaction won't affect the `Bar` instance. If you want both to be transactional you'd need to use two services, or a plugin like `Atomikos` to manage the commit, e.g. with the `Atomikos` plugin.

XA and Two-phase Commit

Grails has no native support for [XA](#) `DataSource`s or [two-phase commit](#), but the [Atomikos plugin](#) makes the simple changes needed in your `DataSource` definitions to reconfigure them as XA `DataSource`s.

3.4 Externalized Configuration

Some deployments require that configuration be sourced from more than one place and be changeable at runtime. In order to support deployment scenarios such as these the configuration can be externalized. This is done by adding a `grails.config.locations` setting in `Config.groovy` to specify the locations of the configuration files that should be used.

```
grails.config.locations = [
    "classpath:${appName}-config.properties",
    "classpath:${appName}-config.groovy",
    "file:${userHome}/.grails/${appName}-config.properties",
    "file:${userHome}/.grails/${appName}-config.groovy" ]
```

In the above example we're loading configuration files (both Java Properties files and [ConfigSlurper](#) config files) from the classpath and files located in `USER_HOME`.

It is also possible to load config by specifying a class that is a config script.

```
grails.config.locations = [com.my.app.MyConfig]
```

This can be useful in situations where the config is either coming from a plugin or some other part of your application re-using configuration provided by plugins across multiple applications.

Ultimately all configuration files get merged into the `config` property of the [GrailsApplication](#) object and

Values that have the same name as previously defined values will overwrite the existing values, and they are loaded in the order in which they are defined.

Config Defaults

The configuration values contained in the locations described by the `grails.config.locations` property defined in your application `Config.groovy` file which may not be what you want. You may want to have defaults that can be overridden in either your application's `Config.groovy` file or in a named config file using the `grails.config.defaults.locations` property.

This property supports the same values as the `grails.config.locations` property (i.e. paths to config files) but the config described by `grails.config.defaults.locations` will be loaded *before* the application config is overridden. Some plugins use this mechanism to supply one or more sets of default configuration for the application config.



Grails also supports the concept of property place holders and property override configurers and for more information on these see the section on [Grails and Spring](#)

3.5 Versioning

Versioning Basics

Grails has built in support for application versioning. The version of the application is set to 0.1 when you run the [create-app](#) command. The version is stored in the application meta data file `application.properties`.

To change the version of your application you can edit the file manually, or run the [set-version](#) command:

```
grails set-version 0.2
```

The version is used in various commands including the [war](#) command which will append the application version to the war file.

Detecting Versions at Runtime

You can detect the application version using Grails' support for application metadata using the [GrailsApplication](#) class. In [Grails controllers](#) there is an implicit [grailsApplication](#) variable that can be used:

```
def version = grailsApplication.metadata['app.version']
```

You can retrieve the version of Grails that is running with:

```
def grailsVersion = grailsApplication.metadata['app.grails.version']
```

or the GrailsUtil class:

```
import grails.util.GrailsUtil
...
def grailsVersion = GrailsUtil.grailsVersion
```

3.6 Project Documentation

Since Grails 1.2, the documentation engine that powers the creation of this documentation has been available.

The documentation engine uses a variation on the [Textile](#) syntax to automatically create project documentation.

Creating project documentation

To use the engine you need to follow a few conventions. First, you need to create a `src/docs/guide` directory where your source files will go. Then, you need to create the source docs themselves. Each chapter should have its sub-sections. You will end up with something like:

```
+ src/docs/guide/introduction.gdoc
+ src/docs/guide/introduction/changes.gdoc
+ src/docs/guide/gettingStarted.gdoc
+ src/docs/guide/configuration.gdoc
+ src/docs/guide/configuration/build.gdoc
+ src/docs/guide/configuration/build/controllers.gdoc
```

Note that you can have all your gdoc files in the top-level directory if you want, but you can also put sub-sections in the parent section - as the above example shows.

Once you have your source files, you still need to tell the documentation engine what the structure of your user guide is. You add a `src/docs/guide/toc.yml` file that contains the structure and titles for each section. This represents the structure of the user guide in tree form. For example, the above files could be represented as

```
introduction:
  title: Introduction
  changes: Change Log
gettingStarted: Getting Started
configuration:
  title: Configuration
  build:
    title: Build Config
    controllers: Specifying Controllers
```

The format is pretty straightforward. Any section that has sub-sections is represented with the section name followed by a colon. The next line should contain `title:` plus the title of the section as seen in the source file. Leaf nodes, i.e. those without any sub-sections, declare their title on the same line as the title.

That's it. You can easily add, remove, and move sections within the `toc.yml` to restructure the generated documentation. All section names, i.e. the gdoc filenames, should be unique since they are used for creating internal links. No worry though, the documentation engine will warn you of duplicate section names.

Creating reference items

Reference items appear in the Quick Reference section of the documentation. Each reference item belongs to a directory located in the `src/docs/ref` directory. For example, suppose you have defined a new controller that belongs to the `Controllers` category so you would create a gdoc text file at the following location:

```
+ src/docs/ref/Controllers/renderPDF.gdoc
```

Configuring Output Properties

There are various properties you can set within your `grails-app/conf/Config.groovy` for the documentation such as:

- **grails.doc.title** - The title of the documentation
- **grails.doc.subtitle** - The subtitle of the documentation
- **grails.doc.authors** - The authors of the documentation
- **grails.doc.license** - The license of the software
- **grails.doc.copyright** - The copyright message to display
- **grails.doc.footer** - The footer to use

Other properties such as the version are pulled from your project itself. If a title is not specified, the application name is used.

Generating Documentation

Once you have created some documentation (refer to the syntax guide in the next chapter) you can generate the documentation using the command:

```
grails doc
```

This command will output an `docs/manual/index.html` which can be opened in a browser to view the generated documentation.

Documentation Syntax

As mentioned the syntax is largely similar to Textile or Confluence style wiki markup. The following are the basics.

Basic Formatting

Monospace: `monospace`

```
@monospace@
```

Italic: *italic*

```
_italic_
```

Bold: **bold**

```
*bold*
```

Image:  **GRAILS**

```
!http://grails.org/images/new/grailslogo_topNav.png!
```

Linking

There are several ways to create links with the documentation generator. A basic external link can either style markup:

```
[SpringSource|http://www.springsource.com/]
```

or

```
"SpringSource":http://www.springsource.com/
```

For links to other sections inside the user guide you can use the `guide:` prefix with the name of the section:

```
[Intro|guide:introduction]
```

The section name comes from the corresponding gdoc filename. The documentation engine will warn you if the section name is broken.

To link to reference items you can use a special syntax:

```
[controllers|renderPDF]
```

In this case the category of the reference item is on the left hand side of the `|` and the name of the reference

Finally, to link to external APIs you can use the `api :` prefix. For example:

```
[String|api:java.lang.String]
```

The documentation engine will automatically create the appropriate javadoc link in this case. To add or configure them in `grails-app/conf/Config.groovy`. For example:

```
grails.doc.api.org.hibernate=  
    "http://docs.jboss.org/hibernate/stable/core/javadocs"
```

The above example configures classes within the `org.hibernate` package to link to the Hibernate web

Lists and Headings

Headings can be created by specifying the letter 'h' followed by a number and then a dot:

```
h3.<space>Heading3  
h4.<space>Heading4
```

Unordered lists are defined with the use of the `*` character:

```
* item 1  
** subitem 1  
** subitem 2  
* item 2
```

Numbered lists can be defined with the `#` character:

```
# item 1
```

Tables can be created using the `table` macro:

Name	Number
Albert	46
Wilma	1348
James	12

```
{table}  
*Name* | *Number*  
Albert | 46  
Wilma | 1348  
James | 12  
{table}
```

Code and Notes

You can define code blocks with the `code` macro:

```
class Book {  
    String title  
}
```

```
{code}  
class Book {  
    String title  
}  
{code}
```

The example above provides syntax highlighting for Java and Groovy code, but you can also highlight XM

```
<hello>world</hello>
```

```
{code:xml}  
<hello>world</hello>  
{code}
```

There are also a couple of macros for displaying notes and warnings:

Note:



This is a note!

```
{note}  
This is a note!  
{note}
```

Warning:



This is a warning!

```
{warning}  
This is a warning!  
{warning}
```

3.7 Dependency Resolution

Grails features a dependency resolution DSL that lets you control how plugins and JAR dependencies are resolved.

You specify a `grails.project.dependency.resolution` property inside the `grails-app/conf` directory that configures how dependencies are resolved:

```
grails.project.dependency.resolution = {
  // config here
}
```

The default configuration looks like the following:

```
grails.project.class.dir = "target/classes"
grails.project.test.class.dir = "target/test-classes"
grails.project.test.reports.dir = "target/test-reports"
//grails.project.war.file = "target/${appName}-${appVersion}.war"

grails.project.dependency.resolution = {
  // inherit Grails' default dependencies
  inherits("global") {
    // uncomment to disable ehcache
    // excludes 'ehcache'
  }
  log "warn"
  repositories {
    grailsPlugins()
    grailsHome()
    grailsCentral()

    // uncomment these to enable remote dependency resolution
    // from public Maven repositories
    //mavenCentral()
    //mavenLocal()
    //mavenRepo "http://snapshots.repository.codehaus.org"
    //mavenRepo "http://repository.codehaus.org"
    //mavenRepo "http://download.java.net/maven/2/"
    //mavenRepo "http://repository.jboss.com/maven2/"
  }
  dependencies {
    // specify dependencies here under either 'build', 'compile',
    // 'runtime', 'test' or 'provided' scopes eg.

    // runtime 'mysql:mysql-connector-java:5.1.16'
  }

  plugins {
    compile ":hibernate:$grailsVersion"
    compile ":jquery:1.6.1.1"
    compile ":resources:1.0"

    build ":tomcat:$grailsVersion"
  }
}
```

The details of the above will be explained in the next few sections.

3.7.1 Configurations and Dependencies

Grails features five dependency resolution configurations (or 'scopes'):

- `build`: Dependencies for the build system only
- `compile`: Dependencies for the compile step
- `runtime`: Dependencies needed at runtime but not for compilation (see above)
- `test`: Dependencies needed for testing but not at runtime (see above)
- `provided`: Dependencies needed at development time, but not during WAR deployment

Within the `dependencies` block you can specify a dependency that falls into one of these configurations. For example, if your application requires the MySQL driver to function at runtime, you can specify that like this:

```
runtime 'com.mysql:mysql-connector-java:5.1.16'
```

This uses the string syntax: `group:name:version`. You can also use a Map-based syntax:

```
runtime group: 'com.mysql',
        name: 'mysql-connector-java',
        version: '5.1.16'
```

In Maven terminology, `group` corresponds to an artifact's `groupId` and `name` corresponds to its `artifactId`.

Multiple dependencies can be specified by passing multiple arguments:

```
runtime 'com.mysql:mysql-connector-java:5.1.16',
        'net.sf.ehcache:ehcache:1.6.1'

// Or

runtime(
    [group: 'com.mysql', name: 'mysql-connector-java', version: '5.1.16'],
    [group: 'net.sf.ehcache', name: 'ehcache', version: '1.6.1']
)
```

Disabling transitive dependency resolution

By default, Grails will not only get the JARs and plugins that you declare, but it will also get their transitive dependencies. However, you may want to exclude some dependencies, or you may want a dependency without all its baggage. In such cases, you can disable transitive dependency resolution on a case-by-case basis:

```
runtime('com.mysql:mysql-connector-java:5.1.16',
        'net.sf.ehcache:ehcache:1.6.1') {
    transitive = false
}

// Or

runtime group: 'com.mysql',
        name: 'mysql-connector-java',
        version: '5.1.16',
        transitive: false
```

Excluding specific transitive dependencies

A far more common scenario is where you want the transitive dependencies, but some of them cause issues that are unnecessary. For example, many Apache projects have 'commons-logging' as a transitive dependency, but our project (we use SLF4J). That's where the `excludes` option comes in:

```
runtime('com.mysql:mysql-connector-java:5.1.16',
        'net.sf.ehcache:ehcache:1.6.1') {
    excludes "xml-apis", "commons-logging"
}

// Or
runtime(group:'com.mysql', name:'mysql-connector-java', version:'5.1.16') {
    excludes([ group: 'xml-apis', name: 'xml-apis'],
             [ group: 'org.apache.httpcomponents',
               [ name: 'commons-logging' ]])
}
```

As you can see, you can either exclude dependencies by their artifact ID (also known as a module name) or by their group ID (if you use the Map notation). You may also come across `exclude` as well, but that can only accept a single artifact ID.

```
runtime('com.mysql:mysql-connector-java:5.1.16',
        'net.sf.ehcache:ehcache:1.6.1') {
    exclude "xml-apis"
}
```

Using Ivy module configurations

If you use Ivy module configurations and wish to depend on a specific configuration, you can use the `dependencyConfiguration` method to specify the configuration to use.

```
provided("my.org:web-service:1.0") {
    dependencyConfiguration "api"
}
```

If the dependency configuration is not explicitly set, the configuration named "default" will be used (which includes all dependencies coming from Maven style repositories).

3.7.2 Dependency Repositories

Remote Repositories

Initially your `BuildConfig.groovy` does not use any remote public Maven repositories. There is a default configuration that locates the JAR files Grails needs from your Grails installation. To use a public repository, specify it in the `repositories` block.

```
repositories {
    mavenCentral()
}
```

In this case the default public Maven repository is specified. To use the SpringSource Enterprise Bund method:

```
repositories {  
    ebr()  
}
```

You can also specify a specific Maven repository to use by URL:

```
repositories {  
    mavenRepo "http://repository.codehaus.org"  
}
```

Controlling Repositories Inherited from Plugins

A plugin you have installed may define a reference to a remote repository just as an application can. By c repository definition when you install the plugin.

If you do not wish to inherit repository definitions from plugins then you can disable repository inheritance

```
repositories {  
    inherit false  
}
```

In this case your application will not inherit any repository definitions from plugins and it is down to internal) repository definitions.

Local Resolvers

If you do not wish to use a public Maven repository you can specify a flat file repository:

```
repositories {  
    flatDir name:'myRepo', dirs:'/path/to/repo'  
}
```

To specify your local Maven cache (~/.m2/repository) as a repository:

```
repositories {  
    mavenLocal()  
}
```

Custom Resolvers

If all else fails since Grails builds on Apache Ivy you can specify an Ivy resolver:

```

/*
 * Configure our resolver.
 */
def libResolver = new org.apache.ivy.plugins.resolver.URLResolver()
['libraries', 'builds'].each {

    libResolver.addArtifactPattern(
        "http://my.repository/${it}/" +
        "[organisation]/[module]/[revision]/[type]s/[artifact].[ext]")

    libResolver.addIvyPattern(
        "http://my.repository/${it}/" +
        "[organisation]/[module]/[revision]/[type]s/[artifact].[ext]")
}

libResolver.name = "my-repository"
libResolver.settings = ivySettings

resolver libResolver

```

It's also possible to pull dependencies from a repository using SSH. Ivy comes with a dedicated resolver your project like so:

```

import org.apache.ivy.plugins.resolver.SshResolver

'''
repositories {
    ...
}

def sshResolver = new SshResolver(
    name: "myRepo",
    user: "username",
    host: "dev.x.com",
    keyFile: new File("/home/username/.ssh/id_rsa"),
    m2compatible: true)

sshResolver.addArtifactPattern(
    "/home/grails/repo/[organisation]/[artifact]/" +
    "[revision]/[artifact]-[revision].[ext]")

sshResolver.latestStrategy =
    new org.apache.ivy.plugins.latest.LatestTimeStrategy()

sshResolver.changingPattern = ".*SNAPSHOT"

sshResolver.setCheckmodified(true)

resolver sshResolver
}

```

Download the [JSch](#) JAR and add it to Grails' classpath to use the SSH resolver. You can do this by passing

```
grails -classpath /path/to/jsch compile|run-app|etc.
```

You can also add its path to the CLASSPATH environment variable but be aware this it affects many Java to create an alias for `grails -classpath ...` so that you don't have to type the extra arguments eac

Authentication

If your repository requires authentication you can configure this using a `credentials` block:

```
credentials {
    realm = ".."
    host = "localhost"
    username = "myuser"
    password = "mypass"
}
```

This can be placed in your `USER_HOME/.grails/settings.groovy` file using the `grails.setting`:

```
grails.project.ivy.authentication = {
    credentials {
        realm = ".."
        host = "localhost"
        username = "myuser"
        password = "mypass"
    }
}
```

3.7.3 Debugging Resolution

If you are having trouble getting a dependency to resolve you can enable more verbose debugging from method:

```
// log level of Ivy resolver, either 'error', 'warn',
// 'info', 'debug' or 'verbose'
log "warn"
```

3.7.4 Inherited Dependencies

By default every Grails application inherits several framework dependencies. This is done through the line

```
inherits "global"
```

Inside the `BuildConfig.groovy` file. To exclude specific inherited dependencies you use the `exclude`

```
inherits("global") {
    excludes "oscache", "ehcache"
}
```

3.7.5 Providing Default Dependencies

Most Grails applications have runtime dependencies on several jar files that are provided by the Grails framework, Spring, Sitemesh, Hibernate etc. When a war file is created, all of these dependencies will be included in the war. To exclude these jar files from the war. This is useful when the jar files will be provided by the container, as Grails applications are deployed to the same container.

The dependency resolution DSL provides a mechanism to express that all of the default dependencies will be done by invoking the `defaultDependenciesProvided` method and passing `true` as an argument:

```
grails.project.dependency.resolution = {
  defaultDependenciesProvided true // all of the default dependencies will
                                // be "provided" by the container

  inherits "global" // inherit Grails' default dependencies

  repositories {
    grailsHome()
    ...
  }
  dependencies {
    ...
  }
}
```



`defaultDependenciesProvided` must come before `inherits`, otherwise the Grails application will not be included in the war.

3.7.6 Dependency Reports

As mentioned in the previous section a Grails application consists of dependencies inherited from the framework and application dependencies itself.

To obtain a report of an application's dependencies you can run the [dependency-report](#) command:

```
grails dependency-report
```

By default this will generate reports in the `target/dependency-report` directory. You can specify a report for by passing an argument containing the configuration name:

```
grails dependency-report runtime
```

3.7.7 Plugin JAR Dependencies

Specifying Plugin JAR dependencies

The way in which you specify dependencies for a [plugin](#) is identical to how you specify dependencies installed into an application the application automatically inherits the dependencies of the plugin.

To define a dependency that is resolved for use with the plugin but not *exported* to the application then you can use the `export` property:

```
test('org.spockframework:spock-core:0.5-groovy-1.8') {
  export = false
}
```

In this case the Spock dependency will be available only to the plugin and not resolved as an application dependency using the Map syntax:

```
test group: 'org.spockframework', name: 'spock-core',  
      version: '0.5-groovy-1.8', export: false
```



You can use `exported = false` instead of `export = false`, but we recommend consistent with the Map argument.

Overriding Plugin JAR Dependencies in Your Application

If a plugin is using a JAR which conflicts with another plugin, or an application dependency then you can override the dependencies inside an application using exclusions. For example:

```
plugins {  
    compile(":hibernate:$grailsVersion") {  
        excludes "javassist"  
    }  
}  
  
dependencies {  
    runtime "javassist:javassist:3.4.GA"  
}
```

In this case the application explicitly declares a dependency on the "hibernate" plugin and specifies an exclusion, effectively excluding the javassist library as a dependency.

3.7.8 Maven Integration

When using the Grails Maven plugin, Grails' dependency resolution mechanics are disabled as it is assumed that dependencies are managed in Maven's `pom.xml` file.

However, if you would like to continue using Grails regular commands like [run-app](#), [test-app](#) and so on then you can load dependencies from the Maven `pom.xml` file instead.

To do so simply add the following line to your `BuildConfig.groovy`:

```
grails.project.dependency.resolution = {  
    pom true  
    ..  
}
```

The line `pom true` tells Grails to parse Maven's `pom.xml` and load dependencies from there.

3.7.9 Deploying to a Maven Repository

If you use Maven to build your Grails project, you can use the standard Maven targets `mvn install` to deploy a Grails project or plugin to a Maven repository using the [maven-publisher](#) plugin.

The plugin provides the ability to publish Grails projects and plugins to local and remote Maven repositories added by the plugin:

- **maven-install** - Installs a Grails project or plugin into your local Maven cache
- **maven-deploy** - Deploys a Grails project or plugin to a remote Maven repository

By default this plugin will automatically generate a valid `pom.xml` for you unless a `pom.xml` is already present, in which case this `pom.xml` file will be used.

maven-install

The `maven-install` command will install the Grails project or plugin artifact into your local Maven cache.

```
grails maven-install
```

In the case of plugins, the plugin zip file will be installed, whilst for application the application WAR file will be installed.

maven-deploy

The `maven-deploy` command will deploy a Grails project or plugin into a remote Maven repository:

```
grails maven-deploy
```

It is assumed that you have specified the necessary `<distributionManagement>` configuration with the `id` of the remote repository to deploy to:

```
grails maven-deploy --repository=myRepo
```

The `repository` argument specifies the 'id' for the repository. Configure the details of the repository in the `grails-app/conf/BuildConfig.groovy` file or in your `$USER_HOME/.grails/settings` file.

```
grails.project.dependency.distribution = {
    localRepository = "/path/to/my/local"
    remoteRepository(id: "myRepo", url: "http://myserver/path/to/repo")
}
```

The syntax for configuring remote repositories matches the syntax from the [remoteRepository](#) element in the following XML:

```
<remoteRepository id="myRepo" url="scp://localhost/www/repository">
    <authentication username="..." privateKey="${user.home}/.ssh/id_dsa"/>
</remoteRepository>
```

Can be expressed as:

```
remoteRepository(id: "myRepo", url: "scp://localhost/www/repository") {  
    authentication username: "...", privateKey: "${userHome}/.ssh/id_dsa"  
}
```

By default the plugin will try to detect the protocol to use from the URL of the repository (ie "http" from different protocol you can do:

```
grails maven-deploy --repository=myRepo --protocol=webdav
```

The available protocols are:

- http
- scp
- scpexe
- ftp
- webdav

Groups, Artifacts and Versions

Maven defines the notion of a 'groupId', 'artifactId' and a 'version'. This plugin pulls this information from the plugin descriptor.

Projects

For applications this plugin will use the Grails application name and version provided by Grails when generating the version you can run the `set-version` command:

```
grails set-version 0.2
```

The Maven groupId will be the same as the project name, unless you specify a different one in `Config.groovy`

```
grails.project.groupId = "com.mycompany"
```

Plugins

With a Grails plugin the groupId and version are taken from the following properties in the `*GrailsPlugin`

```
String groupId = 'myOrg'  
String version = '0.1'
```

The 'artifactId' is taken from the plugin name. For example if you have a plugin called `FeedsGrails` "feeds". If your plugin does not specify a `groupId` then this defaults to "org.grails.plugins".

3.7.10 Plugin Dependencies

As of Grails 1.3 you can declaratively specify plugins as dependencies via the dependency DSL instead of

```
grails.project.dependency.resolution = {
    ...
    repositories {
        ...
    }
    plugins {
        runtime ':hibernate:1.2.1'
    }
    dependencies {
        ...
    }
    ...
}
```

If you don't specify a group id the default plugin group id of `org.grails.plugins` is used. You can specify a particular plugin by using "latest.integration" as the version number:

```
plugins {
    runtime ':hibernate:latest.integration'
}
```

Integration vs. Release

The "latest.integration" version label will also include resolving snapshot versions. To not include snapshot versions use "latest.release" label:

```
plugins {
    runtime ':hibernate:latest.release'
}
```



The "latest.release" label only works with Maven compatible repositories. If you have a registry repository then you should use "latest.integration".

And of course if you use a Maven repository with an alternative group id you can specify a group id:

```
plugins {
    runtime 'mycompany:hibernate:latest.integration'
}
```

Plugin Exclusions

You can control how plugins transitively resolves both plugin and JAR dependencies using exclusions. For

```
plugins {
  runtime(':weceem:0.8') {
    excludes "searchable"
  }
}
```

Here we have defined a dependency on the "weceem" plugin which transitively depends on the "searchable" plugin. You can tell Grails *not* to transitively install the searchable plugin. You can combine this technique with the following:

```
plugins {
  runtime(':weceem:0.8') {
    excludes "searchable" // excludes most recent version
  }
  runtime ':searchable:0.5.4' // specifies a fixed searchable version
}
```

You can also completely disable transitive plugin installs, in which case no transitive dependencies will be

```
plugins {
  runtime(':weceem:0.8') {
    transitive = false
  }
  runtime ':searchable:0.5.4' // specifies a fixed searchable version
}
```

4 The Command Line

Grails' command line system is built on [Gant](#) - a simple Groovy wrapper around [Apache Ant](#).

However, Grails takes it further through the use of convention and the `grails` command. When you type

```
grails [command name]
```

Grails searches in the following directories for Gant scripts to execute:

- `USER_HOME/.grails/scripts`
- `PROJECT_HOME/scripts`
- `PROJECT_HOME/plugins/*/scripts`
- `GRAILS_HOME/scripts`

Grails will also convert command names that are in lower case form such as `run-app` into camel case. So ty

```
grails run-app
```

Results in a search for the following files:

- `USER_HOME/.grails/scripts/RunApp.groovy`
- `PROJECT_HOME/scripts/RunApp.groovy`
- `PLUGINS_HOME/*/scripts/RunApp.groovy`
- `GLOBAL_PLUGINS_HOME/*/scripts/RunApp.groovy`
- `GRAILS_HOME/scripts/RunApp.groovy`

If multiple matches are found Grails will give you a choice of which one to execute.

When Grails executes a Gant script, it invokes the "default" target defined in that script. If there is no defai

To get a list of all commands and some help about the available commands type:

```
grails help
```

which outputs usage instructions and the list of commands Grails is aware of:

```
Usage (optionals marked with *):
grails [environment]* [target] [arguments]*
```

```
Examples:
grails dev run-app
grails create-app books
```

```
Available Targets (type grails help 'target-name' for more info):
grails bootstrap
grails bug-report
grails clean
grails compile
...
```



Refer to the Command Line reference in the Quick Reference menu of the reference guide about individual commands

It's often useful to provide custom arguments to the JVM when running Grails commands, in particular, you may want to set a higher maximum heap size. The Grails command will use any JVM options environment variable, but you can also specify a Grails-specific environment variable too:

```
export GRAILS_OPTS="-Xmx1G -Xms256m -XX:MaxPermSize=256m"
grails run-app
```

non-interactive mode

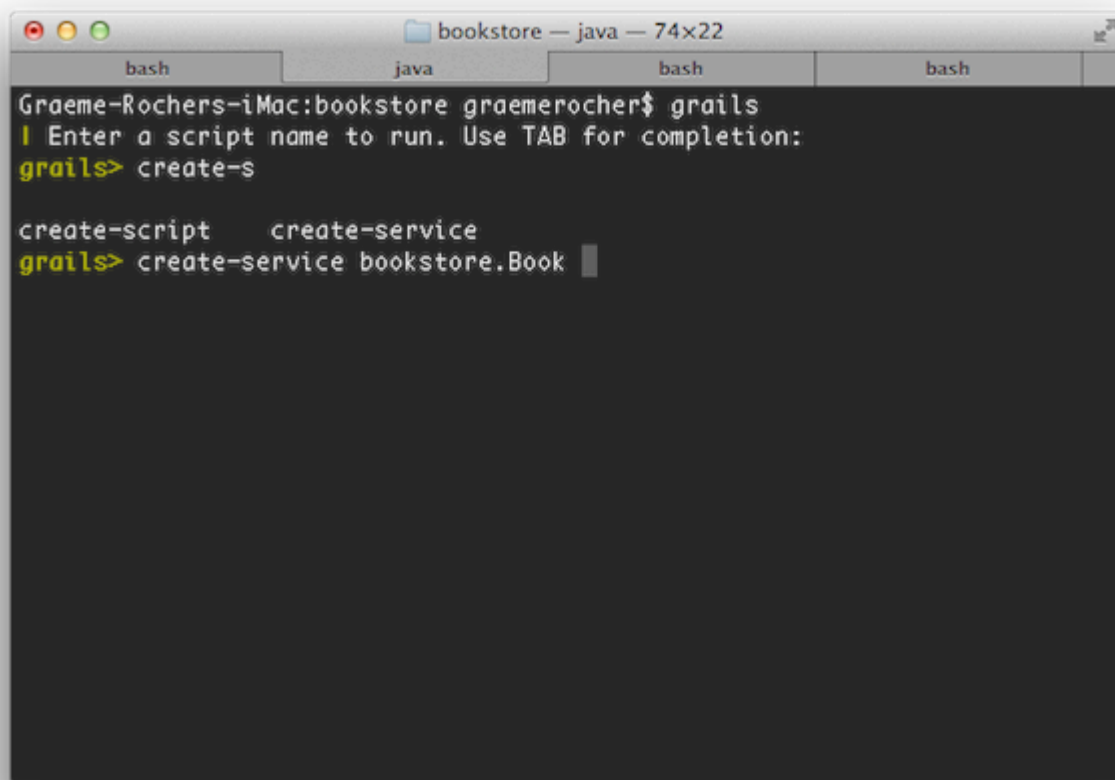
When you run a script manually and it prompts you for information, you can answer the questions and complete the script. When you run a script as part of an automated process, for example a continuous integration build server, there's no way to pass the `--non-interactive` switch to the script command to tell Grails to accept the default answers or whether to install a missing plugin.

For example:

```
grails war --non-interactive
```

4.1 Interactive Mode

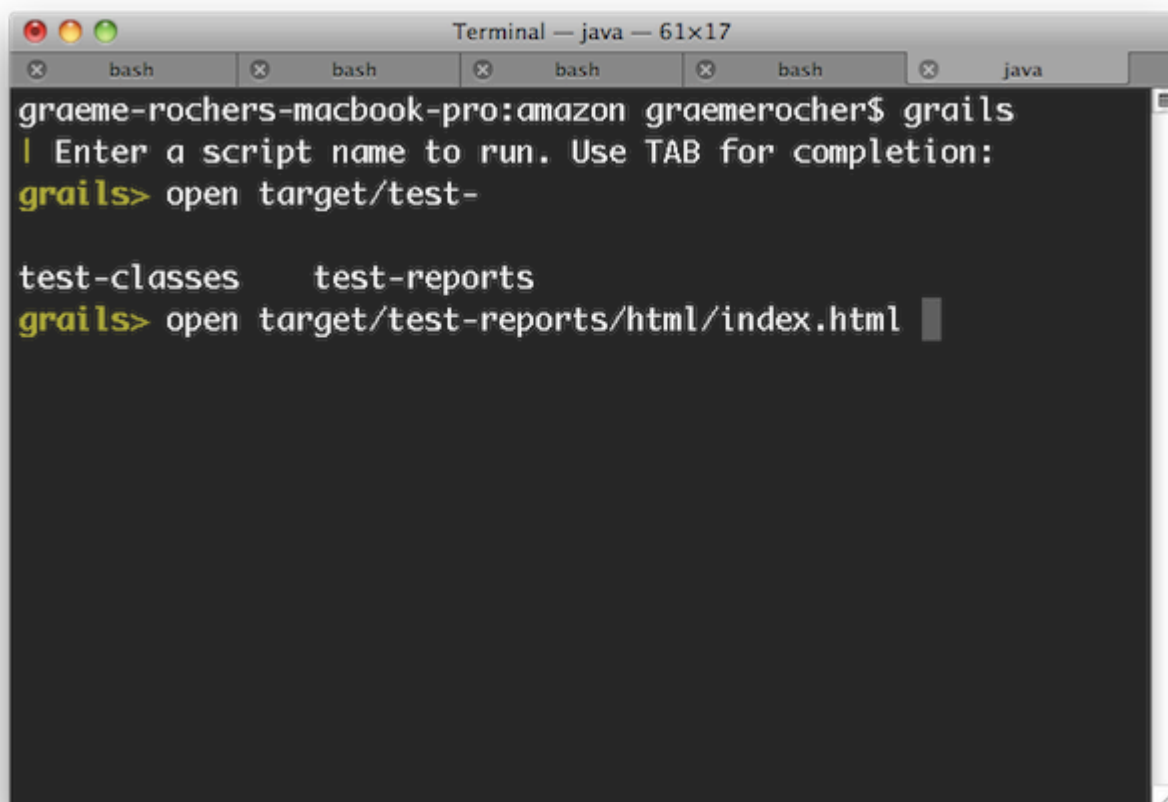
Interactive mode is a feature of the Grails command line which keeps the JVM running and allows you to activate interactive mode by typing 'grails' at the command line and then use TAB completion to get a list of commands.

A terminal window titled 'bookstore — java — 74x22' with tabs for 'bash', 'java', 'bash', and 'bash'. The prompt is 'Graeme-Rochers-iMac:bookstore graemerocher\$'. The user enters 'grails', followed by a prompt 'Enter a script name to run. Use TAB for completion:'. The user then enters 'grails> create-s', which shows suggestions 'create-script' and 'create-service'. Finally, the user enters 'grails> create-service bookstore.Book' with a cursor at the end.

```
Graeme-Rochers-iMac:bookstore graemerocher$ grails
Enter a script name to run. Use TAB for completion:
grails> create-s

create-script    create-service
grails> create-service bookstore.Book
```

If you need to open a file whilst within interactive mode you can use the open command which will TAB

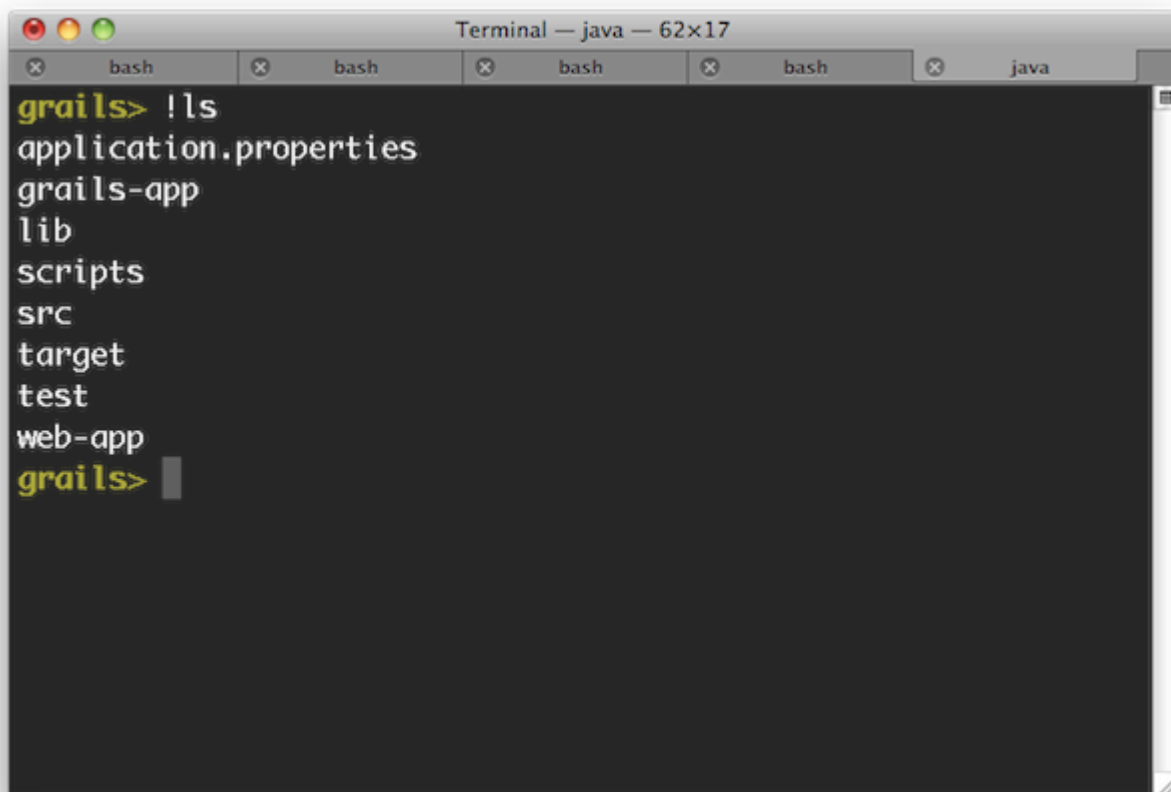
A terminal window titled 'Terminal — java — 61x17' with tabs for 'bash', 'bash', 'bash', 'bash', 'bash', and 'java'. The prompt is 'graeme-rochers-macbook-pro:amazon graemerocher\$'. The user enters 'grails', followed by a prompt 'Enter a script name to run. Use TAB for completion:'. The user then enters 'grails> open target/test-', which shows suggestions 'test-classes' and 'test-reports'. Finally, the user enters 'grails> open target/test-reports/html/index.html' with a cursor at the end.

```
graeme-rochers-macbook-pro:amazon graemerocher$ grails
Enter a script name to run. Use TAB for completion:
grails> open target/test-

test-classes    test-reports
grails> open target/test-reports/html/index.html
```

TAB completion also works for class names after the create-* commands:

If you need to run an external process whilst interactive mode is running you can do so by starting the com

A screenshot of a macOS Terminal window titled "Terminal — java — 62x17". The window has several tabs at the top, all labeled "bash", and a "java" tab on the right. The terminal content shows a prompt "grails>" followed by the command "ls". The output lists the following files and directories: "application.properties", "grails-app", "lib", "scripts", "src", "target", "test", "web-app", and "grails>".

```
grails> ls
application.properties
grails-app
lib
scripts
src
target
test
web-app
grails>
```

4.2 Creating Gant Scripts

You can create your own Gant scripts by running the [create-script](#) command from the root of your project.

```
grails create-script compile-sources
```

Will create a script called `scripts/CompileSources.groovy`. A Gant script itself is similar to supports the concept of "targets" and dependencies between them:

```
target(default:"The default target is the one that gets executed by Grails") {
    depends(clean, compile)
}

target(clean:"Clean out things") {
    ant.delete(dir:"output")
}

target(compile:"Compile some sources") {
    ant.mkdir(dir:"mkdir")
    ant.javac(srcdir:"src/java", destdir:"output")
}
```

As demonstrated in the script above, there is an implicit ant variable (an instance of `groovy.util.Ant` [Apache Ant API](#)).



In previous versions of Grails (1.0.3 and below), the variable was `Ant`, i.e. with a capital first

You can also "depend" on other targets using the `depends` method demonstrated in the `default` target

The default target

In the example above, we specified a target with the explicit name "default". This is one way of defining an alternative approach is to use the `setDefaultTarget()` method:

```
target("clean-compile": "Performs a clean compilation on the app source") {
    depends(clean, compile)
}

target(clean:"Clean out things") {
    ant.delete(dir:"output")
}

target(compile:"Compile some sources") {
    ant.mkdir(dir:"mkdir")
    ant.javac(srcdir:"src/java", destdir:"output")
}

setDefaultTarget("clean-compile")
```

This lets you call the default target directly from other scripts if you wish. Also, although we have put it at the end of the script in this example, it can go anywhere as long as it comes *after* the target it refers to ("clean-compile").

Which approach is better? To be honest, you can use whichever you prefer - there don't seem to be any strong arguments either way. One thing we would say is that if you want to allow other scripts to call your "default" target, you should move it to a separate script. We'll talk some more about this in the next section.

4.3 Re-using Grails scripts

Grails ships with a lot of command line functionality out of the box that you may find useful in your scripts (see the reference in the reference guide for info on all the commands). Of particular use are the [compile](#), [package](#), and [run](#) scripts.

The [bootstrap](#) script for example lets you bootstrap a Spring [ApplicationContext](#) instance to get access to the application context (see the reference in the reference guide for info on all the commands). The [integration tests](#) script uses this:

```
includeTargets << grailsScript("_GrailsBootstrap")

target('default': "Database stuff") {
    depends(configureProxy, packageApp, classpath, loadApp, configureApp)

    Connection c
    try {
        c = appCtx.getBean('dataSource').getConnection()
        // do something with connection
    }
    finally {
        c?.close()
    }
}
```

Pulling in targets from other scripts

Gant lets you pull in all targets (except "default") from another Gant script. You can then depend upon or define targets in the current script. The mechanism for doing this is the `includeTargets` property. Simply use the left-shift operator:

```
includeTargets << new File("/path/to/my/script.groovy")
includeTargets << gant.tools.Ivy
```

Don't worry too much about the syntax using a class, it's quite specialised. If you're interested, look into the `Gant` class.

Core Grails targets

As you saw in the example at the beginning of this section, you use neither the `File`- nor the class-based `includeTargets` method to include core Grails targets. Instead, you should use the special `grailsScript()` method that is provided by Grails (note that this is not available in normal Gant scripts, just Grails ones).

The syntax for the `grailsScript()` method is pretty straightforward: simply pass it the name of the Grails script you want to include. Here is a list of Grails scripts that you could reuse:

Script	Description
<code>_GrailsSettings</code>	You really should include this! Fortunately, it is included automatically by all other Grails scripts, so you usually don't have to include it explicitly.
<code>_GrailsEvents</code>	Include this to fire events. Adds an <code>event(String eventName, List args)</code> target that runs all other Grails scripts.
<code>_GrailsClasspath</code>	Configures compilation, test, and runtime classpaths. If you want to use or play with classpaths, include this script by almost all other Grails scripts.
<code>_GrailsProxy</code>	If you don't have direct access to the internet and use a proxy, include this script to configure the proxy.
<code>_GrailsArgParsing</code>	Provides a <code>parseArguments</code> target that does what it says on the tin: parses the command-line arguments they run your script. Adds them to the <code>argsMap</code> property.
<code>_GrailsTest</code>	Contains all the shared test code. Useful if you want to add any extra tests.
<code>_GrailsRun</code>	Provides all you need to run the application in the configured servlet container, either from a WAR file (<code>runWar/runWarHttps</code>).

There are many more scripts provided by Grails, so it is worth digging into the scripts themselves to find out what they do. Anything that starts with an underscore is designed for reuse.

Script architecture

You maybe wondering what those underscores are doing in the names of the Grails scripts. That is Grails uses `_internal`, or in other words that it has not corresponding "command". So you can't run `"grails _grails-s"` because they don't have a default target.

Internal scripts are all about code sharing and reuse. In fact, we recommend you take a similar approach in your own scripts: create an internal script that can be easily shared, and provide simple command scripts that parse any command-line arguments and call the internal script. For example if you have a script that runs some functional tests, you can split

```

./scripts/FunctionalTests.groovy:
includeTargets << new File("${basedir}/scripts/_FunctionalTests.groovy")
target(default: "Runs the functional tests for this project.") {
    depends(runFunctionalTests)
}

./scripts/_FunctionalTests.groovy:
includeTargets << grailsScript("_GrailsTest")
target(runFunctionalTests: "Run functional tests.") {
    depends(...)
    ...
}

```

Here are a few general guidelines on writing scripts:

- Split scripts into a "command" script and an internal one.
- Put the bulk of the implementation in the internal script.
- Put argument parsing into the "command" script.
- To pass arguments to a target, create some script variables and initialise them before calling the target
- Avoid name clashes by using closures assigned to script variables instead of targets. You can then pas

4.4 Hooking into Events

Grails provides the ability to hook into scripting events. These are events triggered during execution of Gr

The mechanism is deliberately simple and loosely specified. The list of possible events is not fixed in events triggered by plugin scripts, for which there is no equivalent event in the core target scripts.

Defining event handlers

Event handlers are defined in scripts called `_Events.groovy`. Grails searches for these scripts in the fo

- `USER_HOME/.grails/scripts` - user-specific event handlers
- `PROJECT_HOME/scripts` - applicaton-specific event handlers
- `PLUGINS_HOME/*/scripts` - plugin-specific event handlers
- `GLOBAL_PLUGINS_HOME/*/scripts` - event handlers provided by global plugins

Whenever an event is fired, *all* the registered handlers for that event are executed. Note that the automatically by Grails, so you just need to declare them in the relevant `_Events.groovy` file.

Event handlers are blocks defined in `_Events.groovy`, with a name beginning with "event". The `/scripts` directory to demonstrate the feature:

```

eventCreatedArtefact = { type, name ->
    println "Created $type $name"
}

eventStatusUpdate = { msg ->
    println msg
}

eventStatusFinal = { msg ->
    println msg
}

```

You can see here the three handlers `eventCreatedArtefact`, `eventStatusUpdate`, `eventStatusFinal`, which are documented in the command line reference guide. For example the [compile](#) command

- `CompileStart` - Called when compilation starts, passing the kind of compile - source or tests
- `CompileEnd` - Called when compilation is finished, passing the kind of compile - source or tests

Triggering events

To trigger an event simply include the `Init.groovy` script and call the `event()` closure:

```

includeTargets << grailsScript("_GrailsEvents")
event("StatusFinal", ["Super duper plugin action complete!"])

```

Common Events

Below is a table of some of the common events that can be leveraged:

Event	Parameters	Description
StatusUpdate	message	Passed a string indicating current script status/progress
StatusError	message	Passed a string indicating an error message from the current
StatusFinal	message	Passed a string indicating the final script status message, i.e. target does not exit the scripting environment
CreatedArtefact	artefactType,artefactName	Called when a create-xxxx script has completed and created
CreatedFile	fileName	Called whenever a project source file is created, not includi
Exiting	returnCode	Called when the scripting environment is about to exit cleanl
PluginInstalled	pluginName	Called after a plugin has been installed
CompileStart	kind	Called when compilation starts, passing the kind of compile
CompileEnd	kind	Called when compilation is finished, passing the kind of con
DocStart	kind	Called when documentation generation is about to start - jav
DocEnd	kind	Called when documentation generation has ended - javadoc o
SetClasspath	rootLoader	Called during classpath initialization so plugins c rootLoader.addURL(...). Note that this augments the classp you cannot use this to load a class that your event script nee if you load the class by name.
PackagingEnd	none	Called at the end of packaging (which is called prior to the web.xml is generated)

4.5 Customising the build

Grails is most definitely an opinionated framework and it prefers convention to configuration, but this does not mean you cannot configure it. In this section, we look at how you can influence and modify the standard Grails build.

The defaults

The core of the Grails build configuration is the `grails.util.BuildSettings` class, which controls where classes are compiled to, what dependencies the application has, and other such settings.

Here is a selection of the configuration options and their default values:

Property	Config option	Default value
grailsWorkDir	grails.work.dir	\$USER_HOME/.grails/<grailsVersion>
projectWorkDir	grails.project.work.dir	<grailsWorkDir>/projects/<baseDirName>
classesDir	grails.project.class.dir	<projectWorkDir>/classes
testClassesDir	grails.project.test.class.dir	<projectWorkDir>/test-classes
testReportsDir	grails.project.test.reports.dir	<projectWorkDir>/test/reports
resourcesDir	grails.project.resource.dir	<projectWorkDir>/resources
projectPluginsDir	grails.project.plugins.dir	<projectWorkDir>/plugins
globalPluginsDir	grails.global.plugins.dir	<grailsWorkDir>/global-plugins
verboseCompile	grails.project.compile.verbose	false

The `BuildSettings` class has some other properties too, but they should be treated as read-only:

Property	Description
baseDir	The location of the project.
userHome	The user's home directory.
grailsHome	The location of the Grails installation in use (may be null).
grailsVersion	The version of Grails being used by the project.
grailsEnv	The current Grails environment.
compileDependencies	A list of compile-time project dependencies as <code>File</code> instances.
testDependencies	A list of test-time project dependencies as <code>File</code> instances.
runtimeDependencies	A list of runtime-time project dependencies as <code>File</code> instances.

Of course, these properties aren't much good if you can't get hold of them. Fortunately that's easy to do - they're available to your scripts as the `grailsSettings` script variable. You can also access them via the `grails.util.BuildSettingsHolder` class, but this isn't recommended.

Overriding the defaults

All of the properties in the first table can be overridden by a system property or a configuration option - for example, to change the project working directory, you could either run this command:

```
grails -Dgrails.project.work.dir=work compile
```

or add this option to your `grails-app/conf/BuildConfig.groovy` file:

```
grails.project.work.dir = "work"
```

Note that the default values take account of the property values they depend on, so setting the project working directory will relocate the compiled classes, test classes, resources, and plugins.

What happens if you use both a system property and a configuration option? Then the system property wins. The `BuildConfig.groovy` file, which in turn takes precedence over the default values.

The `BuildConfig.groovy` file is a sibling of `grails-app/conf/Config.groovy` - the former is used for build, whereas the latter contains those that affect the application at runtime. It's not limited to the optional build configuration options dotted around the documentation, such as ones for specifying the port that the application listens on for determining what files get packaged in the WAR file.

Available build settings

Name	Description
<code>grails.server.port.http</code>	Port to run the embedded servlet container on ("run-app" and "run-war"). If not set, the default is 8080.
<code>grails.server.port.https</code>	Port to run the embedded servlet container on for HTTPS ("run-app --https"). If not set, the default is 8443.
<code>grails.config.base.webXml</code>	Path to a custom web.xml file to use for the application (alternative to using <code>grails.config.web.xml</code>).
<code>grails.compiler.dependencies</code>	Legacy approach to adding extra dependencies to the compiler classpath. See the documentation for more details. These entries will be processed by an <code>AntBuilder</code> so the corresponding XML elements in an Ant build file, e.g. <code>fileset(dir: "*/**/*.class")</code> .
<code>grails.testing.patterns</code>	A list of Ant path patterns that let you control which files are included in the test case suffix, which is set by the next property.
<code>grails.testing.nameSuffix</code>	By default, tests are assumed to have a suffix of "Tests". You can change this with the <code>grails.testing.nameSuffix</code> property. For example, another common suffix is "Test".
<code>grails.project.war.file</code>	A string containing the file path of the generated WAR file, along with the name of the application. For example, "target/my-app.war".
<code>grails.war.dependencies</code>	A closure containing "fileset()" entries that allows you complete control over the files included in the "WEB-INF/lib" directory.
<code>grails.war.copyToWebApp</code>	A closure containing "fileset()" entries that allows you complete control over the files copied to the web application. It overrides the default behaviour of including everything under "web-app".
<code>grails.war.resources</code>	A closure that takes the location of the staging directory as its first argument and returns a list of files to be removed from the staging directory before packaging into a WAR.
<code>grails.project.web.xml</code>	The location to generate Grails' web.xml to

4.6 Ant and Maven

If all the other projects in your team or company are built using a standard build tool such as Ant or Maven, you can use the Grails command line to build your application. Fortunately, you can easily integrate these main build tools in use today (well, the ones in use in Java projects at least).

Ant Integration

When you create a Grails application with the [create-app](#) command, Grails doesn't automatically create a `build.xml` file. You can generate one with the [integrate-with](#) command:

```
grails integrate-with --ant
```

This creates a `build.xml` file containing the following targets:

- `clean` - Cleans the Grails application
- `compile` - Compiles your application's source code
- `test` - Runs the unit tests
- `run` - Equivalent to "grails run-app"
- `war` - Creates a WAR file
- `deploy` - Empty by default, but can be used to implement automatic deployment

Each of these can be run by Ant, for example:

```
ant war
```

The build file is configured to use [Apache Ivy](#) for dependency management, which means that it will automatically download Grails JAR files and other dependencies on demand. You don't even have to install Grails locally to use it. This makes it suitable for continuous integration systems such as [CruiseControl](#) or [Jenkins](#).

It uses the Grails [Ant task](#) to hook into the existing Grails build system. The task lets you run any Grails script used by the generated build file. To use the task, you must first declare it:

```
<taskdef name="grailsTask"
         classname="grails.ant.GrailsTask"
         classpathref="grails.classpath"/>
```

This raises the question: what should be in "grails.classpath"? The task itself is in the "grails-bootstrap" artifact, so that classpath at least. You should also include the "groovy-all" JAR. With the task defined, you just need to know what attributes are available:

Attribute	Description	Required
home	The location of the Grails installation directory to use for the build.	Yes
classpathref	Classpath to load Grails from. Must include the "grails-bootstrap" artifact and should include "grails-scripts".	Yes
script	The name of the Grails script to run, e.g. "TestApp".	Yes
args	The arguments to pass to the script, e.g. "-unit -xml".	No
environment	The Grails environment to run the script in.	No
includeRuntimeClasspath	Advanced setting: adds the application's runtime classpath to the build classpath if true.	No

The task also supports the following nested elements, all of which are standard Ant path structures:

- `classpath` - The build classpath (used to load Gant and the Grails scripts).
- `compileClasspath` - Classpath used to compile the application's classes.
- `runtimeClasspath` - Classpath used to run the application and package the WAR @`compileClasspath`.
- `testClasspath` - Classpath used to compile and run the tests. Typically includes everything in ru

How you populate these paths is up to you. If you use the `home` attribute and put your own dependencies, you may even need to use any of them. For an example of their use, take a look at the generated Ant build file for ne

Maven Integration

Grails provides integration with [Maven 2](#) with a Maven plugin. The current Maven plugin is based on but who did a great job with the original.

Preparation

In order to use the new plugin, all you need is Maven 2 installed and set up. This is because **you no longer use it with Maven!**



The Maven 2 integration for Grails has been designed and tested for Maven 2.0.9 and above earlier versions.



The default mvn setup DOES NOT supply sufficient memory to run the Grails environment. you add the following environment variable setting to prevent poor performance:

```
export MAVEN_OPTS="-Xmx512m -XX:MaxPermSize=256"
```

Creating a Grails Maven Project

To create a Mavenized Grails project simply run the following command:

```
mvn archetype:generate -DarchetypeGroupId=org.grails \
-DarchetypeArtifactId=grails-maven-archetype \
-DarchetypeVersion=1.3.2 \
-DgroupId=example -DartifactId=my-app
```

Choose whichever grails version, group ID and artifact ID you want for your application, but everything else is a new Maven project with a POM and a couple of other files. What you won't see is anything that looks like the old Grails project structure. The first step is to create the project structure that you're used to. But first, to set target JDK to Java 6, do that now.

```
<plugin>
  <artifactId>maven-compiler-plugin</artifactId>
  <configuration>
    <source>1.5</source>
    <target>1.5</target>
  </configuration>
</plugin>
```

to

```
<plugin>
  <artifactId>maven-compiler-plugin</artifactId>
  <configuration>
    <source>1.6</source>
    <target>1.6</target>
  </configuration>
</plugin>
```

Then you're ready to create the project structure:

```
cd my-app
mvn initialize
```



if you see a message similar to this:

```
Resolving plugin JAR dependencies ...
:: problems summary ::
:::: WARNINGS
      module not found: org.hibernate#hibernate-core;3.3.1.GA
```

you need to add the plugins manually to application.properties:

```
plugins.hibernate=2.0.0
plugins.tomcat=2.0.0
```

then run

```
mvn compile
```

and the hibernate and tomcat plugins will be installed.

Now you have a Grails application all ready to go. The plugin integrates into the standard build cycle, so you can build and package your app: `mvn clean, mvn compile, mvn test, mvn package, mvn install`

You can also use some of the Grails commands that have been wrapped as Maven goals:

- `grails:create-controller` - Calls the [create-controller](#) command
- `grails:create-domain-class` - Calls the [create-domain-class](#) command
- `grails:create-integration-test` - Calls the [create-integration-test](#) command
- `grails:create-pom` - Creates a new Maven POM for an existing Grails project
- `grails:create-script` - Calls the [create-script](#) command
- `grails:create-service` - Calls the [create-service](#) command
- `grails:create-taglib` - Calls the [create-tag-lib](#) command
- `grails:create-unit-test` - Calls the [create-unit-test](#) command
- `grails:exec` - Executes an arbitrary Grails command line script
- `grails:generate-all` - Calls the [generate-all](#) command
- `grails:generate-controller` - Calls the [generate-controller](#) command
- `grails:generate-views` - Calls the [generate-views](#) command
- `grails:install-plugin` - Calls the [install-plugin](#) command
- `grails:install-templates` - Calls the [install-templates](#) command
- `grails:list-plugins` - Calls the [list-plugins](#) command
- `grails:package` - Calls the [package](#) command
- `grails:run-app` - Calls the [run-app](#) command
- `grails:uninstall-plugin` - Calls the [uninstall-plugin](#) command

For a complete, up to date list, run `mvn grails:help`

Mavenizing an existing project

Creating a new project is great way to start, but what if you already have one? You don't want to create a of the old one over. The solution is to create a POM for the existing project using this Maven command (grails version of your existing project):

```
mvn org.grails:grails-maven-plugin:1.3.2:create-pom -DgroupId=com.mycompany
```

When this command has finished, you can immediately start using the standard phases, such as `mvn package` with your group ID when creating the POM.

You may also want to set target JDK to Java 6; see above.

Adding Grails commands to phases

The standard POM created for you by Grails already attaches the appropriate core Grails commands to the standard Maven phases. For example, the `compile` command goes in the "compile" phase and `war` goes in the "package" phase. That doesn't help though if you want to attach a command to a particular phase. The classic example is functional tests. How do you make sure that your functional tests (or whatever you have decided on) are run during the "integration-test" phase?

Fear not: all things are possible. In this case, you can associate the command to a phase using an extra "exec

```
<plugin>
  <groupId>org.grails</groupId>
  <artifactId>grails-maven-plugin</artifactId>
  <version>1.3.2</version>
  <extensions>true</extensions>
  <executions>
    <execution>
      <goals>
        ...
      </goals>
    </execution>
    <!-- Add the "functional-tests" command to the "integration-test" phase -->
    <execution>
      <id>functional-tests</id>
      <phase>integration-test</phase>
      <goals>
        <goal>exec</goal>
      </goals>
      <configuration>
        <command>functional-tests</command>
      </configuration>
    </execution>
  </executions>
</plugin>
```

This also demonstrates the `grails:exec` goal, which can be used to run any Grails command. Simply command system property, and optionally specify the arguments with the `args` property:

```
mvn grails:exec -Dcommand=create-webtest -Dargs=Book
```

Debugging a Grails Maven Project

Maven can be launched in debug mode using the "mvnDebug" command. To launch your Grails applicatio

```
mvnDebug grails:run-app
```

The process will be suspended on startup and listening for a debugger on port 8000.

If you need more control of the debugger, this can be specified using the `MAVEN_OPTS` environment default "mvn" command:

```
MAVEN_OPTS="-Xdebug -Xrunjdpw:transport=dt_socket,server=y,suspend=y,address=5005"
mvn grails:run-app
```

Raising issues

If you come across any problems with the Maven integration, please raise a JIRA issue as a sub-task of [GF](#)

5 Object Relational Mapping (GORM)

Domain classes are core to any business application. They hold state about business processes and hopefully linked together through relationships; one-to-one, one-to-many, or many-to-many.

GORM is Grails' object relational mapping (ORM) implementation. Under the hood it uses Hibernate 3 (an ORM solution) and thanks to the dynamic nature of Groovy with its static and dynamic typing, along with less configuration involved in creating Grails domain classes.

You can also write Grails domain classes in Java. See the section on Hibernate Integration for how to write dynamic persistent methods. Below is a preview of GORM in action:

```
def book = Book.findByTitle("Groovy in Action")

book
  .addToAuthors(name:"Dierk Koenig")
  .addToAuthors(name:"Guillaume LaForge")
  .save()
```

5.1 Quick Start Guide

A domain class can be created with the [create-domain-class](#) command:

```
grails create-domain-class helloworld.Person
```



If no package is specified with the create-domain-class script, Grails automatically uses the package name.

This will create a class at the location `grails-app/domain/helloworld/Person.groovy` such

```
package helloworld

class Person {
}
```



If you have the `dbCreate` property set to "update", "create" or "create-drop" on your [Database](#), Grails will automatically generate/modify the database tables for you.

You can customize the class by adding properties:

```
class Person {
  String name
  Integer age
  Date lastVisit
}
```

Once you have a domain class try and manipulate it with the [shell](#) or [console](#) by typing:

```
grails console
```

This loads an interactive GUI where you can run Groovy commands with access to the Spring Application

5.1.1 Basic CRUD

Try performing some basic CRUD (Create/Read/Update/Delete) operations.

Create

To create a domain class use Map constructor to set its properties and call [save](#):

```
def p = new Person(name: "Fred", age: 40, lastVisit: new Date())
p.save()
```

The [save](#) method will persist your class to the database using the underlying Hibernate ORM layer.

Read

Grails transparently adds an implicit `id` property to your domain class which you can use for retrieval:

```
def p = Person.get(1)
assert 1 == p.id
```

This uses the [get](#) method that expects a database identifier to read the `Person` object back from the data read-only state by using the [read](#) method:

```
def p = Person.read(1)
```

In this case the underlying Hibernate engine will not do any dirty checking and the object will not be persisted. The [save](#) method then the object is placed back into a read-write state.

In addition, you can also load a proxy for an instance by using the [load](#) method:

```
def p = Person.load(1)
```

This incurs no database access until a method other than `getId()` is called. Hibernate then initializes the proxy if no record is found for the specified id.

Update

To update an instance, change some properties and then call [save](#) again:

```
def p = Person.get(1)
p.name = "Bob"
p.save()
```

Delete

To delete an instance use the [delete](#) method:

```
def p = Person.get(1)
p.delete()
```

5.2 Domain Modelling in GORM

When building Grails applications you have to consider the problem domain you are trying to solve. [Amazon](#)-style bookstore you would be thinking about books, authors, customers and publishers to name a

These are modeled in GORM as Groovy classes, so a Book class may have a title, a release date, an ISBN and so on. The following sections show how to model the domain in GORM.

To create a domain class you run the [create-domain-class](#) command as follows:

```
grails create-domain-class org.bookstore.Book
```

The result will be a class at `grails-app/domain/org/bookstore/Book.groovy`:

```
package org.bookstore

class Book {
}
```

This class will map automatically to a table in the database called `book` (the same name as the class). This is the [ORM Domain Specific Language](#).

Now that you have a domain class you can define its properties as Java types. For example:

```
package org.bookstore

class Book {
    String title
    Date releaseDate
    String ISBN
}
```

Each property is mapped to a column in the database, where the convention for column names is all lower case. For example `releaseDate` maps onto a column `release_date`. The SQL types are auto-detected from the property names with [Constraints](#) or the [ORM DSL](#).

5.2.1 Association in GORM

Relationships define how domain classes interact with each other. Unless specified explicitly at both direction it is defined.

5.2.1.1 Many-to-one and one-to-one

A many-to-one relationship is the simplest kind, and is defined with a property of the type of another domain

Example A

```
class Face {  
    Nose nose  
}
```

```
class Nose {  
}
```

In this case we have a unidirectional many-to-one relationship from Face to Nose. To make this relationship as follows:

Example B

```
class Face {  
    Nose nose  
}
```

```
class Nose {  
    static belongsTo = [face:Face]  
}
```

In this case we use the `belongsTo` setting to say that Nose "belongs to" Face. The result of this is that when we save or delete the Face instance, GORM will save or delete the Nose. In other words, the relationship is from Face to the associated Nose:

```
new Face(nose:new Nose()).save()
```

The example above will save both face and nose. Note that the inverse *is not* true and will result in an error:

```
new Nose(face:new Face()).save() // will cause an error
```

Now if we delete the Face instance, the Nose will go too:

```
def f = Face.get(1)
f.delete() // both Face and Nose deleted
```

To make the relationship a true one-to-one, use the `hasOne` property on the owning side, e.g. `Face`:

Example C

```
class Face {
    static hasOne = [nose:Nose]
}
```

```
class Nose {
    Face face
}
```

Note that using this property puts the foreign key on the inverse table to the previous example, so in this case the `nose` table inside a column called `face_id`. Also, `hasOne` only works with bidirectional relationships.

Finally, it's a good idea to add a unique constraint on one side of the one-to-one relationship:

```
class Face {
    static hasOne = [nose:Nose]
    static constraints = {
        nose unique: true
    }
}
```

```
class Nose {
    Face face
}
```


5.2.1.2 One-to-many

A one-to-many relationship is when one class, example `Author`, has many instances of a another class, such a relationship with the `hasMany` setting:

```
class Author {
    static hasMany = [books: Book]
    String name
}
```


```
class Book {
    String title
}
```

In this case we have a unidirectional one-to-many. Grails will, by default, map this kind of relationship with

 The [ORM DSL](#) allows mapping unidirectional relationships using a foreign key association in

Grails will automatically inject a property of type `java.util.Set` into the domain class based on the iterate over the collection:

```
def a = Author.get(1)
for (book in a.books) {
    println book.title
}
```

 The default fetch strategy used by Grails is "lazy", which means that the collection will be lazy access. This can lead to the [n+1 problem](#) if you are not careful.

If you need "eager" fetching you can use the [ORM DSL](#) or specify eager fetching as part of a

The default cascading behaviour is to cascade saves and updates, but not deletes unless a `belongsTo` is a

```
class Author {
    static hasMany = [books: Book]
    String name
}
```

```
class Book {
    static belongsTo = [author: Author]
    String title
}
```

If you have two properties of the same type on the many side of a one-to-many you have to use `mappedBy` mapped:

```
class Airport {
    static hasMany = [flights: Flight]
    static mappedBy = [flights: "departureAirport"]
}
```

```
class Flight {
    Airport departureAirport
    Airport destinationAirport
}
```

This is also true if you have multiple collections that map to different properties on the many side:

```
class Airport {
    static hasMany = [outboundFlights: Flight, inboundFlights: Flight]
    static mappedBy = [outboundFlights: "departureAirport",
                      inboundFlights: "destinationAirport"]
}
```

```
class Flight {
    Airport departureAirport
    Airport destinationAirport
}
```

5.2.1.3 Many-to-many

Grails supports many-to-many relationships by defining a `hasMany` on both sides of the relationship and side of the relationship:

```
class Book {
    static belongsTo = Author
    static hasMany = [authors: Author]
    String title
}
```

```
class Author {
    static hasMany = [books: Book]
    String name
}
```

Grails maps a many-to-many using a join table at the database level. The owning side of the relationship has responsibility for persisting the relationship and is the only side that can cascade saves across.


For example this will work and cascade saves:

```
new Author(name: "Stephen King")
    .addToBooks(new Book(title: "The Stand"))
    .addToBooks(new Book(title: "The Shining"))
    .save()
```

However this will only save the Book and not the authors!

```
new Book(name:"Groovy in Action")
    .addToAuthors(new Author(name:"Dierk Koenig"))
    .addToAuthors(new Author(name:"Guillaume Laforge"))
    .save()
```

This is the expected behaviour as, just like Hibernate, only one side of a many-to-many can take responsib

 Grails' [Scaffolding](#) feature **does not** currently support many-to-many relationship and hence y
to manage the relationship yourself

5.2.1.4 Basic Collection Types

As well as associations between different domain classes, GORM also supports mapping of basic collec
class creates a nicknames association that is a Set of String instances:

```
class Person {
    static hasMany = [nicknames: String]
}
```

GORM will map an association like the above using a join table. You can alter various aspects of ho
joinTable argument:

```
class Person {
    static hasMany = [nicknames: String]
    static mapping = {
        hasMany joinTable: [name: 'bunch_o_nicknames',
                           key: 'person_id',
                           column: 'nickname',
                           type: "text"]
    }
}
```

The example above will map to a table that looks like the following:

bunch_o_nicknames Table

person_id	nickname
1	Fred

5.2.2 Composition in GORM

As well as [association](#), Grails supports the notion of composition. In this case instead of mapping clas
"embedded" within the current table. For example:

```

class Person {
    Address homeAddress
    Address workAddress
    static embedded = ['homeAddress', 'workAddress']
}

class Address {
    String number
    String code
}

```

The resulting mapping would look like this:

Person Table

id	home_address _number	home_address _code	work_address _number	work_address _code
1	47	343432	67	43545



If you define the Address class in a separate Groovy file in the `grails-app/domain` package, you will get an address table. If you don't want this to happen use Groovy's ability to define multiple classes in a single file. Include the Address class below the Person class in the `grails-app/domain/Person.groovy` file.

5.2.3 Inheritance in GORM

GORM supports inheritance both from abstract base classes and concrete persistent GORM entities. For example:

```

class Content {
    String author
}

```

```

class BlogEntry extends Content {
    URL url
}

```

```

class Book extends Content {
    String ISBN
}

```

```

class PodCast extends Content {
    byte[] audioStream
}

```

In the above example we have a parent `Content` class and then various child classes with more specific t

Considerations

At the database level Grails by default uses table-per-hierarchy mapping with a discriminator column `Content`) and its subclasses (`BlogEntry`, `Book` etc.), share the **same** table.

Table-per-hierarchy mapping has a down side in that you **cannot** have non-nullable properties with inheritance table-per-subclass which can be enabled with the [ORM DSL](#)

However, excessive use of inheritance and table-per-subclass can result in poor query performance due to our advice is if you're going to use inheritance, don't abuse it and don't make your inheritance hierarchy too

Polymorphic Queries

The upshot of inheritance is that you get the ability to polymorphically query. For example using the [list](#) will return all subclasses of `Content`:

```
def content = Content.list() // list all blog entries, books and podcasts
content = Content.findAllByAuthor('Joe Bloggs') // find all by author

def podCasts = PodCast.list() // list only podcasts
```

5.2.4 Sets, Lists and Maps

Sets of Objects

By default when you define a relationship with GORM it is a `java.util.Set` which is an unordered collection. In other words when you have:

```
class Author {
    static hasMany = [books: Book]
}
```

The `books` property that GORM injects is a `java.util.Set`. Sets guarantee uniqueness but not order, if you have custom ordering you configure the `Set` as a `SortedSet`:

```
class Author {
    SortedSet books
    static hasMany = [books: Book]
}
```

In this case a `java.util.SortedSet` implementation is used which means you must implement `java.util.SortedSet` class:


```
class Book implements Comparable {
    String title
    Date releaseDate = new Date()
    int compareTo(obj) {
        releaseDate.compareTo(obj.releaseDate)
    }
}
```

The result of the above class is that the Book instances in the books collection of the Author class will be c

Lists of Objects

To keep objects in the order which they were added and to be able to reference them by index like an array a List:

```
class Author {
    List books
    static hasMany = [books: Book]
}
```

In this case when you add new elements to the books collection the order is retained in a sequential list ind

```
author.books[0] // get the first book
```

The way this works at the database level is Hibernate creates a books_idx column where it saves the in retain this order at the database level.

When using a List, elements must be added to the collection before being saved, otherwise I org.hibernate.HibernateException: null index column for collection):

```
// This won't work!
def book = new Book(title: 'The Shining')
book.save()
author.addToBooks(book)

// Do it this way instead.
def book = new Book(title: 'Misery')
author.addToBooks(book)
author.save()
```

Bags of Objects

If ordering and uniqueness aren't a concern (or if you manage these explicitly) then you can use the Hi collections.

The only change required for this is to define the collection type as a Collection:

```
class Author {
  Collection books

  static hasMany = [books: Book]
}
```

Since uniqueness and order aren't managed by Hibernate, adding to or removing from collections maps existing instances from the database, so this approach will perform better and require less memory than using

Maps of Objects

If you want a simple map of string/value pairs GORM can map this with the following:

```
class Author {
  Map books // map of ISBN:book names
}

def a = new Author()
a.books = ["1590597583":"Grails Book"]
a.save()
```

In this case the key and value of the map **MUST** be strings.

If you want a Map of objects then you can do this:

```
class Book {
  Map authors

  static hasMany = [authors: Author]
}

def a = new Author(name:"Stephen King")

def book = new Book()
book.authors = [stephen:a]
book.save()
```

The static hasMany property defines the type of the elements within the Map. The keys for the map **must**

A Note on Collection Types and Performance

The Java Set type doesn't allow duplicates. To ensure uniqueness when adding an entry to a Set association from the database. If you have a large numbers of entries in the association this can be costly in

The same behavior is required for List types, since Hibernate needs to load the entire association recommended that if you anticipate a large numbers of records in the association that you make the association be created on the inverse side. For example consider the following code:

```
def book = new Book(title:"New Grails Book")
def author = Author.get(1)
book.author = author
book.save()
```

In this example the association link is being created by the child (Book) and hence it is not necessary resulting in fewer queries and more efficient code. Given an Author with a large number of associated code like the following you would see an impact on performance:

```
def book = new Book(title:"New Grails Book")
def author = Author.get(1)
author.addToBooks(book)
author.save()
```

You could also model the collection as a Hibernate Bag as described above.

5.3 Persistence Basics

A key thing to remember about Grails is that under the surface Grails is using [Hibernate](#) for persistence. If using [ActiveRecord](#) or [iBatis](#) Hibernate's "session" model may feel a little strange.

Grails automatically binds a Hibernate session to the currently executing request. This lets you use the [save](#) GORM methods transparently.

Transactional Write-Behind

A useful feature of Hibernate over direct JDBC calls and even other frameworks is that when you call perform any SQL operations **at that point**. Hibernate batches up SQL statements and executes them as request when flushing and closing the session. This is typically done for you automatically by Grails, which

Hibernate caches database updates where possible, only actually pushing the changes when it knows that triggered programmatically. One common case where Hibernate will flush cached updates is when information might be included in the query results. But as long as you're doing non-conflicting saves, until the session is flushed. This can be a significant performance boost for applications that do a lot of data

Note that flushing is not the same as committing a transaction. If your actions are performed in the context SQL updates but the database will save the changes in its transaction queue and only finalize the updates when

5.3.1 Saving and Updating

An example of using the [save](#) method can be seen below:

```
def p = Person.get(1)
p.save()
```

This save will be not be pushed to the database immediately - it will be pushed when the next flush occurs. You can want to control when those statements are executed or, in Hibernate terminology, when the session is "flushed" by passing an argument to the save method:

```
def p = Person.get(1)
p.save(flush: true)
```

Note that in this case *all* pending SQL statements including previous saves, deletes, etc. will be synchronized to the database. You can catch any exceptions, which is typically useful in highly concurrent scenarios involving [optimistic locking](#).

```
def p = Person.get(1)
try {
    p.save(flush: true)
}
catch (org.springframework.dao.DataIntegrityViolationException e) {
    // deal with exception
}
```

Another thing to bear in mind is that Grails [validates](#) a domain instance every time you save it. If that validation fails, the instance will not be persisted to the database. By default, `save()` will simply return `null` in this case, but if you would use the `failOnError` argument:

```
def p = Person.get(1)
try {
    p.save(failOnError: true)
}
catch (ValidationException e) {
    // deal with exception
}
```

You can even change the default behaviour with a setting in `Config.groovy`, as described in the [section](#) when you are saving domain instances that have been bound with data provided by the user, the likelihood and you won't want those exceptions propagating to the end user.

You can find out more about the subtleties of saving data in [this article](#) - a must read!

5.3.2 Deleting Objects

An example of the [delete](#) method can be seen below:

```
def p = Person.get(1)
p.delete()
```

As with saves, Hibernate will use transactional write-behind to perform the delete; to perform the delete immediately, you can use the `flush` argument:

```
def p = Person.get(1)
p.delete(flush: true)
```

Using the `flush` argument lets you catch any errors that occur during a delete. A common error that occurs is a `DataIntegrityViolationException`, although this is normally down to a programming or schema error. The following `DataIntegrityViolationException` that is thrown when you violate the database constraints:

```
def p = Person.get(1)

try {
    p.delete(flush: true)
}
catch (org.springframework.dao.DataIntegrityViolationException e) {
    flash.message = "Could not delete person ${p.name}"
    redirect(action: "show", id: p.id)
}
```

Note that Grails does not supply a `deleteAll` method as deleting data is discouraged and can often be a bad idea. If you really need to batch delete data you can use the [executeUpdate](#) method to do batch DML statements

```
Customer.executeUpdate("delete Customer c where c.name = :oldName",
    [oldName: "Fred"])
```

5.3.3 Understanding Cascading Updates and Deletes

It is critical that you understand how cascading updates and deletes work when using GORM. The key setting which controls which class "owns" a relationship.

Whether it is a one-to-one, one-to-many or many-to-many, defining `belongsTo` will result in updates dependant (the other side of the relationship), and for many-/one-to-one and one-to-many relationships delete.

If you *do not* define `belongsTo` then no cascades will happen and you will have to manually save one-to-many, in which case saves will cascade automatically if a new instance is in a `hasMany` collection

Here is an example:

```
class Airport {
    String name
    static hasMany = [flights: Flight]
}
```

```
class Flight {
    String number
    static belongsTo = [airport: Airport]
}
```

If I now create an `Airport` and add some `Flights` to it I can save the `Airport` and have the updates cascade saving the whole object graph:

```
new Airport(name: "Gatwick")
    .addToFlights(new Flight(number: "BA3430"))
    .addToFlights(new Flight(number: "EZ0938"))
    .save()
```

Conversely if I later delete the `Airport` all `Flights` associated with it will also be deleted:

```
def airport = Airport.findByName("Gatwick")
airport.delete()
```

However, if I were to remove `belongsTo` then the above cascading deletion code **would not work**. To summaries below that describe the default behaviour of GORM with regards to specific associations. A series of articles to get a deeper understanding of relationships and cascading.

Bidirectional one-to-many with `belongsTo`

```
class A { static hasMany = [bees: B] }
```

```
class B { static belongsTo = [a: A] }
```

In the case of a bidirectional one-to-many where the many side defines a `belongsTo` then the cascade is `delete` and "NONE" for the many side.

Unidirectional one-to-many

```
class A { static hasMany = [bees: B] }
```

```
class B { }
```

In the case of a unidirectional one-to-many where the many side defines no `belongsTo` then the cascade strategy is `delete`.

Bidirectional one-to-many, no `belongsTo`

```
class A { static hasMany = [bees: B] }
```

```
class B { A a }
```

In the case of a bidirectional one-to-many where the many side does not define a `belongsTo` then the cascade strategy is `delete` for the one side and "NONE" for the many side.

Unidirectional one-to-one with `belongsTo`

```
class A { }
```

```
class B { static belongsTo = [a: A] }
```

In the case of a unidirectional one-to-one association that defines a `belongsTo` then the cascade strategy the relationship (A->B) and "NONE" from the side that defines the `belongsTo` (B->A)

Note that if you need further control over cascading behaviour, you can use the [ORM DSL](#).

5.3.4 Eager and Lazy Fetching

Associations in GORM are by default lazy. This is best explained by example:

```
class Airport {  
    String name  
    static hasMany = [flights: Flight]  
}
```

```
class Flight {  
    String number  
    Location destination  
    static belongsTo = [airport: Airport]  
}
```

```
class Location {  
    String city  
    String country  
}
```

Given the above domain classes and the following code:

```
def airport = Airport.findByName("Gatwick")  
for (flight in airport.flights) {  
    println flight.destination.city  
}
```

GORM will execute a single SQL query to fetch the `Airport` instance, another to get its flights, and then the flights association to get the current flight's destination. In other words you get N+1 queries (if `airport`).

Configuring Eager Fetching

An alternative approach that avoids the N+1 queries is to use eager fetching, which can be specified as follows:

```
class Airport {
    String name
    static hasMany = [flights: Flight]
    static mapping = {
        flights lazy: false
    }
}
```

In this case the `flights` association will be loaded at the same time as its `Airport` instance, although you can fetch the collection. You can also use `fetch: 'join'` instead of `lazy: false`, in which case GORM will fetch the airports and their flights. This works well for single-ended associations, but you need to be careful as you'd expect right up to the moment you add a limit to the number of results you want. At that point, you'll get more than you were expecting. The reason for this is quite technical but ultimately the problem arises from GORM.

So, the recommendation is currently to use `fetch: 'join'` for single-ended associations and `lazy: false`.

Be careful how and where you use eager loading because you could load your entire database into memory. You can find more information on the mapping options in the [section on the ORM DSL](#).

Using Batch Fetching

Although eager fetching is appropriate for some cases, it is not always desirable. If you made everything an eager fetch, you could load your entire database into memory resulting in performance and memory problems. An alternative to eager fetching is to configure Hibernate to lazily fetch results in "batches". For example:

```
class Airport {
    String name
    static hasMany = [flights: Flight]
    static mapping = {
        flights batchSize: 10
    }
}
```

In this case, due to the `batchSize` argument, when you iterate over the `flights` association, Hibernate will fetch them in batches. For example, if you had an `Airport` that had 30 flights, if you didn't configure batch fetching you would get 30 queries to fetch each flight. With batch fetching you get 1 query to fetch the `Airport` and 3 queries to fetch the flights. In other words, batch fetching is an optimization of the lazy fetching strategy. Batch fetching can also be configured for the `Flight` class.

```
class Flight {
    ...
    static mapping = {
        batchSize 10
    }
}
```

Check out [part 3](#) of the GORM Gotchas series for more in-depth coverage of this tricky topic.

5.3.5 Pessimistic and Optimistic Locking

Optimistic Locking

By default GORM classes are configured for optimistic locking. Optimistic locking is a feature of Hibe value in a special `version` column in the database that is incremented after each update.

The `version` column gets read into a `version` property that contains the current versioned state of pers

```
def airport = Airport.get(10)
println airport.version
```

When you perform updates Hibernate will automatically check the `version` property against the version co will throw a [StaleObjectException](#). This will roll back the transaction if one is active.

This is useful as it allows a certain level of atomicity without resorting to pessimistic locking that ha downside is that you have to deal with this exception if you have highly concurrent writes. This requires fl

```
def airport = Airport.get(10)
try {
    airport.name = "Heathrow"
    airport.save(flush: true)
}
catch (org.springframework.dao.OptimisticLockingFailureException e) {
    // deal with exception
}
```

The way you deal with the exception depends on the application. You could attempt a programmatic merg ask them to resolve the conflict.

Alternatively, if it becomes a problem you can resort to pessimistic locking.



The `version` will only be updated after flushing the session.

Pessimistic Locking

Pessimistic locking is equivalent to doing a SQL "SELECT * FOR UPDATE" statement and locking; implication that other read operations will be blocking until the lock is released.

In Grails pessimistic locking is performed on an existing instance with the [lock](#) method:

```
def airport = Airport.get(10)
airport.lock() // lock for update
airport.name = "Heathrow"
airport.save()
```

Grails will automatically deal with releasing the lock for you once the transaction has been committed. H doing is "upgrading" from a regular SELECT to a SELECT..FOR UPDATE and another thread could still call to `get()` and the call to `lock()`.

To get around this problem you can use the static [lock](#) method that takes an id just like [get](#):

```
def airport = Airport.lock(10) // lock for update
airport.name = "Heathrow"
airport.save()
```

In this case only SELECT..FOR UPDATE is issued.

As well as the [lock](#) method you can also obtain a pessimistic locking using queries. For example using a d

```
def airport = Airport.findByName("Heathrow", [lock: true])
```

Or using criteria:

```
def airport = Airport.createCriteria().get {
  eq('name', 'Heathrow')
  lock true
}
```

5.3.6 Modification Checking

Once you have loaded and possibly modified a persistent domain class instance, it isn't straightforward to reload the instance using [get](#) Hibernate will return the current modified instance from its Session cache trigger a flush which could cause problems if your data isn't ready to be flushed yet. So GORM provides values that Hibernate caches when it loads the instance (which it uses for dirty checking).

isDirty

You can use the [isDirty](#) method to check if any field has been modified:

```
def airport = Airport.get(10)
assert !airport.isDirty()

airport.properties = params
if (airport.isDirty()) {
  // do something based on changed state
}
```



`isDirty()` does not currently check collection associations, but it does check all other persistent associations.

You can also check if individual fields have been modified:

```
def airport = Airport.get(10)
assert !airport.isDirty()

airport.properties = params
if (airport.isDirty('name')) {
    // do something based on changed name
}
```

getDirtyPropertyNames

You can use the [getDirtyPropertyNames](#) method to retrieve the names of modified fields; this may be emp

```
def airport = Airport.get(10)
assert !airport.isDirty()

airport.properties = params
def modifiedFieldNames = airport.getDirtyPropertyNames()
for (fieldName in modifiedFieldNames) {
    // do something based on changed value
}
```

getPersistentValue

You can use the [getPersistentValue](#) method to retrieve the value of a modified field:

```
def airport = Airport.get(10)
assert !airport.isDirty()

airport.properties = params
def modifiedFieldNames = airport.getDirtyPropertyNames()
for (fieldName in modifiedFieldNames) {
    def currentValue = airport."$fieldName"
    def originalValue = airport.getPersistentValue(fieldName)
    if (currentValue != originalValue) {
        // do something based on changed value
    }
}
```

5.4 Querying with GORM

GORM supports a number of powerful ways to query from dynamic finders, to criteria to Hibernate's object

Groovy's ability to manipulate collections with [GPath](#) and methods like `sort`, `findAll` and so on combin

However, let's start with the basics.

Listing instances

Use the [list](#) method to obtain all instances of a given class:

```
def books = Book.list()
```

The [list](#) method supports arguments to perform pagination:

```
def books = Book.list(offset:10, max:20)
```

as well as sorting:

```
def books = Book.list(sort:"title", order:"asc")
```

Here, the `sort` argument is the name of the domain class property that you wish to sort on, and the `order` or `desc` for **descending**.

Retrieval by Database Identifier

The second basic form of retrieval is by database identifier using the [get](#) method:

```
def book = Book.get(23)
```

You can also obtain a list of instances for a set of identifiers using [getAll](#):

```
def books = Book.getAll(23, 93, 81)
```

5.4.1 Dynamic Finders

GORM supports the concept of **dynamic finders**. A dynamic finder looks like a static method invocation but actually exist in any form at the code level.

Instead, a method is auto-magically generated using code synthesis at runtime, based on the properties of the `Book` class:

```
class Book {  
    String title  
    Date releaseDate  
    Author author  
}
```

```
class Author {  
    String name  
}
```

The Book class has properties such as title, releaseDate and author. These can be used by the form of "method expressions":

```
def book = Book.findByTitle("The Stand")
book = Book.findByTitleLike("Harry Pot%")
book = Book.findByReleaseDateBetween(firstDate, secondDate)
book = Book.findByReleaseDateGreaterThan(someDate)
book = Book.findByTitleLikeOrReleaseDateLessThan("%Something%", someDate)
```

Method Expressions

A method expression in GORM is made up of the prefix such as [findBy](#) followed by an expression that basic form is:

```
Book.findBy([Property][Comparator][Boolean Operator])?[Property][Comparator]
```

The tokens marked with a '?' are optional. Each comparator changes the nature of the query. For example:

```
def book = Book.findByTitle("The Stand")
book = Book.findByTitleLike("Harry Pot%")
```

In the above example the first query is equivalent to equality whilst the latter, due to the Like comparison expression.

The possible comparators include:

- `InList` - In the list of given values
- `LessThan` - less than a given value
- `LessThanEquals` - less than or equal a given value
- `GreaterThan` - greater than a given value
- `GreaterThanEquals` - greater than or equal a given value
- `Like` - Equivalent to a SQL like expression
- `Ilike` - Similar to a Like, except case insensitive
- `NotEqual` - Negates equality
- `Between` - Between two values (requires two arguments)
- `IsNull` - Not a null value (doesn't take an argument)
- `NotNull` - Is a null value (doesn't take an argument)

Notice that the last three require different numbers of method arguments compared to the rest, as demonstrated

```
def now = new Date()
def lastWeek = now - 7
def book = Book.findByReleaseDateBetween(lastWeek, now)

books = Book.findAllByReleaseDateIsNull()
books = Book.findAllByReleaseDateIsNotNull()
```

Boolean logic (AND/OR)

Method expressions can also use a boolean operator to combine two or more criteria:

```
def books = Book.findAllByTitleLikeAndReleaseDateGreaterThan(
    "%Java%", new Date() - 30)
```

In this case we're using And in the middle of the query to make sure both conditions are satisfied, but you

```
def books = Book.findAllByTitleLikeOrReleaseDateGreaterThan(
    "%Java%", new Date() - 30)
```

You can combine as many criteria as you like, but they must all be combined with And or all Or. If you number of criteria creates a very long method name, just convert the query to a [Criteria](#) or [HQL](#) query.

Querying Associations

Associations can also be used within queries:

```
def author = Author.findByName("Stephen King")
def books = author ? Book.findAllByAuthor(author) : []
```

In this case if the Author instance is not null we use it in a query to obtain all the Book instances for the

Pagination and Sorting

The same pagination and sorting parameters available on the [list](#) method can also be used with dynamic parameter:

```
def books = Book.findAllByTitleLike("Harry Pot%",
    [max: 3, offset: 2, sort: "title", order: "desc"])
```

5.4.2 Criteria

Criteria is a type safe, advanced way to query that uses a Groovy builder to construct potentially complex than building up query strings using a StringBuffer.

Criteria can be used either with the [createCriteria](#) or [withCriteria](#) methods. The builder uses Hibernate's [Criteria](#) class to map the static methods found in the [Restrictions](#) class of the Hibernate Criteria API. For example:

```
def c = Account.createCriteria()
def results = c {
    between("balance", 500, 1000)
    eq("branch", "London")
    or {
        like("holderFirstName", "Fred%")
        like("holderFirstName", "Barney%")
    }
    maxResults(10)
    order("holderLastName", "desc")
}
```

This criteria will select up to 10 Account objects in a List matching the following criteria:

- balance is between 500 and 1000
- branch is 'London'
- holderFirstName starts with 'Fred' or 'Barney'

The results will be sorted in descending order by holderLastName.

If no records are found with the above criteria, an empty List is returned.

Conjunctions and Disjunctions

As demonstrated in the previous example you can group criteria in a logical OR using an `or { }` block:

```
or {
    between("balance", 500, 1000)
    eq("branch", "London")
}
```

This also works with logical AND:

```
and {
    between("balance", 500, 1000)
    eq("branch", "London")
}
```

And you can also negate using logical NOT:

```
not {
    between("balance", 500, 1000)
    eq("branch", "London")
}
```

All top level conditions are implied to be AND'd together.

Querying Associations

Associations can be queried by having a node that matches the property name. For example say the `Account` objects:

```
class Account {  
    ...  
    static hasMany = [transactions: Transaction]  
    ...  
}
```

We can query this association by using the property name `transaction` as a builder node:

```
def c = Account.createCriteria()  
def now = new Date()  
def results = c.list {  
    transactions {  
        between('date', now - 10, now)  
    }  
}
```

The above code will find all the `Account` instances that have performed transactions within the association queries within logical blocks:

```
def c = Account.createCriteria()  
def now = new Date()  
def results = c.list {  
    or {  
        between('created', now - 10, now)  
        transactions {  
            between('date', now - 10, now)  
        }  
    }  
}
```

Here we find all accounts that have either performed transactions in the last 10 days OR have been recently

Querying with Projections

Projections may be used to customise the results. Define a "projections" node within the criteria build equivalent methods within the projections node to the methods found in the Hibernate [Projections](#) class:

```
def c = Account.createCriteria()  
def numberOfBranches = c.get {  
    projections {  
        countDistinct('branch')  
    }  
}
```

When multiple fields are specified in the projection, a List of values will be returned. A single value will b

Using SQL Restrictions

You can access Hibernate's SQL Restrictions capabilities.

```
def c = Person.createCriteria()
def peopleWithShortFirstNames = c.list {
    sqlRestriction "char_length(first_name) <= 4"
}
```



Note that the parameter there is SQL. The `first_name` attribute referenced in the persistence model, not the object model like in HQL queries. The `Person` property named `first_name` maps to the `first_name` column in the database and you must refer to that in the `sqlRestriction` method.

Also note that the SQL used here is not necessarily portable across databases.

Using Scrollable Results

You can use Hibernate's [ScrollableResults](#) feature by calling the `scroll` method:

```
def results = crit.scroll {
    maxResults(10)
}
def f = results.first()
def l = results.last()
def n = results.next()
def p = results.previous()

def future = results.scroll(10)
def accountNumber = results.getLong('number')
```

To quote the documentation of Hibernate `ScrollableResults`:

A result iterator that allows moving around within the results by arbitrary increments. The Query / similar to the JDBC PreparedStatement/ ResultSet pattern and the semantics of methods of this interface named methods on ResultSet.

Contrary to JDBC, columns of results are numbered from zero.

Setting properties in the Criteria instance

If a node within the builder tree doesn't match a particular criterion it will attempt to set a property on the object. This example calls `setMaxResults` and `setFirstResult` on the `Criteria` instance.

```
import org.hibernate.FetchMode as FM
...
def results = c.list {
  maxResults(10)
  firstResult(50)
  fetchMode("aRelationship", FM.JOIN)
}
```

Querying with Eager Fetching

In the section on [Eager and Lazy Fetching](#) we discussed how to declaratively specify fetching to avoid the can also be achieved using a criteria query:

```
def criteria = Task.createCriteria()
def tasks = criteria.list{
  eq "assignee.id", task.assignee.id
  join 'assignee'
  join 'project'
  order 'priority', 'asc'
}
```

Notice the usage of the `join` method: it tells the criteria API to use a `JOIN` to fetch the named association. It is probably best not to use this for one-to-many associations though, because you will most likely end up with 'select' fetch mode:

```
import org.hibernate.FetchMode as FM
...
def results = Airport.withCriteria {
  eq "region", "EMEA"
  fetchMode "flights", FM.SELECT
}
```

Although this approach triggers a second query to get the `flights` association, you will get reliable results.



`fetchMode` and `join` are general settings of the query and can only be specified at the top level. Use them inside projections or association constraints.

An important point to bear in mind is that if you include associations in the query constraints, those associations are loaded. For example, in this query:

```
def results = Airport.withCriteria {
  eq "region", "EMEA"
  flights {
    like "number", "BA%"
  }
}
```

the `flights` collection would be loaded eagerly via a join even though the fetch mode has not been explicitly specified.

Method Reference

If you invoke the builder with no method name such as:

```
c { ... }
```

The build defaults to listing all the results and hence the above is equivalent to:

```
c.list { ... }
```

Method	Description
list	This is the default method. It returns all matching rows.
get	Returns a unique result set, i.e. just one row. The criteria has to be formed that way, that it not to be confused with a limit to just the first row.
scroll	Returns a scrollable result set.
listDistinct	If subqueries or associations are used, one may end up with the same row multiple times in distinct entities and is equivalent to DISTINCT_ROOT_ENTITY of the CriteriaSpecification
count	Returns the number of matching rows.

5.4.3 Hibernate Query Language (HQL)

GORM classes also support Hibernate's query language HQL, a very complete reference for which can be found in the Hibernate documentation.

GORM provides a number of methods that work with HQL including [find](#), [findAll](#) and [executeQuery](#). An

```
def results =  
    Book.findAll("from Book as b where b.title like 'Lord of the%'")
```

Positional and Named Parameters

In this case the value passed to the query is hard coded, however you can equally use positional parameters:

```
def results =  
    Book.findAll("from Book as b where b.title like ?", ["The Shi%"])
```

```
def author = Author.findByName("Stephen King")  
def books = Book.findAll("from Book as book where book.author = ?",  
                        [author])
```

Or even named parameters:

```
def results =
    Book.findAll("from Book as b " +
        "where b.title like :search or b.author like :search",
        [search: "The Shi%"])
```

```
def author = Author.findByName("Stephen King")
def books = Book.findAll("from Book as book where book.author = :author",
    [author: author])
```

Multiline Queries

Use the line continuation character to separate the query across multiple lines:

```
def results = Book.findAll("\
from Book as b, \
    Author as a \
where b.author = a and a.surname = ?", ['Smith'])
```



Triple-quoted Groovy multiline Strings will NOT work with HQL queries.

Pagination and Sorting

You can also perform pagination and sorting whilst using HQL queries. To do so simply specify the pagination method call and include an "ORDER BY" clause in the HQL:

```
def results =
    Book.findAll("from Book as b where " +
        "b.title like 'Lord of the%' " +
        "order by b.title asc",
        [max: 10, offset: 20])
```

5.5 Advanced GORM Features

The following sections cover more advanced usages of GORM including caching, custom mapping and ev

5.5.1 Events and Auto Timestamping

GORM supports the registration of events as methods that get fired when certain events occurs such following is a list of supported events:

- `beforeInsert` - Executed before an object is initially persisted to the database
- `beforeUpdate` - Executed before an object is updated
- `beforeDelete` - Executed before an object is deleted
- `beforeValidate` - Executed before an object is validated
- `afterInsert` - Executed after an object is persisted to the database
- `afterUpdate` - Executed after an object has been updated
- `afterDelete` - Executed after an object has been deleted
- `onLoad` - Executed when an object is loaded from the database

To add an event simply register the relevant closure with your domain class.



Do not attempt to flush the session within an event (such as with `obj.save(flush:true)`). Since flushing this will cause a `StackOverflowError`.

Event types

The `beforeInsert` event

Fired before an object is saved to the database

```
class Person {
  Date dateCreated
  def beforeInsert() {
    dateCreated = new Date()
  }
}
```

The `beforeUpdate` event

Fired before an existing object is updated

```
class Person {
  Date dateCreated
  Date lastUpdated
  def beforeInsert() {
    dateCreated = new Date()
  }
  def beforeUpdate() {
    lastUpdated = new Date()
  }
}
```

The beforeDelete event

Fired before an object is deleted.

```
class Person {
    String name
    Date dateCreated
    Date lastUpdated

    def beforeDelete() {
        ActivityTrace.withNewSession {
            new ActivityTrace(eventName: "Person Deleted", data: name).save()
        }
    }
}
```

Notice the usage of `withNewSession` method above. Since events are triggered whilst Hibernate is flush, `save()` and `delete()` won't result in objects being saved unless you run your operations with a new Session.

Fortunately the `withNewSession` method lets you share the same transactional JDBC connection underlying Session.

The beforeValidate event

Fired before an object is validated.

```
class Person {
    String name

    static constraints = {
        name size: 5..45
    }

    def beforeValidate() {
        name = name?.trim()
    }
}
```

The `beforeValidate` method is run before any validators are run.

GORM supports an overloaded version of `beforeValidate` which accepts a `List` parameter which represents the property names which are about to be validated. This version of `beforeValidate` will be called when the `validate` method is called with a `List` of property names as an argument.

```

class Person {
    String name
    String town
    Integer age

    static constraints = {
        name size: 5..45
        age range: 4..99
    }

    def beforeValidate(List propertiesBeingValidated) {
        // do pre validation work based on propertiesBeingValidated
    }
}

def p = new Person(name: 'Jacob Brown', age: 10)
p.validate(['age', 'name'])

```



Note that when `validate` is triggered indirectly because of a call to the `save` method, the `validate` method is being invoked with no arguments, not a `List` that includes all of the property names.

Either or both versions of `beforeValidate` may be defined in a domain class. GORM will prefer the version with arguments but will fall back on the no-arg version if the `List` version does not exist. Likewise, GORM will prefer the no-arg version but will fall back on the `List` version if the no-arg version does not exist.

The onLoad/beforeLoad event

Fired immediately before an object is loaded from the database:

```

class Person {
    String name
    Date dateCreated
    Date lastUpdated

    def onLoad() {
        log.debug "Loading ${id}"
    }
}

```

`beforeLoad()` is effectively a synonym for `onLoad()`, so only declare one or the other.

The afterLoad event

Fired immediately after an object is loaded from the database:

```
class Person {
    String name
    Date dateCreated
    Date lastUpdated

    def afterLoad() {
        name = "I'm loaded"
    }
}
```

Custom Event Listeners

You can also register event handler classes in an application's `grails-app/conf/spring` `doWithSpring` closure in a plugin descriptor by registering a Spring bean named `hibernateEv` property, `listenerMap` which specifies the listeners to register for various Hibernate events.

The values of the Map are instances of classes that implement one or more Hibernate listener interfaces. all of the required interfaces, or one concrete class per interface, or any combination. The valid Map keys here:

Name	Interface
auto-flush	AutoFlushEventListener
merge	MergeEventListener
create	PersistEventListener
create-onflush	PersistEventListener
delete	DeleteEventListener
dirty-check	DirtyCheckEventListener
evict	EvictEventListener
flush	FlushEventListener
flush-entity	FlushEntityEventListener
load	LoadEventListener
load-collection	InitializeCollectionEventListener
lock	LockEventListener
refresh	RefreshEventListener
replicate	ReplicateEventListener
save-update	SaveOrUpdateEventListener
save	SaveOrUpdateEventListener
update	SaveOrUpdateEventListener
pre-load	PreLoadEventListener
pre-update	PreUpdateEventListener
pre-delete	PreDeleteEventListener
pre-insert	PreInsertEventListener
pre-collection-recreate	PreCollectionRecreateEventListener
pre-collection-remove	PreCollectionRemoveEventListener
pre-collection-update	PreCollectionUpdateEventListener
post-load	PostLoadEventListener
post-update	PostUpdateEventListener
post-delete	PostDeleteEventListener
post-insert	PostInsertEventListener
post-commit-update	PostUpdateEventListener
post-commit-delete	PostDeleteEventListener
post-commit-insert	PostInsertEventListener
post-collection-recreate	PostCollectionRecreateEventListener
post-collection-remove	PostCollectionRemoveEventListener
post-collection-update	PostCollectionUpdateEventListener

For example, you could register a class `AuditEventListener` which implements `PostUpdateEventListener`, and `PostDeleteEventListener` using the following in an applic

```
beans = {
    auditListener(AuditEventListener)

    hibernateEventListeners(HibernateEventListeners) {
        listenerMap = ['post-insert': auditListener,
                       'post-update': auditListener,
                       'post-delete': auditListener]
    }
}
```

or use this in a plugin:

```
def doWithSpring = {
    auditListener(AuditEventListener)

    hibernateEventListeners(HibernateEventListeners) {
        listenerMap = ['post-insert': auditListener,
                       'post-update': auditListener,
                       'post-delete': auditListener]
    }
}
```

Automatic timestamping

The examples above demonstrated using events to update a `lastUpdated` and `dateCreated` property. However, this is actually not necessary. By defining a `lastUpdated` and `dateCreated` property then using GORM.

If this is not the behaviour you want you can disable this feature with:

```
class Person {
    Date dateCreated
    Date lastUpdated
    static mapping = {
        autoTimestamp false
    }
}
```



If you put `nullable: false` constraints on either `dateCreated` or `lastUpdated`, will fail validation - probably not what you want. Leave constraints off these properties unless you want automatic timestamping.

5.5.2 Custom ORM Mapping

Grails domain classes can be mapped onto many legacy schemas with an Object Relational Mapping. The following sections take you through what is possible with the ORM DSL.



None of this is necessary if you are happy to stick to the conventions defined by GORM for names and so on. You only need this functionality if you need to tailor the way GORM maps or configures caching

Custom mappings are defined using a static mapping block defined within your domain class:

```
class Person {  
    ...  
    static mapping = {  
    }  
}
```

5.5.2.1 Table and Column Names

Table names

The database table name which the class maps to can be customized using the `table` method:

```
class Person {  
    ...  
    static mapping = {  
        table 'people'  
    }  
}
```

In this case the class would be mapped to a table called `people` instead of the default name of `person`.

Column names

It is also possible to customize the mapping for individual columns onto the database. For example to change

```
class Person {  
    String firstName  
    static mapping = {  
        table 'people'  
        firstName column: 'First_Name'  
    }  
}
```

Here `firstName` is a dynamic method within the mapping Closure that has a single Map parameter. Since `firstName` is a class persistent field, the parameter values (in this case just "column") are used to configure the mapping

Column type

GORM supports configuration of Hibernate types with the DSL using the `type` attribute. This includes the Hibernate [org.hibernate.usertype.UserType](#) interface, which allows complete customization of how a type is mapped. For example, a `PostCodeType` you could use it as follows:

```
class Address {
    String number
    String postCode

    static mapping = {
        postCode type: PostCodeType
    }
}
```

Alternatively if you just wanted to map it to one of Hibernate's basic types other than the default chosen by

```
class Address {
    String number
    String postCode

    static mapping = {
        postCode type: 'text'
    }
}
```

This would make the `postCode` column map to the default large-text type for the database you're using (See the Hibernate documentation regarding [Basic Types](#) for further information.

Many-to-One/One-to-One Mappings

In the case of associations it is also possible to configure the foreign keys used to map associations. In the association this is exactly the same as any regular column. For example consider the following:

```
class Person {
    String firstName
    Address address

    static mapping = {
        table 'people'
        firstName column: 'First_Name'
        address column: 'Person_Address_Id'
    }
}
```

By default the `address` association would map to a foreign key column called `address_id`. By using the name of the foreign key column to `Person_Address_Id`.

One-to-Many Mapping

With a bidirectional one-to-many you can change the foreign key column used by changing the column name as per the example in the previous section on one-to-one associations. However, with unidirectional associations specified on the association itself. For example given a unidirectional one-to-many relationship between `Person` and `Address`, the following code will change the foreign key in the `address` table:

```

class Person {
  String firstName
  static hasMany = [addresses: Address]
  static mapping = {
    table 'people'
    firstName column: 'First_Name'
    addresses column: 'Person_Address_Id'
  }
}

```

If you don't want the column to be in the address table, but instead some intermediate join table you can

```

class Person {
  String firstName
  static hasMany = [addresses: Address]
  static mapping = {
    table 'people'
    firstName column: 'First_Name'
    addresses joinTable: [name: 'Person_Addresses',
                          key: 'Person_Id',
                          column: 'Address_Id']
  }
}

```

Many-to-Many Mapping

Grails, by default maps a many-to-many association using a join table. For example consider this many-to-

```

class Group {
  ...
  static hasMany = [people: Person]
}

```

```

class Person {
  ...
  static belongsTo = Group
  static hasMany = [groups: Group]
}

```

In this case Grails will create a join table called `group_person` containing foreign keys called `person`, `person_id` and `group` tables. To change the column names you can specify a column within the mappings

```

class Group {
    ...
    static mapping = {
        people column: 'Group_Person_Id'
    }
}
class Person {
    ...
    static mapping = {
        groups column: 'Group_Group_Id'
    }
}

```

You can also specify the name of the join table to use:

```

class Group {
    ...
    static mapping = {
        people column: 'Group_Person_Id',
        joinTable: 'PERSON_GROUP_ASSOCIATIONS'
    }
}
class Person {
    ...
    static mapping = {
        groups column: 'Group_Group_Id',
        joinTable: 'PERSON_GROUP_ASSOCIATIONS'
    }
}

```

5.5.2.2 Caching Strategy

Setting up caching

[Hibernate](#) features a second-level cache with a customizable cache provider. This is configured in the `grails-app/conf/DataSource.groovy` file as follows:

```

hibernate {
    cache.use_second_level_cache=true
    cache.use_query_cache=true
    cache.provider_class='org.hibernate.cache.EhCacheProvider'
}

```

You can customize any of these settings, for example to use a distributed caching mechanism.



For further reading on caching and in particular Hibernate's second-level cache, see the [documentation](#) on the subject.

Caching instances

Call the `cache` method in your mapping block to enable caching with the default settings:

```
class Person {
  ...
  static mapping = {
    table 'people'
    cache true
  }
}
```

This will configure a 'read-write' cache that includes both lazy and non-lazy properties. You can customize

```
class Person {
  ...
  static mapping = {
    table 'people'
    cache usage: 'read-only', include: 'non-lazy'
  }
}
```

Caching associations

As well as the ability to use Hibernate's second level cache to cache instances you can also cache collections. For example:

```
class Person {
  String firstName
  static hasMany = [addresses: Address]
  static mapping = {
    table 'people'
    version false
    addresses column: 'Address', cache: true
  }
}
```

```
class Address {
  String number
  String postCode
}
```

This will enable a 'read-write' caching mechanism on the addresses collection. You can also use:

```
cache: 'read-write' // or 'read-only' or 'transactional'
```

to further configure the cache usage.

Caching Queries

You can cache queries such as dynamic finders and criteria. To do so using a dynamic finder you can pass

```
def person = Person.findByFirstName("Fred", [cache: true])
```



In order for the results of the query to be cached, you must enable caching in your mapping in the previous section.

You can also cache criteria queries:

```
def people = Person.withCriteria {  
    like('firstName', 'Fr%')  
    cache true  
}
```

Cache usages

Below is a description of the different cache settings and their usages:

- `read-only` - If your application needs to read but never modify instances of a persistent class, a read-only cache might be appropriate.
- `read-write` - If the application needs to update data, a read-write cache might be appropriate.
- `nonstrict-read-write` - If the application only occasionally needs to update data (ie. if it is very rare to update the same item simultaneously) and strict transaction isolation is not required, a nonstrict-read-write cache might be appropriate.
- `transactional` - The transactional cache strategy provides support for fully transactional environments using JTA. Such a cache may only be used in a JTA environment. To use this cache, you must set `hibernate.transaction.manager_lookup_class` in the `grails-app/conf/DataSources.groovy` file.

5.5.2.3 Inheritance Strategies

By default GORM classes use `table-per-hierarchy` inheritance mapping. This has the disadvantage of a `NOT-NULL` constraint applied to them at the database level. If you would prefer to use a `table-per-subclass` mapping, you can do so as follows:

```
class Payment {  
    Integer amount  
  
    static mapping = {  
        tablePerHierarchy false  
    }  
}  
  
class CreditCardPayment extends Payment {  
    String cardNumber  
}
```

The mapping of the root `Payment` class specifies that it will not be using `table-per-hierarchy` mapping.

5.5.2.4 Custom Database Identity

You can customize how GORM generates identifiers for the database using the DSL. By default GORM uses the 'identity' generator for generating ids. This is by far the best approach, but there are still many schemas that have different approaches.

To deal with this Hibernate defines the concept of an id generator. You can customize the id generator and

```
class Person {  
    ...  
    static mapping = {  
        table 'people'  
        version false  
        id generator: 'hilo',  
        params: [table: 'hi_value',  
                 column: 'next_value',  
                 max_lo: 100]  
    }  
}
```

In this case we're using one of Hibernate's built in 'hilo' generators that uses a separate table to generate ids.



For more information on the different Hibernate generators refer to the [Hibernate reference documentation](#).

Although you don't typically specify the `id` field (Grails adds it for you) you can still configure its mapping. In the example to customise the column for the id property you can do:

```
class Person {  
    ...  
    static mapping = {  
        table 'people'  
        version false  
        id column: 'person_id'  
    }  
}
```

5.5.2.5 Composite Primary Keys

GORM supports the concept of composite identifiers (identifiers composed from 2 or more properties). It is available to you if you need it:

```

import org.apache.commons.lang.builder.HashCodeBuilder

class Person implements Serializable {

    String firstName
    String lastName

    boolean equals(other) {
        if (!(other instanceof Person)) {
            return false
        }

        other.firstName == firstName && other.lastName == lastName
    }

    int hashCode() {
        def builder = new HashCodeBuilder()
        builder.append firstName
        builder.append lastName
        builder.toHashCode()
    }

    static mapping = {
        id composite: ['firstName', 'lastName']
    }
}

```

The above will create a composite id of the `firstName` and `lastName` properties of the `Person` class. prototype of the object itself:

```

def p = Person.get(new Person(firstName: "Fred", lastName: "Flintstone"))
println p.firstName

```

Domain classes mapped with composite primary keys must implement the `Serializable` interface and methods, using the properties in the composite key for the calculations. The example above uses a `HashCodeBuilder` to implement it yourself.

Another important consideration when using composite primary keys is associations. If for example you the foreign keys are stored in the associated table then 2 columns will be present in the associated table.

For example consider the following domain class:

```

class Address {
    Person person
}

```

In this case the address table will have an additional two columns called `person_first_name` and `person_last_name`. To change the mapping of these columns then you can do so using the following technique:

```
class Address {
    Person person
    static mapping = {
        person {
            column: "FirstName"
            column: "LastName"
        }
    }
}
```

5.5.2.6 Database Indices

To get the best performance out of your queries it is often necessary to tailor the table index definitions. It is also a matter of monitoring usage patterns of your queries. With GORM's DSL you can specify which columns to index.

```
class Person {
    String firstName
    String address
    static mapping = {
        table 'people'
        version false
        id column: 'person_id'
        firstName column: 'First_Name', index: 'Name_Idx'
        address column: 'Address', index: 'Name_Idx,Address_Index'
    }
}
```

Note that you cannot have any spaces in the value of the `index` attribute; in this example `index: 'Name_Idx,Address_Index'` cause an error.

5.5.2.7 Optimistic Locking and Versioning

As discussed in the section on [Optimistic and Pessimistic Locking](#), by default GORM uses optimistic locking by setting the `version` property into every class which is in turn mapped to a `version` column at the database level.

If you're mapping to a legacy schema that doesn't have version columns (or there's some other reason why you can't), you can disable this with the `version` method:

```
class Person {
    ...
    static mapping = {
        table 'people'
        version false
    }
}
```



If you disable optimistic locking you are essentially on your own with regards to concurrent updates. There is the risk of users losing data (due to data overriding) unless you use [pessimistic locking](#).

Version columns types

By default Grails maps the `version` property as a `Long` that gets incremented by one each time an i supports using a `Timestamp`, for example:

```
import java.sql.Timestamp

class Person {

    ...
    Timestamp version

    static mapping = {
        table 'people'
    }
}
```

There's a slight risk that two updates occurring at nearly the same time on a fast server can end up with th very low. One benefit of using a `Timestamp` instead of a `Long` is that you combine the optimistic loc single column.

5.5.2.8 Eager and Lazy Fetching

Lazy Collections

As discussed in the section on [Eager and Lazy fetching](#), GORM collections are lazily loaded by default t the ORM DSL. There are several options available to you, but the most common ones are:

- `lazy: false`
- `fetch: 'join'`

and they're used like this:

```
class Person {
    String firstName
    Pet pet

    static hasMany = [addresses: Address]
    static mapping = {
        addresses lazy: false
        pet fetch: 'join'
    }
}
```

```
class Address {
    String street
    String postCode
}
```

```
class Pet {
    String name
}
```

The first option, `lazy: false`, ensures that when a `Person` instance is loaded, its addresses collection is loaded with a second `SELECT`. The second option is basically the same, except the collection is loaded with a `JOIN` rather than a second `SELECT`. If you want to reduce the number of queries, so `fetch: 'join'` is the more appropriate option. On the other hand, it is a more expensive approach if your domain model and data result in more and larger results than would otherwise be the case.

For more advanced users, the other settings available are:

1. `batchSize: N`
2. `lazy: false, batchSize: N`

where `N` is an integer. These let you fetch results in batches, with one query per batch. As a simple example:

```
class Person {
    String firstName
    Pet pet

    static mapping = {
        pet batchSize: 5
    }
}
```

If a query returns multiple `Person` instances, then when we access the first `pet` property, Hibernate will fetch all the `Pet` instances in the batch. You can get the same behaviour with eager loading by combining `batchSize` with the `lazy: false` option. See these options in the [Hibernate user guide](#) and this [primer on fetching strategies](#). Note that ORM DSL does not have a `fetch` option for fetching strategy.

Lazy Single-Ended Associations

In GORM, one-to-one and many-to-one associations are by default lazy. Non-lazy single ended associations are not recommended for many entities because each non-lazy association will result in an extra `SELECT` statement. If there are many associations, the number of queries grows significantly!

Use the same technique as for lazy collections to make a one-to-one or many-to-one association non-lazy/eager.

```
class Person {
    String firstName
}
```

```
class Address {
    String street
    String postCode

    static belongsTo = [person: Person]

    static mapping = {
        person lazy: false
    }
}
```

Here we configure GORM to load the associated `Person` instance (through the `person` property) when an `Address` instance is loaded.

Lazy Single-Ended Associations and Proxies

Hibernate uses runtime-generated proxies to facilitate single-ended lazy associations; Hibernate dynamically generates the proxy.

Consider the previous example but with a lazily-loaded `person` association: Hibernate will set the proxy as a subclass of `Person`. When you call any of the getters (except for the `id` property) or setters on that proxy, it will load the data from the database.

Unfortunately this technique can produce surprising results. Consider the following example classes:

```
class Pet {  
    String name  
}
```

```
class Dog extends Pet {  
}
```

```
class Person {  
    String name  
    Pet pet  
}
```

and assume that we have a single `Person` instance with a `Dog` as the `pet`. The following code will work:

```
def person = Person.get(1)  
assert person.pet instanceof Dog  
assert Pet.get(person.petId) instanceof Dog
```

But this won't:

```
def person = Person.get(1)  
assert person.pet instanceof Dog  
assert Pet.list()[0] instanceof Dog
```

The second assertion fails, and to add to the confusion, this will work:

```
assert Pet.list()[0] instanceof Dog
```

What's going on here? It's down to a combination of how proxies work and the guarantees that the Hibernate API provides. When you have a `Person` instance, Hibernate creates a proxy for its `pet` relation and attaches it to the session. Once that proxy is attached, if you call `get()` on the proxy, or the `pet` relation *within the same session*, Hibernate gives you the real instance.

Fortunately for us, GORM automatically unwraps the proxy when you use `get()` and `findBy*()`, or `list()` and `findAllBy*()`. That means you don't have to worry at all about proxies in the majority of cases. But GORM doesn't do this for `list()` and `findAllBy*()`. That means you don't have to worry at all about proxies in the majority of cases. But GORM doesn't do this for `list()` and `findAllBy*()`. However, if Hibernate hasn't attached the proxy to the session, it will return a list, such as `list()` and `findAllBy*()`. However, if Hibernate hasn't attached the proxy to the session, it will return a list, such as `list()` and `findAllBy*()`. Hence why the last example works.

You can protect yourself to a degree from this problem by using the `instanceOf` method by GORM:

```
def person = Person.get(1)
assert Pet.list()[0].instanceOf(Dog)
```

However, it won't help here if casting is involved. For example, the following code will throw a `ClassCastException` if the first element in the list is a proxy instance with a class that is neither `Dog` nor a sub-class of `Dog`:

```
def person = Person.get(1)
Dog pet = Pet.list()[0]
```

Of course, it's best not to use static types in this situation. If you use an untyped variable for the pet instead, you can call methods on the instance without any problems.

These days it's rare that you will come across this issue, but it's best to be aware of it just in case. At least you can then occur and be able to work around it.

5.5.2.9 Custom Cascade Behaviour

As described in the section on [cascading updates](#), the primary mechanism to control the way updates and deletions are propagated to another is the static [belongsTo](#) property.

However, the ORM DSL gives you complete access to Hibernate's [transitive persistence](#) capabilities using the `cascade` attribute.

Valid settings for the cascade attribute include:

- `merge` - merges the state of a detached association
- `save-update` - cascades only saves and updates to an association
- `delete` - cascades only deletes to an association
- `lock` - useful if a pessimistic lock should be cascaded to its associations
- `refresh` - cascades refreshes to an association
- `evict` - cascades evictions (equivalent to `discard()` in GORM) to associations if set
- `all` - cascade *all* operations to associations
- `all-delete-orphan` - Applies only to one-to-many associations and indicates that when a child is deleted, the parent should be automatically deleted. Children are also deleted when the parent is.



It is advisable to read the section in the Hibernate documentation on [transitive persistence](#) for a better understanding of the different cascade styles and recommendations for their usage.

To specify the cascade attribute simply define one or more (comma-separated) of the aforementioned settings.

```

class Person {
    String firstName
    static hasMany = [addresses: Address]
    static mapping = {
        addresses cascade: "all-delete-orphan"
    }
}

```

```

class Address {
    String street
    String postCode
}

```

5.5.2.10 Custom Hibernate Types

You saw in an earlier section that you can use composition (with the embedded property) to break down a large domain class into smaller ones. You can also achieve a similar effect with Hibernate's custom user types. These are not domain classes themselves, but these types also has a corresponding "meta-type" class that implements org.hibernate.usertype.UserType.

The [Hibernate reference manual](#) has some information on custom types, but here we will focus on how to take a look at a simple domain class that uses an old-fashioned (pre-Java 1.5) type-safe enum class:

```

class Book {
    String title
    String author
    Rating rating
    static mapping = {
        rating type: RatingUserType
    }
}

```

All we have done is declare the `rating` field the enum type and set the property's type in the custom mapping implementation. That's all you have to do to start using your custom type. If you want, you can also use the `"column"` to change the column name and `"index"` to add it to an index.

Custom types aren't limited to just a single column - they can be mapped to as many columns as you want. You can specify the mapping what columns to use, since Hibernate can only use the property name for a single column. For mapping multiple columns to a property using this syntax:


```

class Book {
    String title
    Name author
    Rating rating

    static mapping = {
        name type: NameUserType, {
            column name: "first_name"
            column name: "last_name"
        }
        rating type: RatingUserType
    }
}

```

The above example will create "first_name" and "last_name" columns for the author property. You'll be using some of the normal column/property mapping attributes in the column definitions. For example:

```

column name: "first_name", index: "my_idx", unique: true

```

The column definitions do *not* support the following attributes: type, cascade, lazy, cache, and join.

One thing to bear in mind with custom types is that they define the *SQL types* for the corresponding database. If you're used to configuring them yourself, but what happens if you have a legacy database that uses a different SQL type? You can override the column's SQL type using the `sqlType` attribute:

```

class Book {
    String title
    Name author
    Rating rating

    static mapping = {
        name type: NameUserType, {
            column name: "first_name", sqlType: "text"
            column name: "last_name", sqlType: "text"
        }
        rating type: RatingUserType, sqlType: "text"
    }
}

```

Mind you, the SQL type you specify needs to still work with the custom type. So overriding a default of "text" with "yes_no" isn't going to work.

5.5.2.11 Derived Properties

A derived property is one that takes its value from a SQL expression, often but not necessarily based on other properties. Consider a Product class like this:

```

class Product {
    Float price
    Float taxRate
    Float tax
}

```

If the `tax` property is derived based on the value of `price` and `taxRate` properties then is probably no SQL used to derive the value of a derived property may be expressed in the ORM DSL like this:

```
class Product {
    Float price
    Float taxRate
    Float tax

    static mapping = {
        tax formula: 'PRICE * TAX_RATE'
    }
}
```

Note that the formula expressed in the ORM DSL is SQL so references to other properties should relate model, which is why the example refers to `PRICE` and `TAX_RATE` instead of `price` and `taxRate`.

With that in place, when a `Product` is retrieved with something like `Product.get(42)`, the SQL that is generated is something like this:

```
select
    product0_.id as id1_0_,
    product0_.version as version1_0_,
    product0_.price as price1_0_,
    product0_.tax_rate as tax4_1_0_,
    product0_.PRICE * product0_.TAX_RATE as formula1_0_
from
    product product0_
where
    product0_.id=?
```

Since the `tax` property is derived at runtime and not stored in the database it might seem that the same method like `getTax()` to the `Product` class that simply returns the product of the `taxRate` and `price` that you would give up the ability query the database based on the value of the `tax` property. Using a derived property to retrieve all `Product` objects that have a `tax` value greater than 21.12 you could execute a query like this

```
Product.findAllByTaxGreaterThan(21.12)
```

Derived properties may be referenced in the Criteria API:

```
Product.withCriteria {
    gt 'tax', 21.12f
}
```

The SQL that is generated to support either of those would look something like this:

```

select
    this_.id as id1_0_,
    this_.version as version1_0_,
    this_.price as price1_0_,
    this_.tax_rate as tax4_1_0_,
    this_.PRICE * this_.TAX_RATE as formula1_0_
from
    product this_
where
    this_.PRICE * this_.TAX_RATE>?

```



Because the value of a derived property is generated in the database and depends on the other derived properties may not have GORM constraints applied to them. If constraints are set on a derived property, they will be ignored.

5.5.2.12 Custom Naming Strategy

By default Grails uses Hibernate's `ImprovedNamingStrategy` to convert domain class names and property names by converting from camel-cased Strings to ones that use underscores as word separators. You can configure this in the mapping closure but if there's a consistent pattern you can specify a different `NamingStrategy`.

Configure the class name to be used in `grails-app/conf/DataSource.groovy` in the `hibernate` block:

```

dataSource {
    pooled = true
    dbCreate = "create-drop"
    ...
}

hibernate {
    cache.use_second_level_cache = true
    ...
    naming_strategy = com.myco.myproj.CustomNamingStrategy
}

```

You can use an existing class or write your own, for example one that prefixes table names and column names.

```

package com.myco.myproj

import org.hibernate.cfg.ImprovedNamingStrategy
import org.hibernate.util.StringHelper

class CustomNamingStrategy extends ImprovedNamingStrategy {

    String classToTableName(String className) {
        "table_" + StringHelper.unqualify(className)
    }

    String propertyToColumnName(String propertyName) {
        "col_" + StringHelper.unqualify(propertyName)
    }
}

```

5.5.3 Default Sort Order

You can sort objects using query arguments such as those found in the [list](#) method:

```
def airports = Airport.list(sort:'name')
```

However, you can also declare the default sort order for a collection in the mapping:

```
class Airport {  
    ...  
    static mapping = {  
        sort "name"  
    }  
}
```

The above means that all collections of `Airports` will by default be sorted by the airport name. If you use `sort` in this syntax:

```
class Airport {  
    ...  
    static mapping = {  
        sort name: "desc"  
    }  
}
```

Finally, you can configure sorting at the association level:

```
class Airport {  
    ...  
    static hasMany = [flights: Flight]  
    static mapping = {  
        flights sort: 'number', order: 'desc'  
    }  
}
```

In this case, the `flights` collection will always be sorted in descending order of flight number.



These mappings will not work for default unidirectional one-to-many or many-to-many relationships that involve a join table. See [this issue](#) for more details. Consider using a `SortedSet` or queries to fetch the data you need.

5.6 Programmatic Transactions

Grails is built on Spring and uses Spring's Transaction abstraction for dealing with programmatic transactions. This abstraction has been enhanced to make this simpler with the [withTransaction](#) method. This method has a single parameter which is a Spring [TransactionStatus](#) instance.

Here's an example of using `withTransaction` in a controller methods:

```

def transferFunds() {
    Account.withTransaction { status ->
        def source = Account.get(params.from)
        def dest = Account.get(params.to)

        def amount = params.amount.toInteger()
        if (source.active) {
            if (dest.active) {
                source.balance -= amount
                dest.amount += amount
            }
            else {
                status.setRollbackOnly()
            }
        }
    }
}

```

In this example we rollback the transaction if the destination account is not active. Also, if an unchecked Exception, even though Groovy doesn't require that you catch checked exceptions) is thrown automatically be rolled back.

You can also use "save points" to rollback a transaction to a particular point in time if you don't want to r be achieved through the use of Spring's [SavePointManager](#) interface.

The withTransaction method deals with the begin/commit/rollback logic for you within the scope of

5.7 GORM and Constraints

Although constraints are covered in the [Validation](#) section, it is important to mention them here as some which the database schema is generated.

Where feasible, Grails uses a domain class's constraints to influence the database columns generate properties.

Consider the following example. Suppose we have a domain model with the following properties:

```

String name
String description

```

By default, in MySQL, Grails would define these columns as

Column	Data Type
name	varchar(255)
description	varchar(255)

But perhaps the business rules for this domain class state that a description can be up to 1000 characters in likely define the column as follows *if* we were creating the table with an SQL script.

Column	Data Type
description	TEXT

Chances are we would also want to have some application-based validation to make sure we don't exceed persist any records. In Grails, we achieve this validation with [constraints](#). We would add the following con

```
static constraints = {  
    description maxSize: 1000  
}
```

This constraint would provide both the application-based validation we want and it would also cause the s
Below is a description of the other constraints that influence schema generation.

Constraints Affecting String Properties

- [inList](#)
- [maxSize](#)
- [size](#)

If either the `maxSize` or the `size` constraint is defined, Grails sets the maximum column length based on

In general, it's not advisable to use both constraints on the same domain class property. However, if both constraint are defined, then Grails sets the column length to the minimum of the `maxSize` constraint and (Grails uses the minimum of the two, because any length that exceeds that minimum will result in a validat

If the `inList` constraint is defined (and the `maxSize` and the `size` constraints are not defined), then C based on the length of the longest string in the list of valid values. For example, given a list including valu would set the column length to 6 (i.e., the number of characters in the string "Groovy").

Constraints Affecting Numeric Properties

- [min](#)
- [max](#)
- [range](#)

If the `max`, `min`, or `range` constraint is defined, Grails attempts to set the column precision based on th attempted influence is largely dependent on how Hibernate interacts with the underlying DBMS.)

In general, it's not advisable to combine the pair `min/max` and `range` constraints together on the same d of these constraints is defined, then Grails uses the minimum precision value from the constraints. (Grails any length that exceeds that minimum precision will result in a validation error.)

- [scale](#)

If the `scale` constraint is defined, then Grails attempts to set the column [scale](#) based on the constraint v point numbers (i.e., `java.lang.Float`, `java.Lang.Double`, `java.lang.BigI` `java.lang.BigDecimal`). The success of this attempted influence is largely dependent on how F DBMS.

The constraints define the minimum/maximum numeric values, and Grails derives the maximum number in mind that specifying only one of `min/max` constraints will not affect schema generation (since there co with `max:100`, for example), unless the specified constraint value requires more digits than default Hi moment). For example:

```
someFloatValue max: 1000000, scale: 3
```

would yield:

```
someFloatValue DECIMAL(19, 3) // precision is default
```

but

```
someFloatValue max: 12345678901234567890, scale: 5
```

would yield:

```
someFloatValue DECIMAL(25, 5) // precision = digits in max + scale
```

and

```
someFloatValue max: 100, min: -100000
```

would yield:

```
someFloatValue DECIMAL(8, 2) // precision = digits in min + default scale
```

6 The Web Layer

6.1 Controllers

A controller handles requests and creates or prepares the response. A controller can generate the response (a controller, simply create a class whose name ends with `Controller` in the `grails-app/controllers` package).

The default [URL Mapping](#) configuration ensures that the first part of your controller name is mapped to your controller maps to URIs within the controller name URI.

6.1.1 Understanding Controllers and Actions

Creating a controller

Controllers can be created with the [create-controller](#) or [generate-controller](#) command. For example try run root of a Grails project:

```
grails create-controller book
```

The command will create a controller at the location `grails-app/controllers/myapp/BookController`.

```
package myapp

class BookController {

    def index() { }
}
```

where "myapp" will be the name of your application, the default package name if one isn't specified.

`BookController` by default maps to the `/book` URI (relative to your application root).



The `create-controller` and `generate-controller` commands are just for convenience; you can as easily create controllers using your favorite text editor or IDE.

Creating Actions

A controller can have multiple public action methods; each one maps to a URI:


```
class BookController {
  def list() {
    // do controller logic
    // create model

    return model
  }
}
```

This example maps to the `/book/list` URI by default thanks to the property being named `list`.

Public Methods as Actions

In earlier versions of Grails actions were implemented with Closures. This is still supported, but the preferred way is to use public methods.

Leveraging methods instead of Closure properties has some advantages:

- Memory efficient
- Allow use of stateless controllers (singleton scope)
- You can override actions from subclasses and call the overridden superclass method with `super.actionName()`
- Methods can be intercepted with standard proxying mechanisms, something that is complicated to do with Closures

If you prefer the Closure syntax or have older controller classes created in earlier versions of Grails using Closures, you can set the `grails.compile.artefacts.closures.convert` property to `true` in `grails.config`.

```
grails.compile.artefacts.closures.convert = true
```

and a compile-time AST transformation will convert your Closures to methods in the generated bytecode.



If a controller class extends some other class which is not defined under the `grails-app` directory, methods inherited from that class are not converted to controller actions. If the inherited methods are converted to controller actions, the methods may be overridden in the subclass and then the subclass will not invoke the method in the super class.

The Default Action

A controller has the concept of a default URI that maps to the root URI of the controller, for example the action that is called when the default URI is requested is dictated by the following rules:

- If there is only one action, it's the default
- If you have an action named `index`, it's the default
- Alternatively you can set it explicitly with the `defaultAction` property:

```
static defaultAction = "list"
```

6.1.2 Controllers and Scopes

Available Scopes

Scopes are hash-like objects where you can store variables. The following scopes are available to controllers:

- [servletContext](#) - Also known as application scope, this scope lets you share state across the entire web application. The object is an instance of [ServletContext](#)
- [session](#) - The session allows associating state with a given user and typically uses cookies to associate the user. The object is an instance of [HttpSession](#)
- [request](#) - The request object allows the storage of objects for the current request only. The object is an instance of [HttpServletRequest](#)
- [params](#) - Mutable map of incoming request query string or POST parameters
- [flash](#) - See below

Accessing Scopes

Scopes can be accessed using the variable names above in combination with Groovy's array index operator. For example, to access the [HttpServletRequest](#) scope:

```
class BookController {
    def find() {
        def findBy = params["findBy"]
        def appContext = request["foo"]
        def loggedUser = session["logged_user"]
    }
}
```

You can also access values within scopes using the de-reference operator, making the syntax even more concise:

```
class BookController {
    def find() {
        def findBy = params.findBy
        def appContext = request.foo
        def loggedUser = session.logged_user
    }
}
```

This is one of the ways that Grails unifies access to the different scopes.

Using Flash Scope

Grails supports the concept of [flash](#) scope as a temporary store to make attributes available for this request. The attributes are cleared. This is useful for setting a message directly before redirecting, for example:

```
def delete() {
  def b = Book.get(params.id)
  if (!b) {
    flash.message = "User not found for id ${params.id}"
    redirect(action:list)
  }
  ... // remaining code
}
```

When the `list` action is requested, the `message` value will be in scope and can be used to display an ir from the `flash` scope after this second request.

Note that the attribute name can be anything you want, and the values are often strings used to display mes

Scoped Controllers

By default, a new controller instance is created for each request. In fact, because the controller is prot each request happens on its own thread.

You can change this behaviour by placing a controller in a particular scope. The supported scopes are:

- `prototype` (default) - A new controller will be created for each request (recommended for actions :
- `session` - One controller is created for the scope of a user session
- `singleton` - Only one instance of the controller ever exists (recommended for actions as methods)

To enable one of the scopes, add a static `scope` property to your class with one of the valid scope values l

```
static scope = "singleton"
```

You can define the default strategy under in `Config.groovy` with the `grails.controllers.def`

```
grails.controllers.defaultScope = "singleton"
```



Use scoped controllers wisely. For instance, we don't recommend having any properties controller since they will be shared for *all* requests. Setting a default scope other than `proto` unexpected behaviors if you have controllers provided by installed plugins that expe `prototype`.

6.1.3 Models and Views

Returning the Model

A model is a `Map` that the view uses when rendering. The keys within that `Map` correspond to variable na couple of ways to return a model. First, you can explicitly return a `Map` instance:

```
def show() {
    [book: Book.get(params.id)]
}
```



The above does *not* reflect what you should use with the scaffolding views - see the [scaffolding](#) details.

If no explicit model is returned the controller's properties will be used as the model, thus allowing you to v

```
class BookController {
    List books
    List authors

    def list() {
        books = Book.list()
        authors = Author.list()
    }
}
```



This is possible due to the fact that controllers are prototype scoped. In other words a new c each request. Otherwise code such as the above would not be thread-safe, and all users would

In the above example the books and authors properties will be available in the view.

A more advanced approach is to return an instance of the Spring [ModelAndView](#) class:

```
import org.springframework.web.servlet.ModelAndView

def index() {
    // get some books just for the index page, perhaps your favorites
    def favoriteBooks = ...

    // forward to the list view to show them
    return new ModelAndView("/book/list", [ bookList : favoriteBooks ])
}
```

One thing to bear in mind is that certain variable names can not be used in your model:

- attributes
- application

Currently, no error will be reported if you do use them, but this will hopefully change in a future version o

Selecting the View

In both of the previous two examples there was no code that specified which [view](#) to render. So how does the answer lie in the conventions. Grails will look for a view at the location `grails-app/views/book/`

```
class BookController {
  def show() {
    [book: Book.get(params.id)]
  }
}
```

To render a different view, use the [render](#) method:

```
def show() {
  def map = [book: Book.get(params.id)]
  render(view: "display", model: map)
}
```

In this case Grails will attempt to render a view at the location `grails-app/views/book/display` qualifies the view location with the `book` directory of the `grails-app/views` directory. This is convenient instead you can use an absolute path instead of a relative one:

```
def show() {
  def map = [book: Book.get(params.id)]
  render(view: "/shared/display", model: map)
}
```

In this case Grails will attempt to render a view at the location `grails-app/views/shared/display`. Grails also supports JSPs as views, so if a GSP isn't found in the expected location but a JSP is, it will be used.

Rendering a Response

Sometimes it's easier (for example with Ajax applications) to render snippets of text or code to the response; the highly flexible `render` method can be used:

```
render "Hello World!"
```

The above code writes the text "Hello World!" to the response. Other examples include:

```
// write some markup
render {
  for (b in books) {
    div(id: b.id, b.title)
  }
}
```

```
// render a specific view
render(view: 'show')
```

```
// render a template for each item in a collection
render(template: 'book_template', collection: Book.list())
```

```
// render some text with encoding and content type
render(text: "<xml>some xml</xml>", contentType: "text/xml", encoding: "UTF-8")
```

If you plan on using Groovy's MarkupBuilder to generate HTML for use with the render method HTML elements and Grails tags, for example:

```
import groovy.xml.MarkupBuilder
...
def login() {
    def writer = new StringWriter()
    def builder = new MarkupBuilder(writer)
    builder.html {
        head {
            title 'Log in'
        }
        body {
            h1 'Hello'
            form {
            }
        }
    }
    def html = writer.toString()
    render html
}
```

This will actually [call the form tag](#) (which will return some text that will be ignored by the MarkupBuilder element, use the following:

```
def login() {
    // ...
    body {
        h1 'Hello'
        builder.form {
        }
    }
    // ...
}
```

6.1.4 Redirects and Chaining

Redirects

Actions can be redirected using the [redirect](#) controller method:

```
class OverviewController {
  def login() {}
  def find() {
    if (!session.user)
      redirect(action: 'login')
    return
  }
  ...
}
```

Internally the [redirect](#) method uses the [HttpServletResponse](#) object's `sendRedirect` method.

The `redirect` method expects one of:

- Another closure within the same controller class:

```
// Call the login action within the same class
redirect(action: login)
```

- The name of an action (and controller name if the redirect isn't to an action in the current controller):

```
// Also redirects to the index action in the home controller
redirect(controller: 'home', action: 'index')
```

- A URI for a resource relative the application context path:

```
// Redirect to an explicit URI
redirect(uri: "/login.html")
```

- Or a full URL:

```
// Redirect to a URL
redirect(url: "http://grails.org")
```

Parameters can optionally be passed from one action to the next using the `params` argument of the method:

```
redirect(action: 'myaction', params: [myparam: "myvalue"])
```

These parameters are made available through the [params](#) dynamic property that accesses request parameters. If a parameter with the same name as a request parameter is provided, the request parameter is overridden and the controller parameter is used.

Since the `params` object is a `Map`, you can use it to pass the current request parameters from one action to the next.

```
redirect(action: "next", params: params)
```

Finally, you can also include a fragment in the target URI:

```
redirect(controller: "test", action: "show", fragment: "profile")
```

which will (depending on the URL mappings) redirect to something like `/myapp/test/show#profile`.

Chaining

Actions can also be chained. Chaining allows the model to be retained from one action to the next. For example:

```
class ExampleChainController {
  def first() {
    chain(action: second, model: [one: 1])
  }
  def second () {
    chain(action: third, model: [two: 2])
  }
  def third() {
    [three: 3]
  }
}
```

results in the model:

```
[one: 1, two: 2, three: 3]
```

The model can be accessed in subsequent controller actions in the chain using the `chainModel` map actions following the call to the `chain` method:

```
class ChainController {
  def nextInChain() {
    def model = chainModel.myModel
    ...
  }
}
```

Like the `redirect` method you can also pass parameters to the `chain` method:

```
chain(action: "action1", model: [one: 1], params: [myparam: "param1"])
```


6.1.5 Controller Interceptors

Often it is useful to intercept processing based on either request, session or application state. This can be achieved by using interceptors. There are currently two types of interceptors: before and after.



If your interceptor is likely to apply to more than one controller, you are almost certainly better off using Filters. Filters can be applied to multiple controllers or URIs without the need to change the logic of each controller.

Before Interception

The `beforeInterceptor` intercepts processing before the action is executed. If it returns `false`, the action is not executed. The interceptor can be defined for all actions in a controller as follows:

```
def beforeInterceptor = {  
    println "Tracing action ${actionUri}"  
}
```

The above is declared inside the body of the controller definition. It will be executed before all actions and a common use case is very simplistic authentication:

```
def beforeInterceptor = [action: this.&auth, except: 'login']  
// defined with private scope, so it's not considered an action  
private auth() {  
    if (!session.user) {  
        redirect(action: 'login')  
        return false  
    }  
}  
  
def login() {  
    // display login page  
}
```

The above code defines a method called `auth`. A private method is used so that it is not exposed as an action. `beforeInterceptor` then defines an interceptor that is used on all actions *except* the login action as the `login` action is referenced using Groovy's method pointer syntax. Within the method it detects whether the user is logged in. If not, it redirects to the login action and returns `false`, causing the intercepted action to not be processed.

After Interception

Use the `afterInterceptor` property to define an interceptor that is executed after an action:

```
def afterInterceptor = { model ->  
    println "Tracing action ${actionUri}"  
}
```

The after interceptor takes the resulting model as an argument and can hence manipulate the model or response.

An after interceptor may also modify the Spring MVC [ModelAndView](#) object prior to rendering. In this ca

```
def afterInterceptor = { model, modelAndView ->
    println "Current view is ${modelAndView.viewName}"
    if (model.someVar) modelAndView.viewName = "/mycontroller/someotherview"
    println "View is now ${modelAndView.viewName}"
}
```

This allows the view to be changed based on the model returned by the current action. Note that the mode being intercepted called `redirect` or `render`.

Interception Conditions

Rails users will be familiar with the authentication example and how the 'except' condition was used when are called 'filters' in Rails; this terminology conflicts with Servlet filter terminology in Java):

```
def beforeInterceptor = [action: this.&auth, except: 'login']
```

This executes the interceptor for all actions except the specified action. A list of actions can also be defined

```
def beforeInterceptor = [action: this.&auth, except: ['login', 'register']]
```

The other supported condition is 'only', this executes the interceptor for only the specified action(s):

```
def beforeInterceptor = [action: this.&auth, only: ['secure']]
```

6.1.6 Data Binding

Data binding is the act of "binding" incoming request parameters onto the properties of an object or array. It should deal with all necessary type conversion since request parameters, which are typically delivered by HTTP, whilst the properties of a Groovy or Java object may well not be.

Grails uses [Spring](#)'s underlying data binding capability to perform data binding.

Binding Request Data to the Model

There are two ways to bind request parameters onto the properties of a domain class. The first involves using

```
def save() {
    def b = new Book(params)
    b.save()
}
```

The data binding happens within the code `new Book(params)`. By passing the [params](#) object to `new`, Grails automatically recognizes that you are trying to bind from request parameters. So if we had an incoming request

```
/book/save?title=The%20Stand&author=Stephen%20King
```

Then the `title` and `author` request parameters would automatically be set on the domain class. You perform data binding onto an existing instance:

```
def save() {  
    def b = Book.get(params.id)  
    b.properties = params  
    b.save()  
}
```

This has the same effect as using the implicit constructor.

Data binding and Single-ended Associations

If you have a `one-to-one` or `many-to-one` association you can use Grails' data binding capability example if you have an incoming request such as:

```
/book/save?author.id=20
```

Grails will automatically detect the `.id` suffix on the request parameter and look up the `Author` instance binding such as:

```
def b = new Book(params)
```

An association property can be set to `null` by passing the literal `String "null"`. For example:

```
/book/save?author.id=null
```

Data Binding and Many-ended Associations

If you have a `one-to-many` or `many-to-many` association there are different techniques for data binding depending on the association type.

If you have a `Set` based association (the default for a `hasMany`) then the simplest way to populate an association is by using the `<g:select>` tag. For example consider the usage of `<g:select>` below:

```
<g:select name="books"  
    from="${Book.list()}"  
    size="5" multiple="yes" optionKey="id"  
    value="${author?.books}" />
```

This produces a select box that lets you select multiple values. In this case if you submit the form Grails will bind the selected values from the select box to populate the `books` association.

However, if you have a scenario where you want to update the properties of the associated objects the this the subscript operator:

```
<g:textField name="books[0].title" value="the Stand" />
<g:textField name="books[1].title" value="the Shining" />
```

However, with Set based association it is critical that you render the mark-up in the same order that you p a Set has no concept of order, so although we're referring to books0 and books1 it is not guaranteed correct on the server side unless you apply some explicit sorting yourself.

This is not a problem if you use List based associations, since a List has a defined order and an inde Map based associations.

Note also that if the association you are binding to has a size of two and you refer to an element that is out:

```
<g:textField name="books[0].title" value="the Stand" />
<g:textField name="books[1].title" value="the Shining" />
<g:textField name="books[2].title" value="Red Madder" />
```

Then Grails will automatically create a new instance for you at the defined position. If you "skipped" a few

```
<g:textField name="books[0].title" value="the Stand" />
<g:textField name="books[1].title" value="the Shining" />
<g:textField name="books[5].title" value="Red Madder" />
```

Then Grails will automatically create instances in between. For example in the above case Grails w association being bound had a size of 2.

You can bind existing instances of the associated type to a List using the same .id syntax as you wou For example:

```
<g:select name="books[0].id" from="${bookList}"
value="${author?.books[0]?.id}" />

<g:select name="books[1].id" from="${bookList}"
value="${author?.books[1]?.id}" />

<g:select name="books[2].id" from="${bookList}"
value="${author?.books[2]?.id}" />
```

Would allow individual entries in the books List to be selected separately.

Entries at particular indexes can be removed in the same way too. For example:

```
<g:select name="books[0].id"
from="${Book.list()}"
value="${author?.books[0]?.id}"
noSelection="['null': '']"/>
```

Will render a select box that will remove the association at books0 if the empty option is chosen.

Binding to a Map property works the same way except that the list index in the parameter name is replaced

```
<g:select name="images[cover].id"
  from="${Image.list()}"
  value="${book?.images[cover]?.id}"
  noSelection="['null': '']"/>
```

This would bind the selected image into the Map property `images` under a key of `"cover"`.

Data binding with Multiple domain classes

It is possible to bind data to multiple domain objects from the [params](#) object.

For example so you have an incoming request to:

```
/book/save?book.title=The%20Stand&author.name=Stephen%20King
```

You'll notice the difference with the above request is that each parameter has a prefix such as `author.` or parameters belong to which type. Grails' `params` object is like a multi-dimensional hash and you can index parameters to bind.

```
def b = new Book(params.book)
```

Notice how we use the prefix before the first dot of the `book.title` parameter to isolate only parameter the same with an `Author` domain class:

```
def a = new Author(params.author)
```

Data Binding and Action Arguments

Controller action arguments are subject to request parameter data binding. There are 2 categories of category is command objects. Complex types are treated as command objects. See the [Command Objec](#) The other category is basic object types. Supported types are the 8 primitives, their corresponding type wrapper behavior is to map request parameters to action arguments by name:

```
class AccountingController {
  // accountNumber will be initialized with the value of params.accountNumber
  // accountType will be initialized with params.accountType
  def displayInvoice(String accountNumber, int accountType) {
    // ...
  }
}
```

For primitive arguments and arguments which are instances of any of the primitive type wrapper classes before the request parameter value can be bound to the action argument. The type conversion happens as shown above, the `params.accountType` request parameter has to be converted to an `int`. If type argument will have its default value per normal Java behavior (null for type wrapper references, false for corresponding error will be added to the `errors` property of the defining controller.

```
/accounting/displayInvoice?accountNumber=B59786&accountType=bogusValue
```

Since "bogusValue" cannot be converted to type `int`, the value of `accountType` will be zero, `controller.errors.errorCount` will be equal to 1 and `controller.errors.getFieldError` contain the corresponding error.

If the argument name does not match the name of the request parameter then the `@grails.web.RequestParameter` applied to an argument to express the name of the request parameter which should be bound to that argument.

```
import grails.web.RequestParameter

class AccountingController {
    // mainAccountNumber will be initialized with the value of params.accountNumber
    // accountType will be initialized with params.accountType
    def displayInvoice(@RequestParameter('accountNumber') String mainAccountNumber
        // ...
    }
}
```

Data binding and type conversion errors

Sometimes when performing data binding it is not possible to convert a particular String into a particular type. Grails will retain type conversion errors inside the [errors](#) property of a Grails domain class.

```
class Book {
    ...
    URL publisherURL
}
```

Here we have a domain class `Book` that uses the `java.net.URL` class to represent URLs. Given an incorrect request:

```
/book/save?publisherURL=a-bad-url
```

it is not possible to bind the string `a-bad-url` to the `publisherURL` property as a type mismatch error. Grails will retain this error:

```
def b = new Book(params)

if (b.hasErrors()) {
    println "The value ${b.errors.getFieldError('publisherURL').rejectedValue} " +
        " is not a valid URL!"
}
```

Although we have not yet covered error codes (for more information see the section on [Validation](#)), for ty message from the `grails-app/i18n/messages.properties` file to use for the error. You can use as:

```
typeMismatch.java.net.URL=The field {0} is not a valid URL
```

Or a more specific one:

```
typeMismatch.Book.publisherURL=The publisher URL you specified is not a valid URL
```

Data Binding and Security concerns

When batch updating properties from request parameters you need to be careful not to allow clients to bind be persisted in the database. You can limit what properties are bound to a given domain class using the `sub`

```
def p = Person.get(1)
p.properties['firstName','lastName'] = params
```

In this case only the `firstName` and `lastName` properties will be bound.

Another way to do this is to use [Command Objects](#) as the target of data binding instead of domain flexible [bindData](#) method.

The `bindData` method allows the same data binding capability, but to arbitrary objects:

```
def p = new Person()
bindData(p, params)
```

The `bindData` method also lets you exclude certain parameters that you don't want updated:

```
def p = new Person()
bindData(p, params, [exclude: 'dateOfBirth'])
```

Or include only certain properties:

```
def p = new Person()
bindData(p, params, [include: ['firstName', 'lastName']])
```



Note that if an empty List is provided as a value for the `include` parameter then all fields are included in binding if they are not explicitly excluded.

6.1.7 XML and JSON Responses

Using the render method to output XML

Grails supports a few different ways to produce XML and JSON responses. The first is the [render](#) method.

The render method can be passed a block of code to do mark-up building in XML:

```
def list() {
  def results = Book.list()
  render(contentType: "text/xml") {
    books {
      for (b in results) {
        book(title: b.title)
      }
    }
  }
}
```

The result of this code would be something like:

```
<books>
  <book title="The Stand" />
  <book title="The Shining" />
</books>
```

Be careful to avoid naming conflicts when using mark-up building. For example this code would produce :

```
def list() {
  def books = Book.list() // naming conflict here
  render(contentType: "text/xml") {
    books {
      for (b in results) {
        book(title: b.title)
      }
    }
  }
}
```

This is because there is local variable `books` which Groovy attempts to invoke as a method.

Using the render method to output JSON

The render method can also be used to output JSON:


```
def list() {
def results = Book.list()
render(contentType: "text/json") {
    books = array {
        for (b in results) {
            book title: b.title
        }
    }
}
}
```

In this case the result would be something along the lines of:

```
[
  {title:"The Stand"},
  {title:"The Shining"}
]
```

The same dangers with naming conflicts described above for XML also apply to JSON building.

Automatic XML Marshalling

Grails also supports automatic marshalling of [domain classes](#) to XML using special converters.

To start off with, import the `grails.converters` package into your controller:

```
import grails.converters.*
```

Now you can use the following highly readable syntax to automatically convert domain classes to XML:

```
render Book.list() as XML
```

The resulting output would look something like the following::

```
<?xml version="1.0" encoding="ISO-8859-1"?>
<list>
  <book id="1">
    <author>Stephen King</author>
    <title>The Stand</title>
  </book>
  <book id="2">
    <author>Stephen King</author>
    <title>The Shining</title>
  </book>
</list>
```

An alternative to using the converters is to use the [codecs](#) feature of Grails. The codecs feature provides methods:

```
def xml = Book.list().encodeAsXML()
render xml
```

For more information on XML marshalling see the section on [REST](#)

Automatic JSON Marshalling

Grails also supports automatic marshalling to JSON using the same mechanism. Simply substitute XML wi

```
render Book.list() as JSON
```

The resulting output would look something like the following:

```
[
  {
    "id":1,
    "class":"Book",
    "author":"Stephen King",
    "title":"The Stand"
  },
  {
    "id":2,
    "class":"Book",
    "author":"Stephen King",
    "releaseDate":new Date(1194127343161),
    "title":"The Shining"
  }
]
```

Again as an alternative you can use the `encodeAsJSON` to achieve the same effect.

6.1.8 More on JSONBuilder

The previous section on on XML and JSON responses covered simplistic examples of rendering XML builder used by Grails is the standard [XmlSlurper](#) found in Groovy, the JSON builder is a custom impleme

JSONBuilder and Grails versions

JSONBuilder behaves different depending on the version of Grails you use. For version below 1.2 the de is used. This section covers the usage of the Grails 1.2 JSONBuilder

For backwards compatibility the old JSONBuilder class is used with the `render` method for older JSONBuilder class set the following in `Config.groovy`:

```
grails.json.legacy.builder = false
```

Rendering Simple Objects

To render a simple JSON object just set properties within the context of the Closure:

```
render(contentType: "text/json") {  
    hello = "world"  
}
```

The above will produce the JSON:

```
{ "hello": "world" }
```

Rendering JSON Arrays

To render a list of objects simple assign a list:

```
render(contentType: "text/json") {  
    categories = [ 'a', 'b', 'c' ]  
}
```

This will produce:

```
{ "categories": [ "a", "b", "c" ] }
```

You can also render lists of complex objects, for example:

```
render(contentType: "text/json") {  
    categories = [ { a = "A" }, { b = "B" } ]  
}
```

This will produce:

```
{ "categories": [ { "a": "A" } , { "b": "B" } ] }
```

Use the special `element` method to return a list as the root:

```
render(contentType: "text/json") {  
    element 1  
    element 2  
    element 3  
}
```

The above code produces:

```
[ 1, 2, 3 ]
```

Rendering Complex Objects

Rendering complex objects can be done with Closures. For example:

```
render(contentType: "text/json") {  
    categories = ['a', 'b', 'c']  
    title = "Hello JSON"  
    information = {  
        pages = 10  
    }  
}
```

The above will produce the JSON:

```
{"categories":["a","b","c"],"title":"Hello JSON","information":{"pages":10}}
```

Arrays of Complex Objects

As mentioned previously you can nest complex objects within arrays using Closures:

```
render(contentType: "text/json") {  
    categories = [ { a = "A" }, { b = "B" } ]  
}
```

You can use the array method to build them up dynamically:

```
def results = Book.list()  
render(contentType: "text/json") {  
    books = array {  
        for (b in results) {  
            book title: b.title  
        }  
    }  
}
```

Direct JSONBuilder API Access

If you don't have access to the render method, but still want to produce JSON you can use the API direc

```

def builder = new JSONBuilder()
def result = builder.build {
  categories = ['a', 'b', 'c']
  title = "Hello JSON"
  information = {
    pages = 10
  }
}

// prints the JSON text
println result.toString()

def sw = new StringWriter()
result.render sw

```

6.1.9 Uploading Files

Programmatic File Uploads

Grails supports file uploads using Spring's [MultipartHttpServletRequest](#) interface. The first step for file like this:

```

Upload Form: <br />
<g:uploadForm action="upload">
  <input type="file" name="myFile" />
  <input type="submit" />
</g:uploadForm>

```

The uploadForm tag conveniently adds the enctype="multipart/form-data" attribute to the s

There are then a number of ways to handle the file upload. One is to work with the Spring [MultipartFile](#) in

```

def upload() {
  def f = request.getFile('myFile')
  if (f.empty) {
    flash.message = 'file cannot be empty'
    render(view: 'uploadForm')
    return
  }

  f.transferTo(new File('/some/local/dir/myfile.txt'))
  response.sendError(200, 'Done')
}

```

This is convenient for doing transfers to other destinations and manipulating the file directly as you can ob the [MultipartFile](#) interface.

File Uploads through Data Binding

File uploads can also be performed using data binding. Consider this Image domain class:

```
class Image {
    byte[] myFile

    static constraints = {
        // Limit upload file size to 2MB
        myFile maxSize: 1024 * 1024 * 2
    }
}
```

If you create an image using the `params` object in the constructor as in the example below, Grails will automatically convert the `byte` to the `myFile` property:

```
def img = new Image(params)
```

It's important that you set the [size](#) or [maxSize](#) constraints, otherwise your database may be created with reasonably sized files. For example, both H2 and MySQL default to a blob size of 255 bytes for `byte` properties.

It is also possible to set the contents of the file as a string by changing the type of the `myFile` property on

```
class Image {
    String myFile
}
```

6.1.10 Command Objects

Grails controllers support the concept of command objects. A command object is similar to a form bean and is useful for populating a subset of the properties needed to update a domain class. Or where there is no database but you need features such as [data binding](#) and [validation](#).

Declaring Command Objects

Command objects are typically declared in the same source file as a controller, directly below the controller class.

```
class UserController {
    ...
}

class LoginCommand {
    String username
    String password

    static constraints = {
        username(blank: false, minSize: 6)
        password(blank: false, minSize: 6)
    }
}
```

As this example shows, you can define [constraints](#) in command objects just like in [domain classes](#).

Using Command Objects

To use command objects, controller actions may optionally specify any number of command object parameters so that Grails knows what objects to create, populate and validate.

Before the controller action is executed Grails will automatically create an instance of the command object binding the request parameters, and validate the command object. For example:

```
class LoginController {
  def login = { LoginCommand cmd ->
    if (cmd.hasErrors()) {
      redirect(action: 'loginForm')
      return
    }

    // work with the command object data
  }
}
```

When using methods instead of Closures for actions, you can specify command objects in arguments:

```
class LoginController {
  def login(LoginCommand cmd) {
    if (cmd.hasErrors()) {
      redirect(action: 'loginForm')
      return
    }

    // work with the command object data
  }
}
```

Command Objects and Dependency Injection

Command objects can participate in dependency injection. This is useful if your command object has some [services](#):

```
class LoginCommand {
  def loginService

  String username
  String password

  static constraints = {
    username validator: { val, obj ->
      obj.loginService.canLogin(obj.username, obj.password)
    }
  }
}
```

In this example the command object interacts with the `loginService` bean which is in `ApplicationContext`.

6.1.11 Handling Duplicate Form Submissions

Grails has built-in support for handling duplicate form submissions using the "Synchronizer Token Pattern" the [form](#) tag:

A common use case is dealing with multiple request parameters of the same name. For example `?name=Bob&name=Judy`.

In this case dealing with one parameter and dealing with many has different semantics since Groovy's `list` iterates over each character. To avoid this problem the [params](#) object provides a `list` method that always returns

```
for (name in params.list('name')) {  
    println name  
}
```

6.1.13 Asynchronous Request Processing

Grails support asynchronous request processing as provided by the Servlet 3.0 specification. To enable the servlet target version to 3.0 in `BuildConfig.groovy`:

```
grails.servlet.version = "3.0"
```

With that done ensure you do a clean re-compile as some async features are enabled at compile time.



With a Servlet target version of 3.0 you can only deploy on Servlet 3.0 containers such as Tomcat 7.0 or Jetty 9.0.

Asynchronous Rendering

You can render content (templates, binary data etc.) in an asynchronous manner by calling the `startAsync` of the Servlet 3.0 `AsyncContext`. Once you have a reference to the `AsyncContext` you can use `render` to render content:

```
def index() {  
    def ctx = startAsync()  
    ctx.start {  
        new Book(title:"The Stand").save()  
        render template:"books", model:[books:Book.list()]  
        ctx.complete()  
    }  
}
```

Note that you must call the `complete()` method to terminate the connection.

Resuming an Async Request

You resume processing of an async request (for example to delegate to view rendering) by using the `AsyncContext` class:

```
def index() {
  def ctx = startAsync()
  ctx.start {
    // do working
    ...
    // render view
    ctx.dispatch()
  }
}
```

6.2 Groovy Server Pages

Groovy Servers Pages (or GSP for short) is Grails' view technology. It is designed to be familiar for users but to be far more flexible and intuitive.

GSPs live in the `grails-app/views` directory and are typically rendered automatically (by convention).

```
render(view: "index")
```

A GSP is typically a mix of mark-up and GSP tags which aid in view rendering.



Although it is possible to have Groovy logic embedded in your GSP and doing this within a view document, the practice is strongly discouraged. Mixing mark-up and code is a **bad** thing. Views should contain no code and needn't do so.

A GSP typically has a "model" which is a set of variables that are used for view rendering. The model is typically provided by the controller. For example consider the following controller action:

```
def show() {
  [book: Book.get(params.id)]
}
```

This action will look up a `Book` instance and create a model that contains a key called `book`. This key is then used in the view using the name `book`:

```
${book.title}
```

6.2.1 GSP Basics

In the next view sections we'll go through the basics of GSP and what is available to you. First off let's cover some basics that should be familiar with.

GSP supports the usage of `<% %>` scriptlet blocks to embed Groovy code (again this is discouraged):

```
<html>
  <body>
    <% out << "Hello GSP!" %>
  </body>
</html>
```

You can also use the `<%= %>` syntax to output values:

```
<html>
  <body>
    <%= "Hello GSP!" %>
  </body>
</html>
```

GSP also supports JSP-style server-side comments (which are not rendered in the HTML response) as the

```
<html>
  <body>
    <!-- This is my comment --%>
    <%= "Hello GSP!" %>
  </body>
</html>
```

6.2.1.1 Variables and Scopes

Within the `<% %>` brackets you can declare variables:

```
<% now = new Date() %>
```

and then access those variables later in the page:

```
<%=now%>
```

Within the scope of a GSP there are a number of pre-defined variables, including:

- application - The [javax.servlet.ServletContext](#) instance
- applicationContext The Spring [ApplicationContext](#) instance
- flash - The [flash](#) object
- grailsApplication - The [GrailsApplication](#) instance
- out - The response writer for writing to the output stream
- params - The [params](#) object for retrieving request parameters
- request - The [HttpServletRequest](#) instance
- response - The [HttpServletResponse](#) instance
- session - The [HttpSession](#) instance
- webRequest - The [GrailsWebRequest](#) instance

6.2.1.2 Logic and Iteration

Using the `<% %>` syntax you can embed loops and so on using this syntax:

```
<html>
  <body>
    <% [1,2,3,4].each { num -> %>
      <p><%= "Hello ${num}!" %></p>
    <%}%>
  </body>
</html>
```

As well as logical branching:

```
<html>
  <body>
    <% if (params.hello == 'true')%>
      <%= "Hello!" %>
    <% else %>
      <%= "Goodbye!" %>
    </body>
</html>
```

6.2.1.3 Page Directives

GSP also supports a few JSP-style page directives.

The import directive lets you import classes into the page. However, it is rarely needed due to Groovy's de

```
<%@ page import="java.awt.*" %>
```

GSP also supports the contentType directive:

```
<%@ page contentType="text/json" %>
```

The contentType directive allows using GSP to render other formats.

6.2.1.4 Expressions

In GSP the `<%= %>` syntax introduced earlier is rarely used due to the support for GSP expressions. An expression or a Groovy GString and takes the form `${expr}`:

```
<html>
  <body>
    Hello ${params.name}
  </body>
</html>
```


However, unlike JSP EL you can have any Groovy expression within the `${..}` block. Variables within default, so any HTML in the variable's string is rendered directly to the page. To reduce the risk of Cross-site scripting, enable automatic HTML escaping with the `grails.views.default.codec` setting in `grails-app`:

```
grails.views.default.codec='html'
```

Other possible values are 'none' (for no default encoding) and 'base64'.

6.2.2 GSP Tags

Now that the less attractive JSP heritage has been set aside, the following sections cover GSP's built-in tags to define GSP pages.

 The section on [Tag Libraries](#) covers how to add your own custom tag libraries.

All built-in GSP tags start with the prefix `g:`. Unlike JSP, you don't specify any tag library imports. If a tag is assumed to be a GSP tag. An example GSP tag would look like:

```
<g:example />
```

GSP tags can also have a body such as:

```
<g:example>
  Hello world
</g:example>
```

Expressions can be passed into GSP tag attributes, if an expression is not used it will be assumed to be a String.

```
<g:example attr="${new Date()}">
  Hello world
</g:example>
```

Maps can also be passed into GSP tag attributes, which are often used for a named parameter style syntax:

```
<g:example attr="${new Date()}" attr2="[one:1, two:2, three:3]">
  Hello world
</g:example>
```

Note that within the values of attributes you must use single quotes for Strings:

```
<g:example attr="${new Date()}" attr2="[one:'one', two:'two']">
  Hello world
</g:example>
```

With the basic syntax out the way, the next sections look at the tags that are built into Grails by default.

6.2.2.1 Variables and Scopes

Variables can be defined within a GSP using the [set](#) tag:

```
<g:set var="now" value="${new Date()}" />
```

Here we assign a variable called `now` to the result of a GSP expression (which simply constructs a new `Date` object). We also use the body of the `<g:set>` tag to define a variable:

```
<g:set var="myHTML">
  Some re-usable code on: ${new Date()}
</g:set>
```

Variables can also be placed in one of the following scopes:

- `page` - Scoped to the current page (default)
- `request` - Scoped to the current request
- `flash` - Placed within [flash](#) scope and hence available for the next request
- `session` - Scoped for the user session
- `application` - Application-wide scope.

To specify the scope, use the `scope` attribute:

```
<g:set var="now" value="${new Date()}" scope="request" />
```

6.2.2.2 Logic and Iteration

GSP also supports logical and iterative tags out of the box. For logic there are [if](#), [else](#) and [elseif](#) tags for us

```
<g:if test="${session.role == 'admin'}">
  <%-- show administrative functions --%>
</g:if>
<g:else>
  <%-- show basic functions --%>
</g:else>
```

Use the [each](#) and [while](#) tags for iteration:

```
<g:each in="${[1,2,3]}" var="num">
  <p>Number ${num}</p>
</g:each>

<g:set var="num" value="${1}" />
<g:while test="${num < 5}">
  <p>Number ${num++}</p>
</g:while>
```

6.2.2.3 Search and Filtering

If you have collections of objects you often need to sort and filter them. Use the [findAll](#) and [grep](#) tags for t

```
Stephen King's Books:
<g:findAll in="${books}" expr="it.author == 'Stephen King'">
  <p>Title: ${it.title}</p>
</g:findAll>
```

The `expr` attribute contains a Groovy expression that can be used as a filter. The [grep](#) tag does a similar j

```
<g:grep in="${books}" filter="NonFictionBooks.class">
  <p>Title: ${it.title}</p>
</g:grep>
```

Or using a regular expression:

```
<g:grep in="${books.title}" filter="~/.*?Groovy.*?/">
  <p>Title: ${it}</p>
</g:grep>
```

The above example is also interesting due to its usage of GPath. GPath is an XPath-like language in Groovy. It can auto-magically iterate the collection, obtain each title, and return a new list!

6.2.2.4 Links and Resources

GSP also features tags to help you manage linking to controllers and actions. The [link](#) tag lets you specify it will automatically work out the link based on the [URL Mappings](#), even if you change them! For example

```
<g:link action="show" id="1">Book 1</g:link>
<g:link action="show" id="${currentBook.id}">${currentBook.name}</g:link>
<g:link controller="book">Book Home</g:link>
<g:link controller="book" action="list">Book List</g:link>
<g:link url="[action: 'list', controller: 'book']">Book List</g:link>
<g:link params="[sort: 'title', order: 'asc', author: currentBook.author]"
  action="list">Book List</g:link>
```

6.2.2.5 Forms and Fields

Form Basics

GSP supports many different tags for working with HTML forms and fields, the most basic of which is the `<g:form>` tag. It is an aware version of the regular HTML form tag. The `url` attribute lets you specify which controller and action attributes apply.

```
<g:form name="myForm" url="[controller:'book',action:'list']">...</g:form>
```

In this case we create a form called `myForm` that submits to the `BookController`'s `list` action. The `url` attribute lets you specify which controller and action attributes apply.

Form Fields

In addition to easy construction of forms, GSP supports custom tags for dealing with different types of field.

- [textField](#) - For input fields of type 'text'
- [passwordField](#) - For input fields of type 'password'
- [checkBox](#) - For input fields of type 'checkbox'
- [radio](#) - For input fields of type 'radio'
- [hiddenField](#) - For input fields of type 'hidden'
- [select](#) - For dealing with HTML select boxes

Each of these allows GSP expressions for the value:

```
<g:textField name="myField" value="${myValue}" />
```

GSP also contains extended helper versions of the above tags such as [radioGroup](#) (for creating groups of [radio](#) buttons) and [timeZoneSelect](#) (for selecting locales, currencies and time zones respectively).

Multiple Submit Buttons

The age old problem of dealing with multiple submit buttons is also handled elegantly with Grails using regular submit, but lets you specify an alternative action to submit to:

```
<g:actionSubmit value="Some update label" action="update" />
```

6.2.2.6 Tags as Method Calls

One major difference between GSP tags and other tagging technologies is that GSP tags can be called as either [controllers](#), [tag libraries](#) or GSP views.

Tags as method calls from GSPs

Tags return their results as a String-like object (a `StreamCharBuffer` which has all of the same methods to the response when called as methods. For example:

```
Static Resource: ${createLinkTo(dir: "images", file: "logo.jpg")}
```

This is particularly useful for using a tag within an attribute:

```

```

In view technologies that don't support this feature you have to nest tags within tags, which becomes the effect of WYSIWIG tools such as Dreamweaver that attempt to render the mark-up as it is not well-formed

```
" />
```

Tags as method calls from Controllers and Tag Libraries

You can also invoke tags from controllers and tag libraries. Tags within the default `g`: [namespace](#) can return a `StreamCharBuffer` result is returned:

```
def imageLocation = createLinkTo(dir:"images", file:"logo.jpg").toString()
```

Prefix the namespace to avoid naming conflicts:

```
def imageLocation = g.createLinkTo(dir:"images", file:"logo.jpg").toString()
```

For tags that use a [custom namespace](#), use that prefix for the method call. For example (from the [FCK Editor](#))

```
def editor = fckeditor.editor(name: "text", width: "100%", height: "400")
```

6.2.3 Views and Templates

Grails also has the concept of templates. These are useful for partitioning your views into maintainable and provide a highly re-usable mechanism for structured views.

Template Basics

Grails uses the convention of placing an underscore before the name of a view to identify it as a template that renders Books located at `grails-app/views/book/_bookTemplate.gsp`:

```
<div class="book" id="${book?.id}">
  <div>Title: ${book?.title}</div>
  <div>Author: ${book?.author?.name}</div>
</div>
```

Use the [render](#) tag to render this template from one of the views in `grails-app/views/book`:

```
<g:render template="bookTemplate" model="[book: myBook]" />
```

Notice how we pass into a model to use using the `model` attribute of the `render` tag. If you have multiple templates for each Book using the `render` tag with a `collection` attribute:

```
<g:render template="bookTemplate" var="book" collection="${bookList}" />
```

Shared Templates

In the previous example we had a template that was specific to the `BookController` and its view. However, you may want to share templates across your application.

In this case you can place them in the root views directory at `grails-app/views` or any subdirectory below it. To use an absolute location starting with `/` instead of a relative location. For example, `grails-app/views/shared/_mySharedTemplate.gsp`, you would reference it as:

```
<g:render template="/shared/mySharedTemplate" />
```

You can also use this technique to reference templates in any directory from any view or controller:

```
<g:render template="/book/bookTemplate" model="[book: myBook]" />
```

The Template Namespace

Since templates are used so frequently there is template namespace, called `tmpl`, available that make example the following usage pattern:

```
<g:render template="bookTemplate" model="[book:myBook]" />
```

This can be expressed with the `tmpl` namespace as follows:

```
<tmpl:bookTemplate book="${myBook}" />
```

Templates in Controllers and Tag Libraries

You can also render templates from controllers using the [render](#) controller method. This is useful for [Ajax](#) HTML or data responses to partially update the current page instead of performing new request:

```
def bookData() {
    def b = Book.get(params.id)
    render(template: "bookTemplate", model:[book:b])
}
```

The [render](#) controller method writes directly to the response, which is the most common behaviour. To in String you can use the [render](#) tag:

```
def bookData() {
    def b = Book.get(params.id)
    String content = g.render(template: "bookTemplate", model:[book:b])
    render content
}
```

Notice the usage of the `g` namespace which tells Grails we want to use the [tag as method call](#) instead of the

6.2.4 Layouts with Sitemesh

Creating Layouts

Grails leverages [Sitemesh](#), a decorator engine, to support view layouts. Layouts are located in the `grail` A typical layout can be seen below:

```

<html>
  <head>
    <title><g:layoutTitle default="An example decorator" /></title>
    <g:layoutHead />
  </head>
  <body onload="\${pageProperty(name:'body.onload')}">
    <div class="menu"><!--my common menu goes here--></menu>
    <div class="body">
      <g:layoutBody />
    </div>
  </div>
</body>
</html>

```

The key elements are the [layoutHead](#), [layoutTitle](#) and [layoutBody](#) tag invocations:

- `layoutTitle` - outputs the target page's title
- `layoutHead` - outputs the target page's head tag contents
- `layoutBody` - outputs the target page's body tag contents

The previous example also demonstrates the [pageProperty](#) tag which can be used to inspect and return aspects of the page.

Triggering Layouts

There are a few ways to trigger a layout. The simplest is to add a meta tag to the view:

```

<html>
  <head>
    <title>An Example Page</title>
    <meta name="layout" content="main" />
  </head>
  <body>This is my content!</body>
</html>

```

In this case a layout called `grails-app/views/layouts/main.gsp` will be used to layout the page. In the previous section the output would resemble this:

```

<html>
  <head>
    <title>An Example Page</title>
  </head>
  <body onload="">
    <div class="menu"><!--my common menu goes here--></div>
    <div class="body">
      This is my content!
    </div>
  </body>
</html>

```

Specifying A Layout In A Controller

Another way to specify a layout is to specify the name of the layout by assigning a value to the `"layout"` property in a controller such as:

```
class BookController {
    static layout = 'customer'

    def list() { ... }
}
```

You can create a layout called `grails-app/views/layouts/customer.gsp` which will `BookController` delegates to. The value of the "layout" property may contain a directory in `grails-app/views/layouts/` directory. For example:

```
class BookController {
    static layout = 'custom/customer'

    def list() { ... }
}
```

Views rendered from that controller would be decorated with the `grails-app/views/layouts/custom/customer.gsp`.

Layout by Convention

Another way to associate layouts is to use "layout by convention". For example, if you have this controller

```
class BookController {
    def list() { ... }
}
```

You can create a layout called `grails-app/views/layouts/book.gsp`, which will be applied to the `list` action delegates to.

Alternatively, you can create a layout called `grails-app/views/layouts/book/list.gsp` within the `BookController`.

If you have both the above mentioned layouts in place the layout specific to the action will take precedence.

If a layout may not be located using any of those conventions, the convention of last resort is to look for `grails-app/views/layouts/application.gsp`. The name of the application default layout is defined in `grails-app/conf/Config.groovy` as follows:

```
grails.sitemesh.default.layout = 'myLayoutName'
```

With that property in place, the application default layout will be `grails-app/views/layouts/myLayoutName.gsp`.

Inline Layouts

Grails' also supports Sitemesh's concept of inline layouts with the [applyLayout](#) tag. This can be used to wrap an arbitrary section of content. This lets you even further modularize your view structure by "decorating" you

Some examples of usage can be seen below:

```
<g:applyLayout name="myLayout" template="bookTemplate" collection="{books}" />
<g:applyLayout name="myLayout" url="http://www.google.com" />
<g:applyLayout name="myLayout">
The content to apply a layout to
</g:applyLayout>
```

Server-Side Includes

While the [applyLayout](#) tag is useful for applying layouts to external content, if you simply want to include you use the [include](#) tag:

```
<g:include controller="book" action="list" />
```

You can even combine the [include](#) tag and the [applyLayout](#) tag for added flexibility:

```
<g:applyLayout name="myLayout">
  <g:include controller="book" action="list" />
</g:applyLayout>
```

Finally, you can also call the [include](#) tag from a controller or tag library as a method:

```
def content = include(controller:"book", action:"list")
```

The resulting content will be provided via the return value of the [include](#) tag.

6.2.5 Static Resources

Grails 2.0 integrates with the [Resources plugin](#) to provide sophisticated static resource management. The Grails applications.

The basic way to include a link to a static resource in your application is to use the [resource](#) tag. This simplifies the file.

However modern applications with dependencies on multiple JavaScript and CSS libraries and frameworks (Grails plugins) require something more powerful.

The issues that the Resources framework tackles are:

- Web application performance tuning is difficult
- Correct ordering of resources, and deferred inclusion of JavaScript
- Resources that depend on others that must be loaded first
- The need for a standard way to expose static resources in plugins and applications
- The need for an extensible processing chain to optimize resources
- Preventing multiple inclusion of the same resource

The plugin achieves this by introducing new artefacts and processing the resources using the server's local file system. It adds artefacts for declaring resources, for declaring "mappers" that can process resources, and a servlet filter. What you get is an incredibly advanced resource system that enables you to easily create highly optimized applications in development and in production.

The Resources plugin documentation provides a more detailed overview of the [concepts](#) which will be a good guide.

6.2.5.1 Including resources using the resource tags

Pulling in resources with `r:require`

To use resources, your GSP page must indicate which resource modules it requires. For example with the "jquery" resource module, to use jQuery in any page on your site you simply add:

```
<html>
  <head>
    <r:require module="jquery" />
    <r:layoutResources />
  </head>
  <body>
    ...
    <r:layoutResources />
  </body>
</html>
```

This will automatically include all resources needed for jQuery, including them at the correct locations in the page, in a disposition to be "head", so they load early in the page.

You can call `r:require` multiple times in a GSP page, and you use the "modules" attribute to provide a

```
<html>
  <head>
    <r:require modules="jquery, main, blueprint, charting" />
    <r:layoutResources />
  </head>
  <body>
    ...
    <r:layoutResources />
  </body>
</html>
```

The above may result in many JavaScript and CSS files being included, in the correct order, with some JavaScript in the body to improve the apparent page load time.

However you cannot use `r:require` in isolation - as per the examples you must have the `<r:layoutResources`

Rendering the links to resources with `r:layoutResources`

When you have declared the resource modules that your GSP page requires, the framework needs to render the resources at the correct time.

To achieve this correctly, you must include the `r:layoutResources` tag twice in your page, or more commonly

```
<html>
  <head>
    <g:layoutTitle/>
    <r:layoutResources/>
  </head>
  <body>
    <g:layoutBody/>
    <r:layoutResources/>
  </body>
</html>
```

This represents the simplest Sitemesh layout you can have that supports Resources.

The Resources framework has the concept of a "disposition" for every resource. This is an indication of whether the resource is included.

The default disposition applied depends on the type of resource. All CSS must be rendered in `<head>` in CSS, and will be rendered by the first `r:layoutResources`. Page load times are improved when JavaScript default for JavaScript files is "defer", which means it is rendered when the second `r:layoutResources` is invoked.

Note that both your GSP page and your Sitemesh layout (as well as any GSP template fragments) can call `r:require`. The only limitation is that you must call `r:require` before the `r:layoutResources` that should render it.

Adding page-specific JavaScript code with `r:script`

Grails has the [javascript](#) tag which is adapted to defer to Resources plugin if installed, but it is recommended when you need to include fragments of JavaScript code.

This lets you write some "inline" JavaScript which is actually **not** rendered inline, but either in the `<head>` disposition.

Given a Sitemesh layout like this:

```
<html>
  <head>
    <g:layoutTitle/>
    <r:layoutResources/>
  </head>
  <body>
    <g:layoutBody/>
    <r:layoutResources/>
  </body>
</html>
```

...in your GSP you can inject some JavaScript code into the head or deferred regions of the page like this:


```

<html>
  <head>
    <title>Testing r:script magic!</title>
  </head>
  <body>
    <r:script disposition="head">
      window.alert('This is at the end of <head>');
    </r:script>
    <r:script disposition="defer">
      window.alert('This is at the end of the body, and the page has loaded.')
    </r:script>
  </body>
</html>

```

The default disposition is "defer", so the disposition in the latter r:script is purely included for demonstration. Note that such r:script code fragments **always** load after any modules that you have used, to ensure that an

Linking to images with r:img

This tag is used to render markup, using the Resources framework to process the resource on the fly (and make it eternally cacheable).

This includes any extra attributes on the tag if the resource has been previously declared in a module.

With this mechanism you can specify the width, height and any other attributes in the resource declaration in as necessary.

Example:

```

<html>
  <head>
    <title>Testing r:img</title>
  </head>
  <body>
    <r:img uri="/images/logo.png"/>
  </body>
</html>

```

Note that Grails has a built-in g:img tag as a shortcut for rendering tags that refer to a static resource. It is Resources-aware and will delegate to r:img if found. However it is recommended that you use r:img directly.

Alongside the regular Grails [resource](#) tag attributes, this also supports the "uri" attribute for increased brevity. See [r:resource documentation](#) for full details.

6.2.5.2 Other resource tags

r:resource

This is equivalent to the Grails [resource](#) tag, returning a link to the processed static resource. Grails' own implementation if found, but if your code requires the Resources plugin, you should use r:resource directly.

Alongside the regular Grails [resource](#) tag attributes, this also supports the "uri" attribute for increased brevity. See [r:resource documentation](#) for full details.

r:external

This is a resource-aware version of Grails [external](#) tag which renders the HTML markup necessary to include CSS, JS or a favicon.

See [r:resource documentation](#) for full details.

6.2.5.3 Declaring resources

A DSL is provided for declaring resources and modules. This can go either in your `Config.groovy` resources, or more commonly in a resources artefact in `grails-app/conf`.

Note that you do not need to declare all your static resources, especially images. However you must declare resource-specific attributes. Any resource that is not declared is called "ad-hoc" and will still be processed.

Consider this example resource configuration file, `grails-app/conf/MyAppResources.groovy`:

```
modules = {
    core {
        dependsOn 'jquery, utils'

        resource url: '/js/core.js', disposition: 'head'
        resource url: '/js/ui.js'
        resource url: '/css/main.css',
        resource url: '/css/branding.css'
        resource url: '/css/print.css', attrs: [media: 'print']
    }

    utils {
        dependsOn 'jquery'

        resource url: '/js/utils.js'
    }

    forms {
        dependsOn 'core,utils'

        resource url: '/css/forms.css'
        resource url: '/js/forms.js'
    }
}
```

This defines three resource modules; 'core', 'utils' and 'forms'. The resources in these modules will be loaded according to the module name, resulting in fewer files. You can override this with `bundle: 'someOtherBundle'` on the module (see [resources plugin documentation](#)).

It declares dependencies between them using `dependsOn`, which controls the load order of the resources.

When you include an `<r:require module="forms"/>` in your GSP, it will pull in all the resources 'jquery', all in the correct order.

You'll also notice the `disposition: 'head'` on the `core.js` file. This tells Resources that while it is part of the body, this one must go into the `<head>`.

The CSS file for print styling adds custom attributes using the `attrs` map option, and these are passed to the engine renders the link to the resource, so you can customize the HTML attributes of the generated link.

There is no limit to the number of modules or `xxxResources.groovy` artefacts you can provide, and plugins can provide applications, which is exactly how the jQuery plugin works.

To define modules like this in your application's Config.groovy, you simply assign the DSL closure to the Config variable.

For full details of the resource DSL please see the [resources plugin documentation](#).

6.2.5.4 Overriding plugin resources

Because a resource module can define the bundle groupings and other attributes of resources, you may correct for your application.

For example, you may wish to bundle jQuery and some other libraries all together in one file. There is a lot often it is the case that you'd like to override some of these settings.

To do this, the DSL supports an "overrides" clause, within which you can change the defaultBundle individual resources that have been declared with a unique id:

```
modules = {
    core {
        dependsOn 'jquery, utils'
        defaultBundle 'monolith'

        resource url: '/js/core.js', disposition: 'head'
        resource url: '/js/ui.js'
        resource url: '/css/main.css',
        resource url: '/css/branding.css'
        resource url: '/css/print.css', attrs: [media: 'print']
    }

    utils {
        dependsOn 'jquery'
        defaultBundle 'monolith'

        resource url: '/js/utils.js'
    }

    forms {
        dependsOn 'core,utils'
        defaultBundle 'monolith'

        resource url: '/css/forms.css'
        resource url: '/js/forms.js'
    }

    overrides {
        jquery {
            defaultBundle 'monolith'
        }
    }
}
```

This will put all code into a single bundle named 'monolith'. Note that this can still result in multiple file head and defer dispositions, and JavaScript and CSS files are bundled separately.

Note that overriding individual resources requires the original declaration to have included a unique id for

For full details of the resource DSL please see the [resources plugin documentation](#).

6.2.5.5 Optimizing your resources

The Resources framework uses "mappers" to mutate the resources into the final format served to the user.

The resource mappers are applied to each static resource once, in a specific order. You can create your own plugins provide some already for zipping, caching and minifying.

Out of the box, the Resources plugin provides bundling of resources into fewer files, which is achieved CSS re-writing to handle when your CSS files are moved into a bundle.

Bundling multiple resources into fewer files

The 'bundle' mapper operates by default on any resource with a "bundle" defined - or inherited from a de Modules have an implicit default bundle name the same as the name of the module.

Files of the same kind will be aggregated into this bundle file. Bundles operate across module boundaries:

```
modules = {
  core {
    dependsOn 'jquery, utils'
    defaultBundle 'common'

    resource url: '/js/core.js', disposition: 'head'
    resource url: '/js/ui.js', bundle: 'ui'
    resource url: '/css/main.css', bundle: 'theme'
    resource url: '/css/branding.css'
    resource url: '/css/print.css', attrs: [media: 'print']
  }

  utils {
    dependsOn 'jquery'

    resource url: '/js/utils.js', bundle: 'common'
  }

  forms {
    dependsOn 'core,utils'

    resource url: '/css/forms.css', bundle: 'ui'
    resource url: '/js/forms.js', bundle: 'ui'
  }
}
```

Here you see that resources are grouped into bundles; 'common', 'ui' and 'theme' - across module boundaries

Note that auto-bundling by module does **not** occur if there is only one resource in the module.

Making resources cache "eternally" in the client browser

Caching resources "eternally" in the client is only viable if the resource has a unique name that changes and requires caching headers to be set on the response.

The [cached-resources](#) plugin provides a mapper that achieves this by hashing your files and renaming them, and setting caching headers on every response for those resources. To use, simply install the cached-resources plugin.

Note that the caching headers can only be set if your resources are being served by your application. If you are serving content from your app (e.g. Apache HTTPD), configure it to send caching headers. Alternatively you can serve resources from your container.

Zippping resources

Returning gzipped resources is another way to reduce page load times and reduce bandwidth.

The [zipped-resources](#) plugin provides a mapper that automatically compresses your content, excluding files of types such as gif, jpeg and png.

Simply install the zipped-resources plugin and it works.

Minifying

There are a number of CSS and JavaScript minifiers available to obfuscate and reduce the size of your publicly released but releases are imminent.

6.2.5.6 Debugging

When your resources are being moved around, renamed and otherwise mutated, it can be hard to debug, especially Safari, Chrome and Firefox have excellent tools that let you view all the resources requested by information about them.

There are several debugging features built in to the Resources framework.

X-Grails-Resources-Original-Src Header

Every resource served in development mode will have the X-Grails-Resources-Original-Src: header added that make up the response.

Adding the debug flag

If you add a query parameter `_debugResources=y` to your URL and request the page, Resources will bypass your original source files.

This also adds a unique timestamp to all your resource URLs, to defeat any caching that browsers may use to see your very latest code when you reload the page.

Turning on debug all the time

You can turn on the aforementioned debug mechanism without requiring a query parameter, but turning it

```
grails.resources.debug = true
```

You can of course set this per-environment.

6.2.5.7 Preventing processing of resources

Sometimes you do not want a resource to be processed in a particular way, or even at all. Occasionally you may want to prevent mapping.

Preventing the application of a specific mapper to an individual resource

All resource declarations support a convention of `noXXXX:true` where XXXX is a mapper name.

So for example to prevent the "hashandcache" mapper from being applied to a resource (which renames relative links written in JavaScript code), you would do this:

```
modules = {
  forms {
    resource url: '/css/forms.css', nohashandcache: true
    resource url: '/js/forms.js', nohashandcache: true
  }
}
```

Excluding/including paths and file types from specific mappers

Mappers have includes/excludes Ant patterns to control whether they apply to a given resource. Mappers their activity, for example the zipped-resources plugin's "zip" mapper is set to exclude images by default.

You can configure this in your `Config.groovy` using the mapper name e.g:

```
// We wouldn't link to .exe files using Resources but for the sake of example:
grails.resources.zip.excludes = ['**/*.zip', '**/*.exe']

// Perhaps for some reason we want to prevent bundling on "less" CSS files:
grails.resources.bundle.excludes = ['**/*.less']
```

There is also an "includes" inverse. Note that settings these replaces the default includes/excludes for that mapper.

Controlling what is treated as an "ad-hoc" (legacy) resource

Ad-hoc resources are those undeclared, but linked to directly in your application **without** using the Grails resource mapping (internal or external).

These may occur with some legacy plugins or code with hardcoded paths in.

There is a `Config.groovy` setting **`grails.resources.adhoc.patterns`** which defines a list of Servlet API controller resources filter will use to detect such "ad-hoc resource" requests.

By default this is set to:

```
grails.resources.adhoc.patterns = ['images/*', '*.js', '*.css']
```

6.2.5.8 Other Resources-aware plugins

At the time of writing, the following plugins include support for the Resources framework:

- [jquery](#)
- [jquery-ui](#)
- [blueprint](#)
- [lesscss-resources](#)
- [zipped-resources](#)
- [cached-resources](#)

6.2.6 Sitemesh Content Blocks

Although it is useful to decorate an entire page sometimes you may find the need to decorate independently can use content blocks. To get started, partition the page to be decorated using the `<content>` tag:

```
<content tag="navbar">
... draw the navbar here...
</content>

<content tag="header">
... draw the header here...
</content>

<content tag="footer">
... draw the footer here...
</content>

<content tag="body">
... draw the body here...
</content>
```

Then within the layout you can reference these components and apply individual layouts to each:

```
<html>
  <body>
    <div id="header">
      <g:applyLayout name="headerLayout">
        <g:pageProperty name="page.header" />
      </g:applyLayout>
    </div>
    <div id="nav">
      <g:applyLayout name="navLayout">
        <g:pageProperty name="page.navbar" />
      </g:applyLayout>
    </div>
    <div id="body">
      <g:applyLayout name="bodyLayout">
        <g:pageProperty name="page.body" />
      </g:applyLayout>
    </div>
    <div id="footer">
      <g:applyLayout name="footerLayout">
        <g:pageProperty name="page.footer" />
      </g:applyLayout>
    </div>
  </body>
</html>
```

6.2.7 Making Changes to a Deployed Application

One of the main issues with deploying a Grails application (or typically any servlet-based one) is that any redeploy your whole application. If all you want to do is fix a typo on a page, or change an image link, it can be a pain. For such simple requirements, Grails does have a solution: the `grails.gsp.view.dir` configuration:

How does this work? The first step is to decide where the GSP files should go. Let's say we have a directory `/var/www/grails/my-app`. We add these two lines to `grails-app/conf/Config.groovy`:

```
grails.gsp.enable.reload = true
grails.gsp.view.dir = "/var/www/grails/my-app/"
```

The first line tells Grails that modified GSP files should be reloaded at runtime. If you don't have this set you like but they won't be reflected in the running application until you restart. The second line tells Grails from.



The trailing slash on the `grails.gsp.view.dir` value is important! Without it, Grails uses the parent directory.

Setting "`grails.gsp.view.dir`" is optional. If it's not specified, you can update files directly to the application. Depending on the application server, these files might get overwritten when the server is restarted. Most "war deployment" which is recommended in this case.

With those settings in place, all you need to do is copy the views from your web application to the external server. It would look something like this:

```
mkdir -p /var/www/grails/my-app/grails-app/views
cp -R grails-app/views/* /var/www/grails/my-app/grails-app/views
```

The key point here is that you must retain the view directory structure, including the `grails-app/views` directory. The path would be `/var/www/grails/my-app/grails-app/views/...`

One thing to bear in mind with this technique is that every time you modify a GSP, it uses up permanent memory. Eventually it will hit "out of permanent memory space" errors unless you restart the server. So this technique is not recommended for production environments.

There are also some System properties to control GSP reloading:

Name	Description
<code>grails.gsp.enable.reload</code>	alternative system property for enabling the GSP reload mode without changing the default
<code>grails.gsp.reload.interval</code>	interval between checking the lastmodified time of the gsp source file, unit is seconds
<code>grails.gsp.reload.granularity</code>	the number of milliseconds leeway to give before deciding a file is out of date. Different roundings usually cause a 1000ms difference in lastmodified times

GSP reloading is supported for precompiled GSPs since Grails 1.3.5 .

6.2.8 GSP Debugging

Viewing the generated source code

- Adding "`?showSource=true`" or "`&showSource=true`" to the url shows the generated Groovy source code. It won't show the source code of included templates. This only works in development mode
- The saving of all generated source code can be activated by setting the property "`grails.views.gsp.keepSource`". It must point to a directory that exists and is writable.
- During "grails war" gsp pre-compilation, the generated source code is stored in `grails.project.generatedSourceDir` (default: `~/grails/(grails_version)/projects/(project name)/gspcompile`).

Debugging GSP code with a debugger

- See [Debugging GSP in STS](#)

Viewing information about templates used to render a single url

GSP templates are reused in large web applications by using the `g:render` taglib. Several small templates might be hard to find out what GSP template actually renders the HTML seen in the result. The debug templates add comments to the output. The comments contain debug information about GSP templates used to render the page.

Usage is simple: append `"?debugTemplates"` or `"&debugTemplates"` to the URL and view the result. `"debugTemplates"` is restricted to development mode. It won't work in production.

Here is an example of comments added by debugTemplates :

```
<!-- GSP #2 START template: /home/.../views/_carousel.gsp
      precompiled: false lastmodified: ... -->
.
.
.
<!-- GSP #2 END template: /home/.../views/_carousel.gsp
      rendering time: 115 ms -->
```

Each comment block has a unique ID so that you can find the start & end of each template call.

6.3 Tag Libraries

Like [Java Server Pages](#) (JSP), GSP supports the concept of custom tag libraries. Unlike JSP, Grails' tag libraries are completely reloadable at runtime.

Quite simply, to create a tag library create a Groovy class that ends with the convention `grails-app/taglib` directory:

```
class SimpleTagLib {
}
```

Now to create a tag create a Closure property that takes two arguments: the tag attributes and the body content.

```
class SimpleTagLib {
    def simple = { attrs, body ->
    }
}
```

The `attrs` argument is a Map of the attributes of the tag, whilst the `body` argument is a Closure that returns the body content.

```
class SimpleTagLib {
    def emoticon = { attrs, body ->
        out << body() << (attrs.happy == 'true' ? " :-)" : " :-(")
    }
}
```

As demonstrated above there is an implicit `out` variable that refers to the output `Writer` which you can use. Then you can reference the tag inside your GSP; no imports are necessary:

```
<g:emoticon happy="true">Hi John</g:emoticon>
```



To help IDEs like SpringSource Tool Suite (STS) and others autocomplete tag attributes, you should add comments to your tag closures with `@attr` descriptions. Since taglibs use Groovy code it can detect all usable attributes.

For example:

```
class SimpleTagLib {
    /**
     * Renders the body with an emoticon.
     *
     * @attr happy whether to show a happy emoticon ('true') or
     * a sad emoticon ('false')
     */
    def emoticon = { attrs, body ->
        out << body() << (attrs.happy == 'true' ? " :-)" : " :-(")
    }
}
```

and any mandatory attributes should include the `REQUIRED` keyword, e.g.

```
class SimpleTagLib {
    /**
     * Creates a new password field.
     *
     * @attr name REQUIRED the field name
     * @attr value the field value
     */
    def passwordField = { attrs ->
        attrs.type = "password"
        attrs.tagName = "passwordField"
        fieldImpl(out, attrs)
    }
}
```

6.3.1 Variables and Scopes

Within the scope of a tag library there are a number of pre-defined variables including:

- `actionName` - The currently executing action name
- `controllerName` - The currently executing controller name
- `flash` - The [flash](#) object
- `grailsApplication` - The [GrailsApplication](#) instance
- `out` - The response writer for writing to the output stream
- `pageScope` - A reference to the [pageScope](#) object used for GSP rendering (i.e. the binding)
- `params` - The [params](#) object for retrieving request parameters
- `pluginContextPath` - The context path to the plugin that contains the tag library
- `request` - The [HttpServletRequest](#) instance
- `response` - The [HttpServletResponse](#) instance
- `servletContext` - The [javax.servlet.ServletContext](#) instance
- `session` - The [HttpSession](#) instance

6.3.2 Simple Tags

As demonstrated in the previous example it is easy to write simple tags that have no body and just call a `dateFormat` style tag:

```
def dateFormat = { attrs, body ->
    out << new java.text.SimpleDateFormat(attrs.format).format(attrs.date)
}
```

The above uses Java's `SimpleDateFormat` class to format a date and then write it to the response. The following:

```
<g:dateFormat format="dd-MM-yyyy" date="{new Date()}" />
```

With simple tags sometimes you need to write HTML mark-up to the response. One approach would be to

```
def formatBook = { attrs, body ->
    out << "<div id='${attrs.book.id}'>"
    out << "Title : ${attrs.book.title}"
    out << "</div>"
}
```

Although this approach may be tempting it is not very clean. A better approach would be to reuse the [render](#)

```
def formatBook = { attrs, body ->
    out << render(template: "bookTemplate", model: [book: attrs.book])
}
```

And then have a separate GSP template that does the actual rendering.

6.3.3 Logical Tags

You can also create logical tags where the body of the tag is only output once a set of conditions have been met. For example, you can create a logical tag that checks if a user is an administrator and only outputs the body content if he/she has the correct permissions:

```
def isAdmin = { attrs, body ->
    def user = attrs.user
    if (user && checkUserPrivs(user)) {
        out << body()
    }
}
```

The tag above checks if the user is an administrator and only outputs the body content if he/she has the correct permissions.

```
<g:isAdmin user="${myUser}">
    // some restricted content
</g:isAdmin>
```

6.3.4 Iterative Tags

Iterative tags are easy too, since you can invoke the body multiple times:

```
def repeat = { attrs, body ->
    attrs.times?.toInteger()?.times { num ->
        out << body(num)
    }
}
```

In this example we check for a `times` attribute and if it exists convert it to a number, then use Groovy's `times` method to repeat the body a number of times:

```
<g:repeat times="3">
<p>Repeat this 3 times! Current repeat = ${it}</p>
</g:repeat>
```

Notice how in this example we use the implicit `it` variable to refer to the current number. This works because the value of `times` is passed in the current value inside the iteration:

```
out << body(num)
```

That value is then passed as the default variable `it` to the tag. However, if you have nested tags this can instead name the variables that the body uses:

```
def repeat = { attrs, body ->
  def var = attrs.var ?: "num"
  attrs.times?.toInteger()?.times { num ->
    out << body((var):num)
  }
}
```

Here we check if there is a `var` attribute and if there is use that as the name to pass into the body invocation.

```
out << body((var):num)
```



Note the usage of the parenthesis around the variable name. If you omit these Groovy assigns the String key and not referring to the variable itself.

Now we can change the usage of the tag as follows:

```
<g:repeat times="3" var="j">
<p>Repeat this 3 times! Current repeat = ${j}</p>
</g:repeat>
```

Notice how we use the `var` attribute to define the name of the variable `j` and then we are able to refer to it in the tag.

6.3.5 Tag Namespaces

By default, tags are added to the default Grails namespace and are used with the `g:` prefix in GSP pages. To use a different namespace by adding a static property to your `TagLib` class:

```
class SimpleTagLib {
  static namespace = "my"
  def example = { attrs ->
    ...
  }
}
```

Here we have specified a namespace of `my` and hence the tags in this tag lib must then be referenced from

```
<my:example name="..." />
```

where the prefix is the same as the value of the static `namespace` property. Namespaces are particularly useful when you have multiple tag libraries.

Tags within namespaces can be invoked as methods using the namespace as a prefix to the method call:

```
out << my.example(name: "foo")
```

This works from GSP, controllers or tag libraries.

6.3.6 Using JSP Tag Libraries

In addition to the simplified tag library mechanism provided by GSP, you can also use JSP tags from Grails. You can use with the `taglib` directive:

```
<%@ taglib prefix="fmt" uri="http://java.sun.com/jsp/jstl/fmt" %>
```

Then you can use it like any other tag:

```
<fmt:formatNumber value="${10}" pattern=".00"/>
```

With the added bonus that you can invoke JSP tags like methods:

```
${fmt.formatNumber(value:10, pattern:".00")}
```

6.3.7 Tag return value

Since Grails 1.2, a tag library call returns an instance of `org.codehaus.groovy.grails.web.util.DefaultTagReturn` by default. This change improves performance by reducing object creation and optimizing buffering during rendering. In previous versions, a `java.lang.String` instance was returned.

Tag libraries can also return direct object values to the caller since Grails 1.2.. Object return is controlled by the `returnObjectForTags` property in the tag library class.

Example:

```
class ObjectReturningTagLib {
    static namespace = "cms"
    static returnObjectForTags = ['content']

    def content = { attrs, body ->
        CmsContent.findByCode(attrs.code)?.content
    }
}
```

6.4 URL Mappings

Throughout the documentation so far the convention used for URLs has been the default of `/controller/action`. This convention is not hard wired into Grails and is in fact controlled by a `UrlMappings` class in `grails-app/conf/UrlMappings.groovy`.

The `UrlMappings` class contains a single property called `mappings` that has been assigned a block of code.

```
class UrlMappings {  
    static mappings = {  
    }  
}
```

6.4.1 Mapping to Controllers and Actions

To create a simple mapping simply use a relative URL as the method name and specify named parameters

```
"/product"(controller: "product", action: "list")
```

In this case we've mapped the URL `/product` to the `list` action of the `ProductController`. O default action of the controller:

```
"/product"(controller: "product")
```

An alternative syntax is to assign the controller and action to use within a block passed to the method:

```
"/product" {  
    controller = "product"  
    action = "list"  
}
```

Which syntax you use is largely dependent on personal preference. To rewrite one URI onto another (expl pair) do something like this:

```
"/hello"(uri: "/hello.dispatch")
```

Rewriting specific URIs is often useful when integrating with other frameworks.

6.4.2 Embedded Variables

Simple Variables

The previous section demonstrated how to map simple URLs with concrete "tokens". In URL mapping characters between each slash, `/`. A concrete token is one which is well defined such as `/product`. don't know what the value of a particular token will be until runtime. In this case you can use variable plac

```
static mappings = {  
    "/product/$id"(controller: "product")  
}
```

In this case by embedding a `$id` variable as the second token Grails will automatically map the second token to the `params` object) called `id`. For example given the URL `/product/MacBook`, the following code will render

```
class ProductController {
    def index() { render params.id }
}
```

You can of course construct more complex examples of mappings. For example the traditional blog URL format

```
static mappings = {
    "/$blog/$year/$month/$day/$id"(controller: "blog", action: "show")
}
```

The above mapping would let you do things like:

```
/graemerocher/2007/01/10/my_funky_blog_entry
```

The individual tokens in the URL would again be mapped into the `params` object with values available for

Dynamic Controller and Action Names

Variables can also be used to dynamically construct the controller and action name. In fact the default Grails

```
static mappings = {
    "/$controller/$action?/$id?"()
}
```

Here the name of the controller, action and id are implicitly obtained from the variables `controller`, `action` and `id` in the URL.

You can also resolve the controller name and action name to execute dynamically using a closure:

```
static mappings = {
    "/$controller" {
        action = { params.goHere }
    }
}
```

Optional Variables

Another characteristic of the default mapping is the ability to append a `?` at the end of a variable to make it optional. For example this technique could be applied to the blog URL mapping to have more flexible linking:


```
static mappings = {
  "$blog/$year?/$month?/$day?/$id?"(controller:"blog", action:"show")
}
```

With this mapping all of these URLs would match with only the relevant parameters being populated in the

```
/graemerocher/2007/01/10/my_funky_blog_entry
/graemerocher/2007/01/10
/graemerocher/2007/01
/graemerocher/2007
/graemerocher
```

Arbitrary Variables

You can also pass arbitrary parameters from the URL mapping into the controller by just setting them in the

```
"/holiday/win" {
  id = "Marrakech"
  year = 2007
}
```

This variables will be available within the [params](#) object passed to the controller.

Dynamically Resolved Variables

The hard coded arbitrary variables are useful, but sometimes you need to calculate the name of the variable possible by assigning a block to the variable name:

```
"/holiday/win" {
  id = { params.id }
  isEligible = { session.user != null } // must be logged in
}
```

In the above case the code within the blocks is resolved when the URL is actually matched and hence can do logic.

6.4.3 Mapping to Views

You can resolve a URL to a view without a controller or action involved. For example to map the root of the rails-app/grails-app/views/index.gsp you could use:

```
static mappings = {
  "/"(view: "/index") // map the root URL
}
```

Alternatively if you need a view that is specific to a given controller you could use:

```
static mappings = {
  "/help"(controller: "site", view: "help") // to a view for a controller
}
```

6.4.4 Mapping to Response Codes

Grails also lets you map HTTP response codes to controllers, actions or views. Just use a method name that you are interested in:

```
static mappings = {
  "403"(controller: "errors", action: "forbidden")
  "404"(controller: "errors", action: "notFound")
  "500"(controller: "errors", action: "serverError")
}
```

Or you can specify custom error pages:

```
static mappings = {
  "403"(view: "/errors/forbidden")
  "404"(view: "/errors/notFound")
  "500"(view: "/errors/serverError")
}
```

Declarative Error Handling

In addition you can configure handlers for individual exceptions:

```
static mappings = {
  "403"(view: "/errors/forbidden")
  "404"(view: "/errors/notFound")
  "500"(controller: "errors", action: "illegalArgument",
    exception: IllegalArgumentException)
  "500"(controller: "errors", action: "nullPointer",
    exception: NullPointerException)
  "500"(controller: "errors", action: "customException",
    exception: MyException)
  "500"(view: "/errors/serverError")
}
```

With this configuration, an `IllegalArgumentException` will be handled by the `illegalArgument` action, a `NullPointerException` will be handled by the `nullPointer` action, and a `MyException` will be handled by the `customException` action. Other exceptions will be handled by the catch-all rule and use the `/errors/serverError` view.

You can access the exception from your custom error handling view or controller action using the request's `exception` property.

```
class ErrorController {
  def handleError() {
    def exception = request.exception
    // perform desired processing to handle the exception
  }
}
```



If your error-handling controller action throws an exception as well, you'll `StackOverflowException`.

6.4.5 Mapping to HTTP methods

URL mappings can also be configured to map based on the HTTP method (GET, POST, PUT or DELETE) and for restricting mappings based on HTTP method.

As an example the following mappings provide a RESTful API URL mappings for the `ProductContro`

```
static mappings = {
  "/product/$id"(controller: "product") {
    action = [GET: "show", PUT: "update", DELETE: "delete", POST: "save"]
  }
}
```

6.4.6 Mapping Wildcards

Grails' URL mappings mechanism also supports wildcard mappings. For example consider the following m

```
static mappings = {
  "/images/*.jpg"(controller: "image")
}
```

This mapping will match all paths to images such as `/image/logo.jpg`. Of course you can achieve the

```
static mappings = {
  "/images/$name.jpg"(controller: "image")
}
```

However, you can also use double wildcards to match more than one level below:

```
static mappings = {
  "/images/**/*.jpg"(controller: "image")
}
```

In this cases the mapping will match `/image/logo.jpg` as well as `/image/other/logo.jpg`. Evn variable:

```
static mappings = {
    // will match /image/logo.jpg and /image/other/logo.jpg
    "/images/$name*.jpg"(controller: "image")
}
```

In this case it will store the path matched by the wildcard inside a name parameter obtainable from the [params](#) object.

```
def name = params.name
println name // prints "logo" or "other/logo"
```

If you use wildcard URL mappings then you may want to exclude certain URIs from Grails' URL mappings. This is done through an `excludes` setting inside the `UrlMappings.groovy` class:

```
class UrlMappings {
    static excludes = ["/images/*", "/css/*"]
    static mappings = {
        ...
    }
}
```

In this case Grails won't attempt to match any URIs that start with `/images` or `/css`.

6.4.7 Automatic Link Re-Writing

Another great feature of URL mappings is that they automatically customize the behaviour of the [link](#) tag to require you to go and change all of your links.

This is done through a URL re-writing technique that reverse engineers the links from the URL mappings defined in one from an earlier section:

```
static mappings = {
    "/$blog/$year?/$month?/$day?/$id?"(controller:"blog", action:"show")
}
```

If you use the link tag as follows:

```
<g:link controller="blog" action="show"
        params="[blog:'fred', year:2007]">
    My Blog
</g:link>

<g:link controller="blog" action="show"
        params="[blog:'fred', year:2007, month:10]">
    My Blog - October 2007 Posts
</g:link>
```

Grails will automatically re-write the URL in the correct format:

```
<a href="/fred/2007">My Blog</a>
<a href="/fred/2007/10">My Blog - October 2007 Posts</a>
```

6.4.8 Applying Constraints

URL Mappings also support Grails' unified [validation constraints](#) mechanism, which lets you further "c" example, if we revisit the blog sample code from earlier, the mapping currently looks like this:

```
static mappings = {
    "/$blog/$year?/$month?/$day?/$id?"(controller:"blog", action:"show")
}
```

This allows URLs such as:

```
/graemerocher/2007/01/10/my_funky_blog_entry
```

However, it would also allow:

```
/graemerocher/not_a_year/not_a_month/not_a_day/my_funky_blog_entry
```

This is problematic as it forces you to do some clever parsing in the controller code. Luckily, URL Mappings also support validation constraints to validate the URL tokens:

```
"/$blog/$year?/$month?/$day?/$id?" {
    controller = "blog"
    action = "show"
    constraints {
        year(matches:/\d{4}/)
        month(matches:/\d{2}/)
        day(matches:/\d{2}/)
    }
}
```

In this case the constraints ensure that the year, month and day parameters match a particular valid pattern. We'll look at this in more detail later on.

6.4.9 Named URL Mappings

URL Mappings also support named mappings, that is are mappings which have a name associated with them. This is useful for generating links when the specific mapping is not known at the time the links are generated.

The syntax for defining a named mapping is as follows:

```
static mappings = {
    name <mapping name>: <url pattern> {
        // ...
    }
}
```

For example:

```
static mappings = {
    name personList: "/showPeople" {
        controller = 'person'
        action = 'list'
    }
    name accountDetails: "/details/$acctNumber" {
        controller = 'product'
        action = 'accountDetails'
    }
}
```

The mapping may be referenced in a link tag in a GSP.

```
<g:link mapping="personList">List People</g:link>
```

That would result in:

```
<a href="/showPeople">List People</a>
```

Parameters may be specified using the params attribute.

```
<g:link mapping="accountDetails" params="[acctNumber:'8675309']">
    Show Account
</g:link>
```

That would result in:

```
<a href="/details/8675309">Show Account</a>
```

Alternatively you may reference a named mapping using the link namespace.

```
<link:personList>List People</link:personList>
```

That would result in:

```
<a href="/showPeople">List People</a>
```

The link namespace approach allows parameters to be specified as attributes.

```
<link:accountDetails acctNumber="8675309">Show Account</link:accountDetails>
```

That would result in:

```
<a href="/details/8675309">Show Account</a>
```

To specify attributes that should be applied to the generated href, specify a Map value to the `attrs` attribute directly to the href, not passed through to be used as request parameters.

```
<link:accountDetails attrs="[class: 'fancy']" acctNumber="8675309">  
  Show Account  
</link:accountDetails>
```

That would result in:

```
<a href="/details/8675309" class="fancy">Show Account</a>
```

6.4.10 Customizing URL Formats

The default URL Mapping mechanism supports camel case names in the URLs. The default URL for `acc` in a controller named `MathHelperController` would be something like `/mathHelper/a` customization of this pattern and provides an implementation which replaces the camel case convention with support URLs like `/math-helper/add-numbers`. To enable hyphenated URLs assign `grails.web.url.converter` property in `grails-app/conf/Config.groovy`.

```
// grails-app/conf/Config.groovy  
grails.web.url.converter = 'hyphenated'
```

Arbitrary strategies may be plugged in by providing a class which implements the [UrlConverter](#) interface the Spring application context with the bean name of `grails.web.UrlConverter.BEAN_NAME`. If that name, it will be used as the default converter and there is no need to assign a value to the `grails.web.url.converter` property.

```
// src/groovy/com/myapplication/MyUrlConverterImpl.groovy
package com.myapplication

class MyUrlConverterImpl implements grails.web.UrlConverter {
    String toUrlElement(String propertyOrClassName) {
        // return some representation of a property or class name that should be
    }
}
```

```
// grails-app/conf/spring/resources.groovy
beans = {
    "${grails.web.UrlConverter.BEAN_NAME}"(com.myapplication.MyUrlConverterImpl)
}
```

6.5 Web Flow

Overview

Grails supports the creation of web flows built on the [Spring Web Flow](#) project. A web flow is a converter that retains state for the scope of the flow. A web flow also has a defined start and end state.

Web flows don't require an HTTP session, but instead store their state in a serialized form, which is then rendered by Grails and passed around as a request parameter. This makes flows far more scalable than other forms of state and its inherent memory and clustering concerns.

Web flow is essentially an advanced state machine that manages the "flow" of execution from one state to the next. You, the user, you don't have to be concerned with ensuring that users enter an action in the middle of some multi-step flow. This makes web flow perfect for use cases such as shopping carts, hotel booking and any application that requires a sequence of actions.



From Grails 1.2 onwards Webflow is no longer in Grails core, so you must install the Webflow plugin. Feature: `grails install-plugin webflow`

Creating a Flow

To create a flow create a regular Grails controller and add an action that ends with the convention `Flow`. For example:

```
class BookController {
    def index() {
        redirect(action: "shoppingCart")
    }
    def shoppingCartFlow = {
        ...
    }
}
```

Notice when redirecting or referring to the flow as an action we omit the `Flow` suffix. In other words the action is `shoppingCart`, not `shoppingCartFlow`.

6.5.1 Start and End States

As mentioned before a flow has a defined start and end state. A start state is the state which is entered w (or flow). The start state of a Grails flow is the first method call that takes a block. For example:

```
class BookController {
  ...
  def shoppingCartFlow = {
    showCart {
      on("checkout").to "enterPersonalDetails"
      on("continueShopping").to "displayCatalogue"
    }
    ...
    displayCatalogue {
      redirect(controller: "catalogue", action: "show")
    }
    displayInvoice()
  }
}
```

Here the showCart node is the start state of the flow. Since the showCart state doesn't define an action (that, by convention, refers to the view `grails-app/views/book/shoppingCart/showCart.gsp`).

Notice that unlike regular controller actions, the views are stored within a directory that `grails-app/views/book/shoppingCart`.

The shoppingCart flow also has two possible end states. The first is displayCatalogue which p controller and action, thus exiting the flow. The second is displayInvoice which is an end state as render a view called `grails-app/views/book/shoppingCart/displayInvoice.gsp` while

Once a flow has ended it can only be resumed from the start state, in this case showCart, and not from a

6.5.2 Action States and View States

View states

A view state is a one that doesn't define an action or a redirect. So for example this is a view state:

```
enterPersonalDetails {
  on("submit").to "enterShipping"
  on("return").to "showCart"
}
```

It will look for a view called `grails-app/views/book/shoppingCart/enterPersonalDet` enterPersonalDetails state defines two events: submit and return. The view is responsible render method to change the view to be rendered:

```
enterPersonalDetails {
  render(view: "enterDetailsView")
  on("submit").to "enterShipping"
  on("return").to "showCart"
}
```

Now it will look for `grails-app/views/book/shoppingCart/enterDetailsView.gsp`. Since it's not found, it will look for a shared view:

```
enterPersonalDetails {
    render(view: "/shared/enterDetailsView")
    on("submit").to "enterShipping"
    on("return").to "showCart"
}
```

Now it will look for `grails-app/views/shared/enterDetailsView.gsp`

Action States

An action state is a state that executes code but does not render a view. The result of the action is used to determine the next state. You define an action to be executed. This is done by calling the `action` method and passing

```
listBooks {
    action {
        [bookList: Book.list()]
    }
    on("success").to "showCatalogue"
    on(Exception).to "handleError"
}
```

As you can see an action looks very similar to a controller action and in fact you can reuse controller actions. When an action successfully returns with no errors the `success` event will be triggered. In this case since we return a list of books and is automatically placed in [flow scope](#).

In addition, in the above example we also use an exception handler to deal with errors on the line:

```
on(Exception).to "handleError"
```

This makes the flow transition to a state called `handleError` in the case of an exception.

You can write more complex actions that interact with the flow request context:

```
processPurchaseOrder {
    action {
        def a = flow.address
        def p = flow.person
        def pd = flow.paymentDetails
        def cartItems = flow.cartItems
        flow.clear()

        def o = new Order(person: p, shippingAddress: a, paymentDetails: pd)
        o.invoiceNumber = new Random().nextInt(9999999)
        for (item in cartItems) { o.addToItems item }
        o.save()
        [order: o]
    }
    on("error").to "confirmPurchase"
    on(Exception).to "confirmPurchase"
    on("success").to "displayInvoice"
}
```

Here is a more complex action that gathers all the information accumulated from the flow scope and creates order as the model. The important thing to note here is the interaction with the request context and "flow scope".

Transition Actions

Another form of action is what is known as a *transition* action. A transition action is executed directly prior to a transition being triggered. A simple example of a transition action can be seen below:

```
enterPersonalDetails {  
  on("submit") {  
    log.trace "Going to enter shipping"  
  }.to "enterShipping"  
  on("return").to "showCart"  
}
```

Notice how we pass a block of the code to the submit event that simply logs the transition. Transition state [validation](#), which is covered in a later section.

6.5.3 Flow Execution Events

In order to *transition* execution of a flow from one state to the next you need some way of triggering an event. Events can be triggered from either view states or action states.

Triggering Events from a View State

As discussed previously the start state of the flow in a previous code listing deals with two possible events: continueShopping and checkout.

```
def shoppingCartFlow = {  
  showCart {  
    on("checkout").to "enterPersonalDetails"  
    on("continueShopping").to "displayCatalogue"  
  }  
  ...  
}
```

Since the showCart event is a view state it will render the view `grails-app/book/shoppingCart/showCart.gsp`. In this view you need to have components that trigger flow execution. On a form this can be done using the [submitButton](#) tag.

```
<g:form action="shoppingCart">  
  <g:submitButton name="continueShopping" value="Continue Shopping" />  
  <g:submitButton name="checkout" value="Checkout" />  
</g:form>
```

The form must submit back to the shoppingCart flow. The name attribute of each [submitButton](#) tag is the name of the event. If you don't have a form you can also trigger an event with the [link](#) tag as follows:

```
<g:link action="shoppingCart" event="checkout" />
```

Triggering Events from an Action

To trigger an event from an action you invoke a method. For example there is the built in `error()` at below triggers the `error()` event on validation failure in a transition action:

```
enterPersonalDetails {
  on("submit") {
    def p = new Person(params)
    flow.person = p
    if (!p.validate()) return error()
  }.to "enterShipping"
  on("return").to "showCart"
}
```

In this case because of the error the transition action will make the flow go back to the `enterPersonal`

With an action state you can also trigger events to redirect flow:

```
shippingNeeded {
  action {
    if (params.shippingRequired) yes()
    else no()
  }
  on("yes").to "enterShipping"
  on("no").to "enterPayment"
}
```

6.5.4 Flow Scopes

Scope Basics

You'll notice from previous examples that we used a special object called `flow` to store objects within different scopes you can utilize:

- `request` - Stores an object for the scope of the current request
- `flash` - Stores the object for the current and next request only
- `flow` - Stores objects for the scope of the flow, removing them when the flow reaches an end state
- `conversation` - Stores objects for the scope of the conversation including the root flow and nested
- `session` - Stores objects in the user's session



Grails service classes can be automatically scoped to a web flow scope. See the documentation information.

Returning a model Map from an action will automatically result in the model being placed in flow scope. you can place objects within `flow` scope as follows:

```
enterPersonalDetails {
  on("submit") {
    [person: new Person(params)]
  }.to "enterShipping"
  on("return").to "showCart"
}
```

Be aware that a new request is always created for each state, so an object placed in request scope is not available in a subsequent view state. Use one of the other scopes to pass objects from one state to another.

1. Moves objects from flash scope to request scope upon transition between states;
2. Merges objects from the flow and conversation scopes into the view model before rendering (so you can reference these objects within a view, e.g. GSP pages).

Flow Scopes and Serialization

When placing objects in flash, flow or conversation scope they must implement `java.io.Serializable`. This has an impact on [domain classes](#) in that domain classes are typically placed within a scope. For example consider the following domain class:

```
class Book {
  String title
}
```

To place an instance of the Book class in a flow scope you will need to modify it as follows:

```
class Book implements Serializable {
  String title
}
```

This also impacts associations and closures you declare within a domain class. For example consider this:

```
class Book implements Serializable {
  String title
  Author author
}
```

Here if the Author association is not Serializable you will also get an error. This also impacts closures like `onLoad`, `onSave` and so on. The following domain class will cause an error if an instance is placed in a flow scope:

```
class Book implements Serializable {
  String title
  def onLoad = {
    println "I'm loading"
  }
}
```

The reason is that the assigned block on the `onLoad` event cannot be serialized. To get around this you sh :

```
class Book implements Serializable {  
    String title  
    transient onLoad = {  
        println "I'm loading"  
    }  
}
```

or as methods:

```
class Book implements Serializable {  
    String title  
    def onLoad() {  
        println "I'm loading"  
    }  
}
```

6.5.5 Data Binding and Validation

In the section on [start and end states](#), the start state in the first example triggered a transition to the `enter` renders a view and waits for the user to enter the required information:

```
enterPersonalDetails {  
    on("submit").to "enterShipping"  
    on("return").to "showCart"  
}
```

The view contains a form with two submit buttons that either trigger the submit event or the return event:

```
<g:form action="shoppingCart">  
    <!-- Other fields -->  
    <g:submitButton name="submit" value="Continue"></g:submitButton>  
    <g:submitButton name="return" value="Back"></g:submitButton>  
</g:form>
```

However, what about the capturing the information submitted by the form? To capture the form info we

```
enterPersonalDetails {  
    on("submit") {  
        flow.person = new Person(params)  
        !flow.person.validate() ? error() : success()  
    }.to "enterShipping"  
    on("return").to "showCart"  
}
```

Notice how we perform data binding from request parameters and place the `Person` instance within `enterPersonalDetails` to perform [validation](#) and invoke the `error()` method if validation fails. This signals to the flow that the `enterPersonalDetails` view so valid entries can be entered by the user, otherwise the transition goes to the `enterShipping` state.

Like regular actions, flow actions also support the notion of [Command Objects](#) by defining the first argument

```
enterPersonalDetails {
  on("submit") { PersonDetailsCommand cmd ->
    flow.personDetails = cmd
    !flow.personDetails.validate() ? error() : success()
  }.to "enterShipping"
  on("return").to "showCart"
}
```

6.5.6 Subflows and Conversations

Grails' Web Flow integration also supports subflows. A subflow is like a flow within a flow. For example

```
def searchFlow = {
  displaySearchForm {
    on("submit").to "executeSearch"
  }
  executeSearch {
    action {
      [results:searchService.executeSearch(params.q)]
    }
    on("success").to "displayResults"
    on("error").to "displaySearchForm"
  }
  displayResults {
    on("searchDeeper").to "extendedSearch"
    on("searchAgain").to "displaySearchForm"
  }
  extendedSearch {
    // Extended search subflow
    subflow(controller: "searchExtensions", action: "extendedSearch")
    on("moreResults").to "displayMoreResults"
    on("noResults").to "displayNoMoreResults"
  }
  displayMoreResults()
  displayNoMoreResults()
}
```

It references a subflow in the `extendedSearch` state. The controller parameter is optional if the subflow is in the same controller as the calling flow.



Prior to 1.3.5, the previous subflow call would look like `subflow(extendedSearch)`. This is deprecated and only supported for backward compatibility. The requirement that the name of the subflow state be the same as the called subflow (minus `Flow`) is deprecated and only supported for backward compatibility.

The subflow is another flow entirely:

```

def extendedSearchFlow = {
  startExtendedSearch {
    on("findMore").to "searchMore"
    on("searchAgain").to "noResults"
  }
  searchMore {
    action {
      def results = searchService.deepSearch(ctx.conversation.query)
      if (!results) return error()
      conversation.extendedResults = results
    }
    on("success").to "moreResults"
    on("error").to "noResults"
  }
  moreResults()
  noResults()
}

```

Notice how it places the `extendedResults` in conversation scope. This scope differs to flow scope whole conversation not just the flow. Also notice that the end state (either `moreResults` or `noResults`) leads back to the main flow:

```

extendedSearch {
  // Extended search subflow
  subflow(controller: "searchExtensions", action: "extendedSearch")
  on("moreResults").to "displayMoreResults"
  on("noResults").to "displayNoMoreResults"
}

```

6.6 Filters

Although Grails [controllers](#) support fine grained interceptors, these are only really useful when applied to manage with larger applications. Filters on the other hand can be applied across a whole group of controllers or actions. Filters are far easier to plugin and maintain completely separately to your main controller logic and concerns such as security, logging, and so on.

6.6.1 Applying Filters

To create a filter create a class that ends with the convention `Filters` in the `grails-app/conf` directory block called `filters` that contains the filter definitions:

```

class ExampleFilters {
  def filters = {
    // your filters here
  }
}

```

Each filter you define within the `filters` block has a name and a scope. The name is the method name and the scope is the arguments. For example to define a filter that applies to all controllers and all actions you can use wildcard

```

sampleFilter(controller: '*', action: '*') {
  // interceptor definitions
}

```


The scope of the filter can be one of the following things:

- A controller and/or action name pairing with optional wildcards
- A URI, with Ant path matching syntax

Filter rule attributes:

- `controller` - controller matching pattern, by default `*` is replaced with `.*` and a regex is compiled
- `controllerExclude` - controller exclusion pattern, by default `*` is replaced with `.*` and a regex is compiled
- `action` - action matching pattern, by default `*` is replaced with `.*` and a regex is compiled
- `actionExclude` - action exclusion pattern, by default `*` is replaced with `.*` and a regex is compiled
- `regex (true/false)` - use regex syntax (don't replace `*` with `.*`)
- `uri` - a uri to match, expressed with as Ant style path (e.g. `/book/**`)
- `uriExclude` - a uri pattern to exclude, expressed with as Ant style path (e.g. `/book/**`)
- `find (true/false)` - rule matches with partial match (see `java.util.regex.Matcher.find`)
- `invert (true/false)` - invert the rule (NOT rule)

Some examples of filters include:

- All controllers and actions

```
all(controller: '*', action: '*') {  
}
```

- Only for the BookController

```
justBook(controller: 'book', action: '*') {  
}
```

- All controllers except the BookController

```
notBook(controller: 'book', invert: true) {  
}
```

- All actions containing 'save' in the action name

```
saveInActionName(action: '*save*', find: true) {  
}
```

- All actions starting with the letter 'b' except for actions beginning with the phrase 'bad*'

```
actionBeginningWithBButNotBad(action: 'b*', actionExclude: 'bad*', find: true) {
}
```

- Applied to a URI space

```
someURIs(uri: '/book/**') {
}
```

- Applied to all URIs

```
allURIs(uri: '/*') {
}
```

In addition, the order in which you define the filters within the `filters` code block dictates the order in which the filters are executed. In addition to the order of execution between `Filters` classes, you can use the `dependsOn` property discussed in [filter dependencies](#).



Note: When exclude patterns are used they take precedence over the matching patterns. For example, if `actionExclude` is 'bad*' then actions like 'best' and 'bien' will have that filter applied but 'badlands' will not.

6.6.2 Filter Types

Within the body of the filter you can then define one or several of the following interceptor types for the filter:

- `before` - Executed before the action. Return `false` to indicate that the response has been handled and the action should not execute.
- `after` - Executed after an action. Takes a first argument as the view model to allow modification of the view model.
- `afterView` - Executed after view rendering. Takes an `Exception` as an argument which will be rendered to the user. Note: this Closure is called before the layout is applied.

For example to fulfill the common simplistic authentication use case you could define a filter as follows:

```

class SecurityFilters {
  def filters = {
    loginCheck(controller: '*', action: '*') {
      before = {
        if (!session.user && !actionName.equals('login')) {
          redirect(action: 'login')
          return false
        }
      }
    }
  }
}

```

Here the `loginCheck` filter uses a `before` interceptor to execute a block of code that checks if a user the login action. Note how returning false ensure that the action itself is not executed.

6.6.3 Variables and Scopes

Filters support all the common properties available to [controllers](#) and [tag libraries](#), plus the application con

- [request](#) - The `HttpServletRequest` object
- [response](#) - The `HttpServletResponse` object
- [session](#) - The `HttpSession` object
- [servletContext](#) - The `ServletContext` object
- [flash](#) - The flash object
- [params](#) - The request parameters object
- [actionName](#) - The action name that is being dispatched to
- [controllerName](#) - The controller name that is being dispatched to
- [grailsApplication](#) - The Grails application currently running
- [applicationContext](#) - The `ApplicationContext` object

However, filters only support a subset of the methods available to controllers and tag libraries. These inclu

- [redirect](#) - For redirects to other controllers and actions
- [render](#) - For rendering custom responses

6.6.4 Filter Dependencies

In a `Filters` class, you can specify any other `Filters` classes that should first be executed using `tl` when a `Filters` class depends on the behavior of another `Filters` class (e.g. setting up the environme and is defined as an array of `Filters` classes.

Take the following example `Filters` classes:

```

class MyFilters {
    def dependsOn = [MyOtherFilters]

    def filters = {
        checkAwesome(uri: "/*") {
            before = {
                if (request.isAwesome) { // do something awesome }
            }
        }

        checkAwesome2(uri: "/*") {
            before = {
                if (request.isAwesome) { // do something else awesome }
            }
        }
    }
}

```

```

class MyOtherFilters {
    def filters = {
        makeAwesome(uri: "/*") {
            before = {
                request.isAwesome = true
            }
        }

        doNothing(uri: "/*") {
            before = {
                // do nothing
            }
        }
    }
}

```

MyFilters specifically dependsOn MyOtherFilters. This will cause all the filters in MyOtherFilters who be executed before those in MyFilters. For a request of "/test", which will match the scope of every filter would be as follows:

- MyOtherFilters - makeAwesome
- MyOtherFilters - doNothing
- MyFilters - checkAwesome
- MyFilters - checkAwesome2

The filters within the MyOtherFilters class are processed in order first, followed by the filters in the MyFilters classes are enabled and the execution order of filters within each Filters class are preserved

If any cyclical dependencies are detected, the filters with cyclical dependencies will be added to the end of the list. Information about any cyclical dependencies that are detected will be written to the logs. Ensure at least `WARN` or configure an appender for the `org.codehaus.groovy.grails.plugins.web.filters.FiltersGrailsPlugin` when

6.7 Ajax

Ajax is the driving force behind the shift to richer web applications. These types of applications in general frameworks written in languages like [Groovy](#) and [Ruby](#) Grails provides support for building Ajax applications. For a full list of these see the Tag Library Reference.

6.7.1 Ajax Support

By default Grails ships with the [jQuery](#) library, but through the [Plugin system](#) provides support for other frameworks like [Google Web Toolkit](#).

This section covers Grails' support for Ajax in general. To get started, add this line to the <head> tag of your layout:

```
<g:javascript library="jquery" />
```

You can replace jQuery with any other library supplied by a plugin you have installed. This works because Grails has a plugin system that supports a number of different Ajax libraries including:

- jQuery
- Prototype
- Dojo
- YUI
- MooTools

6.7.1.1 Remoting Linking

Remote content can be loaded in a number of ways, the most common way is through the [remoteLink](#) tag, which is an anchor tag that performs an asynchronous request and optionally sets the response in an element. The syntax is as follows:

```
<g:remoteLink action="delete" id="1">Delete Book</g:remoteLink>
```

The above link sends an asynchronous request to the `delete` action of the current controller with an id of 1.

6.7.1.2 Updating Content

This is great, but usually you provide feedback to the user about what happened:

```
def delete() {  
    def b = Book.get(params.id)  
    b.delete()  
    render "Book ${b.id} was deleted"  
}
```

GSP code:

```
<div id="message"></div>  
<g:remoteLink action="delete" id="1" update="message">  
Delete Book  
</g:remoteLink>
```

The above example will call the action and set the contents of the message div to the response in this case. This is done by the `update` attribute on the tag, which can also take a Map to indicate what should be updated

```
<div id="message"></div>
<div id="error"></div>
<g:remoteLink update="[success: 'message', failure: 'error']"
               action="delete" id="1">
Delete Book
</g:remoteLink>
```

Here the error div will be updated if the request failed.

6.7.1.3 Remote Form Submission

An HTML form can also be submitted asynchronously in one of two ways. Firstly using the [formRemote](#) tag, those for the [remoteLink](#) tag:

```
<g:formRemote url="[controller: 'book', action: 'delete']"
              update="[success: 'message', failure: 'error']">
  <input type="hidden" name="id" value="1" />
  <input type="submit" value="Delete Book!" />
</g:formRemote >
```

Or alternatively you can use the [submitToRemote](#) tag to create a submit button. This allows some but depending on the action:

```
<form action="delete">
  <input type="hidden" name="id" value="1" />
  <g:submitToRemote action="delete"
                    update="[success: 'message', failure: 'error']" />
</form>
```

6.7.1.4 Ajax Events

Specific JavaScript can be called if certain events occur, all the events start with the "on" prefix and are appropriate, or take other action:

```
<g:remoteLink action="show"
               id="1"
               update="success"
               onLoading="showProgress()"
               onComplete="hideProgress()">Show Book 1</g:remoteLink>
```

The above code will execute the "showProgress()" function which may show a progress bar or whatever is

- `onSuccess` - The JavaScript function to call if successful
- `onFailure` - The JavaScript function to call if the call failed
- `on_ERROR_CODE` - The JavaScript function to call to handle specified error codes (eg `on404="alert"`)
- `onUninitialized` - The JavaScript function to call if the Ajax engine failed to initialise
- `onLoading` - The JavaScript function to call when the remote function is loading the response
- `onLoaded` - The JavaScript function to call when the remote function is completed loading the response
- `onComplete` - The JavaScript function to call when the remote function is complete, including any

If you need a reference to the `XmlHttpRequest` object you can use the implicit event parameter `e` to obtain it.

```
<g:javascript>
  function fireMe(e) {
    alert("XmlHttpRequest = " + e)
  }
</g:javascript>
<g:remoteLink action="example"
  update="success"
  onSuccess="fireMe(e)">Ajax Link</g:remoteLink>
```

6.7.2 Ajax with Prototype

Grails features an external plugin to add [Prototype](#) support to Grails. To install the plugin type the following in a terminal window:

```
grails install-plugin prototype
```

This will download the current supported version of the Prototype plugin and install it into your Grails project. Following reference to the top of your page:

```
<g:javascript library="prototype" />
```

If you require [Scriptaculous](#) too you can do the following instead:

```
<g:javascript library="scriptaculous" />
```

Now all of Grails tags such as [remoteLink](#), [formRemote](#) and [submitToRemote](#) work with Prototype remote functions.

6.7.3 Ajax with Dojo

Grails features an external plugin to add [Dojo](#) support to Grails. To install the plugin type the following in a terminal window:

```
grails install-plugin dojo
```

This will download the current supported version of Dojo and install it into your Grails project. With reference to the top of your page:

```
<g:javascript library="dojo" />
```

Now all of Grails tags such as [remoteLink](#), [formRemote](#) and [submitToRemote](#) work with Dojo remoting.

6.7.4 Ajax with GWT

Grails also features support for the [Google Web Toolkit](#) through a plugin. There is comprehensive [docume](#)

6.7.5 Ajax on the Server

There are a number of different ways to implement Ajax which are typically broken down into:

- Content Centric Ajax - Where you just use the HTML result of a remote call to update the page
- Data Centric Ajax - Where you actually send an XML or JSON response from the server and program
- Script Centric Ajax - Where the server sends down a stream of JavaScript to be evaluated on the fly

Most of the examples in the [Ajax](#) section cover Content Centric Ajax where you are updating the page Centric or Script Centric. This guide covers the different styles of Ajax.

Content Centric Ajax

Just to re-cap, content centric Ajax involves sending some HTML back from the server and is typically implemented using the [render](#) method:

```
def showBook() {  
    def b = Book.get(params.id)  
  
    render(template: "bookTemplate", model: [book: b])  
}
```

Calling this on the client involves using the [remoteLink](#) tag:

```
<g:remoteLink action="showBook" id="${book.id}"  
              update="book${book.id}">Update Book</g:remoteLink>  
  
<div id="book${book.id}">  
    <!--existing book mark-up -->  
</div>
```

Data Centric Ajax with JSON

Data Centric Ajax typically involves evaluating the response on the client and updating programmatically would typically use Grails' [JSON marshalling](#) capability:

```
import grails.converters.JSON

def showBook() {
    def b = Book.get(params.id)

    render b as JSON
}
```

And then on the client parse the incoming JSON request using an Ajax event handler:

```
<g:javascript>
function updateBook(e) {
    var book = eval("(" + e.responseText + ")") // evaluate the JSON
    $("book" + book.id + "_title").innerHTML = book.title
}
</g:javascript>
<g:remoteLink action="test" update="foo" onSuccess="updateBook(e)">
    Update Book
</g:remoteLink>
<g:set var="bookId">book${book.id}</g:set>
<div id="${bookId}">
    <div id="${bookId}_title">The Stand</div>
</div>
```

Data Centric Ajax with XML

On the server side using XML is equally simple:

```
import grails.converters.XML

def showBook() {
    def b = Book.get(params.id)

    render b as XML
}
```

However, since DOM is involved the client gets more complicated:

```
<g:javascript>
function updateBook(e) {
    var xml = e.responseXML
    var id = xml.getElementsByTagName("book").getAttribute("id")
    $("book" + id + "_title") = xml.getElementsByTagName("title")[0].textContent
}
</g:javascript>
<g:remoteLink action="test" update="foo" onSuccess="updateBook(e)">
    Update Book
</g:remoteLink>
<g:set var="bookId">book${book.id}</g:set>
<div id="${bookId}">
    <div id="${bookId}_title">The Stand</div>
</div>
```

Script Centric Ajax with JavaScript

Script centric Ajax involves actually sending JavaScript back that gets evaluated on the client. An example

```
def showBook() {
  def b = Book.get(params.id)

  response.contentType = "text/javascript"
  String title = b.title.encodeAsJavaScript()
  render "$('book${b.id}_title')='${title}'"
}
```

The important thing to remember is to set the `contentType` to `text/javascript`. If you use JavaScript will automatically be evaluated due to this `contentType` setting.

Obviously in this case it is critical that you have an agreed client-side API as you don't want changes on the of the reasons Rails has something like RJS. Although Grails does not currently have a feature such as RJ that offers similar capabilities.

Responding to both Ajax and non-Ajax requests

It's straightforward to have the same Grails controller action handle both Ajax and non-Ajax requests. `HttpServletRequest` which can be used to identify Ajax requests. For example you could render a p requests or the full page for regular HTTP requests:

```
def listBooks() {
  def books = Book.list(params)
  if (request.xhr) {
    render template: "bookTable", model: [books: books]
  } else {
    render view: "list", model: [books: books]
  }
}
```

6.8 Content Negotiation

Grails has built in support for [Content negotiation](#) using either the HTTP Accept header, an explicit for of a mapped URI.

Configuring Mime Types

Before you can start dealing with content negotiation you need to tell Grails what content types you wi: configured with a number of different content types within `grails-app/conf/Config.groov` setting:

```
grails.mime.types = [ xml: ['text/xml', 'application/xml'],
                      text: 'text-plain',
                      js: 'text/javascript',
                      rss: 'application/rss+xml',
                      atom: 'application/atom+xml',
                      css: 'text/css',
                      csv: 'text/csv',
                      all: '*/*',
                      json: 'text/json',
                      html: ['text/html', 'application/xhtml+xml']
                    ]
```

The above bit of configuration allows Grails to detect to format of a request containing either the 'text/x' simply 'xml'. You can add your own types by simply adding new entries into the map.

Content Negotiation using the Accept header

Every incoming HTTP request has a special [Accept](#) header that defines what media types (or mime types) this is typically:

```
*/*
```

Which simply means anything. However, on newer browser something all together more useful is sent such header):

```
text/xml, application/xml, application/xhtml+xml, text/html;q=0.9,
text/plain;q=0.8, image/png, */*;q=0.5
```

Grails parses this incoming format and adds a property to the [response](#) object that outlines the pre example the following assertion would pass:

```
assert 'html' == response.format
```

Why? The `text/html` media type has the highest "quality" rating of 0.9, therefore is the highest priority mentioned previously the result is slightly different:

```
assert 'all' == response.format
```

In this case 'all' possible formats are accepted by the client. To deal with different kinds of requests from the [render](#) method that acts as kind of a switch statement:

```
import grails.converters.XML

class BookController {
  def list() {
    def books = Book.list()
    withFormat {
      html bookList: books
      js { render "alert('hello')" }
      xml { render books as XML }
    }
  }
}
```

If the preferred format is `html` then Grails will execute the `html()` call only. This causes `Grails-app/views/books/list.html.gsp` or `Grails-app/views/books/list.gsp` will be invoked and an XML response rendered.

How do we handle the "all" format? Simply order the content-types within your `withFormat` block so comes first. So in the above example, "all" will trigger the `html` handler.



When using [withFormat](#) make sure it is the last call in your controller action as the `withFormat` method is used by the action to dictate what happens next.

Request format vs. Response format

As of Grails 2.0, there is a separate notion of the *request* format and the *response* format. The request format is used to detect if the incoming request can be parsed into XML or JSON, with the `format` parameter or `ACCEPT` header to attempt to deliver an appropriate response to the client.

The [withFormat](#) available on controllers deals specifically with the response format. If you wish to add support for the request format then you can do so using a separate `withFormat` method available on the request:

```
request.withFormat {
  xml {
    // read XML
  }
  json {
    // read JSON
  }
}
```

Content Negotiation with the format Request Parameter

If fiddling with request headers is not your favorite activity you can override the format used by specifying the `format` parameter in the URL:

```
/book/list?format=xml
```

You can also define this parameter in the [URL Mappings](#) definition:

```
"/book/list"(controller:"book", action:"list") {  
    format = "xml"  
}
```

Content Negotiation with URI Extensions

Grails also supports content negotiation using URI extensions. For example given the following URI:

```
/book/list.xml
```

Grails will remove the extension and map it to `/book/list` instead whilst simultaneously setting the extension. This behaviour is enabled by default, so if you wish to turn it off, you must set the `grails.mime.file.extensions` to `false`:

```
grails.mime.file.extensions = false
```

Testing Content Negotiation

To test content negotiation in a unit or integration test (see the section on [Testing](#)) you can either manipulate

```
void testJavascriptOutput() {  
    def controller = new TestController()  
    controller.request.addHeader "Accept",  
        "text/javascript, text/html, application/xml, text/xml, */*"  
    controller.testAction()  
    assertEquals "alert('hello')", controller.response.contentAsString  
}
```

Or you can set the format parameter to achieve a similar effect:

```
void testJavascriptOutput() {  
    def controller = new TestController()  
    controller.params.format = 'js'  
    controller.testAction()  
    assertEquals "alert('hello')", controller.response.contentAsString  
}
```

7 Validation

Grails validation capability is built on [Spring's Validator API](#) and data binding capabilities. However C provides a unified way to define validation "constraints" with its constraints mechanism.

Constraints in Grails are a way to declaratively specify validation rules. Most commonly they are applied to [Domain Objects](#) and [Command Objects](#) also support constraints.

7.1 Declaring Constraints

Within a domain class [constraints](#) are defined with the constraints property that is assigned a code block:

```
class User {
    String login
    String password
    String email
    Integer age

    static constraints = {
        // ...
    }
}
```

You then use method calls that match the property name for which the constraint applies in combination with the constraints:

```
class User {
    ...

    static constraints = {
        login size: 5..15, blank: false, unique: true
        password size: 5..15, blank: false
        email email: true, blank: false
        age min: 18
    }
}
```

In this example we've declared that the `login` property must be between 5 and 15 characters long, it can also be applied other constraints to the `password`, `email` and `age` properties.



By default, all domain class properties are not nullable (i.e. they have an implicit `nullable: false`). The same is not true for command object properties, which are nullable by default.

A complete reference for the available constraints can be found in the Quick Reference section under the C

A word of warning - referencing domain class properties from constraints

It's very easy to attempt to reference instance variables from the static constraints block, but this isn't legal. You will get a `MissingPropertyException` for your trouble. For example, you may try

```
class Response {
  Survey survey
  Answer answer

  static constraints = {
    survey blank: false
    answer blank: false, inList: survey.answers
  }
}
```

See how the `inList` constraint references the instance property `survey`? That won't work. Instead, use :

```
class Response {
  ...
  static constraints = {
    survey blank: false
    answer blank: false, validator: { val, obj -> val in obj.survey.answers }
  }
}
```

In this example, the `obj` argument to the custom validator is the domain *instance* that is being validated, and return a boolean to indicate whether the new value for the `answer` property, `val`, is valid.

7.2 Validating Constraints

Validation Basics

Call the [validate](#) method to validate a domain class instance:

```
def user = new User(params)

if (user.validate()) {
  // do something with user
}
else {
  user.errors.allErrors.each {
    println it
  }
}
```

The `errors` property on domain classes is an instance of the Spring [Errors](#) interface. The `Errors` interface provides methods to retrieve validation errors and also retrieve the original values.

Validation Phases

Within Grails there are two phases of validation, the first one being [data binding](#) which occurs when creating a new instance such as:

```
def user = new User(params)
```

At this point you may already have errors in the `errors` property due to type conversion (such as converting a string to an integer) and you can access these and obtain the original input value using the `Errors` API:

```
if (user.hasErrors()) {
    if (user.errors.hasFieldErrors("login")) {
        println user.errors.getFieldError("login").rejectedValue
    }
}
```

The second phase of validation happens when you call [validate](#) or [save](#). This is when Grails will validate the constraints you defined. For example, by default the [save](#) method calls `validate` before executing, allowing you to

```
if (user.save()) {
    return user
}
else {
    user.errors.allErrors.each {
        println it
    }
}
```

7.3 Validation on the Client

Displaying Errors

Typically if you get a validation error you redirect back to the view for rendering. Once there you need to render the errors. Grails supports a rich set of tags for dealing with errors. To render the errors as a list you can use [renderErrors](#):

```
<g:renderErrors bean="${user}" />
```

If you need more control you can use [hasErrors](#) and [eachError](#):

```
<g:hasErrors bean="${user}">
    <ul>
        <g:eachError var="err" bean="${user}">
            <li>${err}</li>
        </g:eachError>
    </ul>
</g:hasErrors>
```

Highlighting Errors

It is often useful to highlight using a red box or some indicator when a field has been incorrectly input. This can be done by invoking it as a method. For example:

```
<div class='value ${hasErrors(bean:user,field:'login','errors')}'>
    <input type="text" name="login" value="${fieldValue(bean:user,field:'login')}" />
</div>
```


This code checks if the `login` field of the user bean has any errors and if so it adds an errors CSS CSS rules to highlight the `div`.

Retrieving Input Values

Each error is actually an instance of the [FieldError](#) class in Spring, which retains the original input value the error object to restore the value input by the user using the [fieldValue](#) tag:

```
<input type="text" name="login" value="${fieldValue(bean:user,field:'login')}" />
```

This code will check for an existing `FieldError` in the `User` bean and if there is obtain the originally i

7.4 Validation and Internationalization

Another important thing to note about errors in Grails is that error messages are not hard coded anywhere. messages from message bundles using Grails' [i18n](#) support.

Constraints and Message Codes

The codes themselves are dictated by a convention. For example consider the constraints we looked at earl

```
package com.mycompany.myapp

class User {
    ...

    static constraints = {
        login size: 5..15, blank: false, unique: true
        password size: 5..15, blank: false
        email email: true, blank: false
        age min: 18
    }
}
```

If a constraint is violated Grails will by convention look for a message code of the form:

```
[Class Name].[Property Name].[Constraint Code]
```

In the case of the blank constraint this would be `user.login.blank` so you would need a `messages.properties` file:

```
user.login.blank=Your login name must be specified!
```

The class name is looked for both with and without a package, with the packaged version `com.mycompany.myapp.User.login.blank` will be used before `user.login.blank`. This allows for cases w/ clash with a plugin's.

For a reference on what codes are for which constraints refer to the reference guide for each constraint.

Displaying Messages

The [renderErrors](#) tag will automatically look up messages for you using the [message](#) tag. If you need more control, you can use this yourself:

```
<g:hasErrors bean="${user}">
  <ul>
    <g:eachError var="err" bean="${user}">
      <li><g:message error="${err}" /></li>
    </g:eachError>
  </ul>
</g:hasErrors>
```

In this example within the body of the [eachError](#) tag we use the [message](#) tag in combination with its `error` attribute to display the given error.

7.5 Validation Non Domain and Command Object Classes

[Domain classes](#) and [command objects](#) support validation by default. Other classes may be made validateable by defining a `constraints` property in the class (as described above) and then telling the framework about them. It is important to register the validateable classes with the framework. Simply defining the `constraints` property is not sufficient.

The Validateable Annotation

Classes which define the static `constraints` property and are annotated with `@Validateable` can be made validateable. Consider this example:

```
// src/groovy/com/mycompany/myapp/User.groovy
package com.mycompany.myapp

import grails.validation.Validateable

@Validateable
class User {
    ...

    static constraints = {
        login size: 5..15, blank: false, unique: true
        password size: 5..15, blank: false
        email email: true, blank: false
        age min: 18
    }
}
```

Registering Validateable Classes

If a class is not marked with `Validateable`, it may still be made validateable by registering it. The steps required to do this are to define the static `constraints` property in the class and then telling the framework about the class by assigning a value to the `constraints` property in `Config.groovy`:

```
grails.validateable.classes = [com.mycompany.myapp.User, com.mycompany.dto.Account]
```

8 The Service Layer

GrailsService. GrailsController

Grails,ServiceController

Service

Grails service,[create-service](#):

```
grails create-service helloworld.simple
```



create-servicepackageGrailspackage

serviceservicegrails-app/services/helloworld/SimpleService.groovy.GrailsService

```
package helloworld
class SimpleService {
}
```

8.1 Declarative Transactions

Services--[domain](#):service[withTransaction](#) Spring

GrailsServiceserviceservicetransactional false

```
class CountryService {
    static transactional = false
}
```

trueService



:newnew BookService()service.

Runtime Error[PROPAGATION REQUIRED](#).



Checked. GroovycheckeduncheckedSpring. checkedunchecked

GrailsSpring Transactional



TransactionalGrailsService

listBooksupdateBookdeleteBook

```
import org.springframework.transaction.annotation.Transactional

class BookService {
    @Transactional(readOnly = true)
    def listBooks() {
        Book.list()
    }
    @Transactional
    def updateBook() {
        // ...
    }
    def deleteBook() {
        // ...
    }
}
```

ServiceServiceService Transactional= TRUE

```
import org.springframework.transaction.annotation.Transactional

@Transactional
class BookService {
    def listBooks() {
        Book.list()
    }
    def updateBook() {
        // ...
    }
    def deleteBook() {
        // ...
    }
}
```

, listBooks:

```

import org.springframework.transaction.annotation.Transactional

@Transactional
class BookService {

  @Transactional(readOnly = true)
  def listBooks() {
    Book.list()
  }

  def updateBook() {
    // ...
  }

  def deleteBook() {
    // ...
  }
}

```

updateBook deleteBook

Spring[Using @Transactional](#).

GrailsSpringGrails Transactional

8.1.1 Transactions Rollback and the Session

Understanding Transactions and the Hibernate Session

When using transactions there are important considerations you must take into account with regards to how they are handled by Hibernate. When a transaction is rolled back the Hibernate session used by GORM is cleared. The session becomes detached and accessing uninitialized lazy-loaded collections will lead to LazyInitializationException.

To understand why it is important that the Hibernate session is cleared. Consider the following example:

```

class Author {
  String name
  Integer age

  static hasMany = [books: Book]
}

```

If you were to save two authors using consecutive transactions as follows:

```

Author.withTransaction { status ->
  new Author(name: "Stephen King", age: 40).save()
  status.setRollbackOnly()
}

Author.withTransaction { status ->
  new Author(name: "Stephen King", age: 40).save()
}

```

Only the second author would be saved since the first transaction rolls back the author save() by the end of the first transaction. If the Hibernate session were not cleared then both author instances would be persisted and it would lead to very strange behavior.

It can, however, be frustrating to get `LazyInitializationExceptions` due to the session being closed.

For example, consider the following example:

```
class AuthorService {
  void updateAge(id, int age) {
    def author = Author.get(id)
    author.age = age
    if (author.isTooOld()) {
      throw new AuthorException("too old", author)
    }
  }
}
```

```
class AuthorController {
  def authorService
  def updateAge() {
    try {
      authorService.updateAge(params.id, params.int("age"))
    }
    catch(e) {
      render "Author books ${e.author.books}"
    }
  }
}
```

In the above example the transaction will be rolled back if the Author's age exceeds the maximum value by throwing an `AuthorException`. The `AuthorException` references the author but when the transaction is rolled back a `LazyInitializationException` will be thrown because the underlying Hibernate session has been closed.

To solve this problem you have a number of options. One is to ensure you query eagerly to get the data you need.

```
class AuthorService {
  ...
  void updateAge(id, int age) {
    def author = Author.findById(id, [fetch:[books:"eager"]])
    ...
  }
}
```

In this example the books association will be queried when retrieving the Author.



This is the optimal solution as it requires fewer queries than the following suggested solutions.

Another solution is to redirect the request after a transaction rollback:

```

class AuthorController {
  AuthService authService
  def updateAge() {
    try {
      authService.updateAge(params.id, params.int("age"))
    }
    catch(e) {
      flash.message "Can't update age"
      redirect action:"show", id:params.id
    }
  }
}

```

In this case a new request will deal with retrieving the Author again. And, finally a third solution is to re-render the view to make sure the session remains in the correct state:

```

class AuthorController {
  def authService
  def updateAge() {
    try {
      authService.updateAge(params.id, params.int("age"))
    }
    catch(e) {
      def author = Author.read(params.id)
      render "Author books ${author.books}"
    }
  }
}

```

Validation Errors and Rollback

A common use case is to rollback a transaction if there are validation errors. For example consider this service:

```

import grails.validation.ValidationException

class AuthService {
  void updateAge(id, int age) {
    def author = Author.get(id)
    author.age = age
    if (!author.validate()) {
      throw new ValidationException("Author is not valid", author.errors)
    }
  }
}

```

To re-render the same view that a transaction was rolled back in you can re-associate the errors with a reference to the original author object:


```
import grails.validation.ValidationException

class AuthorController {
    def authorService
    def updateAge() {
        try {
            authorService.updateAge(params.id, params.int("age"))
        }
        catch (ValidationException e) {
            def author = Author.read(params.id)
            author.errors = e.errors
            render view: "edit", model: [author:author]
        }
    }
}
```

8.2 Scoped Services

By default, access to service methods is not synchronised, so nothing prevents concurrent execution of the service. If a service is a singleton and may be used concurrently, you should be very careful about storing state in a service. Never store state in a service.

You can change this behaviour by placing a service in a particular scope. The supported scopes are:

- `prototype` - A new service is created every time it is injected into another class
- `request` - A new service will be created per request
- `flash` - A new service will be created for the current and next request only
- `flow` - In web flows the service will exist for the scope of the flow
- `conversation` - In web flows the service will exist for the scope of the conversation. ie a root flow
- `session` - A service is created for the scope of a user session
- `singleton` (default) - Only one instance of the service ever exists



If your service is `flash`, `flow` or `conversation` scoped it must implement `java.io.Serializable`. It can only be used in the context of a [Web Flow](#).

To enable one of the scopes, add a static `scope` property to your class whose value is one of the above, for example:

```
static scope = "flow"
```

8.3 Dependency Injection and Services

Dependency Injection Basics

A key aspect of Grails services is the ability to use [Spring Framework](#)'s dependency injection features. Grails follows the "Spring convention". In other words, you can use the property name representation of the class name of a service, controllers, tag libraries, and so on.

As an example, given a service called `BookService`, if you define a property called `bookService` in

```
class BookController {
  def bookService
  ...
}
```

In this case, the Spring container will automatically inject an instance of that service based on its configuration done by name. You can also specify the type as follows:

```
class AuthorService {
  BookService bookService
}
```



NOTE: Normally the property name is generated by lower casing the first letter of the class name. An instance of the `BookService` class would map to a property named `bookService`.

To be consistent with standard JavaBean conventions, if the first 2 letters of the class name and the property name is the same as the class name. For example, the property name of the `JDBCHelper` class would be `JDBCHelperService`, not `jDBCHelperService` or `jdbchelperService`.

See section 8.8 of the JavaBean specification for more information on de-capitalization rules.

Dependency Injection and Services

You can inject services in other services with the same technique. If you had an `AuthorService` that declares the `AuthorService` as follows would allow that:

```
class AuthorService {
  def bookService
}
```

Dependency Injection and Domain Classes / Tag Libraries

You can even inject services into domain classes and tag libraries, which can aid in the development of rich

```
class Book {
  ...
  def bookService
  def buyBook() {
    bookService.buyBook(this)
  }
}
```

8.4 Using Services from Java

One of the powerful things about services is that since they encapsulate re-usable logic, you can use them in other classes. There are a couple of ways you can reuse a service from Java. The simplest way is to move your service to the `grails-app/services` directory. The reason this is important is that it is not possible to import a class from a package (the package used when no package declaration is present). So for example the `BookService` below can

```
class BookService {
    void buyBook(Book book) {
        // logic
    }
}
```

However, this can be rectified by placing this class in a package, by moving the class to `grails-app/services/bookstore` and then modifying the package declaration:

```
package bookstore

class BookService {
    void buyBook(Book book) {
        // logic
    }
}
```

An alternative to packages is to instead have an interface within a package that the service implements:

```
package bookstore

interface BookStore {
    void buyBook(Book book)
}
```

And then the service:

```
class BookService implements bookstore.BookStore {
    void buyBook(Book b) {
        // logic
    }
}
```

This latter technique is arguably cleaner, as the Java side only has a reference to the interface and not to the implementation (which is always a good idea to use packages). Either way, the goal of this exercise is to enable Java to statically recompile the code at compile time.

Now that this is done you can create a Java class within the `src/java` directory and add a setter that uses Spring:

```
// src/java/bookstore/BookConsumer.java
package bookstore;

public class BookConsumer {
    private BookStore store;

    public void setBookStore(BookStore storeInstance) {
        this.store = storeInstance;
    }
    ...
}
```

Once this is done you can configure the Java class as a Spring bean in `grails-app/conf/spring/resources.groovy` (for more information see the section on [Grails and Spring](#)):

```
<bean id="bookConsumer" class="bookstore.BookConsumer">
    <property name="bookStore" ref="bookService" />
</bean>
```

or in `grails-app/conf/spring/resources.groovy`:

```
import bookstore.BookConsumer

beans = {
    bookConsumer(BookConsumer) {
        bookStore = ref("bookService")
    }
}
```

9 Testing

Automated testing is a key part of Grails. Hence, Grails provides many ways to making testing easier for functional tests. This section details the different capabilities that Grails offers for testing.



Grails 1.3.x and below used the `grails.test.GrailsUnitTestCase` class hierarchy style. Grails 2.0.x and above deprecates these test harnesses in favour of mixins that can be used for different kinds of tests (JUnit 3, JUnit 4, Spock etc.) without subclassing

The first thing to be aware of is that all of the `create-*` and `generate-*` commands create unit tests. For example if you run the [create-controller](#) command as follows:

```
grails create-controller com.acme.app.simple
```

Grails will create a controller at `grails-app/controllers/com/acme/app/SimpleController` and a unit test at `test/unit/com/acme/app/SimpleControllerTests.groovy`. What Grails won't do however is create a `SimpleControllerTest` class. That is left up to you.



The default class name suffix is `Tests` but as of Grails 1.2.2, the suffix of `Test` is also supported.

Running Tests

Tests are run with the [test-app](#) command:

```
grails test-app
```

Note that you will be able to run unit tests much quicker if you use an IDE or if you use the "interactive mode" which doesn't require the need to stop the JVM:

```
grails
...
test-app
```

The `test-app` command will produce output such as:

```
-----
Running Unit Tests...
Running test FooTests...FAILURE
Unit Tests Completed in 464ms ...
-----

Tests failed: 0 errors, 1 failures
```

Whilst reports will have been written out the `target/test-reports` directory.



You can force a clean before running tests by passing `-clean` to the `test-app` command.

Targeting Tests

You can selectively target the test(s) to be run in different ways. To run all tests for a controller named `SimpleController`

```
grails test-app SimpleController
```

This will run any tests for the class named `SimpleController`. Wildcards can be used...

```
grails test-app *Controller
```

This will test all classes ending in `Controller`. Package names can optionally be specified...

```
grails test-app some.org.*Controller
```

or to run all tests in a package...

```
grails test-app some.org.*
```

or to run all tests in a package including subpackages...

```
grails test-app some.org.**.*
```

You can also target particular test methods...

```
grails test-app SimpleController.testLogin
```

This will run the `testLogin` test in the `SimpleController` tests. You can specify as many patterns :

```
grails test-app some.org.* SimpleController.testLogin BookController
```

Targeting Test Types and/or Phases

In addition to targeting certain tests, you can also target test *types* and/or *phases* by using the `phase:type`



Grails organises tests by phase and by type. A test phase relates to the state of the Grails application and the type relates to the testing mechanism.

Grails comes with support for 4 test phases (`unit`, `integration`, `functional` and `spock`) and 4 test types for the `unit` and `integration` phases. These test types have the same name as the phases.

Testing plugins may provide new test phases or new test types for existing phases. See the [Grails Testing documentation](#).

To execute the JUnit integration tests you can run:

```
grails test-app integration:integration
```

Both phase and type are optional. Their absence acts as a wildcard. The following command will run all unit tests.

```
grails test-app unit:
```

The Grails [Spock Plugin](#) is one plugin that adds new test types to Grails. It adds a `spock` test type to the `functional` phases. To run all spock tests in all phases you would run the following:

```
grails test-app :spock
```

To run all of the spock tests in the `functional` phase you would run...

```
grails test-app functional:spock
```

More than one pattern can be specified...

```
grails test-app unit:spock integration:spock
```

Targeting Tests in Types and/or Phases

Test and type/phase targeting can be applied at the same time:

```
grails test-app integration: unit: some.org.**.*
```

This would run all tests in the `integration` and `unit` phases that are in the package `some.org` or a sub-package.

9.1 Unit Testing

Unit testing are tests at the "unit" level. In other words you are testing individual methods or blocks surrounding infrastructure. Unit tests are typically run without the presence of physical resources like database connections or files. This is to ensure they run as quick as possible since quick feedback is important.

Since Grails 2.0, a collection of unit testing mixins is provided by Grails that lets you enhance the behavior of your tests. The following sections cover the usage of these mixins.



The previous JUnit 3-style `GrailsUnitTestCase` class hierarchy is still present in Grails for compatibility, but is now deprecated. The previous documentation on the subject can be found in the [documentation](#).

9.1.1 Unit Testing Controllers

The Basics

You use the `grails.test.mixin.TestFor` annotation to unit test controllers. Using `TestFor` you can also use `grails.test.mixin.web.ControllerUnitTestMethodMixin` and its associated API. For example:

```
import grails.test.mixin.TestFor

@TestFor(SimpleController)
class SimpleControllerTests {
    void testSomething() {
    }
}
```

Adding the `TestFor` annotation to a controller causes a new `controller` field to be automatically created for you.



The `TestFor` annotation will also automatically annotate any public methods starting with `@Test` annotation. If any of your test methods don't start with "test" just add this manually.

To test the simplest "Hello World"-style example you can do the following:

```
// Test class
class SimpleController {
    def hello() {
        render "hello"
    }
}
```

```
void testHello() {
    controller.hello()

    assert response.text == 'hello'
}
```


The response object is an `org.codehaus.groovy.grails.plugins.testing.GrailsMockHttpServletRequestResponse` and has a number of useful response.

For example to test a redirect you can use the `redirectUrl` property:

```
// Test class
class SimpleController {
    def index() {
        redirect action: 'hello'
    }
    ...
}
```

```
void testIndex() {
    controller.index()

    assert response.redirectedUrl == '/simple/hello'
}
```

Testing View Rendering

To test view rendering you can inspect the state of the controller's `modelAndView` (`org.springframework.web.servlet.ModelAndView`) or you can use the `view` and `model` properties.

```
// Test class
class SimpleController {
    def home() {
        render view: "homePage", model: [title: "Hello World"]
    }
    ...
}
```

```
void testIndex() {
    controller.home()

    assert view == "/simple/homePage"
    assert model.title == "Hello World"
}
```

Testing Template Rendering

Unlike view rendering, template rendering will actually attempt to write the template directly to the `ModelAndView` hence it requires a different approach to testing.

Consider the following controller action:

```
class SimpleController {
    def display() {
        render template:"snippet"
    }
}
```

In this example the controller will look for a template in `grails-app/views/simple/_snippet.gsp`.

```
void testDisplay() {
    controller.display()
    assert response.text == 'contents of template'
}
```

However, you may not want to render the real template, but just test that it was rendered. In this case you can use the `views` property of the `controller` object.

```
void testDisplay() {
    views['/simple/_snippet.gsp'] = 'mock contents'
    controller.display()
    assert response.text == 'mock contents'
}
```

Testing XML and JSON Responses

XML and JSON responses are also written directly to the response. Grails' mocking capabilities provide some support for testing these responses. For example, consider the following action:

```
def renderXml() {
    render(contentType:"text/xml") {
        book(title:"Great")
    }
}
```

This can be tested using the `xml` property of the response:

```
void testRenderXml() {
    controller.renderXml()
    assert "<book title='Great'/>" == response.text
    assert "Great" == response.xml.@title.text()
}
```

The `xml` property is a parsed result from Groovy's [XmlSlurper](#) class which is very convenient for parsing XML.

Testing JSON responses is pretty similar, instead you use the `json` property:

```
// controller action
def renderJson() {
    render(contentType:"text/json") {
        book = "Great"
    }
}
```

```
// test
void testRenderJson() {
    controller.renderJson()
    assert '{"book":"Great"}' == response.text
    assert "Great" == response.json.book
}
```

The `json` property is an instance of `org.codehaus.groovy.grails.web.json.JSONElement` useful for parsing JSON responses.

Testing XML and JSON Requests

Grails provides various convenient ways to automatically parse incoming XML and JSON packets. For example, XML requests using Grails' data binding:

```
def consumeBook() {
    def b = new Book(params['book'])
    render b.title
}
```

To test this Grails provides an easy way to specify an XML or JSON packet via the `xml` or `json` properties. These can be tested by specifying a String containing the XML:

```
void testConsumeBookXml() {
    request.xml = '<book><title>The Shining</title></book>'
    controller.consumeBook()

    assert response.text == 'The Shining'
}
```

Or alternatively a domain instance can be specified and it will be auto-converted into the appropriate XML:

```
void testConsumeBookXml() {
    request.xml = new Book(title:"The Shining")
    controller.consumeBook()

    assert response.text == 'The Shining'
}
```

The same can be done for JSON requests:

```
void testConsumeBookJson() {
    request.json = new Book(title:"The Shining")
    controller.consumeBook()

    assert response.text == 'The Shining'
}
```

If you prefer not to use Grails' data binding but instead manually parse the incoming XML or JSON that c the controller action below:

```
def consume() {
    request.withFormat {
        xml {
            render request.XML.@title
        }
        json {
            render request.JSON.title
        }
    }
}
```

To test the XML request you can specify the XML as a string:

```
void testConsumeXml() {
    request.xml = '<book title="The Stand" />'
    controller.consume()
    assert response.text == 'The Stand'
}
```

And, of course, the same can be done for JSON:

```
void testConsumeJson() {
    request.json = '{title:"The Stand"}'
    controller.consume()

    assert response.text == 'The Stand'
}
```

Testing Spring Beans

When using TestFor only a subset of the Spring beans available to a running Grails application are av beans available you can do so with the defineBeans method of GrailsUnitTestMixin:

```
class SimpleController {
    SimpleService simpleService
    def hello() {
        render simpleService.sayHello()
    }
}
```

```

void testBeanWiring() {
    defineBeans {
        simpleService(SimpleService)
    }
    controller.hello()
    assert response.text == "Hello World"
}

```

The controller is auto-wired by Spring just like in a running Grails application. Autowiring even occurs if the controller:

```

void testAutowiringViaNew() {
    defineBeans {
        simpleService(SimpleService)
    }

    def controller1 = new SimpleController()
    def controller2 = new SimpleController()

    assert controller1.simpleService != null
    assert controller2.simpleService != null
}

```

Testing Mime Type Handling

You can test mime type handling and the `withFormat` method quite simply by setting the response's fo

```

// controller action
def sayHello() {
    def data = [Hello:"World"]
    withFormat {
        xml { render data as XML }
        html data
    }
}

```

```

// test
void testSayHello() {
    response.format = 'xml'
    controller.sayHello()

    String expected = '<?xml version="1.0" encoding="UTF-8"?>' +
        '<map><entry key="Hello">World</entry></map>'

    assert expected == response.text
}

```

Testing Duplicate Form Submissions

Testing duplicate form submissions is a little bit more involved. For example if you have an action that ha

```
def handleForm() {
    withForm {
        render "Good"
    }.invalidToken {
        render "Bad"
    }
}
```

you want to verify the logic that is executed on a good form submission and the logic that is executed on a bad form submission is simple. Just invoke the controller:

```
void testDuplicateFormSubmission() {
    controller.handleForm()
    assert "Bad" == response.text
}
```

Testing the successful submission requires providing an appropriate SynchronizerToken:

```
import org.codehaus.groovy.grails.web.servlet.mvc.SynchronizerToken
...

void testValidFormSubmission() {
    def token = SynchronizerToken.store(session)
    params[SynchronizerToken.KEY] = token.currentToken.toString()

    controller.handleForm()
    assert "Good" == response.text
}
```

If you test both the valid and the invalid request in the same test be sure to reset the response between executions:

```
controller.handleForm() // first execution
...
response.reset()
...
controller.handleForm() // second execution
```

Testing File Upload

You use the `GrailsMockMultipartFile` class to test file uploads. For example consider the following:

```
def uploadFile() {
    MultipartFile file = request.getFile("myFile")
    file.transferTo(new File("/local/disk/myFile"))
}
```

To test this action you can register a `GrailsMockMultipartFile` with the request:

```

void testFileUpload() {
    final file = new GrailsMockMultipartFile("myFile", "foo".bytes)
    request.addFile(file)
    controller.uploadFile()

    assert file.targetFileLocation.path == "/local/disk/myFile"
}

```

The `GrailsMockMultipartFile` constructor arguments are the name and contents of the file. I use the `transferTo` method that simply records the `targetFileLocation` and doesn't write to disk.

Testing Command Objects

Special support exists for testing command object handling with the `mockCommandObject` method on `ControllerUnitTestMethodBuilder`. Here's an example action:

```

def handleCommand(SimpleCommand simple) {
    if (simple.hasErrors()) {
        render "Bad"
    }
    else {
        render "Good"
    }
}

```

To test this you mock the command object, populate it and then validate it as follows:

```

void testInvalidCommand() {
    def cmd = mockCommandObject(SimpleCommand)
    cmd.name = '' // doesn't allow blank names

    cmd.validate()
    controller.handleCommand(cmd)

    assert response.text == 'Bad'
}

```

Testing Calling Tag Libraries

You can test calling tag libraries using `ControllerUnitTestMethodBuilder`, although the mechanism for testing tag libraries is different. For example to test a call to the `message` tag, add a message to the `messageSource`. Consider the following:

```

def showMessage() {
    render g.message(code: "foo.bar")
}

```

This can be tested as follows:

```
void testRenderBasicTemplateWithTags() {
    messageSource.addMessage("foo.bar", request.locale, "Hello World")

    controller.showMessage()

    assert response.text == "Hello World"
}
```

9.1.2 Unit Testing Tag Libraries

The Basics

Tag libraries and GSP pages can be tested with the `grails.test.mixin.web.GroovyPageUnit` declare which tag library is under test with the `TestFor` annotation:

```
@TestFor(SimpleTagLib)
class SimpleTagLibTests {
}
```

Note that if you are testing invocation of a custom tag from a controller you can combine the `Cont` `GroovyPageUnitTestMixin` using the `Mock` annotation:

```
@TestFor(SimpleController)
@Mock(SimpleTagLib)
class GroovyPageUnitTestMixinTests {
}
```

Testing Custom Tags

The core Grails tags don't need to be enabled during testing, however custom tag libraries do. The `G` provides a `mockTagLib()` method that you can use to mock a custom tag library. For example consider

```
class SimpleTagLib {
    static namespace = 's'

    def hello = { attrs, body ->
        out << "Hello ${attrs.name ?: 'World'}"
    }
}
```

You can test this tag library by using `TestFor` and supplying the name of the tag library:


```
@TestFor(SimpleTagLib)
class SimpleTagLibTests {
    void testHelloTag() {
        assert applyTemplate('<s:hello />') == 'Hello World'
        assert applyTemplate('<s:hello name="Fred" />') == 'Hello Fred'
    }
}
```

Alternatively, you can use the `TestMixin` annotation and mock multiple tag libraries using the `mockTagLib`

```
@grails.test.mixin.TestMixin(GroovyPageUnitTestMixin)
class MultipleTagLibraryTests {

    @Test
    void testMuliple() {
        mockTagLib(FirstTagLib)
        mockTagLib(SecondTagLib)

        ...
    }
}
```

The `GroovyPageUnitTestMixin` provides convenience methods for asserting that the template output

```
@grails.test.mixin.TestMixin(GroovyPageUnitTestMixin)
class MultipleTagLibraryTests {

    @Test
    void testMuliple() {
        mockTagLib(FirstTagLib)
        mockTagLib(SecondTagLib)
        assertOutputEquals ('Hello World', '<s:hello />')
        assertOutputMatches (/.Fred./, '<s:hello name="Fred" />')
    }
}
```

Testing View and Template Rendering

You can test rendering of views and templates in `grails-app/views` via the `render` method of `GroovyPageUnitTestMixin`:

```
def result = render(template: "/simple/hello")
assert result == "Hello World"
```

This will attempt to render a template found at the location `grails-app/views/simple/_hello.groovy`. On any custom tag libraries you need to call `mockTagLib` as described in the previous section.

9.1.3 Unit Testing Domains

Overview



The mocking support described here is best used when testing non-domain artifacts that use you focus on testing the artifact without needing a database. But when testing persistence it's tests which configure Hibernate and use a database.

Domain class interaction can be tested without involving a database connection using `DomainClassUnit` mimics the behavior of GORM against an in-memory `ConcurrentHashMap` implementation. Note that GORM implementation. The following features of GORM for Hibernate can only be tested within an integ

- String-based HQL queries
- composite identifiers
- dirty checking methods
- any direct interaction with Hibernate

However a large, commonly-used portion of the GORM API can be mocked using `DomainClassUnit`

- Simple persistence methods like `save()`, `delete()` etc.
- Dynamic Finders
- Named Queries
- Query-by-example
- GORM Events

If something isn't supported then `GrailsUnitTestMixin`'s `mockFor` method can come in handy to you can write an integration test which bootstraps the complete Grails environment at a cost of test executi

The Basics

`DomainClassUnitTestMixin` is typically used in combination with testing either a controller, servi mock collaborator defined by the `Mock` annotation:

```
import grails.test.mixin.*

@TestFor(SimpleController)
@Mock(Simple)
class SimpleControllerTests {
}
```

The example above tests the `SimpleController` class and mocks the behavior of the `Simple` domain typical scaffolded save controller action:

```

class BookController {
  def save() {
    def book = new Book(params)
    if (book.save(flush: true)) {
      flash.message = message(
        code: 'default.created.message',
        args: [message(code: 'book.label',
          default: 'Book'), book.id])
      redirect(action: "show", id: book.id)
    }
    else {
      render(view: "create", model: [bookInstance: book])
    }
  }
}

```

Tests for this action can be written as follows:

```

import grails.test.mixin.*

@TestFor(BookController)
@Mock(Book)
class BookControllerTests {

  void testSaveInvalidBook() {
    controller.save()

    assert model.bookInstance != null
    assert view == '/book/create'
  }

  void testSaveValidBook() {
    params.title = "The Stand"
    params.pages = "500"

    controller.save()

    assert response.redirectedUrl == '/book/show/1'
    assert flash.message != null
    assert Book.count() == 1
  }
}

```

Mock annotation also supports a list of mock collaborators if you have more than one domain to mock:

```

@TestFor(BookController)
@Mock([Book, Author])
class BookControllerTests {
  ...
}

```

Alternatively you can also use the DomainClassUnitTestMethodMixin directly with the TestMixin annotation:

```

@TestFor(BookController)
@TestMixin(DomainClassUnitTestMethodMixin)
class BookControllerTests {
  ...
}

```

And then call the `mockDomain` method to mock domains during your test:

```
void testSave() {
    mockDomain(Author)
    mockDomain(Book)
}
```

The `mockDomain` method also includes an additional parameter that lets you pass a Map of Maps to c fixture-like data:

```
void testSave() {
    mockDomain(Book, [
        [title: "The Stand", pages: 1000],
        [title: "The Shining", pages: 400],
        [title: "Along Came a Spider", pages: 300] ])
}
```

Testing Constraints

Your constraints contain logic and that logic is highly susceptible to bugs - the kind of bugs that can be t default `save()` doesn't throw an exception when it fails). If your answer is that it's too hard or fiddly, `mockForConstraintsTests()` method.

This method is like a much reduced version of the `mockDomain()` method that simply adds a valida All you have to do is mock the class, create an instance with populated data, and then call `validate` property to determine if validation failed. So if all we are doing is mocking the `validate()` method, That is so that we can test the unique constraint as you will soon see.

So, suppose we have a simple domain class:

```
class Book {
    String title
    String author

    static constraints = {
        title blank: false, unique: true
        author blank: false, minSize: 5
    }
}
```

Don't worry about whether the constraints are sensible (they're not!), they are for demonstration only. 7 following:

```

@TestFor(Book)
class BookTests {
  void testConstraints() {

    def existingBook = new Book(
      title: "Misery",
      author: "Stephen King")

    mockForConstraintsTests(Book, [existingBook])

    // validation should fail if both properties are null
    def book = new Book()

    assert !book.validate()
    assert "nullable" == book.errors["title"]
    assert "nullable" == book.errors["author"]

    // So let's demonstrate the unique and minSize constraints

    book = new Book(title: "Misery", author: "JK")
    assert !book.validate()
    assert "unique" == book.errors["title"]
    assert "minSize" == book.errors["author"]

    // Validation should pass!
    book = new Book(title: "The Shining", author: "Stephen King")
    assert book.validate()
  }
}

```

You can probably look at that code and work out what's happening without any further explanation. The `errors` property is used. First, is a real Spring Errors instance, so you can access all the properties and methods. Second, this particular Errors object also has map/property access as shown. Simply specify the name of the constraint that was violated. map/property access will return the name of the constraint that was violated. Note that it is the constraint name (might expect).

That's it for testing constraints. One final thing we would like to say is that testing the constraints in this way is the "constraints" property name! It is currently one of the hardest bugs to track down normally, and you can highlight the problem straight away.

9.1.4 Unit Testing Filters

Unit testing filters is typically a matter of testing a controller where a filter is a mock collaborator. For example:

```

class CancellingFilters {
  def filters = {
    all(controller:"simple", action:"list") {
      before = {
        redirect(controller:"book")
        return false
      }
    }
  }
}

```

This filter intercepts the `list` action of the `simple` controller and redirects to the `book` controller. To test this, you target the `SimpleController` class and add the `CancellingFilters` as a mock collaborator.

```
@TestFor(SimpleController)
@Mock(CancellingFilters)
class SimpleControllerTests {
}
```

You can then implement a test that uses the `withFilters` method to wrap the call to an action in filter c

```
void testInvocationOfListActionIsFiltered() {
    withFilters(action:"list") {
        controller.list()
    }
    assert response.redirectedUrl == '/book'
}
```

Note that the `action` parameter is required because it is unknown what the action to invoke is un
controller parameter is optional and taken from the controller under test. If it is a another controller yo

```
withFilters(controller:"book",action:"list") {
    controller.list()
}
```

9.1.5 Unit Testing URL Mappings

The Basics

Testing URL mappings can be done with the `TestFor` annotation testing a particular URL mappings cla
mappings you can do the following:

```
@TestFor(UrlMappings)
class UrlMappingsTests {
}
```



Note that since the default `UrlMappings` class is in the default package your test must
package

With that done there are a number of useful methods that are defined by the `grails.test.mixin.web`
for testing URL mappings. These include:

- `assertForwardUrlMapping` - Asserts a URL mapping is forwarded for the given controller cl
defined as a mock collaborate for this to work)
- `assertReverseUrlMapping` - Asserts that the given URL is produced when reverse mapping a
- `assertUrlMapping` - Asserts a URL mapping is valid for the given URL. This combines the a
`assertReverseUrlMapping` assertions

Asserting Forward URL Mappings

You use `assertForwardUrlMapping` to assert that a given URL maps to a given controller. For mappings:

```
static mappings = {
    "/action1"(controller: "simple", action: "action1")
    "/action2"(controller: "simple", action: "action2")
}
```

The following test can be written to assert these URL mappings:

```
void testUrlMappings() {
    assertForwardUrlMapping("/action1", controller: 'simple',
                           action: "action1")
    assertForwardUrlMapping("/action2", controller: 'simple',
                           action: "action2")

    shouldFail {
        assertForwardUrlMapping("/action2", controller: 'simple',
                                action: "action1")
    }
}
```

Assert Reverse URL Mappings

You use `assertReverseUrlMapping` to check that correct links are produced for your URL map views. An example test is largely identical to the previous listing except you use `assertReverseUrlMapping` instead of `assertForwardUrlMapping`. Note that you can combine these 2 assertions with `assertUrlMapping`:

Simulating Controller Mapping

In addition to the assertions to check the validity of URL mappings you can also simulate mapping to a controller as a mock collaborator and the `mapURI` method. For example:

```
@TestFor(SimpleController)
@Mock(UrlMappings)
class SimpleControllerTests {
    void testControllerMapping() {
        SimpleController controller = mapURI('/simple/list')
        assert controller != null

        def model = controller.list()
        assert model != null
    }
}
```

9.1.6 Mocking Collaborators

Beyond the specific targeted mocking APIs there is also an all-purpose `mockFor()` method that is an annotation. The signature of `mockFor` is:

```
mockFor(class, loose = false)
```

This is general-purpose mocking that lets you set up either strict or loose demands on a class.

This method is surprisingly intuitive to use. By default it will create a strict mock control object (one for which is important) that you can use to specify demands:

```
def strictControl = mockFor(MyService)
strictControl.demand.someMethod(0..2) { String arg1, int arg2 -> ... }
strictControl.demand.static.aStaticMethod {-> ... }
```

Notice that you can mock static as well as instance methods by using the "static" property. You then specify with an optional range argument. This range determines how many times you expect the method to be called. If it falls outside of that range (either too few or too many) then an assertion error will be thrown. If no range is assumed, i.e. that the method must be called exactly once.

The last part of a demand is a closure representing the implementation of the mock method. The closure takes the types of the mocked method, but otherwise you are free to add whatever you want in the body.

As we mentioned before, call `mockControl.createMock()` to get an actual mock instance of the class. You can call this multiple times to create as many mock instances as you need. And once you have a mock, call `mockControl.verify()` to check that the expected methods were called.

Lastly, the call:

```
def looseControl = mockFor(MyService, true)
```

will create a mock control object that has only loose expectations, i.e. the order that methods are invoked does not matter.

9.2 Integration Testing

Integration tests differ from unit tests in that you have full access to the Grails environment within the test. Grails will create a database for integration tests and clears out all the data from the database between tests.

One thing to bear in mind is that logging is enabled for your application classes, but it is different from logging in a unit test like this:

```
class MyServiceTests extends GroovyTestCase {
    void testSomething() {
        log.info "Starting tests"
        ...
    }
}
```


the "starting tests" message is logged using a different system than the one used by the application. The 1 instance of `java.util.logging.Logger` (inherited from the base class, not injected by Grails), with the `log` property injected into your application artifacts. For example, it doesn't have `debug()` or `trace()` is in fact `warning()`.

Transactions

Integration tests run inside a database transaction by default, which is rolled back at the end of the each test is not persisted to the database. Add a `transactional` property to your test class to check transactional

```
class MyServiceTests extends GroovyTestCase {
    static transactional = false

    void testMyTransactionalServiceMethod() {
        ...
    }
}
```

Be sure to remove any persisted data from a non-transactional test, for example in the `tearDown` method. In standard transactional tests that expect a clean database.

Testing Controllers

To test controllers you first have to understand the Spring Mock Library.

Grails automatically configures each test with a [MockHttpServletRequest](#), [MockHttpServletResponse](#), and [MockServletContext](#). For example consider the following controller:

```
class FooController {
    def text() {
        render "bar"
    }

    def someRedirect() {
        redirect(action: "bar")
    }
}
```

The tests for this would be:

```
class FooControllerTests extends GroovyTestCase {

    void testText() {
        def fc = new FooController()
        fc.text()
        assertEquals "bar", fc.response.contentAsString
    }

    void testSomeRedirect() {
        def fc = new FooController()
        fc.someRedirect()
        assertEquals "/foo/bar", fc.response.redirectedUrl
    }
}
```

In the above case response is an instance of `MockHttpServletResponse` which we can use `contentAsString` (when writing to the response) or the redirected URL. These mocked versions of `response` (unlike the real versions) and hence you can set properties on the request such as the `contextPath` and

Grails **does not** invoke [interceptors](#) or servlet filters when calling actions during integration testing. For isolation, using [functional testing](#) if necessary.

Testing Controllers with Services

If your controller references a service (or other Spring beans), you have to explicitly initialise the service for

Given a controller using a service:

```
class FilmStarsController {
    def popularityService

    def update() {
        // do something with popularityService
    }
}
```

The test for this would be:

```
class FilmStarsTests extends GroovyTestCase {
    def popularityService

    void testInjectedServiceInController () {
        def fsc = new FilmStarsController()
        fsc.popularityService = popularityService
        fsc.update()
    }
}
```

Testing Controller Command Objects

With command objects you just supply parameters to the request and it will automatically do the command. You can test your action with no parameters:

Given a controller using a command object:

```
class AuthenticationController {
    def signup(SignupForm form) {
        ...
    }
}
```

You can then test it like this:

```
def controller = new AuthenticationController()
controller.params.login = "marcpalmer"
controller.params.password = "secret"
controller.params.passwordConfirm = "secret"
controller.signup()
```

Grails auto-magically sees your call to `signup()` as a call to the action and populates the controller parameters. During controller testing, the params are mutable with a mocked request supplied by Grails.

Testing Controllers and the render Method

The [render](#) method lets you render a custom view at any point within the body of an action. For instance, c

```
def save() {
  def book = Book(params)
  if (book.save()) {
    // handle
  }
  else {
    render(view: "create", model:[book:book])
  }
}
```

In the above example the result of the model of the action is not available as the return value, but instead as a property of the controller. The `modelAndView` property is an instance of Spring MVC's [ModelAndView](#) object, the result of an action:

```
def bookController = new BookController()
bookController.save()
def model = bookController.modelAndView.model.book
```

Simulating Request Data

You can use the Spring [MockHttpServletRequest](#) to test an action that requires request data, for example, consider this action which performs data binding from an incoming request:

```
def create() {
  [book: new Book(params.book)]
}
```

To simulate the 'book' parameter as an XML request you could do something like the following:

```
void testCreateWithXML() {
  def controller = new BookController()
  controller.request.contentType = 'text/xml'
  controller.request.content = '''\
    <?xml version="1.0" encoding="ISO-8859-1"?>
    <book>
      <title>The Stand</title>
      ...
    </book>
  '''.stripIndent().getBytes() // note we need the bytes

  def model = controller.create()
  assert model.book
  assertEquals "The Stand", model.book.title
}
```

The same can be achieved with a JSON request:

```
void testCreateWithJSON() {
    def controller = new BookController()
    controller.request.contentType = "text/json"
    controller.request.content =
        '{"id":1,"class":"Book","title":"The Stand"}'.getBytes()

    def model = controller.create()
    assert model.book
    assertEquals "The Stand", model.book.title
}
```



With JSON don't forget the `class` property to specify the name the target type to bind to. It is within the name of the `<book>` node, but this property is required as part of the JSON packet

For more information on the subject of REST web services see the section on [REST](#).

Testing Web Flows

Testing [Web Flows](#) requires a special test harness called `grails.test.WebFlowTestCase` with [AbstractFlowExecutionTests](#) class.



Subclasses of `WebFlowTestCase` **must** be integration tests

For example given this simple flow:

```
class ExampleController {
    def exampleFlow() {
        start {
            on("go") {
                flow.hello = "world"
            }.to "next"
        }
        next {
            on("back").to "start"
            on("go").to "subber"
        }
        subber {
            subflow(action: "sub")
            on("end").to("end")
        }
        end()
    }

    def subFlow() {
        subSubflowState {
            subflow(controller: "other", action: "otherSub")
            on("next").to("next")
        }
        ...
    }
}
```

You need to tell the test harness what to use for the "flow definition". This is done via overriding the abstract

```
import grails.test.WebFlowTestCase

class ExampleFlowTests extends WebFlowTestCase {
    def getFlow() { new ExampleController().exampleFlow }
    ...
}
```

You can specify the flow id by overriding the `getFlowId` method, otherwise the default is `test`:

```
import grails.test.WebFlowTestCase

class ExampleFlowTests extends WebFlowTestCase {
    String getFlowId() { "example" }
    ...
}
```

If the flow under test calls any subflows, these (or mocks) must be registered before the calling the flow:

```
protected void setUp() {
    super.setUp()

    registerFlow("other/otherSub") { // register a simplified mock
        start {
            on("next").to("end")
        }
        end()
    }

    // register the original subflow
    registerFlow("example/sub", new ExampleController().subFlow)
}
```

Then you kick off the flow with the `startFlow` method:

```
void testExampleFlow() {
    def viewSelection = startFlow()
    ...
}
```

Use the `signalEvent` method to trigger an event:

```
void testExampleFlow() {
    ...
    signalEvent("go")
    assert "next" == flowExecution.activeSession.state.id
    assert "world" == flowScope.hello
}
```

Here we have signaled to the flow to execute the event "go" which causes a transition to the "next" state. It also puts a `hello` variable into the flow scope.

Testing Tag Libraries

Testing tag libraries is simple because when a tag is invoked as a method it returns its result as a string (this class implements all of the methods of `String`). So for example if you have a tag library like this:

```
class FooTagLib {
    def bar = { attrs, body ->
        out << "<p>Hello World!</p>"
    }

    def bodyTag = { attrs, body ->
        out << "<${attrs.name}>"
        out << body()
        out << "</${attrs.name}>"
    }
}
```

The tests would look like:

```
class FooTagLibTests extends GroovyTestCase {
    void testBarTag() {
        assertEquals "<p>Hello World!</p>",
            new FooTagLib().bar(null, null).toString()
    }

    void testBodyTag() {
        assertEquals "<p>Hello World!</p>",
            new FooTagLib().bodyTag(name: "p") {
                "Hello World!"
            }.toString()
    }
}
```

Notice that for the second example, `testBodyTag`, we pass a block that returns the body of the tag. This is as a `String`.

Testing Tag Libraries with `GroovyPagesTestCase`

In addition to doing simple testing of tag libraries like in the above examples, you can also use the `grails` class to test tag libraries with integration tests.

The `GroovyPagesTestCase` class is a subclass of the standard `GroovyTestCase` class and adds support for GSP rendering.



`GroovyPagesTestCase` can only be used in an integration test.

For example, consider this date formatting tag library:

```
import java.text.SimpleDateFormat

class FormatTagLib {
    def dateFormat = { attrs, body ->
        out << new SimpleDateFormat(attrs.format) << attrs.date
    }
}
```

This can be easily tested as follows:

```
class FormatTagLibTests extends GroovyPagesTestCase {
    void testDateFormat() {
        def template =
            '<g:dateFormat format="dd-MM-yyyy" date="${myDate}" />'
        def testDate = ... // create the date
        assertEquals('01-01-2008', template, [myDate:testDate])
    }
}
```

You can also obtain the result of a GSP using the `applyTemplate` method of the `GroovyPagesTest`

```
class FormatTagLibTests extends GroovyPagesTestCase {
    void testDateFormat() {
        def template =
            '<g:dateFormat format="dd-MM-yyyy" date="${myDate}" />'
        def testDate = ... // create the date
        def result = applyTemplate(template, [myDate:testDate])
        assertEquals '01-01-2008', result
    }
}
```

Testing Domain Classes

Testing domain classes is typically a simple matter of using the [GORM API](#), but there are a few things you often need to "flush" to ensure the correct state has been persisted to the database. For example

```
void testQuery() {
    def books = [
        new Book(title: "The Stand"),
        new Book(title: "The Shining")]
    books*.save()

    assertEquals 2, Book.list().size()
}
```

This test will fail because calling `save` does not actually persist the `Book` instances when called. Calling `save` at some point in the future these instances should be persisted. To commit changes immediately you "flush" the

```
void testQuery() {
    def books = [
        new Book(title: "The Stand"),
        new Book(title: "The Shining")]
    books*.save(flush: true)

    assertEquals 2, Book.list().size()
}
```

In this case since we're passing the argument `flush` with a value of `true` the updates will be persisted in to the query later on.

9.3 Functional Testing

Functional tests involve making HTTP requests against the running application and verifying the results. Grails has no support for writing functional tests directly, but there are several plugins available for this.

- Canoo Webtest - <http://grails.org/plugin/webtest>
- G-Func - <http://grails.org/plugin/functional-test>
- Geb - <http://grails.org/plugin/geb>
- Selenium-RC - <http://grails.org/plugin/selenium-rc>
- WebDriver - <http://grails.org/plugin/webdriver>

Consult the documentation for each plugin for its capabilities.

Common Options

There are options that are common to all plugins that control how the Grails application is launched, if at a

inline

The `-inline` option specifies that the grails application should be started inline (i.e. like `run-app`).

This option is implicitly set unless the `baseUrl` or `war` options are set

war

The `-war` option specifies that the grails application should be packaged as a war and started. This is production-like state, but it has a longer startup time than the `-inline` option. It also runs the war in a way that allows access to any internal application objects.

```
grails test-app functional: -war
```

Note that the same build/config options for the [run-war](#) command apply to functional testing against the W

https

The `-https` option results in the application being able to receive https requests as well as http requests. It can be used in conjunction with `-inline` and `-war` options.


```
grails test-app functional: -https
```

Note that this does not change the test *base url* to be https, it will still be http unless the `-httpsBaseUr`:

httpsBaseUrl

The `-httpsBaseUrl` causes the implicit base url to be used for tests to be a https url.

```
grails test-app functional: -httpsBaseUrl
```

This option is ignored if the `-baseUrl` option is specified.

baseUrl

The `baseUrl` option allows the base url for tests to be specified.

```
grails test-app functional: -baseUrl=http://mycompany.com/grailsapp
```

This option will prevent the local grails application being started unless `-inline` or `-war` are given as test against the local Grails application you **must** specify one of either the `-inline` or `-war` options.

10 Internationalization

Grails supports Internationalization (i18n) out of the box by leveraging the underlying Spring MVC internals. Grails is able to customize the text that appears in a view based on the user's Locale. To quote the javadoc for the `Locale` class:

A Locale object represents a specific geographical, political, or cultural region. An operation that is called locale-sensitive and uses the Locale to tailor information for the user. For example, a locale-sensitive operation--the number should be formatted according to the customs/conventions of the region, or culture.

A Locale is made up of a [language code](#) and a [country code](#). For example "en_US" is the code for US English.

10.1 Understanding Message Bundles

Now that you have an idea of locales, to use them in Grails you create a message bundle file containing translations. Message bundles in Grails are located inside the `grails-app/i18n` directory and are simple Java properties files.

Each bundle starts with the name `messages` by convention and ends with the locale. Grails ships with a range of languages within the `grails-app/i18n` directory. For example:

```
messages.properties
messages_da.properties
messages_de.properties
messages_es.properties
messages_fr.properties
...
```

By default Grails looks in `messages.properties` for messages unless the user has specified a locale. You can create a new properties file that ends with the locale you are interested in. For example `messages_en_GB.properties` for British English.

10.2 Changing Locales

By default the user locale is detected from the incoming `Accept-Language` header. However, you can change the locale by simply passing a parameter called `lang` to Grails as a request parameter:

```
/book/list?lang=es
```

Grails will automatically switch the user's locale and store it in a cookie so subsequent requests will have the same locale.

10.3 Reading Messages

Reading Messages in the View

The most common place that you need messages is inside the view. Use the [message](#) tag for this:

```
<g:message code="my.localized.content" />
```

As long as you have a key in your `messages.properties` (with appropriate locale suffix) such as the message:

```
my.localized.content=Hola, Me llamo John. Hoy es domingo.
```

Messages can also include arguments, for example:

```
<g:message code="my.localized.content" args="${ ['Juan', 'lunes'] }" />
```

The message declaration specifies positional parameters which are dynamically specified:

```
my.localized.content=Hola, Me llamo {0}. Hoy es {1}.
```

Reading Messages in Controllers and Tag Libraries

It's simple to read messages in a controller since you can invoke tags as methods:

```
def show() {  
    def msg = message(code: "my.localized.content", args: ['Juan', 'lunes'])  
}
```

The same technique can be used in [tag libraries](#), but if your tag library uses a custom [namespace](#) then you

```
def myTag = { attrs, body ->  
    def msg = g.message(code: "my.localized.content", args: ['Juan', 'lunes'])  
}
```

10.4 Scaffolding and i18n

Grails [scaffolding](#) templates for controllers and views are fully i18n-aware. The GSPs use the [message](#) tag and flash messages use i18n to resolve locale-specific messages.

11 Security

Grails is no more or less secure than Java Servlets. However, Java servlets (and hence Grails) are extremely vulnerable to common buffer overrun and malformed URL exploits due to the nature of the Java Virtual Machine underlying them.

Web security problems typically occur due to developer naivety or mistakes, and there is a little Grails magic that makes writing secure applications easier to write.

What Grails Automatically Does

Grails has a few built in safety mechanisms by default.

1. All standard database access via [GORM](#) domain objects is automatically SQL escaped to prevent SQL injection.
2. The default [scaffolding](#) templates HTML escape all data fields when displayed
3. Grails link creating tags ([link](#), [form](#), [createLink](#), [createLinkTo](#) and others) all use appropriate escaping
4. Grails provides [codecs](#) to let you trivially escape data when rendered as HTML, JavaScript and URLs

11.1 Securing Against Attacks

SQL injection

Hibernate, which is the technology underlying GORM domain classes, automatically escapes data when creating SQL queries. However it is still possible to write bad dynamic HQL code that uses unchecked request parameters and is vulnerable to HQL injection attacks:

```
def vulnerable() {
    def books = Book.find("from Book as b where b.title = '" + params.title + "'")
}
```

or the analogous call using a GString:

```
def vulnerable() {
    def books = Book.find("from Book as b where b.title = '${params.title}'")
}
```

Do **not** do this. Use named or positional parameters instead to pass in parameters:

```
def safe() {
    def books = Book.find("from Book as b where b.title = ?",
                          [params.title])
}
```

or

```
def safe() {  
    def books = Book.find("from Book as b where b.title = :title",  
                          [title: params.title])  
}
```

Phishing

This really a public relations issue in terms of avoiding hijacking of your branding and a declared comm Customers need to know how to identify valid emails.

XSS - cross-site scripting injection

It is important that your application verifies as much as possible that incoming requests were originate another site. Ticketing and page flow systems can help this and Grails' support for [Spring Web Flow](#) inclu

It is also important to ensure that all data values rendered into views are escaped correctly. For example you must call [encodeAsHTML](#) on every object to ensure that people cannot maliciously inject JavaScript viewed by others. Grails supplies several [Dynamic Encoding Methods](#) for this purpose and if your output can easily write your own codec.

You must also avoid the use of request parameters or data fields for determining the next URL to redirect parameter for example to determine where to redirect a user to after a successful login, attackers can imitate site, and then redirect the user back to their own site once logged in, potentially allowing JavaScript code to the site.

Cross-site request forgery

CSRF involves unauthorized commands being transmitted from a user that a website trusts. A typical embedding a link to perform an action on your website if the user is still authenticated.

The best way to decrease risk against these types of attacks is to use the `useToken` attribute on your [Submissions](#) for more information on how to use it. An additional measure would be to not use remember-

HTML/URL injection

This is where bad data is supplied such that when it is later used to create a link in a page, clicking it will may redirect to another site or alter request parameters.

HTML/URL injection is easily handled with the [codecs](#) supplied by Grails, and the tag libraries supplied appropriate. If you create your own tags that generate URLs you will need to be mindful of doing this too.

Denial of service

Load balancers and other appliances are more likely to be useful here, but there are also issues relating to link is created by an attacker to set the maximum value of a result set so that a query could exceed the n system down. The solution here is to always sanitize request parameters before passing them to dynamic fi

```
def safeMax = Math.max(params.max?.toInteger(), 100) // limit to 100 results  
return Book.list(max:safeMax)
```

Guessable IDs

Many applications use the last part of the URL as an "id" of some object to retrieve from GORM or else these are easily guessable as they are typically sequential integers.

Therefore you must assert that the requesting user is allowed to view the object with the requested id before.

Not doing this is "security through obscurity" which is inevitably breached, just like having a default password.

You must assume that every unprotected URL is publicly accessible one way or another.

11.2 Encoding and Decoding Objects

Grails supports the concept of dynamic encode/decode methods. A set of standard codecs are bundled with the framework and a mechanism for developers to contribute their own codecs that will be recognized at runtime.

Codec Classes

A Grails codec class is one that may contain an encode closure, a decode closure or both. When a Grails application is started, the framework dynamically loads codecs from the `grails-app/utils/` directory.

The framework looks under `grails-app/utils/` for class names that end with the convention `Codec`. The first codec that ships with Grails is `HTMLCodec`.

If a codec contains an `encode` closure Grails will create a dynamic `encode` method and add that method to the `Object` class, representing the codec that defined the `encode` closure. For example, the `HTMLCodec` class defines an `encodeAsHTML` method with the name `encodeAsHTML`.

The `HTMLCodec` and `URLCodec` classes also define a `decode` closure, so Grails attaches those methods to the `Object` class as `decodeURL` respectively. Dynamic codec methods may be invoked from anywhere in a Grails application. In a GSP, a report contains a property called 'description' which may contain special characters that must be escaped in an HTML document. One way to deal with that in a GSP is to encode the description property using the dynamic `encodeAsHTML` method.

```
${report.description.encodeAsHTML( )}
```

Decoding is performed using `value.decodeHTML()` syntax.

Standard Codecs

HTMLCodec

This codec performs HTML escaping and unescaping, so that values can be rendered safely in an HTML document without damaging the page layout. For example, given a value "Don't you know that 2 > 1?" you wouldn't be able to click on a link on the page because the `>` will look like it closes a tag, which is especially bad if you render this data within an input field.

Example of usage:

```
<input name="comment.message" value="${comment.message.encodeAsHTML( )}" />
```



Note that the HTML encoding does not re-encode apostrophe/single quote so you must attribute values to avoid text with apostrophes affecting your page.

URLCodec

URL encoding is required when creating URLs in links or form actions, or any time data is used to create from getting into the URL and changing its meaning, for example "Apple & Blackberry" is not going to work as the ampersand will break parameter parsing.

Example of usage:

```
<a href="/mycontroller/find?searchKey=${lastSearch.encodeAsURL()} ">
Repeat last search
</a>
```

Base64Codec

Performs Base64 encode/decode functions. Example of usage:

```
Your registration code is: ${user.registrationCode.encodeAsBase64()}
```

JavaScriptCodec

Escapes Strings so they can be used as valid JavaScript strings. For example:

```
Element.update('${elementId}',
    '${render(template: "/common/message").encodeAsJavaScript()}')
```

HexCodec

Encodes byte arrays or lists of integers to lowercase hexadecimal strings, and can decode hexadecimal strings.

```
Selected colour: #${[255,127,255].encodeAsHex()}
```

MD5Codec

Uses the MD5 algorithm to digest byte arrays or lists of integers, or the bytes of a string (in default system usage). Example of usage:

```
Your API Key: ${user.uniqueID.encodeAsMD5()}
```

MD5BytesCodec

Uses the MD5 algorithm to digest byte arrays or lists of integers, or the bytes of a string (in default system usage).

```
byte[] passwordHash = params.password.encodeAsMD5Bytes()
```

SHA1Codec

Uses the SHA1 algorithm to digest byte arrays or lists of integers, or the bytes of a string (in default system encoding) and return a hexadecimal string. Example of usage:

```
Your API Key: ${user.uniqueID.encodeAsSHA1()}
```

SHA1BytesCodec

Uses the SHA1 algorithm to digest byte arrays or lists of integers, or the bytes of a string (in default system encoding) and return a byte array. Example of usage:

```
byte[] passwordHash = params.password.encodeAsSHA1Bytes()
```

SHA256Codec

Uses the SHA256 algorithm to digest byte arrays or lists of integers, or the bytes of a string (in default system encoding) and return a hexadecimal string. Example of usage:

```
Your API Key: ${user.uniqueID.encodeAsSHA256()}
```

SHA256BytesCodec

Uses the SHA256 algorithm to digest byte arrays or lists of integers, or the bytes of a string (in default system encoding) and return a byte array. Example of usage:

```
byte[] passwordHash = params.password.encodeAsSHA256Bytes()
```

Custom Codecs

Applications may define their own codecs and Grails will load them along with the standard codecs. A custom codec must be placed in the `grails-app/utils/` directory and the class name must end with `Codec`. The codec may contain a `encode` closure or both. The closure must accept a single argument which will be the object that the codec is encoding. Example:

```
class PigLatinCodec {
    static encode = { str ->
        // convert the string to pig latin and return the result
    }
}
```


With the above codec in place an application could do something like this:

```
${lastName.encodeAsPigLatin()}
```

11.3 Authentication

Grails has no default mechanism for authentication as it is possible to implement authentication in many ways. You can implement a simple authentication mechanism using either [interceptors](#) or [filters](#). This is sufficient for simple applications. To use an established security framework, for example by using the [Spring Security](#) or the [Shiro](#) plugin.

Filters let you apply authentication across all controllers or across a URI space. For example you can create `grails-app/conf/SecurityFilters.groovy` by running:

```
grails create-filters security
```

and implement your interception logic there:

```
class SecurityFilters {
    def filters = {
        loginCheck(controller: '*', action: '*') {
            before = {
                if (!session.user && actionName != "login") {
                    redirect(controller: "user", action: "login")
                    return false
                }
            }
        }
    }
}
```

Here the `loginCheck` filter intercepts execution *before* all actions except `login` are executed, and redirect to the `login` action.

The `login` action itself is simple too:

```
def login() {
    if (request.get) {
        return // render the login view
    }
}

def u = User.findByLogin(params.login)
if (u) {
    if (u.password == params.password) {
        session.user = u
        redirect(action: "home")
    }
    else {
        render(view: "login", model: [message: "Password incorrect"])
    }
}
else {
    render(view: "login", model: [message: "User not found"])
}
}
```

11.4 Security Plugins

If you need more advanced functionality beyond simple authentication such as authorization, roles etc. the available security plugins.

11.4.1 Spring Security

The Spring Security plugins are built on the [Spring Security](#) project which provides a flexible, extensible authentication and authorization schemes. The plugins are modular so you can install just the functional. The Spring Security plugins are the official security plugins for Grails and are actively maintained and supported.

There is a [Core plugin](#) which supports form-based authentication, encrypted/salted passwords, HTTP Basic dependent plugins provide alternate functionality such as [OpenID authentication](#), [ACL support](#), [servlet authentication](#), [Kerberos authentication](#), and a plugin providing [user interface extensions](#) and security workflow.

See the [Core plugin page](#) for basic information and the [user guide](#) for detailed information.

11.4.2 Shiro

[Shiro](#) is a Java POJO-oriented security framework that provides a default domain model that models real world security. In Shiro you extend a controller base class called `JSecAuthBase` in each controller you add an `accessControl` block to setup the roles. An example below:

```
class ExampleController extends JSecAuthBase {
    static accessControl = {
        // All actions require the 'Observer' role.
        role(name: 'Observer')

        // The 'edit' action requires the 'Administrator' role.
        role(name: 'Administrator', action: 'edit')

        // Alternatively, several actions can be specified.
        role(name: 'Administrator', only: [ 'create', 'edit', 'save', 'update' ])
    }
    ...
}
```

For more information on the Shiro plugin refer to the [documentation](#).

12 Plugins

Grails is first and foremost a web application framework, but it is also a platform. By exposing a number of things from the command line interface to the runtime configuration engine, Grails can be customised. On this platform, all you need to do is create a plugin.

Extending the platform may sound complicated, but plugins can range from trivially simple to incredibly complex. In a Grails application, you'll know how to create a plugin for [sharing a data model](#) or some static resources.

12.1 Creating and Installing Plugins

Creating Plugins

Creating a Grails plugin is a simple matter of running the command:

```
grails create-plugin [PLUGIN NAME]
```

This will create a plugin project for the name you specify. For example running `grails create-plugin example` will create a plugin project called `example`.

The structure of a Grails plugin is very nearly the same as a Grails application project's except that in the root of the plugin directory you will find a plugin Groovy file called the "plugin descriptor".

Being a regular Grails project has a number of benefits in that you can immediately test your plugin by running the application.

```
grails run-app
```

The plugin descriptor name ends with the convention `GrailsPlugin` and is found in the root of the plugin directory.

```
class ExampleGrailsPlugin {
    def version = "0.1"
    ...
}
```

All plugins must have this class in the root of their directory structure. The plugin class defines the version of the plugin and optionally various hooks into plugin extension points (covered shortly).

You can also provide additional information about your plugin using several special properties:

- `title` - short one-sentence description of your plugin
- `version` - The version of your plugin. Valid values include example "0.1", "0.2-SNAPSHOT", "1.1"
- `grailsVersion` - The version of version range of Grails that the plugin supports. eg. "1.2 > *" (in
- `author` - plugin author's name
- `authorEmail` - plugin author's contact e-mail
- `description` - full multi-line description of plugin's features
- `documentation` - URL of the plugin's documentation

Here is an example from the [Quartz Grails plugin](#):

```
class QuartzGrailsPlugin {
    def version = "0.1"
    def grailsVersion = "1.1 > *"
    def author = "Sergey Nebolsin"
    def authorEmail = "nebolsin@gmail.com"
    def title = "Quartz Plugin"
    def description = '''\
The Quartz plugin allows your Grails application to schedule jobs\
to be executed using a specified interval or cron expression. The\
underlying system uses the Quartz Enterprise Job Scheduler configured\
via Spring, but is made simpler by the coding by convention paradigm.\
'''
    def documentation = "http://grails.org/plugin/quartz"

    ...
}
```

Installing and Distributing Plugins

To distribute a plugin you navigate to its root directory in a console and run:

```
grails package-plugin
```

This will create a zip file of the plugin starting with `grails-` then the plugin name and version. For ex earlier this would be `grails-example-0.1.zip`. The `package-plugin` command will also contains machine-readable information about plugin's name, version, author, and so on.

Once you have a plugin distribution file you can navigate to a Grails project and run:

```
grails install-plugin /path/to/grails-example-0.1.zip
```

If the plugin is hosted on an HTTP server you can install it with:

```
grails install-plugin http://myserver.com/plugins/grails-example-0.1.zip
```

Notes on excluded Artefacts

Although the [create-plugin](#) command creates certain files for you so that the plugin can be run as a Grails application, the following is a list of artefacts created, but not included by [package](#):

- `grails-app/conf/BootStrap.groovy`
- `grails-app/conf/BuildConfig.groovy` (although it is used to generate dependencies)
- `grails-app/conf/Config.groovy`
- `grails-app/conf/DataSource.groovy` (and any other `*DataSource.groovy`)
- `grails-app/conf/UrlMappings.groovy`
- `grails-app/conf/spring/resources.groovy`
- Everything within `/web-app/WEB-INF`
- Everything within `/web-app/plugins/**`
- Everything within `/test/**`
- SCM management files within `**/.svn/**` and `**/CVS/**`

If you need artefacts within `WEB-INF` it is recommended you use the `_Install.groovy` script (once the plugin is installed, to provide such artefacts. In addition, although `UrlMappings.groovy` is excluded, you can provide a `UrlMappings` definition with a different name, such as `MyPluginUrlMappings.groovy`.

Specifying Plugin Locations

An application can load plugins from anywhere on the file system, even if they have not been installed. This is done by specifying the plugin location in the application's `grails-app/conf/BuildConfig.groovy` file:

```
// Useful to test plugins you are developing.
grails.plugin.location.shiro =
    "/home/dilbert/dev/plugins/grails-shiro"

// Useful for modular applications where all plugins and
// applications are in the same directory.
grails.plugin.location.'grails-ui' = "../grails-grails-ui"
```

This is particularly useful in two cases:

- You are developing a plugin and want to test it in a real application without packaging and installing it
- You have split an application into a set of plugins and an application, all in the same "super-project" context

Global plugins

Plugins can also be installed globally for all applications for a particular version of Grails using the `-global` option:

```
grails install-plugin webtest -global
```

The default location is `$USER_HOME/.grails/<grailsVersion>/global-plugins` but this `grails.global.plugins.dir` setting in `BuildConfig.groovy`.

12.2 Plugin Repositories

Distributing Plugins in the Grails Central Plugins Repository

The preferred way to distribute plugin is to publish to the official Grails Plugins Repository. This will make the `grails list-plugins` command:

```
grails list-plugins
```

which lists all plugins in the Grails Plugin repository, and also the [plugin-info](#) command:

```
grails plugin-info [plugin-name]
```

which outputs more information based on the meta info entered into the plugin descriptor.



If you have created a Grails plugin and want it to be hosted in the central repository take a look at the [release-plugin](#) command which details how to release your plugin.

When you have access to the Grails Plugin repository, execute the `release-plugin` command to release your plugin.

```
grails release-plugin
```

This will automatically commit changes to SVN, create tags, and make your changes available to the [list-plugins](#) command.

Configuring Additional Repositories

The process for configuring repositories in Grails differs between versions. For version of Grails 1.2 and below see the [documentation](#) on the subject. The following sections cover Grails 1.3 and above.

Grails 1.3 and above use Ivy under the hood to resolve plugin dependencies. The mechanism for defining the same as [defining repositories for JAR dependencies](#). For example you can define a remote Maven repository using the following syntax in `grails-app/conf/BuildConfig.groovy`:

```
repositories {  
    mavenRepo "http://repository.codehaus.org"  
}
```

You can also define a SVN-based Grails repository (such as the one hosted at <http://plugins.grails.org>) using the `svnRepo` method.

```
repositories {
    grailsRepo "http://myserver/mygrailsrepo"
}
```

There is a shortcut to setup the Grails central repository:

```
repositories {
    grailsCentral()
}
```

The order in which plugins are resolved is based on the ordering of the repositories. So in this case the Grails central repository is the last:

```
repositories {
    grailsRepo "http://myserver/mygrailsrepo"
    grailsCentral()
}
```

All of the above examples use HTTP; however you can specify any [Ivy resolver](#) to resolve plugins with. For example:

```
def sshResolver = new SshResolver(user:"myuser", host:"myhost.com")
sshResolver.addArtifactPattern(
    "/path/to/repo/grails-[artifact]/tags/" +
    "LATEST_RELEASE/grails-[artifact]-[revision].[ext]")
sshResolver.latestStrategy =
    new org.apache.ivy.plugins.latest.LatestTimeStrategy()

sshResolver.changingPattern = ".*SNAPSHOT"
sshResolver.setCheckmodified(true)
```

The above example defines an artifact pattern which tells Ivy how to resolve a plugin zip file. For a more complete example, see the [relevant section](#) in the Ivy user guide.

Publishing to Maven Compatible Repositories

In general it is recommended for Grails 1.3 and above to use standard Maven-style repositories to self-manage the ability for existing tooling and repository managers to interpret the structure of a Maven repository. Grails repositories are not tied to SVN as Grails repositories are.

You use the Maven publisher plugin to publish a plugin to a Maven repository. Please refer to the section on publishing to the subject.

Publishing to Grails Compatible Repositories

Specify the `grails.plugin.repos.distribution.myRepository` setting within the `grails.config` file to publish a Grails plugin to a Grails-compatible repository:

```
grails.plugin.repos.distribution.myRepository =  
    "https://svn.codehaus.org/grails/trunk/grails-test-plugin-repo"
```

You can also provide this settings in the `$USER_HOME/.grails/settings.groovy` file if you prefer to store settings in a file.

Once this is done use the `repository` argument of the `release-plugin` command to specify the repository.

```
grails release-plugin -repository = myRepository
```

12.3 Understanding a Plugin's Structure

As mentioned previously, a plugin is basically a regular Grails application with a plugin descriptor. However, a plugin differs slightly. For example, take a look at this plugin directory structure:

```
+ grails-app  
  + controllers  
  + domain  
  + taglib  
  etc.  
+ lib  
+ src  
  + java  
  + groovy  
+ web-app  
  + js  
  + css
```

When a plugin is installed the contents of the `grails-app` directory will go into `plugins/example-1.0/grails-app`. They **will not** be copied into the main source tree. A plugin is a primary source tree.

Dealing with static resources is slightly different. When developing a plugin, just like an application, a plugin has a `resources` directory. You can then link to static resources just like in an application. This example links to a JavaScript resource.

```
<g:resource dir="js" file="mycode.js" />
```

When you run the plugin in development mode the link to the resource will resolve to something like `/plugins/example-1.0/js/mycode.js`. When the plugin is installed into an application the path will automatically change to something like `/plugin/example-1.0/js/mycode.js`. Grails will deal with making sure the resources are in the right place.

There is a special `pluginContextPath` variable that can be used whilst both developing the plugin and running the application to find out what the correct path to the plugin is.

At runtime the `pluginContextPath` variable will either evaluate to an empty string or `/plugins` if the plugin is running standalone or has been installed in an application.

Java and Groovy code that the plugin provides within the `lib` and `src/java` and `src/groovy` directories will be copied to the project's `web-app/WEB-INF/classes` directory so that they are made available at runtime.

12.4 Providing Basic Artefacts

Adding a new Script

A plugin can add a new script simply by providing the relevant Gant script in its scripts directory:

```
+ MyPlugin.groovy
+ scripts      <-- additional scripts here
+ grails-app
  + controllers
  + services
  + etc.
+ lib
```

Adding a new grails-app artifact (Controller, Tag Library, Service, etc.)

A plugin can add new artifacts by creating the relevant file within the `grails-app` tree. Note that the plugin artifacts are not copied into the main application tree.

```
+ ExamplePlugin.groovy
+ scripts
+ grails-app
  + controllers <-- additional controllers here
  + services <-- additional services here
  + etc. <-- additional XXX here
+ lib
```

Providing Views, Templates and View resolution

When a plugin provides a controller it may also provide default views to be rendered. This is an excellent way to provide views through plugins. Grails' view resolution mechanism will first look for the view in the application it is installed in, then look for the view within the plugin. This means that you can override views provided by a plugin by providing views in the application's `grails-app/views` directory.

For example, consider a controller called `BookController` that's provided by an 'amazon' plugin. If the controller calls `render` it will first look for a view called `grails-app/views/book/list.gsp` then if that fails it will look for a view in the plugin.

However if the view uses templates that are also provided by the plugin then the following syntax may be used:

```
<g:render template="fooTemplate" plugin="amazon"/>
```

Note the usage of the `plugin` attribute, which contains the name of the plugin where the template resides. Grails will then look for the template relative to the application.

Excluded Artefacts

By default Grails excludes the following files during the packaging process:

- `grails-app/conf/BootStrap.groovy`
- `grails-app/conf/BuildConfig.groovy` (although it is used to generate dependencies)
- `grails-app/conf/Config.groovy`
- `grails-app/conf/DataSource.groovy` (and any other `*DataSource.groovy`)
- `grails-app/conf/UrlMappings.groovy`
- `grails-app/conf/spring/resources.groovy`
- Everything within `/web-app/WEB-INF`
- Everything within `/web-app/plugins/**`
- Everything within `/test/**`
- SCM management files within `**/.svn/**` and `**/CVS/**`

If your plugin requires files under the `web-app/WEB-INF` directory it is recommended to use the `scripts/_Install.groovy` Gant script to install these artefacts into the target project's directory tree.

In addition, the default `UrlMappings.groovy` file is excluded to avoid naming conflicts, however a custom definition under a different name which **will** be included. For example a file called `grails-app/conf/CustomUrlMappings.groovy` will be included.

The list of excludes is extensible with the `pluginExcludes` property:

```
// resources that are excluded from plugin packaging
def pluginExcludes = [
    "grails-app/views/error.gsp"
]
```

This is useful for example to include demo or test resources in the plugin repository, but not include them in the final application.

12.5 Evaluating Conventions

Before looking at providing runtime configuration based on conventions you first need to understand how the `GrailsApplication` interface works. Every plugin has an implicit `application` variable which is an instance of the [GrailsApplication](#) interface.

The `GrailsApplication` interface provides methods to evaluate the conventions within the project's artifact classes within your application.

Artifacts implement the [GrailsClass](#) interface, which represents a Grails resource such as a controller or service. Here are some of the things you can do with `GrailsClass` instances you can do:

```
for (grailsClass in application.allClasses) {
    println grailsClass.name
}
```

`GrailsApplication` has a few "magic" properties to narrow the type of artefact you are interested in. Here are some of the things you can use:

```
for (controllerClass in application.controllerClasses) {
    println controllerClass.name
}
```

The dynamic method conventions are as follows:

- `*Classes` - Retrieves all the classes for a particular artefact name. For example `application.controllerClasses`
- `get*Class` - Retrieves a named class for a particular application. `application.getControllerClass("PersonController")`
- `is*Class` - Returns true if the given class is of the given application. `application.isControllerClass(PersonController)`

The `GrailsClass` interface has a number of useful methods that let you further evaluate and work with

- `getPropertyValue` - Gets the initial value of the given property on the class
- `hasProperty` - Returns true if the class has the specified property
- `newInstance` - Creates a new instance of this class.
- `getName` - Returns the logical name of the class in the application without the trailing convention postfix
- `getShortName` - Returns the short name of the class without package prefix
- `getFullName` - Returns the full name of the class in the application with the trailing convention postfix
- `getPropertyName` - Returns the name of the class as a property name
- `getLogicalPropertyName` - Returns the logical property name of the class in the application applicable
- `getNaturalName` - Returns the name of the property in natural terms (eg. 'lastName' becomes 'Last Name')
- `getPackageName` - Returns the package name

For a full reference refer to the [javadoc API](#).

12.6 Hooking into Build Events

Post-Install Configuration and Participating in Upgrades

Grails plugins can do post-install configuration and participate in application upgrade process (the [upgrade](#) two specially named scripts under the `scripts` directory of the plugin - `_Install.groovy` and `_Upgrade.groovy`). `_Install.groovy` is executed after the plugin has been installed and `_Upgrade.groovy` is executed on application (but not the plugin) with [upgrade](#) command.

These scripts are [Gant](#) scripts, so you can use the full power of Gant. An addition to the standard `pluginBasedir` variable which points at the plugin installation basedir.

As an example this `_Install.groovy` script will create a new directory type under the `grails-app` template:

```
ant.mkdir(dir: "${basedir}/grails-app/jobs")
ant.copy(file: "${pluginBasedir}/src/samples/SamplePluginConfig.groovy",
        todir: "${basedir}/grails-app/conf")
```

Scripting events

It is also possible to hook into command line scripting events. These are events triggered during execution. For example, you can hook into status update output (i.e. "Tests passed", "Server running") and the creation of a plugin. A plugin just has to provide an `_Events.groovy` script to listen to the required events. Refer the documentation for further information.

12.7 Hooking into Runtime Configuration

Grails provides a number of hooks to leverage the different parts of the system and perform runtime configuration.

Hooking into the Grails Spring configuration

First, you can hook in Grails runtime configuration by providing a property called `doWithSpring` via the `GrailsPlugin` interface. For example the following snippet is from one of the core Grails plugins that provides [i18n](#) support:

```
import org.springframework.web.servlet.i18n.CookieLocaleResolver
import org.springframework.web.servlet.i18n.LocaleChangeInterceptor
import org.springframework.context.support.ReloadableResourceBundleMessageSource

class I18nGrailsPlugin {
    def version = "0.1"

    def doWithSpring = {
        messageSource(ReloadableResourceBundleMessageSource) {
            basename = "WEB-INF/grails-app/i18n/messages"
        }
        localeChangeInterceptor(LocaleChangeInterceptor) {
            paramName = "lang"
        }
        localeResolver(CookieLocaleResolver)
    }
}
```

This plugin configures the Grails `messageSource` bean and a couple of other beans to manage Locale. See the [Spring Bean Builder](#) syntax to do so.

Participating in web.xml Generation

Grails generates the `WEB-INF/web.xml` file at load time, and although plugins cannot change this file, they can participate in the generation of the file. A plugin can provide a `doWithWebDescriptor` property that is assigned to the `web.xml` as an `XmlSlurper` `GPathResult`.

Add servlet and servlet-mapping

Consider this example from the `ControllersPlugin`:

```

def doWithWebDescriptor = { webXml ->
def mappingElement = webXml.'servlet-mapping'
def lastMapping = mappingElement[mappingElement.size() - 1]
  lastMapping + {
    'servlet-mapping' {
      'servlet-name' ("grails")
      'url-pattern' ("*.dispatch")
    }
  }
}

```

Here the plugin gets a reference to the last `<servlet-mapping>` element and appends Grails' service programmatically modify XML using closures and blocks.

Add filter and filter-mapping

Adding a filter with its mapping works a little differently. The location of the `<filter>` element doesn't it's simplest to insert your custom filter definition immediately after the last `<context-param>` element the usual approach is to add it immediately after the last `<filter>` element like so:

```

def doWithWebDescriptor = { webXml ->
def contextParam = webXml.'context-param'
contextParam[contextParam.size() - 1] + {
  'filter' {
    'filter-name' ('springSecurityFilterChain')
    'filter-class' (DelegatingFilterProxy.name)
  }
}

def filter = webXml.'filter'
filter[filter.size() - 1] + {
  'filter-mapping' {
    'filter-name' ('springSecurityFilterChain')
    'url-pattern' ('/*')
  }
}
}

```

In some cases you need to ensure that your filter comes after one of the standard Grails filters, such as the SiteMesh filter. Fortunately you can insert filter mappings immediately after the standard ones (more a web.xml file) like so:

```

def doWithWebDescriptor = { webXml ->
    ...

    // Insert the Spring Security filter after the Spring
    // character encoding filter.
    def filter = webXml.'filter-mapping'.find {
        it.'filter-name'.text() == "charEncodingFilter"
    }

    filter + {
        'filter-mapping' {
            'filter-name'('springSecurityFilterChain')
            'url-pattern'('/')
        }
    }
}

```

Doing Post Initialisation Configuration

Sometimes it is useful to be able to do some runtime configuration after the Spring [ApplicationContext](#) has been created. This can be done using the `doWithApplicationContext` closure property.

```

class SimplePlugin {
    def name = "simple"
    def version = "1.1"

    def doWithApplicationContext = { appCtx ->
        def sessionFactory = appCtx.sessionFactory
        // do something here with session factory
    }
}

```

12.8 Adding Dynamic Methods at Runtime

The Basics

Grails plugins let you register dynamic methods with any Grails-managed or other class at `doWithDynamicMethods` closure.

For Grails-managed classes like controllers, tag libraries and so forth you can add methods, constructors and other dynamic behavior by accessing each controller's [MetaClass](#):

```

class ExamplePlugin {
    def doWithDynamicMethods = { applicationContext ->
        for (controllerClass in application.controllerClasses) {
            controllerClass.metaClass.myNewMethod = {-> println "hello world" }
        }
    }
}

```

In this case we use the implicit application object to get a reference to all of the controller classes' `MetaClass` and add a `myNewMethod` to each controller. If you know beforehand the class you wish to add a method to, you can use the `metaClass` property.

For example we can add a new method `swapCase` to `java.lang.String`:

```
class ExamplePlugin {
  def doWithDynamicMethods = { applicationContext ->
    String.metaClass.swapCase = {->
      def sb = new StringBuilder()
      delegate.each {
        sb << (Character.isUpperCase(it as char) ?
              Character.toLowerCase(it as char) :
              Character.toUpperCase(it as char))
      }
      sb.toString()
    }
  }
  assert "UpAndDown" == "uPaNDdOWN".swapCase()
}
```

Interacting with the ApplicationContext

The `doWithDynamicMethods` closure gets passed the Spring `ApplicationContext` instance. The objects within it. For example if you were implementing a method to interact with Hibernate you could use a combination with a `HibernateTemplate`:

```
import org.springframework.orm.hibernate3.HibernateTemplate

class ExampleHibernatePlugin {
  def doWithDynamicMethods = { applicationContext ->
    for (domainClass in application.domainClasses) {
      domainClass.metaClass.static.load = { Long id->
        def sf = applicationContext.sessionFactory
        def template = new HibernateTemplate(sf)
        template.load(delegate, id)
      }
    }
  }
}
```

Also because of the autowiring and dependency injection capability of the Spring container you can replace constructors that use the application context to wire dependencies into your object at runtime:

```
class MyConstructorPlugin {
  def doWithDynamicMethods = { applicationContext ->
    for (domainClass in application.domainClasses) {
      domainClass.metaClass.constructor = {->
        return applicationContext.getBean(domainClass.name)
      }
    }
  }
}
```

Here we actually replace the default constructor with one that looks up prototyped Spring beans instead!

12.9 Participating in Auto Reload Events

Monitoring Resources for Changes

Often it is valuable to monitor resources for changes and perform some action when they occur. This is often done by reloading of application state at runtime. For example, consider this simplified snippet from the Grails Services plugin:

```
class ServicesGrailsPlugin {
    ...
    def watchedResources = "file:./grails-app/services/*Service.groovy"
    ...
    def onChange = { event ->
        if (event.source) {
            def serviceClass = application.addServiceClass(event.source)
            def serviceName = "${serviceClass.propertyName}"
            def beans = beans {
                "$serviceName"(serviceClass.getClass()) { bean ->
                    bean.autowire = true
                }
            }
            if (event.ctx) {
                event.ctx.registerBeanDefinition(
                    serviceName,
                    beans.getBeanDefinition(serviceName))
            }
        }
    }
}
```

First it defines `watchedResources` as either a String or a List of strings that contain either the ref to watch. If the watched resources specify a Groovy file, when it is changed it will automatically be reloaded. The `onChange` closure in the event object.

The event object defines a number of useful properties:

- `event.source` - The source of the event, either the reloaded Class or a Spring Resource
- `event.ctx` - The Spring ApplicationContext instance
- `event.plugin` - The plugin object that manages the resource (usually this)
- `event.application` - The GrailsApplication instance
- `event.manager` - The GrailsPluginManager instance

These objects are available to help you apply the appropriate changes based on what changed. In the "Services" plugin, the bean is re-registered with the ApplicationContext when one of the service classes changes.

Influencing Other Plugins

In addition to reacting to changes, sometimes a plugin needs to "influence" another.

Take for example the Services and Controllers plugins. When a service is reloaded, unless you reload the controllers, you will get an exception when you try to auto-wire the reloaded service into an older controller Class.

To get around this, you can specify which plugins another plugin "influences". This means that when one plugin is reloaded, it will also reload its influenced plugins. For example consider this snippet from the ServicesGrailsPlugin:


```
def influences = ['controllers']
```

Observing other plugins

If there is a particular plugin that you would like to observe for changes but not necessary watch the res "observe" property:

```
def observe = ["controllers"]
```

In this case when a controller is changed you will also receive the event chained from the controllers plugin.

It is also possible for a plugin to observe all loaded plugins by using a wildcard:

```
def observe = ["*"]
```

The Logging plugin does exactly this so that it can add the `log` property back to *any* artefact that changes.

12.10 Understanding Plugin Load Order

Controlling Plugin Dependencies

Plugins often depend on the presence of other plugins and can adapt depending on the presence of properties. The first is called `dependsOn`. For example, take a look at this snippet from the Hibernate plugin.

```
class HibernateGrailsPlugin {  
  def version = "1.0"  
  def dependsOn = [dataSource: "1.0",  
                   domainClass: "1.0",  
                   i18n: "1.0",  
                   core: "1.0"]  
}
```

The Hibernate plugin is dependent on the presence of four plugins: the `dataSource`, `domainClass`, `i18n`, and `core`.

The dependencies will be loaded before the Hibernate plugin and if all dependencies do not load, then the plugin will not load.

The `dependsOn` property also supports a mini expression language for specifying version ranges. A few examples are shown below:

```
def dependsOn = [foo: "* > 1.0"]  
def dependsOn = [foo: "1.0 > 1.1"]  
def dependsOn = [foo: "1.0 > *"]
```

When the wildcard `*` character is used it denotes "any" version. The expression syntax also excludes any version less than the specified version. So for example the expression `"1.0 > 1.1"` would match any of the following versions:

- 1.1
- 1.0
- 1.0.1
- 1.0.3-SNAPSHOT
- 1.1-BETA2

Controlling Load Order

Using `dependsOn` establishes a "hard" dependency in that if the dependency is not resolved, the plugin will fail. It is also possible though to have a weaker dependency using the `loadAfter` property:

```
def loadAfter = ['controllers']
```

Here the plugin will be loaded after the `controllers` plugin if it exists, otherwise it will just be loaded. The presence of the other plugin, for example the Hibernate plugin has this code in its `doWithSpring` closure:

```
if (manager?.hasGrailsPlugin("controllers")) {
    openSessionInViewInterceptor(OpenSessionInViewInterceptor) {
        flushMode = HibernateAccessor.FLUSH_MANUAL
        sessionFactory = sessionFactory
    }
    grailsUrlHandlerMapping.interceptors << openSessionInViewInterceptor
}
```

Here the Hibernate plugin will only register an `OpenSessionInViewInterceptor` if the `controllerManager` variable is an instance of the [GrailsPluginManager](#) interface and it provides methods to interact with the plugin.

Scopes and Environments

It's not only plugin load order that you can control. You can also specify which environments your plugin applies to (stages of a build). Simply declare one or both of these properties in your plugin descriptor:

```
def environments = ['development', 'test', 'myCustomEnv']
def scopes = [excludes:'war']
```

In this example, the plugin will only load in the 'development' and 'test' environments. Nor will it be packaged for the 'war' phase. This allows development-only plugins to not be packaged for production.

The full list of available scopes are defined by the enum [BuildScope](#), but here's a summary:

- `test` - when running tests
- `functional-test` - when running functional tests
- `run` - for `run-app` and `run-war`
- `war` - when packaging the application as a WAR file
- `all` - plugin applies to all scopes (default)

Both properties can be one of:

- a string - a sole inclusion
- a list - a list of environments or scopes to include
- a map - for full control, with 'includes' and/or 'excludes' keys that can have string or list values

For example,

```
def environments = "test"
```

will only include the plugin in the test environment, whereas

```
def environments = ["development", "test"]
```

will include it in both the development *and* test environments. Finally,

```
def environments = [includes: ["development", "test"]]
```

will do the same thing.

12.11 The Artefact API

You should by now understand that Grails has the concept of artefacts: special types of classes that it knows normal Groovy and Java classes, for example by enhancing them with extra properties and methods. classes and controllers. What you may not be aware of is that Grails allows application and plugin infrastructure for artefacts, which means you can find out what artefacts are available and even enhance your own custom artefact types.

12.11.1 Asking About Available Artefacts

As a plugin developer, it can be important for you to find out about what domain classes, controllers, or other application. For example, the [Searchable plugin](#) needs to know what domain classes exist so it can check and index the appropriate ones. So how does it do it? The answer lies with the `grailsApplication` object that's available automatically in controllers and GSPs and can be [injected](#) everywhere else.

The `grailsApplication` object has several important properties and methods for querying artefacts. that gives you all the classes of a particular artefact type:

```
for (cls in grailsApplication.<artefactType>Classes) {  
    ...  
}
```

In this case, `artefactType` is the property name form of the artefact type. With core Grails you have:

- `domain`
- `controller`
- `tagLib`
- `service`
- `codec`
- `bootstrap`
- `urlMappings`

So for example, if you want to iterate over all the domain classes, you use:

```
for (cls in grailsApplication.domainClasses) {  
    ...  
}
```

and for URL mappings:

```
for (cls in grailsApplication.urlMappingsClasses) {  
    ...  
}
```

You need to be aware that the objects returned by these properties are not instances of [Class](#). Instead, they are instances of `GrailsClass`, which has some particularly useful properties and methods, including one for the underlying `Class`:

- `shortName` - the class name of the artefact without the package (equivalent of `Class.simpleName`)
- `logicalPropertyName` - the artefact name in property form without the 'type' suffix. So `MyGreeter` would be `greeter`.
- `isAbstract()` - a boolean indicating whether the artefact class is abstract or not.
- `getPropertyValue(name)` - returns the value of the given property, whether it's a static or a non-static property. If the property is initialised on declaration, e.g. `static transactional = true`.

The artefact API also allows you to fetch classes by name and check whether a class is an artefact:

- `get<type>Class(String name)`
- `is<type>Class(Class clazz)`

The first method will retrieve the `GrailsClass` instance for the given name, e.g. `'MyGreatController'`. is a particular type of artefact. For example, `grailsApplication.isControllerClass(org.example.MyGreatController)` to check in fact a controller.

12.11.2 Adding Your Own Artefact Types

Plugins can easily provide their own artefacts so that they can easily find out what implementations are available. What you need to do is create an `ArtefactHandler` implementation and register it in your main plugin class

```
class MyGrailsPlugin {
    def artefacts = [ org.somewhere.MyArtefactHandler ]
    ...
}
```

The `artefacts` list can contain either handler classes (as above) or instances of handlers.

So, what does an artefact handler look like? Well, put simply it is an implementation of the [ArtefactHandler](#) interface. There is a skeleton implementation that can readily be extended: [ArtefactHandlerAdapter](#).

In addition to the handler itself, every new artefact needs a corresponding wrapper class that implements `GrailsClass`. Implementations are available such as [AbstractInjectableGrailsClass](#), which is particularly useful as it is auto-wired, just like controllers and services.

The best way to understand how both the handler and wrapper classes work is to look at the Quartz plugin:

- [GrailsJobClass](#)
- [DefaultGrailsJobClass](#)
- [JobArtefactHandler](#)

Another example is the [Shiro plugin](#) which adds a realm artefact.

12.12 Binary Plugins

Regular Grails plugins are packaged as zip files containing the full source of the plugin. This has some disadvantages (anyone can see the source), in addition to avoiding problems with the source compatibility

As of Grails 2.0 you can pre-compile Grails plugins into regular JAR files known as "binary plugins". This has some disadvantages (as discussed in the advantages of source plugins above) including:

- Binary plugins can be published as standard JAR files to a Maven repository
- Binary plugins can be declared like any other JAR dependency
- Commercial plugins are more viable since the source isn't published
- IDEs have a better understanding since binary plugins are regular JAR files containing classes

Packaging

To package a plugin in binary form you can use the `package-plugin` command and the `--binary` flag:

```
grails package-plugin --binary
```

Supported artefacts include:

- Grails artifact classes such as controllers, domain classes and so on
- I18n Message bundles
- GSP Views, layouts and templates

You can also specify the packaging in the plugin descriptor:

```
def packaging = "binary"
```

in which case the packaging will default to binary.

Using Binary Plugins

The packaging process creates a JAR file in the `target` directory of the plugin, for example `target` two ways to incorporate a binary plugin into an application.

One is simply placing the plugin JAR file in your application's `lib` directory. The other is to publish to a repository and declare it as a dependency in `grails-app/conf/BuildConfig.groovy`:

```
dependencies {  
    compile "mycompany:myplugin:0.1"  
}
```



Since binary plugins are packaged as JAR files, they are declared as dependencies in the `dependencies` block *not* in the `plugins` block as you may be naturally inclined to do. The `plugins` block is for traditional source plugins packaged as zip files.

13 Web Services

Web services are all about providing a web API onto your web application and are typically implemented :

13.1 REST

REST is not really a technology in itself, but more an architectural pattern. REST is very simple and just a communication medium, combined with URL patterns that are "representational" of the underlying system. PUT, POST and DELETE.

Each HTTP method maps to an action type. For example GET for retrieving data, PUT for creating data, POST for updating data and DELETE for deleting data. In this sense REST fits quite well with [CRUD](#).

URL patterns

The first step to implementing REST with Grails is to provide RESTful [URL mappings](#):

```
static mappings = {  
    "/product/$id?"(resource:"product")  
}
```

This maps the URI `/product` onto a `ProductController`. Each HTTP method such as GET, PUT, POST and DELETE is mapped to an action within the controller as outlined by the table below:

Method	Action
GET	show
PUT	update
POST	save
DELETE	delete

In addition, Grails provides automatic XML or JSON marshalling for you.

You can alter how HTTP methods are handled by using URL Mappings to [map to HTTP methods](#):

```
"/product/$id"(controller: "product") {  
    action = [GET: "show", PUT: "update", DELETE: "delete", POST: "save"]  
}
```

However, unlike the `resource` argument used previously, in this case Grails will not provide automatic marshalling. You must specify the `parseRequest` argument:

```
"/product/$id"(controller: "product", parseRequest: true) {  
    action = [GET: "show", PUT: "update", DELETE: "delete", POST: "save"]  
}
```

HTTP Methods

In the previous section you saw how you can easily define URL mappings that map specific HTTP methods. Writing a REST client that then sends a specific HTTP method is then easy (example in Groovy's HTTPBuilder).

```
import groovyx.net.http.*
import static groovyx.net.http.ContentType.JSON

def http = new HTTPBuilder("http://localhost:8080/amazon")

http.request(Method.GET, JSON) {
    url.path = '/book/list'
    response.success = { resp, json ->
        for (book in json.books) {
            println book.title
        }
    }
}
```

Issuing a request with a method other than GET or POST from a regular browser is not possible without some [form](#) you can specify an alternative method such as DELETE:

```
<g:form controller="book" method="DELETE">
    ..
</g:form>
```

Grails will send a hidden parameter called `_method`, which will be used as the request's HTTP method. For non-browser clients is to use the X-HTTP-Method-Override to specify the alternative method.

XML Marshalling - Reading

The controller can use Grails' [XML marshalling](#) support to implement the GET method:

```
import grails.converters.XML

class ProductController {
    def show() {
        if (params.id && Product.exists(params.id)) {
            def p = Product.findByName(params.id)
            render p as XML
        }
        else {
            def all = Product.list()
            render all as XML
        }
    }
    ..
}
```

If there is an `id` we search for the `Product` by name and return it, otherwise we return all `Products`. This way, if we go to `/product/MacBook` we only get a `MacBook`.

XML Marshalling - Updating

To support updates such as PUT and POST you can use the [params](#) object which Grails enhances with packet. Given an incoming XML packet of:

```
<?xml version="1.0" encoding="ISO-8859-1"?>
<product>
  <name>MacBook</name>
  <vendor id="12">
    <name>Apple</name>
  </vendor>
</product>
```

you can read this XML packet using the same techniques described in the [Data Binding](#) section, using the :

```
def save() {
  def p = new Product(params.product)

  if (p.save()) {
    render p as XML
  }
  else {
    render p.errors
  }
}
```

In this example by indexing into the `params` object using the `product` key we can automatically `Product` constructor. An interesting aspect of the line:

```
def p = new Product(params.product)
```

is that it requires no code changes to deal with a form submission that submits form data, or an XML request.



If you require different responses to different clients (REST, HTML etc.) you can use [content negotiation](#).

The `Product` object is then saved and rendered as XML, otherwise an error message is produced using the `render` method:

```
<error>
  <message>The property 'title' of class 'Person' must be specified</message>
</error>
```

REST with JAX-RS

There also is a [JAX-RS Plugin](#) which can be used to build web services based on the Java API for RESTful Web Services.

13.2 SOAP

There are several plugins that add SOAP support to Grails depending on your preferred approach. For C/S, the [Spring WS](#) plugin, whilst if you want to generate a SOAP API from Grails services there are several plugins available.

- [CXF](#) plugin which uses the [CXF](#) SOAP stack
- [Axis2](#) plugin which uses [Axis2](#)
- [Metro](#) plugin which uses the [Metro](#) framework (and can also be used for [Contract First](#))

Most of the SOAP integrations integrate with Grails [services](#) via the `exposes` static property. This example

```
class BookService {
    static expose = ['cxf']

    Book[] getBooks() {
        Book.list() as Book[]
    }
}
```

The WSDL can then be accessed at the location: `http://127.0.0.1:8080/your_grails_app/`

For more information on the CXF plugin refer to [the documentation](#) on the wiki.

13.3 RSS and Atom

No direct support is provided for RSS or Atom within Grails. You could construct RSS or ATOM capability. There is however a [Feeds plugin](#) available for Grails that provides a RSS and Atom builder example of its usage can be seen below:

```
def feed() {
    render(feedType: "rss", feedVersion: "2.0") {
        title = "My test feed"
        link = "http://your.test.server/yourController/feed"

        for (article in Article.list()) {
            entry(article.title) {
                link = "http://your.test.server/article/${article.id}"
                article.content // return the content
            }
        }
    }
}
```

14 Grails and Spring

This section is for advanced users and those who are interested in how Grails integrates with and builds on for [plugin developers](#) considering doing runtime configuration Grails.

14.1 The Underpinnings of Grails

Grails is actually a [Spring MVC](#) application in disguise. Spring MVC is the Spring framework's built Although Spring MVC suffers from some of the same difficulties as frameworks like Struts in terms of its architected and was, for Grails, the perfect framework to build another framework on top of.

Grails leverages Spring MVC in the following areas:

- Basic controller logic - Grails subclasses Spring's [DispatcherServlet](#) and uses it to delegate to Grails [controllers](#)
- Data Binding and Validation - Grails' [validation](#) and [data binding](#) capabilities are built on those provided by Spring
- Runtime configuration - Grails' entire runtime convention based system is wired together by a Spring [ApplicationContext](#)
- Transactions - Grails uses Spring's transaction management in [GORM](#)

In other words Grails has Spring embedded running all the way through it.

The Grails ApplicationContext

Spring developers are often keen to understand how the Grails ApplicationContext instance is constructed.

- Grails constructs a parent ApplicationContext from the web-app/WEB-INF/applicationContext.xml. This ApplicationContext configures the [GrailsApplication](#) instance and the [GrailsPluginManager](#).
- Using this ApplicationContext as a parent Grails' analyses the conventions with the GrailsPluginManager and creates a child ApplicationContext that is used as the root ApplicationContext of the web application.

Configured Spring Beans

Most of Grails' configuration happens at runtime. Each [plugin](#) may configure Spring beans that are registered with the Spring container. For a reference as to which beans are configured, refer to the reference guide which describes each of the beans.

14.2 Configuring Additional Beans

Using the Spring Bean DSL

You can easily register new (or override existing) beans by configuring them in `grails-app/conf/spring` files. These files use the Grails [Spring DSL](#). Beans are defined inside a `beans` property (a Closure):

```
beans = {  
    // beans here  
}
```

As a simple example you can configure a bean with the following syntax:

```
import my.company.MyBeanImpl

beans = {
    myBean(MyBeanImpl) {
        someProperty = 42
        otherProperty = "blue"
    }
}
```

Once configured, the bean can be auto-wired into Grails artifacts and other classes that support `Bootstrap.groovy` and integration tests) by declaring a public field whose name is your bean's name

```
class ExampleController {
    def myBean
    ...
}
```

Using the DSL has the advantage that you can mix bean declarations and logic, for example based on the [environment](#)

```
import grails.util.Environment
import my.company.mock.MockImpl
import my.company.MyBeanImpl

beans = {
    switch(Environment.current) {
        case Environment.PRODUCTION:
            myBean(MyBeanImpl) {
                someProperty = 42
                otherProperty = "blue"
            }
            break
        case Environment.DEVELOPMENT:
            myBean(MockImpl) {
                someProperty = 42
                otherProperty = "blue"
            }
            break
    }
}
```

The `GrailsApplication` object can be accessed with the `application` variable and can be used (amongst other things):

```

import grails.util.Environment
import my.company.mock.MockImpl
import my.company.MyBeanImpl

beans = {
    if (application.config.my.company.mockService) {
        myBean(MockImpl) {
            someProperty = 42
            otherProperty = "blue"
        }
    } else {
        myBean(MyBeanImpl) {
            someProperty = 42
            otherProperty = "blue"
        }
    }
}

```



If you define a bean in `resources.groovy` with the same name as one previously registered, your bean will replace the previous registration. This is a convenient way without resorting to editing plugin code or other approaches that would affect maintainability.

Using XML

Beans can also be configured using a `grails-app/conf/spring/resources.xml`. In earlier versions of Grails, this file was automatically generated for you by the `run-app` script, but the DSL in `resources.groovy` is now automatically generated. But it is still supported - you just need to create it yourself.

This file is typical Spring XML file and the Spring documentation has an [excellent reference](#) on how to create it. The `myBean` bean that we configured using the DSL would be configured with this syntax in the XML file:

```

<bean id="myBean" class="my.company.MyBeanImpl">
    <property name="someProperty" value="42" />
    <property name="otherProperty" value="blue" />
</bean>

```

Like the other bean it can be auto-wired into any class that supports dependency injection:

```

class ExampleController {
    def myBean
}

```

Referencing Existing Beans

Beans declared in `resources.groovy` or `resources.xml` can reference other beans by their name. For example, if the `BookService` class has a bean named `bookService`, your bean would reference it like this:

```
beans = {
    myBean(MyBeanImpl) {
        someProperty = 42
        otherProperty = "blue"
        bookService = ref("bookService")
    }
}
```

or like this in XML:

```
<bean id="myBean" class="my.company.MyBeanImpl">
    <property name="someProperty" value="42" />
    <property name="otherProperty" value="blue" />
    <property name="bookService" ref="bookService" />
</bean>
```

The bean needs a public setter for the bean reference (and also the two simple properties), which in Groovy:

```
package my.company

class MyBeanImpl {
    Integer someProperty
    String otherProperty
    BookService bookService // or just "def bookService"
}
```

or in Java like this:

```
package my.company;

class MyBeanImpl {

    private BookService bookService;
    private Integer someProperty;
    private String otherProperty;

    public void setBookService(BookService theBookService) {
        this.bookService = theBookService;
    }

    public void setSomeProperty(Integer someProperty) {
        this.someProperty = someProperty;
    }

    public void setOtherProperty(String otherProperty) {
        this.otherProperty = otherProperty;
    }
}
```

Using `ref` (in XML or the DSL) is very powerful since it configures a runtime reference, so the reference is in place when the final application context configuration occurs, everything will be resolved correctly.

For a full reference of the available beans see the plugin reference in the reference guide.

14.3 Runtime Spring with the Beans DSL

This Bean builder in Grails aims to provide a simplified way of wiring together dependencies that uses Spring.

In addition, Spring's regular way of configuration (via XML and annotations) is static and difficult to modify. Grails' [BeanBuilder](#) class changes programmatic XML creation which is both error prone and verbose. Grails' [BeanBuilder](#) change programmatically wire together components at runtime, allowing you to adapt the logic based on system parameters.

This enables the code to adapt to its environment and avoids unnecessary duplication of code (having development and production environments).

The BeanBuilder class

Grails provides a [grails.spring.BeanBuilder](#) class that uses dynamic Groovy to construct bean definitions.

```
import org.apache.commons.dbcp.BasicDataSource
import org.codehaus.groovy.grails.orm.hibernate.ConfigurableLocalSessionFactoryBean
import org.springframework.context.ApplicationContext
import grails.spring.BeanBuilder

def bb = new BeanBuilder()

bb.beans {
  dataSource(BasicDataSource) {
    driverClassName = "org.h2.Driver"
    url = "jdbc:h2:mem:grailsDB"
    username = "sa"
    password = ""
  }

  sessionFactory(ConfigurableLocalSessionFactoryBean) {
    dataSource = ref('dataSource')
    hibernateProperties = [ "hibernate.hbm2ddl.auto": "create-drop",
                          "hibernate.show_sql": "true" ]
  }
}

ApplicationContext appContext = bb.createApplicationContext()
```



Within [plugins](#) and the [grails-app/conf/spring/resources.groovy](#) file you don't need to create a [BeanBuilder](#). Instead the DSL is implicitly available inside the `doWithSpring` and `doWithSpringBlock` methods respectively.

This example shows how you would configure Hibernate with a data source with the [BeanBuilder](#) class.

Each method call (in this case `dataSource` and `sessionFactory` calls) maps to the name of the bean. The first argument of the method is the bean's class, whilst the last argument is a block. Within the body of the block you can set properties using Groovy syntax.

Bean references are resolved automatically using the name of the bean. This can be seen in the `sessionFactory` bean which resolves the `dataSource` reference.

Certain special properties related to bean management can also be set by the builder, as seen in the following example.

```

 sessionFactory(ConfigurableLocalSessionFactoryBean) { bean ->
    // Autowiring behaviour. The other option is 'byType'. [autowire]
    bean.autowire = 'byName'
    // Sets the initialisation method to 'init'. [init-method]
    bean.initMethod = 'init'
    // Sets the destruction method to 'destroy'. [destroy-method]
    bean.destroyMethod = 'destroy'
    // Sets the scope of the bean. [scope]
    bean.scope = 'request'
    dataSource = ref('dataSource')
    hibernateProperties = [ "hibernate.hbm2ddl.auto": "create-drop",
                          "hibernate.show_sql":      "true" ]
  }

```

The strings in square brackets are the names of the equivalent bean attributes in Spring's XML definition.

Using BeanBuilder with Spring MVC

Include the `grails-spring-<version>.jar` file in your classpath to use BeanBuilder in a regular following `<context-param>` values to your `/WEB-INF/web.xml` file:

```

<context-param>
  <param-name>contextConfigLocation</param-name>
  <param-value>/WEB-INF/applicationContext.groovy</param-value>
</context-param>

<context-param>
  <param-name>contextClass</param-name>
  <param-value>
    org.codehaus.groovy.grails.commons.spring.GrailsWebApplicationContext
  </param-value>
</context-param>

```

Then create a `/WEB-INF/applicationContext.groovy` file that does the rest:

```

import org.apache.commons.dbcp.BasicDataSource

beans {
  dataSource(BasicDataSource) {
    driverClassName = "org.h2.Driver"
    url = "jdbc:h2:mem:grailsDB"
    username = "sa"
    password = ""
  }
}

```

Loading Bean Definitions from the File System

You can use the BeanBuilder class to load external Groovy scripts that define beans using the same example:

```

def bb = new BeanBuilder()
bb.loadBeans("classpath:*SpringBeans.groovy")

def applicationContext = bb.createApplicationContext()

```


Here the BeanBuilder loads all Groovy files on the classpath ending with `SpringBeans.groovy`. An example script can be seen below:

```
import org.apache.commons.dbcp.BasicDataSource
import org.codehaus.groovy.grails.orm.hibernate.ConfigurableLocalSessionFactoryBean

beans {
    dataSource(BasicDataSource) {
        driverClassName = "org.h2.Driver"
        url = "jdbc:h2:mem:grailsDB"
        username = "sa"
        password = ""
    }
    sessionFactory(ConfigurableLocalSessionFactoryBean) {
        dataSource = dataSource
        hibernateProperties = [ "hibernate.hbm2ddl.auto": "create-drop",
                             "hibernate.show_sql": "true" ]
    }
}
```

Adding Variables to the Binding (Context)

If you're loading beans from a script you can set the binding to use by creating a Groovy Binding:

```
def binding = new Binding()
binding.maxSize = 10000
binding.productGroup = 'finance'

def bb = new BeanBuilder()
bb.binding = binding
bb.loadBeans("classpath:*SpringBeans.groovy")

def ctx = bb.createApplicationContext()
```

Then you can access the `maxSize` and `productGroup` properties in your DSL files.

14.4 The BeanBuilder DSL Explained

Using Constructor Arguments

Constructor arguments can be defined using parameters to each bean-defining method. Put them after the f

```
bb.beans {
    exampleBean(MyExampleBean, "firstArgument", 2) {
        someProperty = [1, 2, 3]
    }
}
```

This configuration corresponds to a `MyExampleBean` with a constructor that looks like this:

```
MyExampleBean(String foo, int bar) {
    ...
}
```

Configuring the BeanDefinition (Using factory methods)

The first argument to the closure is a reference to the bean configuration instance, which you can use to any method on the [AbstractBeanDefinition](#) class:

```
bb.beans {
  exampleBean(MyExampleBean) { bean ->
    bean.factoryMethod = "getInstance"
    bean.singleton = false
    someProperty = [1, 2, 3]
  }
}
```

As an alternative you can also use the return value of the bean defining method to configure the bean:

```
bb.beans {
  def example = exampleBean(MyExampleBean) {
    someProperty = [1, 2, 3]
  }
  example.factoryMethod = "getInstance"
}
```

Using Factory beans

Spring defines the concept of factory beans and often a bean is created not directly from a new instance of a class, but from a factory bean. In this case the bean has no Class argument and instead you must pass the name of the factory bean.

```
bb.beans {
  myFactory(ExampleFactoryBean) {
    someProperty = [1, 2, 3]
  }
  myBean(myFactory) {
    name = "blah"
  }
}
```

Another common approach is provide the name of the factory method to call on the factory bean. The parameter syntax:

```
bb.beans {
  myFactory(ExampleFactoryBean) {
    someProperty = [1, 2, 3]
  }
  myBean(myFactory: "getInstance") {
    name = "blah"
  }
}
```

Here the `getInstance` method on the `ExampleFactoryBean` bean will be called to create the `myBean`.

Creating Bean References at Runtime

Sometimes you don't know the name of the bean to be created until runtime. In this case you can use `ref` by defining method dynamically:

```
def beanName = "example"
bb.beans {
  "${beanName}Bean" (MyExampleBean) {
    someProperty = [1, 2, 3]
  }
}
```

In this case the `beanName` variable defined earlier is used when invoking a bean defining method. This would work just as well with a name that is generated programmatically based on configuration, system properties, etc.

Furthermore, because sometimes bean names are not known until runtime you may need to reference the beans, in this case using the `ref` method:

```
def beanName = "example"
bb.beans {
  "${beanName}Bean" (MyExampleBean) {
    someProperty = [1, 2, 3]
  }
  anotherBean(AnotherBean) {
    example = ref("${beanName}Bean")
  }
}
```

Here the `example` property of `AnotherBean` is set using a runtime reference to the `exampleBean`. This is done by passing the bean name to `ref` from a parent `ApplicationContext` that is provided in the constructor of the `BeanBuilder`.

```
ApplicationContext parent = ...//
def bb = new BeanBuilder(parent)
bb.beans {
  anotherBean(AnotherBean) {
    example = ref("${beanName}Bean", true)
  }
}
```

Here the second parameter `true` specifies that the reference will look for the bean in the parent context.

Using Anonymous (Inner) Beans

You can use anonymous inner beans by setting a property of the bean to a block that takes an argument that

```

bb.beans {
  marge(Person) {
    name = "Marge"
    husband = { Person p ->
      name = "Homer"
      age = 45
      props = [overweight: true, height: "1.8m"]
    }
    children = [bart, lisa]
  }

  bart(Person) {
    name = "Bart"
    age = 11
  }

  lisa(Person) {
    name = "Lisa"
    age = 9
  }
}

```

In the above example we set the marge bean's husband property to a block that creates an inner bean factory bean you can omit the type and just use the specified bean definition instead to setup the factory:

```

bb.beans {
  personFactory(PersonFactory)
  marge(Person) {
    name = "Marge"
    husband = { bean ->
      bean.factoryBean = "personFactory"
      bean.factoryMethod = "newInstance"
      name = "Homer"
      age = 45
      props = [overweight: true, height: "1.8m"]
    }
    children = [bart, lisa]
  }
}

```

Abstract Beans and Parent Bean Definitions

To create an abstract bean definition define a bean without a Class parameter:

```

class HolyGrailQuest {
  def start() { println "lets begin" }
}

```

```

class KnightOfTheRoundTable {
  String name
  String leader
  HolyGrailQuest quest

  KnightOfTheRoundTable(String name) {
    this.name = name
  }

  def embarkOnQuest() {
    quest.start()
  }
}

```

```

import grails.spring.BeanBuilder

def bb = new BeanBuilder()
bb.beans {
  abstractBean {
    leader = "Lancelot"
  }
  ...
}

```

Here we define an abstract bean that has a `leader` property with the value of "Lancelot". To use the child bean:

```

bb.beans {
  ...
  quest(HolyGrailQuest)

  knights(KnightOfTheRoundTable, "Camelot") { bean ->
    bean.parent = abstractBean
    quest = ref('quest')
  }
}

```



When using a parent bean you must set the parent property of the bean before setting any other property of the bean!

If you want an abstract bean that has a `Class` specified you can do it this way:

```
import grails.spring.BeanBuilder

def bb = new BeanBuilder()
bb.beans {

    abstractBean(KnightOfTheRoundTable) { bean ->
        bean.'abstract' = true
        leader = "Lancelot"
    }

    quest(HolyGrailQuest)

    knights("Camelot") { bean ->
        bean.parent = abstractBean
        quest = quest
    }
}
```

In this example we create an abstract bean of type `KnightOfTheRoundTable` and use the bean argument to create a knights bean that has no Class defined, but inherits the Class from the parent bean.

Using Spring Namespaces

Since Spring 2.0, users of Spring have had easier access to key features via XML namespaces. You can use them by declaring it with this syntax:

```
xmlns context:"http://www.springframework.org/schema/context"
```

and then invoking a method that matches the names of the Spring namespace tag and its associated attributes:

```
context.'component-scan'('base-package': "my.company.domain")
```

You can do some useful things with Spring namespaces, such as looking up a JNDI resource:

```
xmlns jee:"http://www.springframework.org/schema/jee"
jee.'jndi-lookup'(id: "dataSource", 'jndi-name': "java:comp/env/myDataSource")
```

This example will create a Spring bean with the identifier `dataSource` by performing a JNDI lookup. In addition to using namespaces you also get full access to all of the powerful AOP support in Spring from BeanBuilder. For example:

```
class Person {
    int age
    String name

    void birthday() {
        ++age;
    }
}
```

```

class BirthdayCardSender {
    List peopleSentCards = []
    void onBirthday(Person person) {
        peopleSentCards << person
    }
}

```

You can define an aspect that uses a pointcut to detect whenever the `birthday()` method is called:

```

xmlns aop:"http://www.springframework.org/schema/aop"

fred(Person) {
    name = "Fred"
    age = 45
}

birthdayCardSenderAspect(BirthdayCardSender)

aop {
    config("proxy-target-class": true) {
        aspect(id: "sendBirthdayCard", ref: "birthdayCardSenderAspect") {
            after method: "onBirthday",
            pointcut: "execution(void ..Person.birthday()) and this(person)"
        }
    }
}

```

14.5 Property Placeholder Configuration

Grails supports the notion of property placeholder configuration through an extended version of Spring's is typically useful in combination with [externalized configuration](#).

Settings defined in either [ConfigSlurper](#) scripts or Java properties files can be used as placeholder in `grails-app/conf/spring/resources.xml`. For example given the following entries in `grails` an externalized config):

```

database.driver="com.mysql.jdbc.Driver"
database.dbname="mysql:mydb"

```

You can then specify placeholders in `resources.xml` as follows using the familiar `${..}` syntax:

```

<bean id="dataSource"
    class="org.springframework.jdbc.datasource.DriverManagerDataSource">
    <property name="driverClassName">
        <value>${database.driver}</value>
    </property>
    <property name="url">
        <value>jdbc:${database.dbname}</value>
    </property>
</bean>

```

14.6 Property Override Configuration

Grails supports setting of bean properties via [configuration](#). This is often useful when used in combination

You define a beans block with the names of beans and their values:

```
beans {
    bookService {
        webserviceURL = "http://www.amazon.com"
    }
}
```

The general format is:

```
[bean name].[property name] = [value]
```

The same configuration in a Java properties file would be:

```
beans.bookService.webserviceURL=http://www.amazon.com
```


15 Grails and Hibernate

If [GORM](#) (Grails Object Relational Mapping) is not flexible enough for your liking you can alternate with Hibernate, either with XML mapping files or JPA annotations. You will be able to map Grails domain classes and have more flexibility in the creation of your database schema. Best of all, you will still be able to use all the query methods provided by GORM!

15.1 Using Hibernate XML Mapping Files

Mapping your domain classes with XML is pretty straightforward. Simply create a `hibernate` directory in `grails-app/conf`, either manually or with the [create-hibernate-cfg-xml](#) command:

```
<?xml version='1.0' encoding='UTF-8'?>
<!DOCTYPE hibernate-configuration PUBLIC
    "-//Hibernate/Hibernate Configuration DTD 3.0//EN"
    "http://hibernate.sourceforge.net/hibernate-configuration-3.0.dtd">
<hibernate-configuration>
    <session-factory>
        <!-- Example mapping file inclusion -->
        <mapping resource="org.example.Book.hbm.xml"/>
        ...
    </session-factory>
</hibernate-configuration>
```

The individual mapping files, like `org.example.Book.hbm.xml` in the above example, also go into the `hibernate` directory. To find out how to map domain classes with XML, check out the [Hibernate manual](#).

If the default location of the `hibernate.cfg.xml` file doesn't suit you, you can change it by editing `grails-app/conf/DataSource.groovy`:

```
hibernate {
    config.location = "file:/path/to/my/hibernate.cfg.xml"
}
```

or even a list of locations:

```
hibernate {
    config.location = [ "file:/path/to/one/hibernate.cfg.xml",
                       "file:/path/to/two/hibernate.cfg.xml" ]
}
```

Grails also lets you write your domain model in Java or reuse an existing one that already has Hibernate mapping files into `grails-app/conf/hibernate` and either put the Java files in `src/java` or the classpath. If the domain model is packaged as a JAR. You still need the `hibernate.cfg.xml` though!

15.2 Mapping with Hibernate Annotations

To map a domain class with annotations, create a new class in `src/java` and use the annotations to define the domain. For more info on this see the [Hibernate Annotations Docs](#):

```

package com.books;

import javax.persistence.Entity;
import javax.persistence.GeneratedValue;
import javax.persistence.Id;

@Entity
public class Book {
    private Long id;
    private String title;
    private String description;
    private Date date;

    @Id
    @GeneratedValue
    public Long getId() {
        return id;
    }

    public void setId(Long id) {
        this.id = id;
    }

    public String getTitle() {
        return title;
    }

    public void setTitle(String title) {
        this.title = title;
    }

    public String getDescription() {
        return description;
    }

    public void setDescription(String description) {
        this.description = description;
    }
}

```

Then register the class with the Hibernate sessionFactory by adding `grails-app/conf/hibernate/hibernate.cfg.xml` file as follows:

```

<!DOCTYPE hibernate-configuration SYSTEM
    "http://hibernate.sourceforge.net/hibernate-configuration-3.0.dtd">
<hibernate-configuration>
    <session-factory>
        <mapping package="com.books" />
        <mapping class="com.books.Book" />
    </session-factory>
</hibernate-configuration>

```

See the previous section for more information on the `hibernate.cfg.xml` file.

When Grails loads it will register the necessary dynamic methods with the class. To see what else you can do, see the section on [Scaffolding](#).

15.3 Adding Constraints

You can still use GORM validation even if you use a Java domain model. Grails lets you define constraints in the `src/java` directory. The script must be in a directory that matches the package of the corresponding domain class. For example, if you had a domain class `org.example.Book`, the constraints script would be `src/java/org/example/BookConstraints.groovy`.

Add a standard GORM constraints block to the script:

```
constraints = {  
    title blank: false  
    author blank: false  
}
```

Once this is in place you can validate instances of your domain class! .

16 Scaffolding

Scaffolding lets you auto-generate a whole application for a given domain class including:

- The necessary [views](#)
- Controller actions for create/read/update/delete (CRUD) operations

Dynamic Scaffolding

The simplest way to get started with scaffolding is to enable it with the `scaffold` property. Set the `scaffold` property to `true` for the `Book` domain class:

```
class BookController {  
    static scaffold = true  
}
```

This works because the `BookController` follows the same naming convention as the `Book` domain class. We could reference the class directly in the `scaffold` property:

```
class SomeController {  
    static scaffold = Author  
}
```

With this configured, when you start your application the actions and views will be auto-generated and dynamically implemented by default by the runtime scaffolding mechanism:

- list
- show
- edit
- delete
- create
- save
- update

A CRUD interface will also be generated. To access this open `http://localhost:8080/app/book`

If you prefer to keep your domain model in Java and [mapped with Hibernate](#) you can still use scaffolding, its name as the `scaffold` argument.

You can add new actions to a scaffolded controller, for example:

```
class BookController {
  static scaffold = Book

  def changeAuthor() {
    def b = Book.get(params.id)
    b.author = Author.get(params["author.id"])
    b.save()

    // redirect to a scaffolded action
    redirect(action:show)
  }
}
```

You can also override the scaffolded actions:

```
class BookController {
  static scaffold = Book

  // overrides scaffolded action to return both authors and books
  def list() {
    [bookInstanceList: Book.list(),
     bookInstanceTotal: Book.count(),
     authorInstanceList: Author.list()]
  }

  def show() {
    def book = Book.get(params.id)
    log.error(book)
    [bookInstance : book]
  }
}
```

All of this is what is known as "dynamic scaffolding" where the CRUD interface is generated dynamically



By default, the size of text areas in scaffolded views is defined in the CSS, so adding 'rows' attributes have no effect.

Also, the standard scaffold views expect model variables of the form <propertyName> collections and <propertyName>Instance for single instances. It's tempting to use properties like 'book', but those won't work.

Customizing the Generated Views

The views adapt to [Validation constraints](#). For example you can change the order that fields appear in the builder:

```
def constraints = {
  title()
  releaseDate()
}
```

You can also get the generator to generate lists instead of text inputs if you use the `inList` constraint:

```
def constraints = {
  title()
  category(inList: ["Fiction", "Non-fiction", "Biography"])
  releaseDate()
}
```

Or if you use the `range` constraint on a number:

```
def constraints = {
  age(range:18..65)
}
```

Restricting the size with a constraint also effects how many characters can be entered in the generated view

```
def constraints = {
  name(size:0..30)
}
```

Static Scaffolding

Grails also supports "static" scaffolding.

The above scaffolding features are useful but in real world situations it's likely that you will want to cus you generate a controller and the views used to create the above interface from the command line. To gene

```
grails generate-controller Book
```

or to generate the views:

```
grails generate-views Book
```

or to generate everything:

```
grails generate-all Book
```

If you have a domain class in a package or are generating from a [Hibernate mapped class](#) remember to incl

```
grails generate-all com.bookstore.Book
```

Customizing the Scaffolding templates

The templates used by Grails to generate the controller and views can be customized by installing the `grails-templates` plugin. The command `grails install-plugin grails-templates` will install the plugin and the templates will be available in the `grails-templates` directory.

17 Deployment

Grails applications can be deployed in a number of ways, each of which has its pros and cons.

"grails run-app"

You should be very familiar with this approach by now, since it is the most common method of running phase. An embedded Tomcat server is launched that loads the web application from the development source changes to application files.

This approach is not recommended at all for production deployment because the performance is poor. The sizable overhead on the server. Having said that, `grails prod run-app` removes the per-request recompilation, so frequently the regular check takes place.

Setting the system property "disable.auto.recompile" to `true` disables this regular check completely, which controls the frequency. This latter property should be set to the number of seconds you want between each check.

"grails run-war"

This is very similar to the previous option, but Tomcat runs against the packaged WAR file rather than the development source, so you get good performance without the hassle of having to deploy the WAR file elsewhere.

WAR file

When it comes down to it, current java infrastructures almost mandate that web applications are deployed as WAR files. This is the common approach to Grails application deployment in production. Creating a WAR file is as simple as executing the following command:

```
grails war
```

There are also many ways in which you can customise the WAR file that is created. For example, you can specify (via a relative path) to the command that instructs it where to place the file and what name to give it:

```
grails war /opt/java/tomcat-5.5.24/foobar.war
```

Alternatively, you can add a line to `grails-app/conf/BuildConfig.groovy` that changes the default WAR file name:

```
grails.project.war.file = "foobar-prod.war"
```

Any command line argument that you provide overrides this setting.

It is also possible to control what libraries are included in the WAR file, for example to avoid conflicts v default behavior is to include in the WAR file all libraries required by Grails, plus any libraries contain libraries contained in the application's "lib" directory. As an alternative to the default behavior you can libraries to include in the WAR file by setting the property `grails.war.dependencies` in BuildCon patterns or closures containing AntBuilder syntax. Closures are invoked from within an Ant "copy" step included, whereas each item in a pattern list is included. Any closure or pattern assigned to the latter pr `grails.war.dependencies`.

Be careful with these properties: if any of the libraries Grails depends on are missing, the application example that includes a small subset of the standard Grails dependencies:

```
def deps = [
    "hibernate3.jar",
    "groovy-all-*.jar",
    "standard-${servletVersion}.jar",
    "jstl-${servletVersion}.jar",
    "oscache-*.jar",
    "commons-logging-*.jar",
    "sitemesh-*.jar",
    "spring-*.jar",
    "log4j-*.jar",
    "ognl-*.jar",
    "commons-*.jar",
    "xstream-1.2.1.jar",
    "xpp3_min-1.1.3.4.O.jar" ]

grails.war.dependencies = {
    fileset(dir: "libs") {
        for (pattern in deps) {
            include(name: pattern)
        }
    }
}
```

This example only exists to demonstrate the syntax for the properties. If you attempt to use it as is in you probably not work. You can find a list of dependencies required by Grails in the "dependencies.txt" file distribution. You can also find a list of the default dependencies included in WAR generation in `DEFAULT_DEPS` and `DEFAULT_J5_DEPS` variables.

The remaining two configuration options available to you are `grails.war.copyToWebApp` and `grails.war.resources` these lets you customise what files are included in the WAR file from the "web-app" directory. The secur want before the WAR file is finally created.

```
// This closure is passed the command line arguments used to start the
// war process.
grails.war.copyToWebApp = { args ->
    fileset(dir: "web-app") {
        include(name: "js/**")
        include(name: "css/**")
        include(name: "WEB-INF/**")
    }
}

// This closure is passed the location of the staging directory that
// is zipped up to make the WAR file, and the command line arguments.
// Here we override the standard web.xml with our own.
grails.war.resources = { stagingDir, args ->
    copy(file: "grails-app/conf/custom-web.xml",
        tofile: "${stagingDir}/WEB-INF/web.xml")
}
```

Application servers

Ideally you should be able to simply drop a WAR file created by Grails into any application server and i things are rarely ever this simple. The [Grails website](#) contains an up-to-date list of application servers that any additional steps required to get a Grails WAR file working.

Copies of this document may be made for your own use and for distribution to others, provided that you do not charge any fee for such copies and further provided that each copy contains this Copyright Notice, whether distributed in print or electronically. Sponsored by [SpringSource](#)